

SAMS
**Teach
Yourself**

THIRD EDITION
Updated to cover
Perl 5.8

templating systems

CGI scripts

server-side includes

Perl

Clinton Pierce

**in 24
Hours**



SAMS Teach Yourself Perl in 24 Hours THIRD EDITION

By [Clinton Pierce](#)

.....
Publisher: **Sams Publishing**

Pub Date: **June 15, 2005**

ISBN: **0-672-32793-7**

Pages: **480**

[Table of Contents](#) | [Index](#)

Overview

Learn Perl programming quickly and easily with 24 one-hour lessons in *Sams Teach Yourself Perl in 24 Hours, 3/e*. These step-by-step lessons will teach you the basics of Perl and how to apply it in web development and system administration. Plus, the third edition has been updated to include 5 chapters on new technologies, information on the latest version of Perl and a look ahead to Perl 6. You will focus on real-world development, including how to:

- Effectively use Perl for large development projects using Perl Modules.
- Use Perl for data processing.
- Utilize Perl as a "glue" language with other programming languages.
- Use Perl as a web development language.

The concise, modular format of the *Sams Teach Yourself Perl in 24 Hours, 3/e* is the ideal way to learn for today's working professional.



SAMS Teach Yourself Perl in 24 Hours THIRD EDITION

By [Clinton Pierce](#)

Publisher: **Sams Publishing**

Pub Date: **June 15, 2005**

ISBN: **0-672-32793-7**

Pages: **480**

[Table of Contents](#) | [Index](#)

[Copyright](#)

[About the Author](#)

[Acknowledgments](#)

[We Want to Hear from You!](#)

[Reader Services](#)

[Introduction](#)

[How to Use This Book](#)

[Conventions Used in This Book](#)

[Part I. Perl Fundamentals](#)

[Hour 1. Getting Started with Perl](#)

[Installing Perl](#)

[Documentation](#)

[Your First Program](#)

[Summary](#)

[Q&A](#)

[Workshop](#)

[Hour 2. Perl's Building Blocks: Numbers and Strings](#)

[Literals](#)

[Scalar Variables](#)

[Expressions and Operators](#)

[More Operators](#)

[Exercise: Interest Calculator](#)

[Summary](#)

[Q&A](#)

[Workshop](#)

[Hour 3. Controlling the Program's Flow](#)

[Blocks](#)

[The if Statement](#)

[Looping](#)

- Other Flow Control Tools
 - Exercise: Finding Primes
 - Summary
 - Q&A
 - Workshop
- Hour 4. Stacking Building Blocks: Lists and Arrays
 - Putting Things into Lists and Arrays
 - Getting Elements Out of an Array
 - Manipulating Arrays
 - Exercise: Playing a Little Game
 - Summary
 - Q&A
 - Workshop
- Hour 5. Working with Files
 - Opening Files
 - Reading
 - Writing
 - Free Files, Testing Files, and Binary Data
 - Summary
 - Q&A
 - Workshop
- Hour 6. Pattern Matching
 - Simple Patterns
 - The Metacharacters
 - Substitution
 - Exercise: Cleaning Up Input Data
 - Pattern Matching Odds and Ends
 - Summary
 - Q&A
 - Workshop
- Hour 7. Hashes
 - Filling Your Hash
 - Getting Data Out of a Hash
 - Lists and Hashes
 - Hash Odds and Ends
 - Useful Things to Do with a Hash
 - Exercise: Creating a Simple Customer Database with Perl
 - Summary
 - Q&A
 - Workshop
- Hour 8. Functions

- Creating and Calling Subroutines
- Scope
- Exercise: Statistics
- Function Footnotes
- Summary
- Q&A
- Workshop

— Part II. Advanced Features

- Hour 9. More Functions and Operators
 - Searching Scalars
 - Transliteration, Not Substitution
 - A Better Way to print
 - Exercise: A Formatted Report
 - New Ways with Arrays
 - Summary
 - Q&A
 - Workshop
- Hour 10. Files and Directories
 - Getting a Directory Listing
 - Exercise: The Unix grep
 - Directories
 - Unix Stuff
 - Everything You Ever Wanted to Know About THAT File
 - Exercise: Renaming Files En Masse
 - Summary
 - Q&A
 - Workshop
- Hour 11. System Interaction
 - The system() Function
 - Capturing Output
 - Pipes
 - First Lesson in Portability
 - Summary
 - Q&A
 - Workshop
- Hour 12. Using Perl's Command-Line Tools
 - What Is the Debugger?
 - Exercise: Finding the Bug
 - Other Command-Line Stuff
 - Summary
 - Q&A

- Workshop
- Hour 13. References and Structures
 - Reference Basics
 - Recipes for Structures
 - Exercise: Another Game, Maze
 - Summary
 - Q&A
- Workshop
- Hour 14. Using Modules
 - A Gentle Introduction
 - A Quick Tour
 - Full List of Standard Modules
 - Summary
 - Q&A
- Workshop
- Hour 15. Finding Permanence
 - DBM Files
 - Exercise: A Free-Form Memo Pad
 - Text Files as Databases
 - Random File Access
 - Locking
 - Summary
 - Q&A
- Workshop
- Hour 16. The Perl Community
 - What's Perl All About, Anyway?
 - The Comprehensive Perl Archive Network (CPAN)
 - Your Next Steps
 - Other Resources
 - Summary
 - Q&A
- Workshop

Part III. Applying Perl

- Hour 17. Writing Modules
 - Building a Module
 - Scoping Revisited
 - Another statement: our
 - Example: A Module to Handle Common File Information Requests
 - Summary
 - Q&A
- Workshop

- Hour 18. Object Primer
 - Classes, Properties, and Methods
 - Example: Implementation of Car in Perl
 - Example: File Information Class
 - Summary
 - Q&A
 - Workshop
- Hour 19. Data Processing
 - How to Look at Data
 - Dealing with Table Data
 - XML Data
 - Summary
 - Q&A
 - Workshop
- Hour 20. Perl as a Glue Language
 - Weather Station
 - Presenting Data as PDF
 - Reading and Writing Excel Spreadsheets
 - Summary
 - Q&A
 - Workshop
- Hour 21. Introduction to CGI
 - Browsing the Web
 - Don't Skip This Section
 - Your First CGI Program
 - What to Do When Your CGI Program Doesn't Work
 - Summary
 - Q&A
 - Workshop
- Hour 22. Basic Forms
 - How Forms Work
 - Passing Information to Your CGI Program
 - Web Security 101
 - A Guestbook
 - Summary
 - Q&A
 - Workshop
- Hour 23. Complex Forms
 - The Stateless Web
 - Hidden Fields
 - A Multipage Survey

[Summary](#)

[Q&A](#)

[Workshop](#)

[Hour 24. Manipulating HTTP and Cookies](#)

[The HTTP Conversation](#)

[More Details on Calling CGI Programs](#)

[Cookies](#)

[Long Term Cookies](#)

[Problems with Cookies](#)

[Summary](#)

[Q&A](#)

[Workshop](#)

[Part IV. Appendixes](#)

[Appendix A. Installing Modules](#)

[Picking the Right Module](#)

[Installing the Modules Under...](#)

[What to Do When You're Not Allowed to Install Modules](#)

[Index](#)

[PREV](#)

< Day Day Up >

[NEXT](#)

Copyright

Copyright © 2005 by Sams Publishing

All rights reserved. No part of this book shall be reproduced, stored in a retrieval system, or transmitted by any means, electronic, mechanical, photocopying, recording, or otherwise, without written permission from the publisher. No patent liability is assumed with respect to the use of the information contained herein. Although every precaution has been taken in the preparation of this book, the publisher and author assume no responsibility for errors or omissions. Nor is any liability assumed for damages resulting from the use of the information contained herein.

Library of Congress Catalog Card Number: 2004195543

Printed in the United States of America

First Printing: June 2005

08 07 06 05 4 3 2 1

Trademarks

All terms mentioned in this book that are known to be trademarks or service marks have been appropriately capitalized. Sams Publishing cannot attest to the accuracy of this information. Use of a term in this book should not be regarded as affecting the validity of any trademark or service mark.

Warning and Disclaimer

Every effort has been made to make this book as complete and as accurate as possible, but no warranty or fitness is implied. The information provided is on an "as is" basis. The author and the publisher shall have neither liability nor responsibility to any person or entity with respect to any loss or damages arising from the information contained in this book.

Bulk Sales

Sams Publishing offers excellent discounts on this book when ordered in quantity for bulk purchases or special sales. For more information, please contact

U.S. Corporate and Government Sales

1-800-382-3419

corpsales@pearsontechgroup.com

For sales outside of the U.S., please contact

International Sales
international@pearsoned.com

Credits

Senior Editor

Scott Meyers

Managing Editor

Charlotte Clapp

Project Editor

Andy Beaster

Copy Editor

Kate Givens

Indexer

Chris Barrick

Proofreader

Heather Arle

Technical Editor

Andy Lester

Publishing Coordinator

Vanessa Evans

Book Designer

Gary Adair

Page Layout

Kelly Maish

Dedication

To Heidi and Calvin, for not killing me this summer and supporting me during this adventure. Without your support, I'd have lost even more of my marbles.

 PREV

< Day Day Up >

NEXT 

About the Author

Clinton Pierce is a software engineer, freelance programmer, and instructor. He has been answering questions about Perl on USENET for many years and has been writing courseware and teaching Perl to his co-workers and anyone else who will listen for about as long. He is a software engineer for a payroll company, who, when not designing middleware software to integrate legacy systems to the Web, navigating the intricacies of payroll taxes, teaching UNIX and Perl, writing books and articles, or writing programs at home Just For The Fun Of It, harbors secret dreams of being abducted by wood nymphs and living in the deep forest without technology. You can visit his Web site at <http://www.geeksalad.org> for updates and corrections, or just to say hello.

Acknowledgments

"If I have seen further, it is by standing on the shoulders of Giants" —Isaac Newton.

Any book on Perl should first recognize the Giant who gave us Perl in the first place. Thanks, Larry.

I not only stood on the shoulders of Giants, but also was led by hoards of others. To make sure this book came out as error-free as possible, I invited perfect strangers to watch, criticize, and correct me when necessary. This has been a humbling experience. In no particular order, the people who annoyed me the most and who deserve the most thanks are Abigail, Greg Bacon, Sean Burke, Ken Fox, Kevin Meltzer, Tom Phoenix, and Randal Schwartz. Also, thanks to Michael Schwern, Tom Grydeland, Matt Bielanski, Mark Jason-Dominus, Jeff Pinyan, Gary Ross, Andrew Chen, and John Bell for finding bugs and offering suggestions. Over the years helpful readers have been instrumental in weeding out small problems with the text and helping with the explanations, notably: Lisa M, Robert B, Qinglin L, Ihor P, David K, Dan G, and Laura C.

Some small credit should also go to #perl for letting me sound my ideas off them and for giving me their honest—and sometimes brutal—opinions.

Somewhat responsible for this book are Bill Crawford and Donna Hinkle for getting me into this whole training mess in the first place. Thanks, I think.

Of course, I'd like to thank the people at Sams Publishing who put up with a new author and all the pain and agony that goes with the territory. Without people like Randi Roger, Scott Meyers, Kate Givens, and everyone behind the scenes that I didn't get to talk to but who are just as important, I'd never have put together anything like this.

We Want to Hear from You!

As the reader of this book, *you* are our most important critic and commentator. We value your opinion and want to know what we're doing right, what we could do better, what areas you'd like to see us publish in, and any other words of wisdom you're willing to pass our way.

You can email or write me directly to let me know what you did or didn't like about this book—as well as what we can do to make our books stronger.

Please note that I cannot help you with technical problems related to the topic of this book, and that due to the high volume of mail I receive, I might not be able to reply to every message.

When you write, please be sure to include this book's title and author as well as your name and phone or email address. I will carefully review your comments and share them with the author and editors who worked on the book.

E-mail: opensource@samspublishing.com

Mail: Mark Taber
Associate Publisher
Sams Publishing
800 East 96th Street
Indianapolis, IN 46240 USA

Reader Services

For more information about this book or another Sams title, visit our website at www.samspublishing.com. Type the ISBN (excluding hyphens) or the title of a book in the Search field to find the page you're looking for.

Introduction

"Any sufficiently advanced technology is virtually indistinguishable from magic."

—Arthur C. Clarke

Remember this well: *There's nothing magical about programming a computer.*

Like anything that seems magical, there's always a trick to it—and programming is no different. All you need are some analytical thinking skills, the desire to learn, and some time to learn Perl. The best way to learn how to program a computer—or any task really—is to have a goal. Your goal could be to spice up your Web pages, convert a program you already have to Perl, or to satisfy idle curiosity—it doesn't really matter.

Now, given that you've established a goal and have some thinking skills, what does this book have for you?

This book will teach you the basics of the Perl programming language. You'll learn just enough Perl to do something useful. You will not be overwhelmed with details that, although interesting, will only get in your way and that you'll probably never use again. Each new concept in this book is demonstrated by lots of working code examples. Go ahead—flip through the book and see.

But why Perl? Almost every company that uses programming of some kind uses Perl. Perl is found in finance, manufacturing, genetics, the military, and every other application known to humankind. And, of course, Perl is used on the Internet and the World Wide Web. Perl isn't going away any time soon, which is why learning it is an investment in your time that will pay off for years to come.

It's also possible to create really stunning Perl programs with just a small bit of code. Using a little bit of Perl to glue together other applications, languages, and technologies, you'll be creating useful Perl programs in no time.

How to Use This Book

This book is divided into 24 segments that take roughly an hour to complete. You can work through the lessons in the space of a day (if you don't plan to eat or sleep), or you can take your time and work through the hour lessons at your own pace.

At the end of each hour, you'll be able to accomplish a new set of tasks. The lessons contain clear explanations of the language features and how they work. In addition, each hour provides you with the opportunity for hands-on training, simply by following the steps described.

Conventions Used in This Book

Sams Teach Yourself Perl in 24 Hours uses a number of conventions that are consistent throughout this book:

- Each hour begins with an overview of what you will learn.
- Step-by-step instructions are preceded by a To Do icon.
- Every hour ends with a summary and a series of commonly asked questions and answers; hopefully, you'll find the answers to your questions among them.

In addition, these elements appear throughout the book:

By the Way

Notes provide you with comments and asides about the topic at hand.

Did you Know?

Tips offer shortcuts and hints on getting the task done.

Watch Out!

Cautions explain roadblocks you might encounter when you work with Perl and tell you how to avoid them.

New terms are emphasized by being placed in bold for your easy reference.

Part I: Perl Fundamentals

[HOUR 1](#) Getting Started with Perl

[HOUR 2](#) Perl's Building Blocks: Numbers and Strings

[HOUR 3](#) Controlling the Program's Flow

[HOUR 4](#) Stacking Building Blocks: Lists and Arrays

[HOUR 5](#) Working with Files

[HOUR 6](#) Pattern Matching

[HOUR 7](#) Hashes

[HOUR 8](#) Functions

Hour 1. Getting Started with Perl

What You'll Learn in This Hour

- Installing Perl
- Accessing Perl's internal documentation
- Writing your first Perl script

Perl is a general-purpose programming language. It can be used for anything that any other programming language can be used for. It has been used in every industry imaginable for almost any task you can think of. It's used on the stock market; in manufacturing, design, bioinformatics, customer support, quality control, systems programming, payroll, and inventory; and, of course, on the Web.

Perl is used in so many places because Perl is a **glue language**. A glue language is used to bind things together. You probably wouldn't want to write a word processor in Perl—although you could—because good word processors are already available. Writing a database, a spreadsheet, an operating system, or a full-featured Web server in Perl would be silly—but again, possible.

What Perl is good at is tying these elements together. Perl can take your database, convert it into a spreadsheet-ready file, and, during the processing, fix the data if you want. Perl can also take your word processing documents and convert them to HTML for display on the Web.

As a side effect of being a language designed to glue elements together, Perl is very adaptable. It runs under, at last count, about two dozen operating systems—and probably more. Perl's programming style is very flexible, so you can do the same things in many ways. Your Perl programs may look nothing like mine, but if they both work, that's okay. Perl can be a strict language when it needs to be, but it can also be forgiving to new programmers if you want. It's all up to you.

Let me just clear up a few points. The name of the programming language is *Perl*. The name of the program that runs your programs, the interpreter, is *perl*. The distinction usually isn't very important to you—except when you're trying to start your programs; then it's always *perl*. Sometimes—but not here—you will see Perl written as *PERL*, probably because one explanation of Perl's name is as an acronym for *Practical Extraction and Report Language*. Nobody ever really called it PERL; it's too pretentious. It's just Perl to its friends.

By the Way

Many of Perl's features were "borrowed" from other languages. This borrowing gave rise, early on, to Perl's other acronym-expansion: *Pathologically Eclectic Rubbish Lister*.

Installing Perl

To play with Perl, you first have to install it. Perl's installation is designed to be easy and error free. In fact, as part of the installation steps, Perl should do a self-test to make sure that it's okay. The installation procedures vary widely depending on what operating system you have. So, to get things moving, pick which operating system you have from the sections here, and follow along.

Stop! Wait! Maybe You Already Have Perl

Before you go through all the trouble to install Perl on your system, you should check to see whether you already have it. Some Unix vendors ship Perl with the operating system. Windows NT comes with Perl as part of the Windows NT Resource Kit (but that version is a bit old). To see whether you have Perl properly installed on your operating system, you need to get to a command prompt.

Under Unix, simply log in to the system. If you have a graphical environment, you need to open a terminal window. After you've logged in or opened your window, you may see a prompt like this:

\$

It might be a % or it might look like `bash%`, but either way, it is called the **shell prompt** or **command prompt**. For the first few hours, you'll need to be at this prompt to interact with Perl.

To see whether Perl is properly installed on your system, type the following (but not the \$ prompt):

```
$ perl -v
```

Either the system responds with an error message such as `command not found`, or Perl responds and prints its version number. If Perl prints its version number, it's installed. You probably do not have to reinstall it.

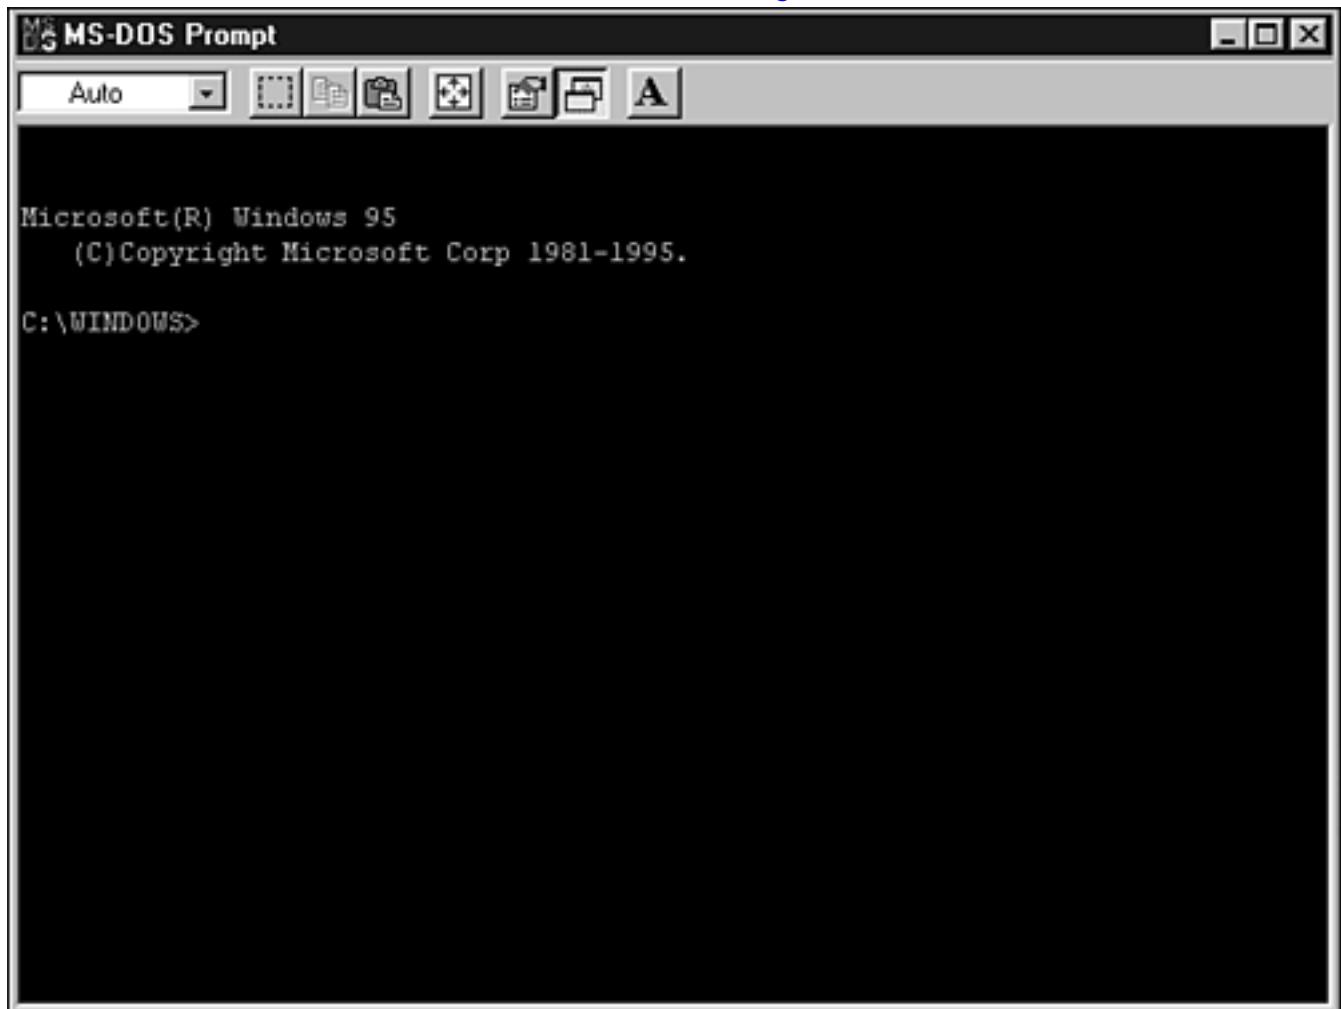
Watch Out!

The version number reported should be at *least* 5—perhaps 5.004, 5.005, 5.6, 5.8, and so on—but no lower. If Perl responds with a version of 4.x, you must install a new copy. Perl version 4 is old, buggy, and no longer maintained; few of the examples in this book will work with Perl version 4. At the time of this writing, 5.8.6 is considered the current version of Perl.

To see whether Perl is properly installed on a Windows machine you need to get to an MS-DOS prompt, similar to [Figure 1.1](#).

Figure 1.1. You can check your version of Perl at this DOS prompt.

[\[View full size image\]](#)



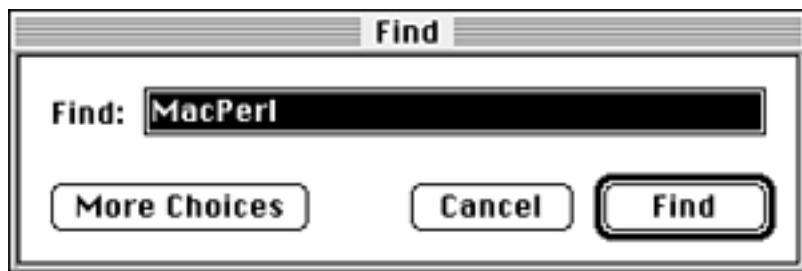
At that prompt, type the following (but don't repeat the prompt):

```
C:\> perl -v
```

If Perl is properly installed, it answers with its version number. As noted in the previous caution, it should be at least version 5. If MS-DOS replies with `Bad command or file name`, you need to install Perl properly.

On the Macintosh (running MacOS), you can check to see whether Perl is installed by running File Find (Command-f) for 'MacPerl' as shown in [Figure 1.2](#). If the application is found, open it and look at the "About MacPerl" option under the Apple menu. You should have at least Version 5.2.0 Patchlevel 5.004; otherwise, you should install a new version of MacPerl.

Figure 1.2. Looking for Perl on the Macintosh.



Newer Macintosh systems run Mac OS X. Underneath the point-and-click interface, there is a Unix operating system. To get a shell prompt (so that you can follow the Unix instructions) go to Finder, Applications, Utilities and select Terminal. You should see a window similar to [Figure 1.3](#).

Figure 1.3. Mac OS X Terminal, with Perl version output.

```
Last login: Thu Feb 10 14:05:08 on ttys01
Welcome to Darwin!
support_mac: ~ $ perl -v

This is perl, v5.8.1-RC3 built for darwin-thread-multi-2level
(with 1 registered patch, see perl -V for more detail)

Copyright 1987-2003, Larry Wall

Perl may be copied only under the terms of either the Artistic License or the
GNU General Public License, which may be found in the Perl 5 source kit.

Complete documentation for Perl, including FAQ lists, should be found on
this system using 'man perl' or 'perldoc perl'. If you have access to the
Internet, point your browser at http://www.perl.com/, the Perl Home Page.

support_mac: ~ $
```

Generally for Mac OS X, follow the Unix instructions in this book. There are extensions particular to Mac OS X that you can download from CPAN.

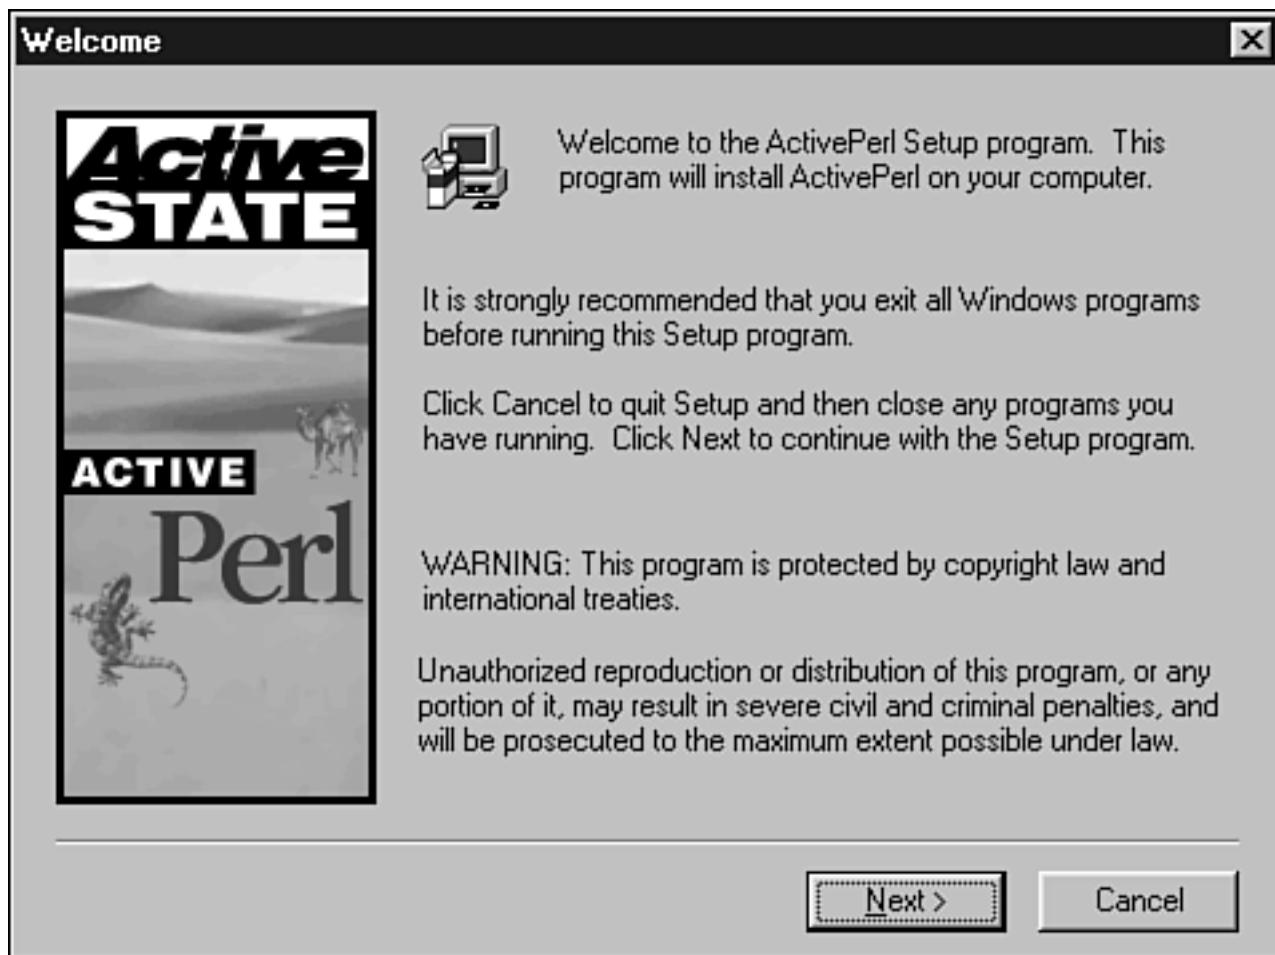
Installing Perl on Windows

To install Perl under Windows, keep in mind that, as with many things in life, you can take the Easy Way or

the Hard Way. If you're proficient with a C compiler and the tools needed in a development environment—makefiles, shells, and so on—you can take the Hard Way and build your own Perl from scratch. The source to the `perl` interpreter is free for you to look at, modify, and change to suit your needs; see Hour 16, "The Perl Community," for details. Building Perl from scratch under Windows is not easy and probably more hassle than it's worth for most people.

Installing Perl the Easy Way really is easy. ActiveState Tool Corp. provides a self-installing Perl distribution, and installation works as it does for any other Windows application, as shown in Figure 1.4. This Perl is distributed under the ActiveState Community License, which you should read. The URL is <http://www.ActiveState.com>.

Figure 1.4. Installing Perl under Windows with ActiveState.



This is exactly the same Perl you would get if you had built it yourself. ActiveState has simply done the hard parts for you and wrapped it up with an installation program. ActiveState also offers commercial support for Perl, if you require it, and offers add-on products such as debuggers and other development tools and libraries.

Did you Know?

You can go to ActiveState's Web site to get the latest version of Perl.

Installing Perl on Unix

To install Perl on Unix, you need a couple of things. First, you need a copy of the Perl source bundle. You can always download the latest bundle from the Downloads area of <http://www.perl.com>. You can find multiple versions there, but the one you want is always labeled "Stable" or "Production." You also need an ANSI C compiler. Don't worry if you don't know what that means. Perl's configuration program checks for one, and if the compiler is not there, you can install a prebuilt version.

By the Way

If your Unix has a system for installing prebuilt packages, you may be able to install a prebuilt version of Perl. Linux, Solaris, AIX, and other Unixes have prebuilt bundled versions of Perl that can be installed easily. Consult your documentation on where to obtain these packages.

After you have the Perl source bundle—which comes in a file named something like stable.tar.gz—then you need to unpack and install it. To do so, enter these commands:

```
$ gunzip stable.tar.gz  
$ tar xf stable.tar
```

These commands take a while to run. If you don't have the gunzip decompression program, you can obtain a copy from <http://www.gnu.org>. The program bundle is called gzip. When you're finished unpacking, type the following command at a prompt:

```
$ sh Configure
```

The `Configure` program starts and asks you a *lot* of questions. If you don't know the answers to most of them, that's fine; just press Enter. The default answer is usually the best one. Perl can be built without problems on almost any Unix system imaginable. When that's all done, type this command:

```
$ make
```

Building Perl takes quite a while. Get some coffee. If you have a slow system, get lunch. When the build is complete, type two more commands:

```
$ make test
```

```
# make install
```

The `make test` command ensures that Perl is 100 percent okay and ready to run. To run `make install`, you might need to be logged in as root—that's why the prompt is `#`, root's prompt, in the example—because it wants to install Perl into system directories.

When `make install` is correct, you can test Perl's installation by typing the following at the prompt again:

```
$ perl -v
```

If this command works, congratulations!

Did you Know?

You can go to <http://www.perl.com> for the latest version of Perl.

Installing Perl on Mac OS

By the Way

If your Macintosh is running the new operating system Mac OS X, it is essentially a Unix machine and you should generally follow the instructions for Unix in this book. Mac OS X ships with a recent version of Perl and should already be installed.

The latest release of the Macintosh Perl—called MacPerl—is available from the CPAN ports directory. To get the release, you need to go to <http://www.perl.com/CPAN/ports/mac> and download the installation file from there. You should download the most recent version of MacPerl appl.bin from that directory. Install it by using StuffIt Expander to extract the MacPerl installation program from the downloaded file, and then running the installation program.

When you're finished, you will want to set up a helper for the Perl documentation reader, Shuck, which was installed with MacPerl. MacOS 8 users can do this in the Internet Control Panel by selecting File Mapping from the Advanced menu and adding a file extension mapping for `.pod` to the Shuck application. This will give you easier access to the documentation. You might also want to set up mappings for `.ph`, `.pl`, `.plx`, `.pm`, `.cgi`, and `.xs` (all extensions used by Perl) to the MacPerl application. Be sure to set the file type to

'TEXT'.

MacOS 7 users will have to use the InternetConfig utility to perform similar mappings. In the InternetConfig, select Helpers, and add a new helper application [shuck](#) for [.pod](#). Also add helpers for the other extensions mentioned previously to the MacPerl application.



< Day Day Up >



Documentation

This point is important, so please pay close attention: With *every* Perl installation, you receive a *full* copy of the current documentation for the Perl language and the interpreter.

That's right. The entire set of documentation available for Perl comes with the installation bundle. You get it for free. The 5.8 distribution includes more than 1,700 pages of documentation. This documentation includes reference material, tutorials, FAQs, history, and even notes on Perl's internals.

You can access this documentation using a variety of methods. On Windows and Unix, a utility called `perldoc` is installed with Perl. You can use the `perldoc` program to search the documentation and provide formatted output for the manuals. To run `perldoc`, you need to be at a command prompt. The following example uses a Unix prompt, but a DOS command prompt is fine:

```
$ perldoc perl

PERL(1)          User Contributed Perl Documentation          PERL(1)

NAME
    perl - Practical Extraction and Report Language

SYNOPSIS
    perl [ -sTuU ]      [ -hv ]      [ -V[:configvar] ]
    [ -cw ]      [ -d[:debugger] ]  [ -D[number/list] ]
    [ -pna ]     [ -Fpattern ]   [ -l[octal] ]  [ -O[octal] ]
    [ -I[dir] ]   [ -m[-]module ] [ -M[-]'module...' ]
    [ -P ]        [ -S ]        [ -x[dir] ]
    [ -i[extension] ] [ -e 'command' ] [ ? ]
    [ programfile ] [ argument ]...
```

For ease of access, the Perl manual has been split up into a number of sections:

:

The manual sections are divided into pieces with names like `perlfunc` (Perl Functions), `perlop` (Perl Operators), and `perlfaq` (Perl FAQ). To access the `perlfunc` manual page, you enter the command `perldoc perlfunc`. The names of all the manual sections are listed in the `perldoc perl` manual page.

To search the manual for a function name, you can run the `perldoc` utility with a `-tf` switch. The following example finds the manual page for Perl's `print` function:

```
$ perldoc -tf print
```

The FAQs are Frequently Asked Questions about Perl. They are questions that people learning Perl ask over and over again. To save people time and trouble, the questions have been collected into files called FAQs. To search an FAQ for a keyword, you should use the `-q` switch along with a word that might appear in the FAQ's title. For example, if you want to know about Perl support, you could use the following query:

```
$ perldoc -q support
```

In this case, the entry for the FAQ questions "Who supports Perl? Who develops it? Why is it free?" is displayed.

Some Special Documentation Cases

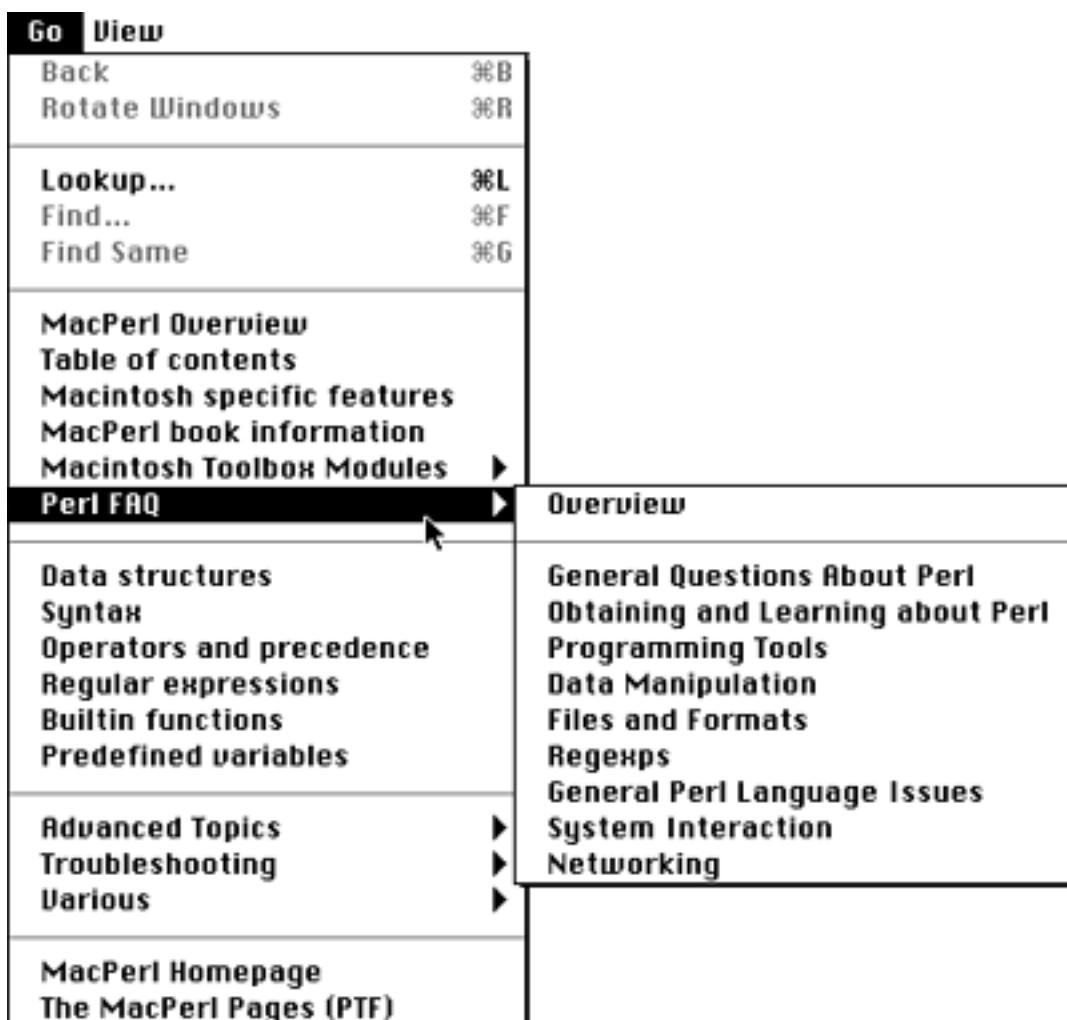
When Perl is installed on a Unix system, the installer is given the option to install manual pages in traditional "man" format. If the installer chooses yes, the standard Perl documentation is converted into man format and stored in an appropriate place. To access the Perl documentation, you can use either the `perldoc` program or the `man` program as you normally would with Unix:

```
$ man perl
```

When ActiveState's Perl distribution is installed on a Microsoft Windows system, the manual pages are converted to HTML format and can be accessed with a Web browser. If you want to read the manuals, point your frame-capable Web browser to the local directory `C:\Perl\html` if you used the standard installation directory; if you didn't, use your selected directory instead.

For the Macintosh (not Mac OS X), MacPerl comes with a utility called `Shuck` that is located in the MacPerl folder. You can use it to read and search the Perl documentation as shown in Figure 1.5.

Figure 1.5. The MacPerl Shuck documentation reader.



What If You Can't Find the Documentation?

If you can't find the documentation, only two things can be wrong. The first is you're not looking in the right places. Perhaps the `perldoc` utility is installed in a directory that's not in your shell's search path. Or perhaps your search path is overly restrictive. Hunt around for the `perldoc` utility, and add that directory to your shell's `PATH` environment variable.

The second reason is the documentation was removed, either by accident or malice. Perl's installation includes documentation. You can't install Perl without it. If the documentation isn't there, a good argument can be made that Perl wasn't properly installed or that it has since been corrupted. Perhaps you—or the system administrator—should consider reinstalling Perl. The documentation is an integral part of the Perl development environment, and without it some pieces of Perl will not function.

If all else fails, and you can't get a local copy of the documentation, you can fall back to the Web. On Perl's primary distribution site (<http://www.perl.com>), you can access the standard set of documentation. Having the actual documentation that came with your version of Perl would be better—it's tailor-made to your particular version and installation—but this online documentation will do in a pinch.

Your First Program

To write your Perl programs, you need a utility called a **text editor**. A text editor is a program that allows you to enter plain text, without any formatting, into a file. Microsoft Windows Notepad and the MS-DOS EDIT. EXE program are both examples of text editors. For Unix, vi, Emacs, and pico are all text editors. At least one of them will be available on your system.

Under the Mac, the MacPerl application contains a rudimentary text editor; to open a new program, select New under the File menu. For Mac OS X use theTextEdit program. This is located in the Finder under Applications. Be sure to set the preferences to "plain text" or select "Make Plain Text" under the Format menu.

You should *not* use a word processor to type in your Perl programs. Word processors—such as Microsoft Word, WordPad, WordPerfect, and so on—embed formatting codes by default when they save documents, even if the documents contain no boldface, italic, font changes, or other special formatting. These formatting codes will confuse Perl, and your programs will not work correctly. If you need to use a word processor, be sure to save your programs as plain text.

Typing Your First Program

Open your text editor, and type the following Perl program exactly as shown:

```
#!/usr/bin/perl  
  
print "Hello, World!\n";
```

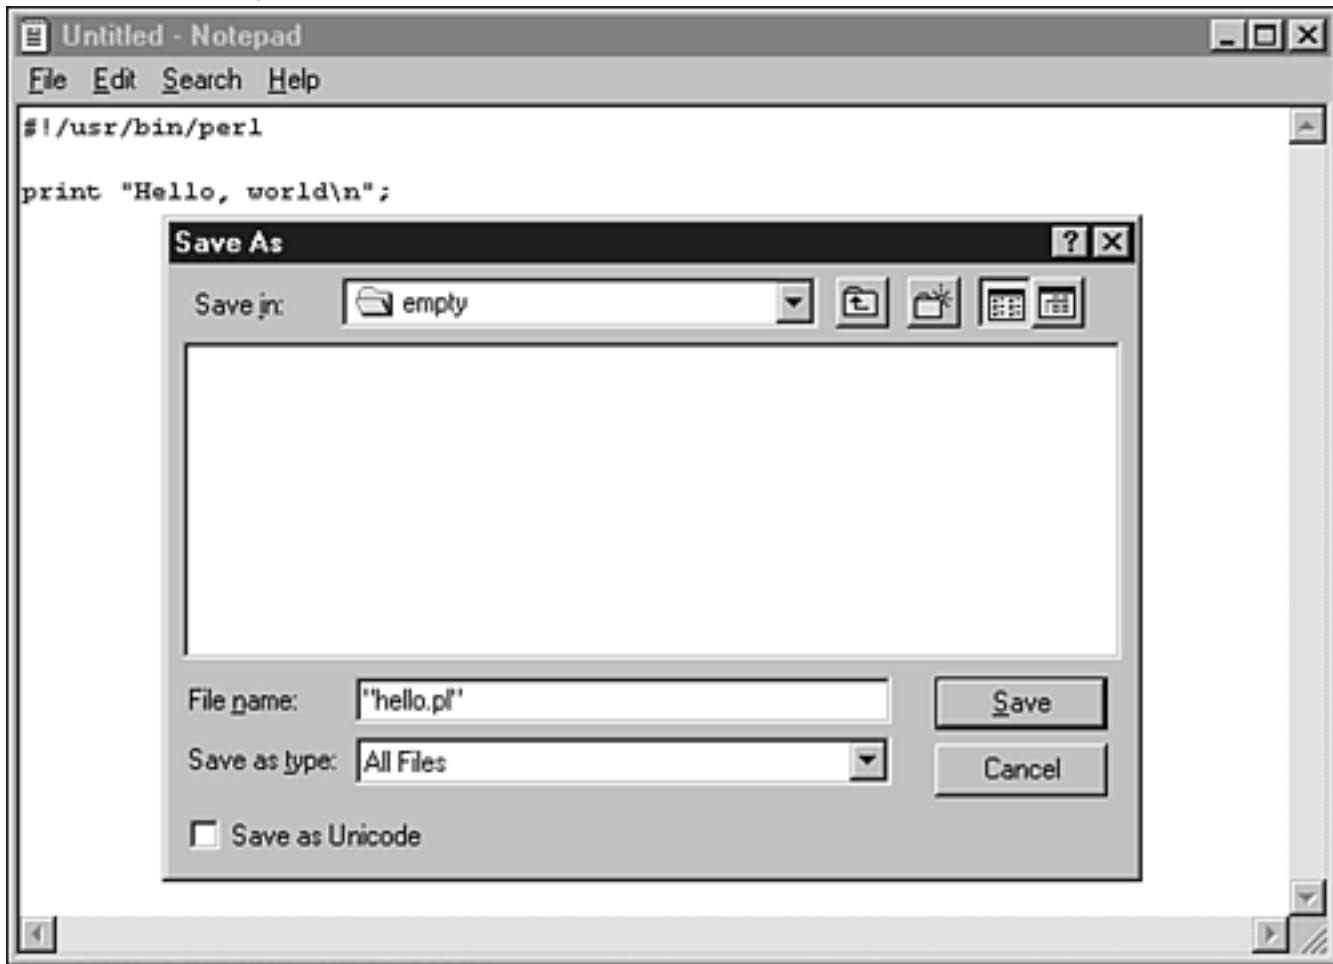
The `#!` line should be the first line in the file.

After you type this program into your text editor, save it in a file named `hello`. You don't need an extension on the filename, but Perl permits you to put one there. Some Windows and Macintosh utilities use extensions to tell what kind of file this is. If you need or want to use an extension, `.pl` or `.plx` is common; for example, you can use `hello.pl`.

Watch Out!

One of the most common problems here besets users of Microsoft Notepad. When you save the file, Notepad always appends a `.txt` extension to the file, whether or not you asked it to. If you want a `.pl` extension, you have to specify the filename in quotes, like this: "`hello.pl`", as shown in the following figure. This forces Notepad to save the file without the `.txt` extension.

[\[View full size image\]](#)



Running the Program

To run the program, for now, you need to get to a command prompt. In Unix, open a terminal window or log in. On a Microsoft Windows machine, open an MS-DOS prompt. Under Mac OS X, use the Finder, select Applications, Utilities and run Terminal. You should also change into the directory where you stored the hello program by using your shell's `cd` command.

When you're at the prompt, type the following. (A DOS prompt is shown here; Unix has a slightly different prompt.)

```
C:\PROGRAMS> perl hello
```

If all goes well, Perl should respond with the following message:

```
Hello, World!
```

If typing this command worked, congratulations! Remember how you ran this program because this is how you will start your programs for the rest of the book. (You can use some other ways, which will be covered shortly.)

If it did not work, check for the following:

- If you receive the error message `Bad command or file name or perl: command not found`, the Perl program is not in your execution path. You need to figure out where the Perl program is installed and add that directory to the `PATH` variable in your shell.
- If you receive the error message `Can't open perl script hello: A file or directory does not exist`, you're probably not in the same directory as the file named `hello` that you saved earlier, or you saved the file under a different name. To change into the proper directory under Unix or Windows, use the `cd` command. For example, if you saved the file on your desktop in Windows 98, the command would be:

```
c:\> cd \Windows\Desktop  
c:\> perl hello
```

Remember, if you used an extension on the filename (such as `.pl` or `.plx`), you need to specify that extension when you run the program.

- If you receive an error such as `syntax error`, Perl starts normally but cannot figure out what is in the file named `hello`. You've either mistyped the file's contents, or you've used a word processor that applied formatting to the saved file. The `cat` command (Unix) or the `type` command (DOS) can be used to verify what is in the file. If you mistyped, check everything; the quotation marks and punctuation are important.

If you're using MacPerl, simply select `Run "hello"` from the `Script` menu to run your first Perl program. If you're not using MacPerl's built-in editor to write your program, use the `Open` command from the `File` menu to open your program in MacPerl, and then select `Run`.

It Worked! So What Happened?

When you typed the command `perl hello`, a program named `perl` was started on your computer. This program is called the **Perl interpreter**. The Perl interpreter is the heart and soul of Perl. Its job is to take the file it has been given (`hello`, in this case), find the program inside it, and run it.

When it runs a program, the interpreter "first checks the statements, functions, operators, math, and everything else that makes up the Perl program to make sure the syntax is correct, and then executes the statements one at a time."

When the Perl interpreter is all done reading your program from disk, it begins to run the program and

continues running that program until it's all done. When it's finished with your program, the Perl interpreter exits and returns control to your operating system.

Now let's look at how the `hello` program is "run."

Perl Play-by-Play

The first line of the `hello` program is

```
#!/usr/bin/perl
```

To Perl, anything following a `#` on a line is considered a comment line. A **comment** is something that Perl ignores. The `#!` on the first line of a program is different in some cases. The pathname that follows—`/usr/bin/perl`—is the path to the Perl interpreter. If a Unix file begins with `#!` followed by the path of an interpreter and is executable, Unix knows that this is a program (rather than a shell script) and how to interpret it. See the "[Q&A](#)" section at the end of the hour for an explanation of how to run programs.

Some Web servers that can execute Perl programs—Apache, for example—also pay attention to the `#!` line and can call the programs without having to use an explicit `perl` command.

For now, just consider `#!` a comment line.

The next line of the program is

```
print "Hello, World!\n";
```

A whole lot is going on here. This line constitutes a Perl statement; it marks off for Perl one thing to do.

First, this line contains a function called `print`. The `print` function takes whatever follows it and displays it to your screen by default. What the `print` function is supposed to print lasts all the way up to the semicolon (`;`).

The semicolon in Perl is a **statement separator**. You should put a semicolon between statements in your Perl program to show where one statement ends and the next one begins.

In this case, the `print` function displays the phrase `Hello, World!`. The `\n` at the end of the line tells Perl to insert a new blank line after it prints the phrase. The quotation marks around the phrase and the `\n` tell Perl that this is a literal string, and not another function. Strings are discussed at great length next hour.

Something You Should Know

Perl is called a **free-form** programming language. This means that Perl statements aren't very picky about how they're written. You can insert spaces, tabs, and even carriage returns—called **whitespace**—nearly anywhere in a Perl statement, and it really doesn't matter.

The only places that you really can't insert whitespace willy-nilly are places you'd expect restrictions to be. For example, you can't insert spaces in the middle of a function name; `pr int` is not a valid function. Nor can you insert them in numbers, as in `25 61`. Also, whitespace inside literal strings like `"Hello, World!"` shows up as, well, whitespace. Almost everywhere else it's valid. You could write the sample Perl program like this:

```
#!/usr/bin/perl

print
    "Hello, World!\n"
;

;
```

It would have functioned identically to the original. This free-form nature gives you a lot of range in "style" in your Perl programs. Feel free to express yourself with form. Just remember, someday other users may have to look at your programs. Be kind to them.

The style used in the examples in this book is fairly conservative. Sometimes statements are broken across several lines for clarity or to save space because Perl statements can get quite long. The Perl documentation even has a style guide, if you want to browse it for suggestions. You should search for the document named `perlstyle`.

By the Way

Style in Perl programs can be taken to extremes. Valid Perl programs can be written as poetry, even haiku. Some memorable Perl programs actually look like pictures but do something useful. *The Perl Journal* (<http://www.tpj.com>) used to run a contest that was all about writing obscure-looking Perl programs: the Obfuscated Perl Contest. You should not take style lessons from these entries.

Summary

In this hour, you learned a little about what Perl is and how it works. As you go through this book, you'll pick up more details here and there. You also learned how to install Perl on your system and verify that it's working properly and that all the documentation is in place. Finally, you typed in—and, I hope, ran—your first Perl program. Afterward, you dissected that program and learned a little bit about how Perl works.

Q&A

Q1: *These things that Perl runs, are they called Perl scripts or Perl programs?*

A1: The name really doesn't matter. Traditionally, *programs* are compiled into machine code and stored that way, and the machine code can be run many times. On the other hand, *scripts* are fed to an external program that translates them into actions every time they're run. Larry Wall, creator of Perl, has said that a "script is what you give an actor, a program is what you give an audience." Take that any way you'd like. For the remainder of this book, I'll call them *Perl programs*, and if you learn well, you can then be known as a *Perl programmer*.

Q2: *I found my program under Windows using the Explorer and double-clicked on it. A DOS window opened, printed something, and then immediately closed. What's wrong?*

A2: Nothing! The Perl interpreter is associated with the `.pl` extension as a DOS-mode program. When you clicked on it, a DOS window was opened and the interpreter ran your program and then immediately exited. (It *was* a short program.) If you actually need to see your Perl program run, open a window and run it manually.

Q3: *In the "Running the Program" section, you hinted at an easier way to run Perl programs under Unix. How?*

A3: First, you must make sure that the `#!` line of the program is correct, and that the pathname there really points to a perl interpreter; `/usr/bin/perl` is the usual place for it, or `/usr/local/bin/perl` on some machines. Next, you must make the program executable by using the `chmod` command. For the `hello` program, the Unix shell command would be `chmod 755 hello`. After that's done, you can run the Perl program by typing `hello` or `./hello`. Hint: Don't ever name your program "test" under Unix. The Unix shells have a command called `test`, and you'll be very frustrated when the wrong program runs. See your shell's documentation for other names you should avoid.

Workshop

Quiz

1: *Perl* is the name of the language; *perl* is the name of

a.

The language also

b.

The interpreter

c.

A DOS command

2: Where can you always find a copy of the Perl documentation?

a.

<http://www.microsoft.com>

b.

<http://www.perl.com>

c.

<http://www.perl.net>

3: In which manual page can you find a description of Perl's syntax?

a.

perlsyn

b.

perlop

c.

perlfaq

Answers

A1: b. After it's installed, `perl` is a valid command in a DOS shell, so c is acceptable also.

A2: b. It's installed on your system as well.

A3: a. Unless you actually run `perldoc perl`, there's no way you could have known for sure, now is there?

Activities

- Browse through the FAQ. Even if you don't understand everything in there, you should get a feel for what kind of information is in the FAQ.

If you prefer reading it on a browser, surf to <http://www.perl.com> and read it there, but do peruse it.

Hour 2. Perl's Building Blocks: Numbers and Strings

What You'll Learn in This Hour:

- Literal numbers and strings
- Scalar variables
- Operators

Every programming language—and every human language—has a similar beginning: You have to have something to talk about. In Perl, numbers and strings are the basic unit of conversation, and these are called **scalars**.

Every hour in this book deals in some way with scalars—increasing, decreasing, querying, testing, collecting, clearing, separating, folding, sorting, saving, loading, printing, and deleting them. Scalars are Perl's singular nouns; they can represent a word, a record, a document, a line of text, or a character.

Some scalars in Perl can represent information that does not change over the life of the program. Some programming languages call these values **constants** or **literals**. Literal data is used for things that simply do not change, such as the value of π , the rate of acceleration of a body falling to Earth, and the name of the 15th President of the United States. If these values are needed by a Perl program, they would, at some point in the program, be represented by a scalar literal.

The other kinds of scalars in Perl are those that change, called **scalar variables**. Variables hold data while you manipulate it. You can change the contents of variables because they merely act as handles for the data they represent. Variables are given names—convenient and easy-to-remember names, ideally—that enable you to refer to the data you are manipulating.

This hour also introduces Perl's operators. They are one kind of verb in the Perl language. Operators take Perl's nouns and do the actual manipulations you need to write programs that perform useful tasks.

Literals

Perl has two different types of scalar constants called *literals*: numeric literals and string literals.

Numbers

Numeric literals are numbers, and Perl accepts several different ways of writing numbers. All the examples shown in [Table 2.1](#) are valid numeric literals in Perl.

Table 2.1. Samples of Numeric Literals

Number	Type of Literal
6	An integer
12.5	A floating point number
15.	Another floating point number
.7320508	Yet another floating-point number
1e10	Scientific notation
6.67E - 33	Scientific notation (e or E is acceptable)
4_294_296	A large number with underscores instead of commas

Numbers are expressed as you think they would be. Integers are just groups of consecutive digits. Floating-point decimal numbers contain a decimal point in the correct position, even if there are no digits to the right of it. A floating-point number can be expressed in scientific notation as an exponent preceded by the letter e (or E) and a decimal number called the mantissa. The value of the scientific-notation literal is 10 raised to the power indicated by the exponent, multiplied by the mantissa; for example, 6.5536E4 = 65,536.0. (Strictly speaking, a mantissa is the decimal part of a logarithm, where it serves the same purpose as it does here.)

You cannot put commas into a numeric literal to improve readability, but you can use underscores where commas would normally appear. Perl removes the underscores when using the value.

Watch Out!

Do not use a leading zero in front of a number, such as 010. To Perl, a leading zero means that the literal represents an octal number—base 8. Perl also allows you to use literal hexadecimal numbers—base 16—and binary numbers—base 2. More information on these is in the online documentation in the [perldata](#) section.

Strings

String literals in Perl are sequences of characters, such as `Hello, World`. They can contain as much data as you want; strings have no real limit on their size except the amount of virtual memory in your computer. Strings can also contain any kind of data—simple ASCII text, ASCII with the high bits on, even binary data. Strings can even be empty.

In Perl you must enclose string literals, with very few exceptions, in quotation marks. This process is called **quoting** the string. The two primary ways of quoting strings are to use single quotation marks (`' '`) or double quotation marks (`" "`). The following are some sample string literals:

```
"foo!"  
'Fourscore and seven years ago'  
"One fish,\nTwo fish,\nRed fish,\nBlue fish\n"  
"  
"Frankly, my dear, I don't give a hoot.\n"
```

What if you need to put another quotation mark inside your string literal? For example, the following string literal would make no sense to Perl:

```
"Then I said to him, "Go ahead, make my day""
```

Here, the quotation mark in front of the word `Go` marks the end of the string literal started by the first quotation mark, leaving the phrase `Go ahead, make my day` outside the string literal—in which case it would have to be valid Perl code, which it isn't. To prevent this situation, you must use a **backslash** (`\`) in front of the quotation marks that are inside the string. A backslash inside of a string literal tells Perl that the character that follows should not be treated as Perl would normally treat it—in this case, it should be handled as a character in the string literal rather than as the marker for the end of the string literal. Put a backslash character in front of each quotation mark that you want Perl to treat simply as a character, as shown here:

```
"Then I said to him, \"Go ahead, make my day.\""
```

The backslashes let Perl know that the quotation mark that follows is not the match for the first quotation

mark that started the string literal. This rule applies to single quotation marks as well as double quotation marks, as you can see here:

```
'The doctors\'s stethoscope was cold.'
```

The primary difference between double-quoting and single-quoting a string is that single-quoted strings are quite literal; every character in a single-quoted string (except the sequence `\'`) means exactly what is there. In a double-quoted string, on the other hand, Perl checks to see whether variable names or **escape sequences** are present and translates them if they are. Escape sequences are special strings that allow you to embed characters in strings when just typing the characters would cause problems. [Table 2.2](#) shows a short list of Perl's escape sequences.

Table 2.2. Sample String Escape Sequences

Sequence Represents

<code>\n</code>	Newline
<code>\r</code>	Carriage return
<code>\t</code>	Tab
<code>\b</code>	Backspace
<code>\u</code>	Change next character to uppercase
<code>\l</code>	Change next character to lowercase
<code>\\"</code>	A literal backslash character
<code>\'</code>	A literal ' inside of a string surrounded by single-quotation marks ("').
<code>\"</code>	A literal " inside of a string surrounded by quotation marks.

By the Way

You can find the full list of escape sequences in the online manual. As I indicated in [Hour 1](#), "Getting Started with Perl," you can find the entire Perl language documentation by using the `perldoc` utility included with the Perl distribution. The escape sequences are listed in the "perlop" manual page under the heading "Quote and Quote-like Operators."

Having many quotation marks embedded in a string can make typing the string error-prone and difficult

because each embedded quote mark has to be escaped, as shown here:

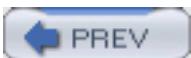
```
"I said, \"Go then,\\"", and he said \"I'm gone.\"."
```

Perl provides another quoting mechanism: the `qq` and `q` operators. To use `qq`, you can surround the string literal with `qq()` instead of quotation marks:

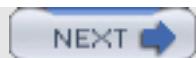
```
qq(I said, "Go then," and he said "I'm gone")
```

`qq` replaces the double quotation marks; this mechanism behaves exactly like double quotes in almost all respects.

You can use the `q` operator to surround text instead of



< Day Day Up >



Scalar Variables

To store scalar data in Perl, you must use a scalar variable. In Perl, you indicate a scalar variable with a dollar sign followed by the name of the variable. The following are some examples:

\$a

\$total

\$Date

\$serial_number

\$cat450

The dollar sign—called a **type identifier**—indicates to Perl that the variable contains scalar data. Other variable types (hashes and arrays) use a different identifier or no identifier at all (filehandles). Variable names in Perl—whether for hashes, arrays, filehandles, or scalars—must conform to the following rules:

- Variable names can contain alphabetic (a to z, A to Z) characters, numbers, or an underscore character (_) after the type identifier. The first character of a variable name can't be a number, though.
- Variable names are case sensitive. This means that upper- and lowercase are significant in variable names. Thus, each of the following represents a different scalar variable:

\$value

\$VALUE

\$Value

\$valuE

Perl reserves to itself single-character variable names that do not start with an alphabetic character or underscore. Variables such as `$_`, `$ "`, `$ /`, `$2`, and `$$` are **special variables** and should not be used as normal variables in your Perl programs. The purpose of these special variables will be covered later.

Scalar variables in Perl—in contrast to some other languages—do not have to be declared or initialized in any way before you can use them. To create a scalar variable, just use it. Perl uses a default value when an uninitialized variable is used. If it's used as a number (such as in a math operation), Perl will use the value `0` (zero); if it's used like a string (almost everywhere else), Perl will use the value `" "`, the empty string.

Watch Out!

Using a variable before it has been initialized with a value is a bad practice. Perl can be made to warn you when this situation happens. If your program has `-w` on the `#!` line at the beginning, the directive "use warnings" appears in your program (Perl 5.6). If you invoke Perl with the `-w` switch on the command line, and then try to use the value of a variable that you haven't previously set, Perl responds with this error message when your program runs and you try to use the value: `Use of uninitialized value.`

The Special Variable `$_`

Perl has a special variable, `$_`, whose value is used as a "default" by many operators and functions. For example, if you simply state `print` by itself—without specifying a scalar variable or string literal to print—Perl will print the current value of `$_`:

```
$_="Dark Side of the Moon";  
  
print;      # Prints the value of $_, "Dark side..."
```

Using `$_` like this can understandably cause some confusion. It's not really apparent what `print` is actually printing, especially if the assignment to `$_` occurs higher up in the program.

Some operators and functions are actually easier to use with the `$_` variable, notably the pattern-matching operators discussed in [Hour 6, "Pattern Matching."](#) In this book, however, I'll keep the use of `$_` to a minimum so that the lessons are easier to follow.

Expressions and Operators

Now that you've learned what scalar data is and know how to use scalar variables, you can start doing something useful with Perl. Perl programs are just collections of expressions and statements executed in order from the top of your Perl program to the bottom (unless you specify otherwise with flow-control statements, covered in [Hour 3](#), "Controlling the Program's Flow").

[Listing 2.1](#) shows a valid Perl program.

Listing 2.1. A Simple Perl Program

```
1:  #!/usr/bin/perl -w
2:
3:  $radius=50;
4:
5:  $area=3.14159*($radius ** 2);
6:  print $area;
```

Line 1: This line is the path to the Perl interpreter, as explained in [Hour 1](#). The `-w` switch tells Perl to inform you of any warnings encountered.

Line 3: This line is an **assignment**. The numeric scalar data value `50` is stored in the scalar variable `$radius`.

Line 5: This line is another assignment. On the right side of the assignment operator is an expression. This expression contains the scalar variable `$radius`, operators (`*` and `**`,

explained later), and a numeric scalar (2). The expression's value is computed and assigned to `$area`.

Line 6: This line prints the result of the calculation stored in `$area`.

An expression in Perl is simply something that has a value. For example, 2 is a valid expression. So are `54*$r`, "Java", `sin($pi*8)`, and `$t=6`. The values of expressions are computed when your program is run. The program evaluates the functions, operators, and scalar constants in the expression and reduces it to a value. You can use these expressions with assignments, as part of other expressions, or as part of other Perl statements.

Basic Operators

As you saw in Listing 2.1, to assign scalar data to a scalar variable, you use the assignment operator, =. The assignment operator takes the value on the right side and puts it in the variable on the left:

```
$title="Gone With the Wind";  
$pi=3.14159;
```

The operand on the left side of the assignment operator must be something that a value can be assigned to—namely, a variable. The operand on the right side can be any kind of expression. The entire assignment itself is an expression; its value is that of the right-hand expression. This means that, in the following snippet, `$a`, `$b`, and `$c` are all set to 42:

```
$a=$b=$c=42;
```

Here, `$c` is first set to 42. `$b` is set to the value of the expression `$c=42` (which is 42). `$a` is then set to the value of the expression `$b=42`.

The variable being assigned to can even appear on the right side of the assignment operator, as shown here:

```
$a=89*$a;  
$count=$count+1;
```

The right side of the assignment operator is evaluated using the old value of `$a` or `$count`, and then the result is assigned to the left side as the new value. The second example has a special name in Perl; it's called an **increment**. You'll read more about incrementing values later.

Numeric Operators

Perl has many operators to manipulate numeric expressions. Some of these operators are familiar to you already; some you will be meeting for the first time. The first kind of operator you should already know—the arithmetic operators. [Table 2.3](#) shows a list of these operators.

 PREV

NEXT 

More Operators

Perl has so many operators that there isn't enough space in this book to document them all fully. (You can read about most of them in the "perlop" manual page.) But for the remainder of this hour, you'll read about the most commonly used operators and functions in Perl.

One-Operand (Unary) Operators

Until now, all the operators you've read about take two **operands**—that is, two expressions, from whose values the operators make a new value in some way. For example, division (6/3) requires a dividend (6) and a divisor (3), multiplication (5*2) requires a multiplicand (5) and a multiplier (2), and so on. Another kind of operator, called a **unary** operator, takes only one operand. You're probably familiar with one example of these already—the unary minus (-). The unary minus returns the value of its operand with the sign reversed: negative if the operand is positive, positive if the operand is a negative number:

```
6;      # Six
-6;     # Negative six.
-(-5);  # Positive five, not negative five.
```

Many of Perl's unary operators are actually **named operators**; that is, instead of a symbol—like the - for unary minus—they use a word. The parentheses around the operand with named unary operators are optional but are shown in [Table 2.4](#) for clarity. Named unary operators in Perl look and act like functions (discussed in [Hour 8](#), "Functions"), and their operands are called **arguments**—the term used for the expressions that Perl functions act on.

Table 2.4. Some Named Unary Operators

Operator	Sample Usage	Result
int	int(5.6234)	Returns the integer portion of its argument (5).
length	length("nose")	Returns the length of its string argument (4).
lc	lc("ME TOO")	Returns its argument shifted to lowercase letters ("me too").
uc	uc("hal 9000")	Returns its argument shifted to uppercase letters ("HAL 9000").
cos	cos(50)	Returns the cosine of 50 in radians (.964966).

`rand`

`rand(5)`

Returns a random number from 0 to less than its argument. If the argument is omitted, a number between 0 and 1 is returned.

Some named unary operators are listed briefly in [Table 2.4](#).

You can find the full list of named operators in the online manual. As discussed in [Hour 1](#), you can find the entire Perl language documentation by using the `perldoc` utility included with the Perl distribution. All of the operators are listed in the "perlop" manual and in the "perlfunc" manual. More operators will be presented in later lessons, as they are needed.

Increment and Decrement

In the "[Numeric Operators](#)" section, you read about a special type of assignment called an increment, which looks like the following:

```
$counter=$counter+1;
```

An increment is typically used to count things, such as the number of records read, or to generate sequence numbers, such as numbering the items in a list. It's such a common idiom in Perl that you can use a special operator called an **autoincrement operator** (`++`). The autoincrement operator adds 1 to its operand:

```
$counter++;
```

When this code is executed, `$counter` is increased by 1.

Perl also offers a shortcut for decreasing a variable's value called, not surprisingly, the **autodecrement operator** (`--`). You use the autodecrement exactly the same way you use the autoincrement:

```
$countdown=10;  
  
$countdown--;      # decrease to 9
```

Let me add a final note on the autoincrement operator: When the operator is applied to a text string, and the text string starts with an alphabetic character and is followed by alphabetic characters or numbers, this operator becomes magical. The last (rightmost) character of the string is increased. If it's an alphabetic character, it becomes the next letter in sequence; if it's numeric, the number increases by 1. You can carry

across alphabetic and numeric columns as follows:

```
$a="999";  
  
$a++;  
  
print $a;          # prints 1000, as you'd expect  
  
$a="c9";  
  
$a++;  
  
print $a;          # prints d0.  9+1=10, carry 1 to the c.  
  
$a="zzz";  
  
$a++;  
  
print $a;          # prints "aaaa".
```

The autodecrement operator does not decrement strings like this.

Angle Operator (<>)

The angle operator (<>), sometimes called a diamond operator, is primarily used for reading and writing files; it will be covered fully in [Hour 5](#), "Working with Files." However, a brief introduction now will make the exercises more interesting, and when you begin [Hour 5](#), the operator will look somewhat familiar to you.

Until then, you can use the angle operator in its simplest form: <STDIN>. This form indicates to Perl that a line of input should be read from the standard input device—usually the keyboard. The <STDIN> expression returns the line read from the keyboard:

```
print "What size is your shoe? ";  
  
$size=<STDIN>;  
  
print "Your shoe size is $size. Thank you!";
```

The preceding code, when executed (assuming you type `9.5` as your shoe size), would print this on the screen:

What is size is your shoe? **9.5**

Your shoe size is 9.5

. Thank you!

The `<STDIN>` expression reads from the keyboard until the user presses the Enter key. The entire line of input is returned and is placed in `$size`. The line of text returned by `<STDIN>` also includes the newline character that the user typed by pressing Enter. That's why the period and "Thank you!" appear on a new line in the preceding display—the newline was part of the value of `$size`. Often, as here, you don't want the newline character at the end of the string—just the text. To remove it, you can use the `chomp` function as follows:

```
print "What size is your shoe?" ;  
  
$size=<STDIN>;  
  
chomp $size;  
  
print "Your shoe size is $size. Thank you!\n" ;
```

`chomp` removes any trailing newline character at the end of its argument

More Assignment Operators

Earlier, you learned that to assign a value to a scalar variable you use the assignment operator (`=`). Perl actually has an entire set of operators for doing assignments. Every Perl arithmetic operator and quite a few others can be combined to do an assignment and an operation all at the same time. The general rule for building an assignment with an operator is as follows:

variable operator=expression

This form of the assignment produces the same result as the following:

variable=variable operator expression

Using combined assignments generally doesn't make your programs any more readable but can make them more concise. Following this rule, the statement

```
$a=$a+3;
```

can be reduced to

```
$a+=3;
```

The following are some more examples of assignments:

```
$line.=", at the end";           # ", at the end" is appended to $line  
$y*=$x                         # same as $y=$y*$x  
$r%=67;                          # Divide by 67, put remainder in $r
```

A Few Words on Strings and Numbers

For the most part, Perl allows you to use numbers and strings interchangeably; the representation it uses depends on what Perl is looking for in that situation.

- If something looks like a number, Perl can use it as a number when it needs a number:

```
$a=42;                      # A number  
  
print $a+18;      # displays 60.  
  
$b="50";  
  
print $b-10;      # Displays 40.
```

- If something looks like a number, when Perl needs a string, it uses the string representation of the number:

```
$a=42/3;  
  
$a=$a . "Hello"; # Using a number like a string.  
  
print $a          # displays "14Hello"
```

- If something doesn't look like a number, but you've used it where a number was expected, Perl simply uses the value 0 in its place:

```
$a="Hello, World!";  
  
print $a+6;      # displays the number 6
```

If you have warnings enabled, Perl emits a warning if you do this, however.

All these uses are in keeping with Perl's philosophy of Least Surprise. Even when given nonsense—as in the last example here—Perl tries to do something sensible with it. If you have warnings enabled in your Perl program—by putting a `-w` on the `#!` line or invoking the Perl interpreter with the `-w` option—Perl warns you that you're doing something nonsensical by giving you the following message: `Argument X isn't numeric.`



< Day Day Up >



Exercise: Interest Calculator

For this exercise, you're going to perform a compound interest calculation. This program will calculate the interest on a savings account, given some information about interest rates, deposits, and time. The formula you're going to use is as follows:

$$\text{accrued} = \text{payment} \left(\frac{\left((1 + \text{MonthlyInterest})^{\text{NumberOfDeposits}} \right) - 1}{\text{MonthlyInterest}} \right)$$

Using your text editor, type the program from [Listing 2.2](#) and save it as `Interest`. Do not type in the line numbers. Make the program executable according to the instructions you learned in [Hour 1](#).

When you're done, try running the program by typing the following at a command line:

```
perl Interest
```

[Listing 2.3](#) shows a sample of the `Interest` program's output.

Listing 2.2. The Complete Source of the Interest Program

```
1:  #!/usr/bin/perl -w
2:
3:  print "Monthly deposit amount? ";
4:  $pmt=<STDIN>;
5:  chomp $pmt;
6:
7:  print "Annual Interest rate? (ex. 7% is .07) ";
8:  $interest=<STDIN>;
9:  chomp $interest;
10:
```

```
11: print "Number of months to deposit? ";
12: $mons=<STDIN>;
13: chomp $mons;
14:
15: # Formula requires a monthly interest
16: $interest/=12;
17:
18: $total=$pmt * ( ( ( 1 + $interest) ** $mons ) -1 )/ $interest;
19:
20: print "After $mons months, at $interest monthly you\n";
21: print "will have $total.\n";
```

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The user is prompted for an amount.

Line 4: `$pmt` is read from the standard input device (the keyboard).

Line 5: The newline character is removed from the end of `$pmt`.

Lines 7–9: `$interest` is read in from the keyboard, and the newline is removed.

Lines 11–13: `$mons` is read in from the keyboard, and the newline is removed.

Line 16: `$interest` is divided by 12 and stored back in `$interest`.

Line 18: The interest calculation is performed, and the result is stored in `$total`.

Lines 20–21: The results are printed.

Listing 2.3. Output from the Interest Program

1: Monthly deposit amount? **180**

2: Annual Interest rate? (ex. 6% is .06) **.06**

3: Number of months to deposit? **120**

4: After 120 months, at 0.005 monthly you

5: will have 29498.2824251624.

 PREV

< Day Day Up >

NEXT 

Summary

In this hour, you learned that Perl's most basic type of data is a scalar. Scalars can consist of almost any kind of data. Scalar values can be represented by string or numeric literals. Numeric literals can have many different formats representing integers and floating-point numbers. String literals are sequences of characters surrounded by either double or single quotation marks. Scalar values are stored in scalar variables, which have a dollar sign \$ in front of their names. Perl provides operators for performing string manipulation and basic arithmetic.

Q&A

Q1: *The output of the Interest program looks sloppy. How do I control how many digits are displayed?*

A1: The easiest way to control the number of decimal digits is to use the `printf` function, which is covered in [Hour 9](#), "More Functions and Operators."

Q2: *Does Perl have a function for rounding?*

A2: The `printf` function usually does what you want for rounding when displaying numbers. If you really need the `round` function, check out the POSIX module, which has this function and many more.

Q3: *How large (or small) of a number will Perl let me manipulate?*

A3:

Workshop

Quiz

1: Variables are interpolated inside `qq` quotes.

a.

True

b.

False

2: What value is stored in `$c` after running the following code?

```
$a=6;
```

```
$a++;
```

```
$b=$a;
```

```
$b--;
```

```
$c=$b;
```

a.

6

b.

7

c.

8

3: Concatenation can be performed only with the concatenation operator (.).

a.

True

b.

False

Answers

A1: a. True. `qq` behaves in every way like a pair of double quotation marks. This means it can interpolate variables.

A2: a. `$a` is set to 6. `$a` is then incremented to 7, and assigned to `$b`. `$b` is decremented to 6, and assigned to `$c`.

A3: b. False. A motto in the Perl community is "There Is More Than One Way To Do It" (TMTOWTDI). Concatenation can be performed by including two (or more) scalars in a double-quoted string, as follows:

```
qq($a$b$c);
```

Activities

- Write a short program that prompts the user for a Fahrenheit temperature and prints the temperature in Celsius. Converting Fahrenheit temperature to Celsius can be accomplished by taking the Fahrenheit temperature and subtracting 32, and then multiplying by 5/9. For example, 75 degrees Fahrenheit is 23.8 degrees Celsius.
- Modify the Interest program in Listing 2.3 to print the amount with no more than two decimal places. You can do so without `printf` with clever use of the `int` operator, multiplication, and division.

Hour 3. Controlling the Program's Flow

What You'll Learn in This Hour:

- Block statements
- Operators
- Looping
- Labels
- Exiting Perl after a program execution

In [Hour 2](#), "Perl's Building Blocks: Numbers and Strings," you learned about statements, operators, and expressions. All the examples in that hour had one thing in common: All the statements were executed in order from top to bottom and were executed only once.

One of the reasons that you use computers is that computers are very good at performing repetitive tasks—over and over again—without getting tired or bored and without developing carpal tunnel syndrome. So far, you haven't had any way of telling Perl to "do this task X times" or to "repeat this task until it's done." In this hour, you will learn about Perl's control structures. Using them, you can group statements into something called a **statement block** and run the group of statements repeatedly until they've done what you want.

The other chore that computers excel at is making decisions quickly. It would be tiresome—not to mention silly—if a computer had to ask you every time it made a decision. The very act of retrieving and reading your email causes your computer to make millions of decisions that you really don't want to deal with: how to assemble network traffic, which colors to make each pixel on your screen, how your incoming mail should be pulled apart and displayed, what should be done when your mouse cursor moves even a tiny bit, and countless others. All these decisions are made up of other decisions, and some of them are made thousands of times per second. In this hour, you will learn about **conditional statements**. Using these statements, you can write blocks of code that will be executed or not, depending on decisions made in your Perl program.

Blocks

The fundamental way to group statements in Perl is the **block**. To group statements in a block, just surround the statements with a matched set of curly braces, as shown here:

```
{  
    statement_a;  
  
    statement_b;  
  
    statement_c;  
}
```

Within the block, statements execute from the top down, as they have until now. You can have other, smaller blocks of statements nested within a block, as you can see here:

```
{  
    statement_a;  
  
    {  
        statement_x;  
  
        statement_y;  
    }  
}
```

The format of the block, like the rest of Perl, is free-form. The statements and curly braces can be on one line or on several lines, as shown here, and with any kind of alignment you want, as long as you always have a matched set of curly braces:

```
{ statement; { another_statement; }
```

```
{ last_statement; } }
```

Although you can arrange blocks any way you would like, programs can be hard to read if they're just thrown together. Good indenting, although not necessary for Perl to understand your code, makes for human-readable Perl. It can help you keep track of your program's logic.

Blocks that occur by themselves within a program are called **bare blocks** or **naked blocks**. Most of the time, however, you will encounter blocks attached to other Perl statements.

[PREV](#)[< Day Day Up >](#)[NEXT](#)

The if Statement

To control whether statements are executed based on a condition in a Perl program, you usually use an `if` statement. The syntax of an `if` statement is as follows:

```
if (expression) block
```

The statement works like this: If the *expression* evaluates to true, the *block* of code is run. If the *expression* is false, the *block* of code is not run. Remember that the block includes the braces. Consider this example:

```
if ( $r == 5 ) {  
  
    print 'The value of $r is equal to 5.';  
  
}
```

The expression being tested is `$r == 5`. The `==` symbol is an equality operator. If the two operands on either side—`$r` and `5`—are numerically equal to one another, the expression is considered to be true, and the `print` statement is executed. If `$r` is not equal to `5`, the `print` statement is not executed.

The `if` statement can also run one set of statements if a condition is true and another set of statements if it's not. That structure is called an `if-else` statement. The syntax looks like this:

```
if (expression)      #If expression is true...  
  
    block1          # ...this block of code is run.  
  
else  
  
    block2          # Otherwise this block is run.
```

The first block, `block1`, is run only if the *expression* is true; if the *expression* is not true, `block2`, following the `else`, is run. Now consider this example:

```
$r=<STDIN>;  
  
chomp $r;  
  
if ($r == 10) {  
  
    print '$r is 10', "\n";  
  
} else {  
  
    print '$r is something other than 10...', "\n";  
  
$r=10;  
  
    print '$r has been set to 10', "\n";  
  
}
```

Watch Out!

In the preceding example, notice that to assign a value to `$r`, I used the assignment operator, `=`. To *test* the value of `$r`, I used the numeric equality test operator, `==`. Do not confuse them in your programs because debugging can be very difficult. Remember that `=` assigns a value and `==` tests for equality. If you use the `-w` option to turn warnings on, Perl can sometimes warn you if you have made this error.

Yet another way of structuring an `if` statement is to check multiple expressions and run code depending on which expressions are true:

```
if (expression1)      # If expression1 is true ...  
  
    block1           # ...run this block of code.  
  
elsif (expression2)  # Otherwise, if expression2 is true...  
  
    block2           # ...Run this block of code.  
  
else  
  
    block3           # If neither expression was true, run this.
```

You can read the preceding block like this: If the expression labeled *expression1* is true, the block *block1* is run. Otherwise, control falls to the `elsif` and *expression2* is tested; if it's true, *block2* is run. If neither *expression1* nor *expression2* is true, *block3* is run. The following is an example of real Perl code that demonstrates this syntax:

```
$r=10;

if ($r==10) {
    print '$r is 10!';
} elsif ($r == 20) {
    print '$r is 20!';
} else {
    print '$r is neither 10 nor 20';
}
```

The Other Relational Operators

So far, you've been comparing numeric quantities in your `if` statements with the equality operator, `==`. Perl actually has quite a few operators for comparing numeric values, most of which are listed in [Table 3.1](#).

Table 3.1. Numeric Relational Operators

Operator	Example	Explanation
<code>==</code>	<code>\$x == \$y</code>	True if <code>\$x</code> equals <code>\$y</code>
<code>></code>	<code>\$x > \$y</code>	True if <code>\$x</code> is greater than <code>\$y</code>
<code><</code>	<code>\$x < \$y</code>	True if <code>\$x</code> is less than <code>\$y</code>
<code>>=</code>	<code>\$x >= \$y</code>	True if <code>\$x</code> is greater than or equal to <code>\$y</code>
<code><=</code>	<code>\$x <= \$y</code>	True if <code>\$x</code> is less than or equal to <code>\$y</code>
<code>!=</code>	<code>\$x != \$y</code>	True if <code>\$x</code> is not equal to <code>\$y</code>

To use these operators, you can simply put them in anywhere that your program needs to test relations between numeric values. An example of the use of these operators in an `if` statement is shown in [Listing 3.1](#), which you can type in and run (do not type the line numbers—"1:" and so forth). The various parts of the program work as follows:

Listing 3.1. A Small Number Guessing Game

```
1:  #!/usr/bin/perl -w
2:
3:  $im_thinking_of=int(rand 10);
4:  print "Pick a number:";
5:  $guess=<STDIN>;
6:  chomp $guess;    # Don't forget to remove the newline!
7:
8:  if ($guess>$im_thinking_of) {
9:      print "You guessed too high!\n";
10: } elsif ($guess < $im_thinking_of) {
11:     print "You guessed too low!\n";
12: } else {
13:     print "You got it right!\n";
14: }
```

Line 1: This line is the standard first line of a Perl program; it indicates the interpreter you want to run and the `-w` switch to enable warnings. See [Hour 1](#), "Getting Started with Perl"; your first line may need to look slightly different.

Line 3: The `(rand 10)` function picks a number between 0 and 10, and the `int()` function truncates it so that only integers 0 to 9 are assigned to `$im_thinking_of`.

Lines 4–6: These lines ask the user for the guess, assign it to `$guess`, and remove the trailing newline character.

Lines 8–9: If `$guess` is greater than the number in `$im_thinking_of`, these lines print an appropriate message.

Lines 10–11: Otherwise, if `$guess` is less than the number in `$im_thinking_of`, these lines print that message.

Lines 12–13: The only choice left is that the user guessed the number.

The operators in [Table 3.1](#) are used only for testing numeric values. Using them to test nonalphabetic data results in behavior that you probably don't want. Consider this example:

```
$first="Simon";  
  
$last="simple";  
  
if ($first == $last) {      # == is not what you want!  
    print "The words are the same!\n";  
  
}
```

The two values `$first` and `$last` actually test *equal* to each other. The reason was explained in [Hour 2](#): If nonnumeric strings are used when Perl is expecting numeric values, the strings evaluate to zero. So the preceding `if` expression looks something like this to Perl: `if (0 == 0)`. This expression evaluates to true, and that's probably not what you wanted.

By the Way

If warnings are turned on, trying to test two alphabetic values (`simple` and `Simon` in the preceding snippet) with `==` will generate a warning message when the program runs to alert you to this problem.

If you want to test nonnumeric values, you can use another set of Perl operators, which are listed in [Table 3.2](#).

Table 3.2. Alphanumeric Relational Operators

Operator	Example	Explanation
----------	---------	-------------

Eq	<code>\$s eq \$t</code>	True if <code>\$s</code> is equal to <code>\$t</code>
Gt	<code>\$s gt \$t</code>	True if <code>\$s</code> is greater than <code>\$t</code>
Lt	<code>\$s lt \$t</code>	True if <code>\$s</code> is less than <code>\$t</code>
Ge	<code>\$s ge \$t</code>	True if <code>\$s</code> is greater than or equal to <code>\$t</code>
Le	<code>\$s le \$t</code>	True if <code>\$s</code> is less than or equal to <code>\$t</code>
Ne	<code>\$s ne \$t</code>	True if <code>\$s</code> is not equal to <code>\$t</code>

These operators decide "greater than" and "less than" by examining each character left to right and comparing them in ASCII order. This means that strings sort in ascending order: most punctuation first, and then numbers, uppercase, and finally lowercase. For example, `1506` compares less than `Happy`, which compares less than `happy`.

What Truth Means to Perl

Up to this point, you've been reading about "if this expression is true..." or "...evaluates to true..." but you haven't seen any formal definition of what Perl thinks "true" is. Perl has a few short rules about what is true and what is not true, and the rules actually make sense when you think about them for a bit. The rules are as follows:

- The number 0 is false.
- The empty string (" ") and the string "0" are false.
- The undefined value `undef` is false.
- Everything else is true.

Make sense? The only other point to remember is that when you're testing an expression to see whether it's true or false, the expression is simplified—functions are called, operators are applied, math expressions are reduced, and so on—and then converted to a scalar value for evaluation to determine whether it is true or false.

Think about these rules, and then take a look at [Table 3.3](#). Try to guess whether the expression is true or false before you look at the answer.

Table 3.3. True or False Examples

Expression	True or False?
0	False. The number 0 is false.
10	True. It is a nonzero number and therefore true.
9>8	True. Relational operators return true or false, as you would expect.

-5+5	False. This expression is evaluated and reduced to 0, and 0 is false.
0.00	False. This number is another representation of 0, as are 0x0, 00, 0b0, and 0e00.
""	False. This expression is explicitly mentioned in the rules as false.
" "	True. There's a space between the quotes, which means they're not entirely empty.
"0.00"	True. Surprise! It's already a string, but not "0" or "". Therefore, it is true.
"00"	True also, for the same reason as "0.00"
"0.00" + 0	False. In this expression, 0.00+0 is evaluated, the result is 0, and that is false.

Until now, you've seen only expressions with relational operators as the conditions in `if` statements. Actually, you can use *any* expression that will evaluate to true or false the way you would want:

```
# The scalar variable $a is evaluated for true/false
if ($a) { ... }

# Checks the length of $b.  If nonzero, the test is true.
if (length($b)) { .... }
```

Recall from [Hour 2](#) that the assignment operator `=` returns a value—the value that was assigned. That value, of course, is also true or false:

```
$a = 1;
$b = 2;

print qq(The statement "$a = $b" is );
if ($a = $b) { # value is 2, therefore true
    print "true";
} else {
    print "false";
}
```

This code prints `The statement "1 = 2" is true`. Now you see why using `=` when you meant `==` is such a pitfall.

The value `undef` is a special value in Perl. Variables that have not yet been set have the value of `undef`, and some functions return `undef` on failure. It's not 0, and it's not a regular scalar value. It's kind of special. In a test for truth, `undef` always evaluates to false. If you try to use the `undef` value in a math expression, it's treated as though it were 0.

Using variables that haven't been set yet is usually a sign of a programming error. If you're running your Perl programs with warnings enabled, the value `undef` in an expression or as an argument to some functions causes Perl to generate the warning `Use of uninitialized value`.

Logical Operators

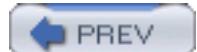
When you're writing programs, you sometimes need to code something like the following: Do this if `$x` is true and if `$y` is true, but not if `$z` is true. You can code this example into a series of `if` statements, but it's not pretty:

```
if ($x) {  
    if ($y) {  
        if ($z) {  
            # do nothing  
        } else {  
            print "All conditions met.\n";  
        }  
    }  
}
```

Perl has a whole class of operators for connecting together true and false statements like this, called **logical operators**. The logical operators are shown in [Table 3.4](#).

Table 3.4. Logical Operators

Operator	Alternative Name	Example	Analysis
<code>&&</code>	and	<code>\$s && \$t</code>	True only if <code>\$s</code> and <code>\$t</code> are true
		<code>\$q and \$p</code>	True only if <code>\$q</code> and <code>\$p</code> are true
<code> </code>	or	<code>\$a \$b</code>	True if <code>\$a</code> is true or <code>\$b</code> is true
		<code>\$c or \$d</code>	True if <code>\$c</code> is true or <code>\$d</code> is true
<code>!</code>	not		

 PREV

NEXT 

Looping

As you read in the Introduction, sometimes just making decisions and running code conditionally are not enough. Often you need to run pieces of code over and over again. The exercise presented in Listing 3.1 wasn't much fun because you could take only one guess (well, that and because it's a pointless game). If you want to be able to take multiple guesses, you need to be able to repeat sections of code conditionally, and that's what looping is all about.

Looping with `while`

The simplest kind of loop is a `while` loop. A `while` loop repeats a block of code as long as an expression is true. The syntax for a `while` loop looks like this:

```
while (expression) block
```

When Perl encounters the `while` statement, it evaluates the *expression*. If the *expression* is true, the *block* of code is run. Then, when the end of the block is reached, the *expression* is re-evaluated. If it's still true, the *block* is repeated, as in the following snippet:

Listing 3.2. Sample `while` Loop

```
1: $counter=0;  
2: while ($counter < 10 ) {  
3:     print "Still counting...$counter\n";  
4:     $counter++;  
5: }
```

Line 1: `$counter` is initialized to zero.

Line 2: The expression `$counter < 10` is evaluated. If it's true, the code in the block is run.

Line 4: The value of `$counter` is incremented by 1.

Line 5: The `}` marks the end of the block started on line 2 with a `{`. At this point, Perl returns

to the top of the `while` loop and re-evaluates the conditional expression.

Looping with `for`

The `for` statement is the most complicated and versatile of Perl's looping constructs. The syntax looks like this:

```
for ( initialization; test; increment ) block
```

The three sections of the `for` statement—*initialization*, *test*, and *increment*—are separated by semicolons. When Perl encounters a `for` loop, the following sequence takes place:

1. The *initialization* expression is evaluated.
2. The *test* expression is evaluated; if it's true, the *block* of code is run.
3. After the *block* is executed, the *increment* is performed, and the *test* is evaluated again. If the *test* is still true, the *block* is run again. This process continues until the *test* expression is false.

The following is an example of a `for` loop:

```
for( $a=0; $a<10; $a=$a+2 ) {  
    print "a is now $a\n";  
}
```

In this snippet, `$a` is set to 0, and the test `$a<10` is performed and found to be true. The body of the loop prints a message. The increment is then run—`$a=$a+2`—which increases the value of `$a` by 2. The test is performed again, and the loop repeats. This particular loop repeats until the value of `$a` is 10, when the test will be false and the program will continue running after the `for` loop.

You don't have to use the increment in the `for` statement for counting; it simply iterates until the test is false. In fact, you should be aware that each of the three parts of the `for` statement is *optional*, although the two semicolons are required. The following `for` statement is missing some pieces but is still perfectly valid:

```
$i=10;          # initialization
```

```
for( ; $i>-1; ) {  
  
    print "$i..";  
  
    $i--;          # actually, a decrement.  
  
}  
  
print "Blast off!\n";
```

Omitting the test portion means that you have to have some other way to exit the loop, or it will loop forever.

 PREV

< Day Day Up >

NEXT 

Other Flow Control Tools

Controlling the way your program executes with loops and condition statements is fine, but other flow control statements are needed to make readable programs. For example, Perl has statements to exit a `while` loop early, to skip certain portions of a `for` loop, to exit an `if` statement before the end of a block, or even to exit your program without falling off the end. Using some of the constructs explained in this section can make your Perl programs more concise and easier to read.

Odd Arrangements

The `if` statements have one more possible syntax. If you have only one expression inside the `if` block, the expression can actually precede the `if` statements. So, instead of writing

```
if (test_expression) {  
    expression;  
}
```

you can write

```
expression if (test_expression);
```

The following are a couple of examples of this variation of the syntax:

```
$correct=1 if ($guess == $question);  
  
print "No pi for you!" if ( $ratio != 3.14159);
```

You usually use this syntax in Perl code for clarity; sometimes reading the code is easier if you see the action before the condition. The expression preceding the `if` must be a single expression. The `if` statement must also be followed by a semicolon.

Fine-Grained Control

In addition to blocks, `for`, `while`, `if`, and other flow-control statements that control blocks of code, you can use Perl statements to control the flow *within* the blocks.

The simplest statement that gives you this control is `last`. The `last` statement causes the innermost currently running loop block to be exited. Consider this example:

```
while($i<15) {  
    last if ($i==5);  
  
    $i++;  
}
```

The `last` statement causes the `while` loop to exit when the value of `$i` is 5, instead of normally when the `while` test is false. When you have multiple nested loop statements, `last` exits the loop currently running.

Listing 3.3. Example of the `last` Statement

```
1:   for($i=0; $i<100; $i++) {  
2:     for($j=0; $j<100; $j++) {  
3:       if ($i * $j == 140) {  
4:         print "The product of $i and $j is 140\n";  
5:         last;  
6:       }  
7:     }  
8:   }
```

The set of nested loops in Listing 3.3 finds all the whole numbers less than 100 whose products are 140—2 and 70, 4 and 35, and so on—rather inefficiently. The point to note here is the `last` statement. When a product is found, the result is printed, and the inner loop (the loop iterating over `$j`) is exited. The outer loop continues executing (by incrementing `$i`) and reruns the inner loop.

The `next` statement causes control to be passed back to the top of the loop and the next iteration of the loop

to begin, if the loop isn't finished:

```
for($i=0; $i<100; $i++) {  
  
    next if (not $i % 2);  
  
    print "An odd number=$i\n";  
  
}
```

This loop prints all the even numbers from 0 to 98. The `next` statement causes the loop to go through its next iteration if `$i` is not even; the `$i % 2` expression is the remainder of `$i` divided by 2. In this case, the `print` statement is skipped. (A much more efficient way to write this loop would be simply to increase `$i` by 2, but that wouldn't demonstrate `next`, would it?)

The `redo` statement is similar to `next`, except that the condition isn't re-evaluated. Perl resumes execution back at the beginning of the block and doesn't check whether the termination condition has been met yet.

Labels

Perl allows blocks and some loop statements (`for`, `while`) to be *labeled*. That is, you can place an identifier in front of the block or statement:

```
MYBLOCK: {  
  
}
```

The preceding block is labeled as `MYBLOCK`. Label names follow the same conventions as variable names, with one small exception: Label names do not have an identifying character—`%`, `$`, `@`—as variables do. It's important to make sure that label names do not clash with Perl's built-in keywords. As a matter of style, it's best if label names are all uppercase. You should not have any conflicts with any current or future Perl keywords that way. The `for` and `while` statements can all have labels as well.

```
OUTER: while($expr) {  
  
    INNER: while($expr) {  
  
        statement;  
  
    }  
}
```

```
}
```

The `last`, `redo`, and `next` statements can each take a label as an argument. You therefore can exit a specific block. The code in [Listing 3.2](#) found two factors of 140 by using a nested pair of `for` loops. Suppose you wanted to exit the loops as soon as a factor is found. Without labels, you would need a complex arrangement of flag variables (variables whose only purpose is to store a true or false value for program flow control) and `if` statements between the two loops because you cannot exit the outer loop from within the inner loop. Labels solve this problem:

```
OUTER: for($i=0; $i<100; $i++) {  
    for($j=0; $j<100; $j++) {  
        if ($i * $j == 140) {  
            print "The product of $i and $j is 140\n";  
            last OUTER;  
        }  
    }  
}
```

Now the last statement can specify *which* loop it wants to exit—in this case, the `OUTER` loop. This snippet prints only the first pair of factors of 140 that it finds.

Leaving Perl

The `exit` statement is the ultimate flow-control tool. When Perl encounters an `exit` statement, the program stops executing, and an `exit status` is returned by Perl to the operating system. This exit status is usually used to indicate successful completion of the program. You'll learn more about exit statuses in [Hour 11](#), "System Interaction." For now, an exit status of zero means everything went okay. The following is an example of `exit`:

```
if ($user_response eq 'quit') {  
    print "Good Bye!\n";  
    exit 0;          # Exit with a status of 0.
```

}

By the Way

The `exit` statement has some side effects that are important to your operating system. When an `exit` is performed, any open files are closed, file locks are released, memory allocated by Perl is released to the system, and the Perl interpreter performs a clean shutdown.

 PREV

< Day Day Up >

NEXT 

Exercise: Finding Primes

What computer language primer would be complete without this little gem? In this exercise, you will examine a small program to find and print prime numbers. Prime numbers are divisible only by 1 and themselves; for example, 2 is prime, 3 is prime, 4 is not (because it is divisible by 1, 4, and 2), and so on. The list of primes is infinite, and they take a lot of computer power to find.

Using your text editor, type the program from Listing 3.4 and save it as `Primes`. Do not type in the line numbers. Make the program executable according to the instructions you learned in Hour 1.

When you're done, try running the program by typing the following at a command line:

```
perl -w Primes
```

Listing 3.4. The Co

Summary

In this hour, you learned about Perl's many flow control constructs. Some constructs, such as `if` and the logical operators, are used to control whether portions of the program run, depending on true or false values. Other constructs, such as `while`, `until`, and `for`, are used for looping over pieces of code as many times as necessary. You also learned what Perl's particular idea of truth is, which is used by virtually all test conditions in Perl.

Q&A

Q1: I'm familiar with another programming language, C, which has a `switch` (or `case`) statement. Where is Perl's `switch` statement?

A1: Perl doesn't have one! Perl provides such a variety of tests that figuring out the best syntax for a `switch` statement is nightmarish. The simplest way to emulate a `switch` statement is as follows:

```
if ($variable_to_test == $value1) {  
  
    statement1;  
  
} elsif ($variable_to_test == $value2) {  
  
    statement2;  
  
} else {  
  
    default_statement;  
  
}
```

The online syntax manual page—which you can view by typing `perldoc perlsyn` at a command prompt—contains many clever examples of how to emulate a `switch` statement in Perl, some with very `switch`-like syntax.

Q2: How many `for` (`while`, `if`) blocks can I nest inside each other?

A2: As many as you like, within memory restrictions of your system. Usually, however, if you have deeply nested loops, it is a sign that you should approach the problem differently.

Q3: Help! Perl is giving me the message `Unmatched right bracket` (or `Missing right bracket`). The line number reported is the end of the file!

A3: Somewhere in your program, you've used an open brace (`{`) without a close brace (`}`), or vice versa. Perl can sometimes guess where the typo is in your program, but sometimes not. Because control structures can nest arbitrarily deeply, Perl doesn't know you've made a mistake until it unexpectedly reaches the End of File without finding the balancing brace. A good program editor (such as vi, Emacs, or UltraEdit) has features to help you find mismatched braces. Use one.



< Day Day Up >



Workshop

Quiz

1: The `while` statement loops as long as a condition is true. What statement loops as long as a condition is false?

a.

```
if (not) {}
```

b.

```
while (!condition) {}
```

2: Is the following expression true or false?

```
(0 and 5) || ( ("0" or 0 or "") and (6 and "Hello")) or 1
```

a.

True

b.

False

3: What is the value of `$i` after this loop is run?

```
for($i=0; $i<=10; $i++) { }
```

a.

10

b.

9

c.

11

Answers

A1: b. The `while (! condition) { }` syntax loops until the condition is false.

A2: a. The expression can be reduced in these steps:

`(false) || ((false) and (true)) or true`

`false || false or true`

`true`

A3: c. The test is `$i<=10`, so when the test is finally false, `$i` must be 11. If you got this one wrong, don't worry. It's such a common mistake that it even has a special name among programmers: a **fence post** error or an **off-by-one** error.

Activities

- Modify Listing 3.1 to keep playing the game until a successful guess is made.
- Listing 3.3, as it is written, is actually quite inefficient at finding primes. For example, it analyzes all the even numbers above 2—which cannot possibly be prime. Make additions to the algorithm to make the Primes program more efficient.

 PREV

< Day Day Up >

NEXT 

Hour 4. Stacking Building Blocks: Lists and Arrays

What You'll Learn in This Hour:

- How to fill and empty arrays
- How to examine arrays element by element
- How to sort and print arrays
- How to split scalars into arrays and join arrays back into scalars

Scalars are Perl's singular nouns. They can represent any one thing—a word, a record, a document, a line of text, or a character. Often, though, you need to talk about collections of things—many words, a few records, two documents, fifty lines of text, or a dozen characters.

When you need to talk about many things in Perl, you use *list data*. You can represent list data in three different ways: by using lists, arrays, and hashes.

Lists are the simplest representation of list data. A list is simply a group of scalars. Sometimes they're written with a set of parentheses encasing the scalars, which are separated by commas. For example, `(2, 5, $a, "Bob")` is a list that contains two numbers, a scalar variable `$a`, and the string `"Bob"`. Each scalar in a list is called a *list element*. In keeping with the philosophy of Least Surprise (see [Hour 2](#), "Perl's Building Blocks: Numbers and Strings"), Perl's lists can contain as many scalar elements as you like. Because scalars can also be arbitrarily large, a list can hold quite a lot of data.

To store list data so that you can refer to it throughout your program, you need an array variable. Array variables are represented in Perl with an "at" sign (@) as the type identifier followed by a valid variable name (as discussed in [Hour 2](#), "Perl's Building Blocks: Numbers and Strings"). For example, `@foo` is a valid array variable in Perl. You can have the same name for an array variable as a scalar variable; for example, `$names` and `@names` refer to different things—`$names` to a scalar variable, and `@names` to an array. The two variables have nothing to do with each other.

Individual items in an array are called *array elements*. Individual array elements are referred to by their position within the array, called an *index*. That is, we can refer to the third array element of the array `@foo`, the fifth array element of the array `@names`, and so on.

The other list type, a *hash*, is similar to an array. Hashes will be discussed further in [Hour 7](#), "Hashes."

Putting Things into Lists and Arrays

Putting things into a literal list is easy. As you just saw, the syntax for a literal list is a set of parentheses enclosing scalar values. The following is an example:

```
(5, 'apple', $x, 3.14159)
```

This example creates a four-element list containing the numbers 5, the string 'apple', whatever happens to be in the scalar variable \$x, and pi.

If the list contains only simple strings, and putting single quotation marks around each string gets to be too much for you, Perl provides a shortcut—the `qw` operator. An example of `qw` follows:

```
qw( apples oranges 45.6 $x )
```

This example creates a four-element list. Each element of the list is separated from the others by whitespace (spaces, tabs, or newlines). If you have list elements that have embedded whitespace, you cannot use the `qw` operator. This code works just as though you had written the following:

```
('apples', 'oranges', '45.6', '$x')
```

Notice that the \$x is encased in single quotation marks. The `qw` operator does not do variable interpolation on elements that look like variables; they are treated as though you wanted them that way literally. So '\$x' is not converted to whatever the value of the scalar variable \$x is; it's left alone as a string containing a dollar sign and the letter x.

Perl also has a useful operator that works in literal lists; it's called the *range operator*. The range operator is designated by a pair of periods (...). The following is an example of this operator:

```
(1..10)
```

The range operator takes the left operand (the 1) and the right operand (the 10) and constructs a list consisting

of all the numbers between 1 and 10, inclusive. If you need several ranges in a list, you can simply use multiple operators:

```
(1..10, 20..30);
```

The preceding example creates a list of 21 elements: 1 through 10 and 20 through 30. Giving the range operator a right operand less than the left, such as `(10..1)`, produces an empty list.

The range operator works on strings as well as numbers. The range `(a..z)` generates a list of all 26 lowercase letters. The range `(aa..zz)` generates a much larger list of 676 letter pairs starting with `aa`, `ab`, `ac`, `ad` and ending with `zx`, `zy`, `zz`.

Arrays

Literal lists are usually used to initialize some other structure: an array or a hash. To create an array in Perl, you can simply put something into it. With Perl, unlike other languages, you don't have to tell it ahead of time that you're creating an array or how large the array is going to be. To create a new array and populate it with a list of items, you could do the following:

```
@boys=qw( Greg Peter Bobby );
```

This example, called an *array assignment*, uses the array assignment operator—the equals sign, just as in a scalar assignment. After that code runs, the array `@boys` contains three elements: `Greg`, `Peter`, and `Bobby`. Notice also that the code uses the `qw` operator; using this operator saves you from having to type six quotation marks and two commas.

Array assignments can also involve other arrays or even empty lists, as shown in the following examples:

```
@copy=@original;  
@clean=();
```

Here, all the elements of `@original` are copied into a new array called `@copy`. If `@copy` already had elements before the assignment, they are now lost. After the second statement is executed, `@clean` is empty. Assigning an empty list (or an empty array) to an array variable removes all the elements from the array.

If a literal list contains other lists, arrays, or hashes, these lists are all flattened into one large list. Observe this snippet of code:

```
@boys=qw( Greg Peter Bobby );
@girls=qw( Marcia Jan Cindy );
@kids=(@girls, @boys);
@family=(@kids, ('Mike', 'Carol'), 'Alice');
```

The list (@`girls`, @`boys`) is flattened by Perl to a simple list containing first all the girls' names and then all the boys' names before the values are assigned to @`kids`. On the next line, the array @`kids` is flattened, and the list ('Mike', 'Carol') is flattened into one long list; then that list is assigned to @`family`. The original structures of @`boys`, @`girls`, @`kids`, and the list ('Mike', 'Carol') are not preserved in @`family`—only the individual elements from `Greg` through `Alice`. In other words, the preceding snippet for building @`family` is equivalent to this assignment:

```
@family=qw(Marcia Jan Cindy Greg Peter Bobby Mike Carol Alice
);
```

The left side of an array assignment can be a list if it contains only variable names. The array assignment initializes the variables on that list. Consider this example:

```
($a, $b, $c)=qw(apples oranges bananas);
```

Here, \$`a` is initialized to 'apples', \$`b` to 'oranges', and \$`c` to 'bananas'.

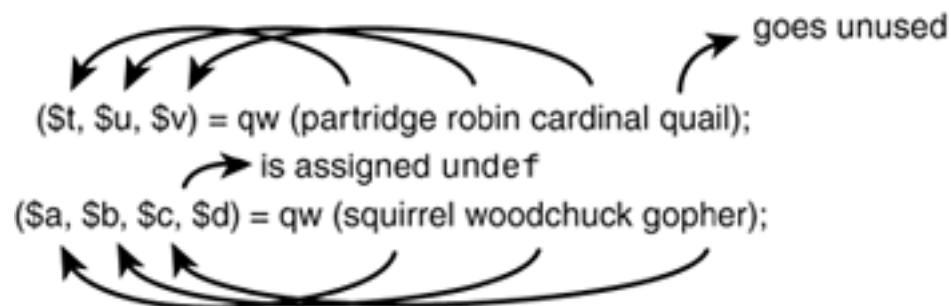
If the list on the left contains an array, that array receives *all* the remaining values from the right side, no matter where it is in the list. The reason is that an array can contain an indefinite number of elements. Observe the following:

```
($a, @fruit, $c) = qw (peaches mangoes grapes cherries);
```

In this example, \$`a` is set to 'peaches'. The remaining fruits in the list on the right are assigned to @`fruit` on the left. No elements are left for \$`c` to receive a value (because the array on the left side of an

assignment absorbs all the remaining values from the right), so `$c` is set to `undef`.

It's also important to note that if the left side contains more variables than it has elements, the leftover variables receive the value `undef`. If the right side has more variables than the list on the left has elements, the extra elements on the right are simply ignored. Here is another example to help understand that concept:



In the first line, `$t`, `$u`, and `$v` all receive a value from the right side. The extra right-side element ('`quail`') is simply not used for this expression. In the second line, `$a`, `$b`, and `$c` all receive a value from the right. `$d`, however, has nothing to get from the right (`$c` takes the last value, '`gopher`'), so `$d` is set to `undef`.

[PREV](#)

< Day Day Up >

[NEXT](#)

Getting Elements Out of an Array

So far in this hour, you've been slinging around whole arrays and lists and putting information into arrays. How can you get that information back out?

One way to get the contents of the entire array is to put the array variable in double quotation marks:

```
print "@array";
```

An array in double quotes is interpolated, and its elements are returned separated by spaces. This example prints the elements of `@array` with a space separating each element.

Many times, though, you need to get to individual elements of arrays. You may need to search for an element, change the value of an element, or to add or remove individual elements in an array.

Individual elements in an array are accessed by a numeric index. The index for array elements starts at the number 0 and increases by 1 for each additional element. Each element of the array has an index value, as shown here:

@trees			
0	1	2	3
oak	cedar	maple	apple

The number of elements in an array is limited only by your system's memory. To access an element, you use the syntax

```
$array[index]
```

where `array` is the array name and `index` is the index of the element you want (also called a *subscript*). The array doesn't have to exist before you refer to individual elements; if it does not already exist, it just automagically springs into existence. Some examples of accessing array elements follow:

```
@trees=qw(oak cedar maple apple);  
  
print $trees[0];           # Prints "oak"  
  
print $trees[3];           # Prints "apple".  
  
$trees[4]='pine';
```

Notice that to talk about an individual element of `@trees`, the code uses a `$`. "I thought the `$` marker was usually reserved for scalars; what's going on?" you might ask. The answer is that the `$` in `$trees[3]` does refer to a scalar: one scalar value within `@trees`. (Scalars are also indicated by a dollar sign because they're singular as well. You should notice a pattern here.)

At the beginning of this hour, you discovered that scalars and arrays can have the same variable names and remain unrelated. Perl can tell the difference between `$trees`, a scalar variable that has nothing to do with the `@trees` array, and `$trees[0]`, the first element in the `@trees` array, because of the square brackets in `$trees[0]`. Perl knows that you're taking about the first element of `@trees` and not talking about `$trees` at all.

You can also talk about a subgroup within an array, called a *slice*. To take a slice of an array, you use both the `@` type identifier—to indicate that you're talking about a group of things—and square brackets—to indicate you're talking about individual elements of an array, as shown here:

```
@trees=qw(oak cedar maple apple cherry pine peach fir);  
  
@trees[3,4,6];           # Just the fruit trees  
  
@conifers=@trees[5,7];   # Just the conifers
```

Finding the End of an Array

Sometimes you need to find the end of the array—for example, to see how many trees are in the `@trees` array or to cut some trees out of the `@trees` array. Perl provides a couple of mechanisms for finding the end. The first is a special variable in the form `$#arrayname`. It returns the number of the last valid index of the array. Check out this example:

```
@trees=qw(oak cedar maple apple cherry pine peach fir);
```

```
print $#trees;
```

This example contains eight elements, but you must remember that arrays are numbered starting at 0. So the preceding example prints the number 7. Modifying the value of `$#trees` changes the length of the array. Making it smaller truncates the array at whatever index you specify, and making it larger gives the array more elements. The newly added elements all have their values set to `undef`.

The other method of finding the size of an array is to use the array variable in a place where a scalar is expected:

```
$size=@array;
```

This puts the number of elements in `@array` into `$size`. This takes advantage of a Perl concept called *context*, explained in the next section.

By the Way

You can also specify negative indexes for arrays. Negative index numbers start counting from the end of the array and work backward. For example, `$array[-1]` is the last element of `@array`, `$array[-2]` is the next to the last element, and so on.

Learning More about Context

What is *context*? Context means the things that surround an item of interest to help define what that item means. For example, seeing a man in surgical scrubs can have different meanings depending on where he is: In a hospital, the fact that the man is wearing scrubs might mean that he's a doctor; at a Halloween party, he could be just another party guest in costume.

Human language uses context to help determine the meaning of words. For example, the word *level* can have several different meanings depending on how it's used and what context it's in:

- The carpenter used a level to hang the door straight.
- The moderator spoke in a level tone.
- The water in the pool was at waist level.

It's the same word each time, but the meaning has changed. It becomes a noun, an adjective, and a different kind of noun depending on how it's used in a sentence.

Perl is also sensitive to context. Functions and operators in Perl can behave differently depending on what context they're used in. The two most important contexts in Perl are *list context* and *scalar context*.

As you've seen, you can use one operator—the equals sign—to perform assignment with both arrays and scalars. The type of expression (list or scalar) on the left side of the assignment operator determines what context the things on the right side are evaluated in, as shown in the following lines of code:

```
$a=$b;      # Scalar on the left: this is scalar context.  
  
@foo=@bar;  # Array on the left: this is list context  
  
($a)=@foo;  # List on the left: this is also list context.  
  
$b=@bar;    # Scalar on the left: this is scalar context.
```

The last line is interesting, because it puts an array into scalar context. As was stated in the previous section, evaluating an array in a scalar context returns the number of elements in the array.

More about the Size and End of an Array

Observe \$a and \$b in the following few lines of code; they do almost the same thing:

```
@foo=qw( water cola juice lemonade );  
  
$a=@foo;  
  
$b=$#foo;  
  
print "$a\n";  
  
print "$b\n";
```

At the end of this code, \$a contains the number 4, and \$b contains the number 3. Why the difference? \$a is @foo evaluated in a scalar context, and it contains the *number of elements*. \$b, on the other hand, is set to the *index of the last element*, and indexes start counting at 0.

Because arrays in a scalar context return the number of elements in the array, testing whether an array contains elements becomes this simple:

```
@mydata=qw( oats peas beans barley );  
  
if (@mydata) {  
  
    print "The array has elements!\n";  
  
}
```

Here, the array `@mydata` is evaluated as a scalar, and it returns the number of elements—in this case, 4. The number 4 evaluates to true in an `if` statement, and the body of the `if` block is run.

By the Way

Actually, `@mydata` here is used in a special kind of scalar context called a *Boolean* context, but it behaves the same way. Boolean context occurs when Perl expects a true or false value, such as in an `if` statement's test expression. One other context, called *void* context, will be explained in [Hour 9, "More Functions and Operators."](#)

Context with Operators and Functions

Many of Perl's operators and functions force their arguments to be either scalar context or list context. Sometimes the operators or functions behave differently depending on what context they're in. Some functions you've already encountered have these properties; however, this fact hasn't been important until now, because they have only had scalars to work on.

The `print` function expects a list as an argument. It doesn't particularly matter what context the list is evaluated in, though. So printing an array with `print` like this causes the array to be evaluated in a list context, yielding the elements of `@foo`:

```
print @foo;
```

You can use a special pseudofunction called `scalar` to force something into a scalar

context:

```
print scalar(@foo);
```

This example prints the number of elements in `@foo`. The `scalar` function forces `@foo` to be evaluated in a scalar context, so `@foo` returns the number of elements in `@foo`. Then the `print` function simply prints the number returned.

The `chomp` function you learned about in [Hour 2](#), "Perl's Building Blocks: Numbers and Strings," takes either an array or a scalar as an argument. If `chomp` is presented with a scalar, it removes the record separator from the end of the scalar. If it is presented with an array, it removes the record separator from the end of each scalar in the array.

 PREV

NEXT 

Manipulating Arrays

Now that you've learned the basic rules for building arrays, it's time to learn some tools to help you manipulate those arrays to perform useful tasks.

Stepping Through an Array

In [Hour 3](#), "Controlling the Program's Flow," you learned about making loops with `while`, `for`, and other constructs. Many tasks you'll want to perform involve examining each element of an array. This process is called *iterating* over the array. One way you could do so is to use a `for` loop, as follows:

```
@flavors=qw( chocolate vanilla strawberry mint sherbet );
for($index=0; $index<@flavors; $index++) {
    print "My favorite flavor is $flavors[$index] and..."
}
print "many others.\n";
```

The first line initializes the array with ice cream flavors, using the `qw` operator for clarity. (If I had included a two-word flavor such as Rocky Road), I would have needed a regular, single-quoted list. The second line does most of the work. `$index` is initialized to 0 and incremented by 1 until `@flavors` is reached. Because it is being compared to the scalar `$index`, `@flavors` is evaluated in a scalar context, so it evaluates to 5—the number of elements in `@flavors`.

The preceding example seems like an awful lot of work for just iterating over an array. Usually in Perl, if something seems like an awful lot of work, you can find an easier way to do it. This is no exception. Perl has another loop statement that wasn't mentioned in [Hour 3](#), called `foreach`. The `foreach` statement sets an index variable, called an *iterator*, equal to each element of a list in turn. Consider this example:

```
foreach $cone (@flavors) {
    print "I'd like a cone of $cone\n";
}
```

Here, the variable `$cone` is set to each value in `@flavors`. As `$cone` is set to each value in `@flavors`,

the body of the loop is executed, printing the message for each value in `@flavors`.

In a `foreach` loop, the iterator isn't just a variable that is assigned the value of each element in the list; it actually refers to that list element itself. If you modify the iterator, the corresponding element in the list will remain modified after the loop is done. Check out this example:

```
foreach $flavor (@flavors) {  
  
    print "I'd like a bowl of $flavor ice cream, please.\n";  
  
    $flavor = "$flavor (I've had some)";  
  
}  
  
print "The available flavors are\n";  
  
foreach $flavor (@flavors) {  
  
    print "$flavor\n";  
  
}
```

In the first loop, the second line prints "I'd like a bowl of chocolate ice cream, please." continuing with `vanilla`, `strawberry`, and so on. The third line, however, modifies `$flavor`, and therefore the corresponding element of `@flavors`, by appending (I've had some) on the end. After the first loop finishes, the second loop lists the flavors, showing for each one the fact that I've had some.

By the Way

In Perl, the `foreach` and `for` loop statements are actually synonyms; they can be used interchangeably. For clarity, throughout this book, you'll find that I use the `foreach()` loop statement to iterate over arrays and the `for()` loop statement for the kind of `for` loops presented in [Hour 3](#), which did not involve arrays. Keep in mind that they are interchangeable.

Converting Between Arrays and Scalars

Perl doesn't have one general rule about converting between scalars and arrays. Rather, Perl provides many functions and operators for converting between the two types.

One method to convert a scalar into an array is the `split` function. The `split` function takes a pattern and a scalar, uses the pattern to split the scalar apart, and returns a list of the pieces. The first argument is the pattern (here surrounded by slashes), and the second argument is the scalar to split apart:

```
@words=split(/ /, "The quick brown fox");
```

After you run this code, `@words` contains each of the words `The`, `quick`, `brown`, and `fox`—without the spaces. If you don't specify a second argument, the variable `$_` is split. If you don't specify a pattern or a string, whitespace is used to split apart the variable `$_`. One special pattern, `//` (the null pattern), splits apart the scalar into individual characters, as shown here:

```
while(<STDIN>) {  
    ($firstchar)=split(//, $_);  
  
    print "The first character was $firstchar\n";  
  
}
```

The first line reads from the terminal one line at a time, setting `$_` equal to that line. The second line splits `$_` apart using the null pattern. The `split` function returns a list of each character from the line in `$_`. That list is assigned to the list on the left side, and the first element of the list is assigned to `$firstchar`; the rest are discarded.

By the Way

The patterns used by `split` are actually *regular expressions*. Regular expressions are a complex pattern-matching language introduced in Hour 6, "Pattern Matching." For now, the examples will use simple patterns such as spaces, colons, commas, and such. After you've learned about regular expressions, I will give examples that use more complex patterns to pull apart scalars with `split`.

This method of splitting a scalar into a list of scalar variables is common in Perl. When you're splitting apart a scalar in which each piece is a distinct element—such as fields in a record—it's easier to figure out which piece is what when you name each piece as it's split. Observe the following:

```
@Music=('White Album', Beatles,  
        'Graceland', Paul Simon,  
        'A Boy Named Sue', Goo Goo Dolls');
```

```
foreach $record (@Music) {  
    ($record_name, $artist)=split(',',$record);  
}
```

When you split directly into a list with named scalars, you can clearly see which fields represent what. The first field is a record name, and the second field is the artist. Had the code split into an array, the distinction between fields might not have been as clear.

To create scalars out of arrays—the reverse of `split`—you can use the Perl `join` function. `join` takes a string and a list, joins the elements of the list together using the string as a separator, and then returns the resulting string. Consider this example:

```
$numbers=join(',', (1..10));
```

This example assigns the string `1,2,3,4,5,6,7,8,9,10` to `$numbers`.

In Perl the output (return value) of one function can be used as an input value (argument) in another function. You can use `split` and `join` to pull a string apart and put it back together all at the same time, as seen here:

```
$message="Elvis was here";  
  
print "The string \"$message\" consists of:",  
  
      join('-', split('//', $message));
```

In this example, the `$message` is split into a list by `split`. That list is used by the `join` function and put back together with dashes. The result is the following message:

The string "Elvis was here" consists of: E-l-v-i-s- -w-a-s- -h-e-r-e

Reordering Your Array

When you're building arrays, often you might want them to come out in a different order than you built them.

For example, if your Perl program reads a list of customers in from a file, printing that customer list in alphabetical order would be reasonable. For sorting data, Perl provides the `sort` function. The `sort` function takes as its argument a list and sorts it in (roughly speaking) alphabetical order; the function then returns a new list in sorted order. The original array remains untouched, as you can see in this example:

```
@Chiefs=qw(Clinton Bush Reagan Carter Ford Nixon);  
  
print join(' ', sort @Chiefs), "\n";  
  
print join(' ', @Chiefs), "\n";
```

This example prints the sorted list of presidents (`Bush Carter Clinton Ford Nixon Reagan`) and then prints the list again in its original order.

Be forewarned that the default sort order is ASCII order. This means that all words that start with uppercase characters sort before words that begin in lowercase letters. Numbers do not sort in ASCII order the way you would expect. They don't sort by value. For example, 11 sorts higher than 100. In cases like this, you need to sort by something other than the default order.

The `sort` function allows you to sort in whatever order you want by using a block of code (or a subroutine name, discussed in [Hour 8](#), "Functions") as the second argument. Inside the block (or subroutine), two variables, `$a` and `$b`, are set to two elements of the list. The block's task is to return `-1`, `0`, or `1` depending on whether the `$a` is less than `$b`, equal to `$b`, or greater than `$b`, respectively. The following is an example of the hard way to do a numeric sort, assuming that `@numbers` is full of numeric values:

```
@sorted=sort { return(1) if ($a>$b);  
  
              return(0) if ($a==$b);  
  
              return(-1) if ($a<$b); } @numbers;
```

(spaceship)> (spaceship operator)>>The preceding example certainly sorts `@numbers` numerically. But the code looks far too complicated for such a common task. As you might suspect for anything this cumbersome, Perl has a shortcut: the "spaceship" operator, `<=>`. The spaceship operator gets its name because it somewhat resembles a flying saucer, seen from the side. It returns `-1` if its left operand is less than the right, `0` if the two operands are equal, and `1` if the left operand is greater than the right:

```
@sorted=sort { $a<=>$b; } @numbers;
```

This code is much cleaner, easier to look at, and more straightforward. You should use the spaceship operator only to compare numeric values.

To compare alphabetic strings, use the `cmp` operator, which works exactly the same way. You can put together more complex sorting arrangements by simply making a more sophisticated sort routine. Section 4 of the Perl Frequently Asked Questions (FAQ) has some more sophisticated examples of this if you need them.

The final function for this hour is an easy function, `reverse`. The `reverse` function, when given a scalar value in a scalar context, reverses the string's characters and returns the reversed string. The call `reverse ("Perl")` in a scalar context, for example, returns `lreP`. When given a list in a list context, `reverse` returns the elements of the list in reverse order, as in this example:

```
@lines=qw(I do not like gree
```

 PREV

< Day Day Up >

NEXT 

Exercise: Playing a Little Game

This hour has really been full of Catch-22's: familiar operators behaving differently depending on context, a handful of new operators and functions, and quite a few new rules to remember about syntax. To keep you from developing any hang-ups, I've added this exercise, which puts your knowledge of arrays and lists to good use—a game.

Using your text editor, type the program from Listing 4.1 and save it as [Hangman](#). As always, don't type the line numbers or their following colons ("1:", etc.) Be sure to make the program executable according to the instructions you learned in Hour 1, "Introduction to the Perl Language."

When you're done, try running the program by typing the following at a command line:

Hangman

or, if your system does not allow making the program executable,

perl -w Hangman

Listing 4.1. Complete Listing of the Hangman Program

```
1:  #!/usr/bin/perl -w
2:
3:  @words=qw( internet answers printer program );
4:  @guesses=();
5:  $wrong=0;
6:
7:  $choice=$words[rand @words];
8:  $hangman=" 0 - | --< ";
9:
10: @letters=split(//, $choice);
```

```
11: @hangman=split(//, $hangman);
12: @blankword=(0) x scalar(@letters);
13: OUTER:
14:     while ($wrong<@hangman) {
15:         foreach $i (0..$#letters) {
16:             if ($blankword[$i]) {
17:                 print $blankword[$i];
18:             } else {
19:                 print "-";
20:             }
21:         }
22:         print "\n";
23:         if ($wrong) {
24:             print @hangman[0..$wrong-1]
25:         }
26:         print "\n Your Guess: ";
27:         $guess=<STDIN>; chomp $guess;
28:         foreach(@guesses) {
29:             next OUTER if ($_ eq $guess);
30:         }
31:         $guesses[@guesses]=$guess;
32:         $right=0;
33:         for ($i=0; $i<@letters; $i++) {
34:             if ($letters[$i] eq $guess) {
35:                 $blankword[$i]=$guess;
```

```
36:                     $right=1;
37:                 }
38:             }
39:             $wrong++ if (not $right);
40:             if (join('', @blankword) eq $choice) {
41:                 print "You got it right!\n";
42:                 exit;
43:             }
44:         }
45:     print "$hangman\nSorry, the word was $choice.\n";
```

[Listing 4.2](#) shows a sample of the Hangman program's output.

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The array `@words` is initialized with the list of possible words that the game can use.

Lines 4–5: Some variables are initialized. `@guesses` is used to hold a list of all past guesses the player has made. `$wrong` holds the number of wrong guesses so far.

Line 7: A word is chosen at random from the array `@words` and assigned to `$choice`. The `rand()` function expects a scalar argument, and because `@words` is being treated as a scalar, it returns the number of elements (in this case, 4). The `rand` function returns a number between 0 and 3—but not including 0 or 4. As it turns out, when you use a decimal number as an array index, the decimal portion is dropped.

Line 8: The hangman is defined. He's not pretty, but he gets the point across.

Line 10: The mystery word in `$choice` is split into individual letters in `@letters`.

Line 11: The hangman scalar is split into pieces in `@hangman`. The head is `$hangman[0]`, the neck is `$hangman[1]`, and so on.

Line 12: The array `@blankword` is used to mark which letters the player has guessed successfully. `(0) x scalar(@letters)` creates a list that is as long as the number of elements in `@letters`, which is stored in `@blankword`. As letters are guessed, these `0`s are changed to letters in line 35—this will mark the positions of the correctly guessed letters.

Lines 13–14: The loop containing most of the program is set up. It has a label—`OUTER`—so that inside the loop you can have some fine-grained control over it. It continues looping until the number of wrong guesses is the same as the length of the hangman.

Lines 15–21: This `foreach` loop iterates over the array `@blankword` for each letter in the puzzle. If `@blankword` doesn't contain a letter in that particular element, a dash is printed; otherwise, the letter is printed.

Lines 23–25: `$wrong` contains the number of wrong guesses. If that number is at least one, line 24 uses a slice to print the hangman array from position 0 up to the number of wrong guesses (less 1).

Lines 26–27: These lines get the guess from the player. `chomp()` removes the trailing newline.

Lines 28–30: These lines search `@guesses` to see whether the player has already guessed that letter. If he or she has, you restart the loop at line 13. The player isn't penalized for duplicate wrong guesses.

Line 31: The letter guessed by the player is recorded in `@guesses`. Using `@guesses` in a scalar context causes the number of elements to be substituted in the brackets `[]`, so each time this statement is run, the index one beyond the last index is set to `$guess`.

Lines 32–38: The meat of the program! The array `@letters`, which contains the puzzle, is searched. If the guess is found in the puzzle, the corresponding element of `@blankword` is set to the letter. The array `@blankword` contains either a correctly guessed letter or `undef` at any particular element. A flag called `$right` is set to 1 to indicate that at least one letter was successfully found.

Line 39: `$wrong` is incremented, unless the player correctly guessed a letter.

Lines 40–43: The elements of the array `@blankword` are joined together to form a string and compared to the original puzzle. If they match, the player has guessed all the letters.

Line 45: The player was unable to guess the puzzle, and the interpreter has dropped out of the loop started at line 13. This line prints a conciliatory message and exits the game.

This game exercises most of the concepts from this hour—literal lists, arrays, `split`, `join`, context, and `foreach` loops. You can implement this small game of Hangman in an almost unlimited number of ways, but I hope that you've caught on to some of the possibilities that arrays provide.

Listing 4.2 shows a sample of the Hangman program's output.

Listing 4.2. Sample Output from Hangman

Your Guess: **t**

----t-

Your Guess: **s**

----t-

O

Your Guess: **e**

----te-

O

Your Guess:

 PREV

< Day Day Up >

NEXT 

Summary

Arrays and lists are Perl's collective variables. You can use them to hold an almost unlimited number of scalars, and you can either manipulate them as a whole or manipulate their individual elements. Perl provides easy mechanisms for copying arrays, sorting arrays, combining arrays, and converting data back and forth between scalars and arrays. Also, many of Perl's operators and functions are sensitive to the context that they appear in, and they behave differently depending on whether they are in a scalar or list context.

Q&A

Q1: *Can you suggest a quick way to find a particular string in an array element?*

A1: Iterating over the array and checking each element are the usual ways to do so. If you frequently find yourself searching an array to see whether an element is in the array, you probably didn't want to store the data in an array in the first place. A much more efficient structure for random-element access is a hash, which is covered in [Hour 7](#).

Q2: *How can I eliminate duplicate elements from an array?*

How can I count the number of unique elements in an array?

How can I see whether two arrays contain the same or different elements?

A2: The answer to all these questions is the same: Use a hash. Hashes allow you to do some rather interesting manipulations on arrays quickly and efficiently. All these questions are answered in [Hour 7](#).

Workshop

Quiz

1: What's an efficient way to swap the values contained in two scalar variables \$a and \$b?

a.

`$a=$b;`

b.

`($a,$b)=($b,$a);`

c.

`$c=$a; $a=$b; $b=$c;`

2: What does the statement `$a=scalar(@array);` assign to the variable \$a?

a.

The number of elements in `@array`

b.

The index of the last element of `@array`

c.

That syntax is not valid

3: What does the statement `$a=@array;` assign to the variable `$a`?

a.

The number of elements in `@array`

b.

The index of the last element of `@array`

c.

That syntax is not valid

Answers

A1: b. The first choice clearly will not work; the value contained in `$a` is destroyed. Choice c answers the question but requires a third variable to hold the data during the swap. Choice b swaps the data correctly, using no extra variables, and is fairly clear code.

A2: a. Using an array in a scalar context returns the number of elements in the array. `$#array` would have returned the last index of the array. The use of `scalar()` in this example is unnecessary; having a scalar on the left side of the assignment operator is enough to put `@array` in a scalar context.

A3: a. Trick question, it's actually the same as #2, except that the context is implied by the assignment instead of being explicitly given with the scalar operator.

Activities

- Modify the Hangman game to print the hangman in an upright position.

Hour 5. Working with Files

What You'll Learn in This Hour:

- How to open and close files
- How to write data to files
- How to read data from files
- How to write Perl defensively so that your programs are robust

Until now, your Perl programs have been self-contained. They have been unable to communicate with the outside world other than to provide messages to the user and receive input from the keyboard. All of that is about to change.

Perl is an outstanding language for reading from and writing to files on disk or elsewhere. Perl's scalars can stretch to hold the longest possible record in a file, and Perl's arrays can stretch to hold the entire contents of files—as long as enough memory is available, of course. When the data is contained within Perl's scalars and arrays, you can perform endless manipulations on that data and write new files.

Perl tries very hard not to get in your way while reading or writing files. In some places, Perl's built-in statements are even optimized for performing common types of file input/output (I/O) operations.

In this hour, you will learn how Perl can give you access to all the data available to you in files.

Opening Files

To read or write files in Perl, you need to open a *filehandle*. Filehandles in Perl are yet another kind of variable. They act as convenient references (handles, if you will) between your program and the operating system about a particular file. They contain information about how the file was opened and how far along you are in reading (or writing) the file; they also contain user-definable attributes about how the file is to be read or written.

From previous hours you're already familiar with one filehandle: [STDIN](#). This filehandle is given to you automatically by Perl when your program starts, and it's usually connected to the keyboard device (you'll learn more details about [STDIN](#) later). The format for filehandle names is the same as that for variable names outlined in [Hour 2](#), "Perl's Building Blocks: Numbers and Strings," except that no type identifier appears in front of the name ([\\$](#), [@](#)). For this reason, it's recommended that filehandle names be in uppercase so that they do not clash with Perl's current or future reserved words: [foreach](#), [else](#), [if](#), and so on.

By the Way

You can also use a string scalar or anything that returns a string—such as a function—as a filehandle name. This type is called an *indirect filehandle*. Describing their use is a bit confusing for a primer in Perl. For more information on indirect filehandles, see the online documentation on the [open](#) function in the perlfunc manual page.

Any time you need to access a file on your disk, you need to create a new filehandle and prepare it by opening the filehandle. You open filehandles, not surprisingly, with the [open](#) function. The syntax of the [open](#) function is as follows:

```
open(filehandle, pathname)
```

The [open](#) function takes a filehandle as its first argument and a pathname as the second argument. The pathname indicates which file you want to open, so if you don't specify a full pathname—such as [c:/windows/system/](#)—open will try to open the file in the current directory. If the [open](#) function succeeds, it returns a nonzero value. If the [open](#) function fails, it returns [undef](#) (false):

```
if (open(MYFILE, "mydatafile")) {  
    # Run this if the open succeeds  
}  
else {  
    print "Cannot open mydatafile!\n";  
    exit 1;  
}
```

In the preceding snippet, if [open](#) succeeds, it evaluates to a true value, and the [if](#) block is run with the open filehandle called MYFILE which is now open for input. Otherwise, the file cannot be opened, and the [else](#) portion of the code is run, indicating an error. In many Perl programs, this "[open or fail](#)" syntax is written using the [die](#) function. The [die](#) function stops execution of your Perl program and prints an error message:

Died at *scriptname* line *xxx*

Here, `scriptname` is the name of the Perl program, and `xxx` is the line number where the `die` was encountered. The `die` and `open` functions are frequently seen together in this form:

```
open(MYTEXT, "novel.txt") || die;
```

This line is read as "open or die," which sums up how you will usually want your program to handle the situation when a file can't be opened. As described in Hour 3, "Controlling the Program's Flow," if the `open` does not succeed—if it returns false—then the logical OR (`||`) needs to evaluate the right-hand argument (the `die`). If the `open` succeeds—if it returns true—then the `die` is never evaluated. This idiom is also written with the other symbol for logical OR, `or`.

When you are done with a file, it is good programming practice to close the filehandle. Closing notifies the operating system that the filehandle is available for reuse and that any unwritten data for the filehandle can now be written to disk. Also, your operating system may allow you to open only a fixed number of filehandles; after that limit is exceeded, you cannot open more filehandles until you close some. To close filehandles, you use the `close` function as follows:

```
close(MYTEXT);
```

If a filehandle name is reused—that is, if another file is opened with the same filehandle name—the original filehandle is first closed and then reopened.

Pathnames

Until now, you've opened only files with simple names like `novel.txt` that did not include a path. When you try to open a filename that doesn't specify a directory name, Perl assumes the file is in the current directory. To open a file that's in another directory, you must use a *pathname*. The pathname describes the path that Perl must take to find the file on your system.

You specify the pathname in the manner in which your operating system expects it, as shown in the following examples:

```
open(MYFILE, "DISK5:[USER.PIERCE.NOVEL]") || die;      # VMS  
open(MYFILE, "Drive:folder:file") || die;      # Macintosh  
open(MYFILE, "/usr/pierce/novel") || die;      # Unix.
```

Under Windows and MS-DOS systems, pathnames contain backslashes as separators—for example, `\Windows\users\pierce\novel.txt`. The only catch is that when you use backslash-separated pathnames in a double-quoted string in Perl, the backslash character sequence gets translated to a special character. Consider this example:

```
open(MYFILE, "\Windows\users\pierce\novel.txt") || die;      # WRONG
```

This example will probably fail, because `\n` in a double-quoted string is a newline character—not the letter *n*—and all the other backslashes will get quietly removed by Perl. As you might guess from Hour 2, "Perl's Building Blocks: Numbers and Strings," one correct way to open the file is by escaping each backslash with another backslash, as follows:

```
open(MYFILE, "C:\\Windows\\users\\pierce\\novel.txt") || die;      # Right, but messy.
```

You can get rid of the double slashes by using the `qq` function as well. However, you can also use forward slash (`/`) in Perl, even under Windows and MS-DOS, to separate the elements of the path. Perl interprets them just fine, as you can see here:

```
open(MYFILE, "C:/Windows/users/pierce/novel.txt") || die; # Much nicer
```

The pathnames you specify can be absolute pathnames—for example, `/home/foo` in UNIX or `c:/windows/win.ini` in Windows—or they can be relative pathnames—`../junkfile` in UNIX or `../bobdir/bobsfile.txt` in Windows. The `open` function can also accept pathnames that are Universal Naming Convention (UNC) pathnames under Microsoft Windows. UNC pathnames are formatted like this:

```
\\\machinename\sharename
```

Perl accepts UNC pathnames with either backslashes or forward slashes and opens files on remote systems if your operating system's networking and file sharing are otherwise set up correctly, as you can see here:

```
open(REMOTE, "//fileserver/common/foofile") || die;
```

On the Macintosh, pathnames are specified by volume, folder, and then file, separated by colons, as shown in [Table 5.1](#).



Reading

You can read from Perl's filehandles in a couple of different ways. The most common method is to use the *file input operator*, also called the *angle operator* (`<>`). To read a filehandle, simply put the filehandle name inside the angle operator and assign the value to a variable:

```
open(MYFILE, "myfile") || die "Can't open myfile: $!";  
  
$line=<MYFILE>;           # Reading the filehandle
```

The angle operator in a scalar context reads one line of input from the file. When called after the entire file has been read, the angle operator returns the value `undef`.

By the Way

A "line of input" is usually considered to be a text stream until the first end-of-line sequence is found. In Unix, that end-of-line sequence is a newline character (ASCII 10); in DOS and Windows, it's the sequence of carriage return and newline characters (ASCII 13,10). This default end-of-line value can be manipulated by Perl to achieve some interesting results. This topic will be covered in [Hour 12](#), "Using Perl's Command-Line Tools."

To read and print the entire file, you can use the following if `MYFILE` is an open filehandle:

```
while(defined($a=<MYFILE>)) {  
  
    print $a;  
  
}
```

As it turns out, a shortcut for reading the filehandle is to use a `while` loop. If the angle operators are the *only* elements inside the conditional expression of a `while` loop, Perl automatically assigns the input line to the special variable `$_` (described in [Hour 2](#), "Perl's Building Blocks: Numbers and Strings") and repeats the loop until the input is exhausted:

```
while(<MYFILE>) {
```

```
print $_;  
}
```

The `while` takes care of assigning the input line to `$_` and making sure the data in the file hasn't been exhausted (called *end of file*). This magic behavior happens only with a `while` loop and only if the angle operators are the only characters in the conditional expression.

Watch Out!

Remember that every line of data read in with a filehandle in Perl contains the end-of-line characters in addition to the text from the line. If you want just the text, use `chomp` on the input line to get rid of the end-of-line characters.

In a list context, the angle operators read in the entire file and assign it to the list. Each line of the file is assigned to each element of the list or array, as shown here:

```
open(MYFILE, "novel.txt") || die "$!";  
  
@contents=<MYFILE>;  
  
close(MYFILE);
```

In the preceding snippet, the remaining data in the filehandle `MYFILE` is read and assigned to `@contents`. The first line of the file `novel.txt` is assigned to the first element in `@contents`: `$contents[0]`. The second line is assigned to `$contents[1]`, and so on.

In most cases, reading an entire file into an array (if it isn't too large) is an easy way for you to deal with the file's data. You can go back and forth through the array, manipulate the array elements, and deal with the array's contents with all the array and scalar operators without worrying because you're actually working with just a copy of the file in the array. Listing 5.1 shows some of the manipulations possible on in-memory files.

Listing 5.1. Reversing a File

```
1:  #!/usr/bin/perl -w  
2:  
3:  open(MYFILE, "testfile") || die "opening testfile: $!";
```

```
4:     @stuff=<MYFILE>;
5:     close(MYFILE);
6:     # Actually, any manipulation can be done now.
7:     foreach(reverse(@stuff)) {
8:         print scalar(reverse($_));
9:     }
```

If the file testfile contains the text

I am the very model of

a modern major-general.

the program in [Listing 5.1](#) would produce the output

```
.lareneg-rojam nredom a
fo ledom yrev eht ma I
```

Line 1: This line contains the path to the interpreter (change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The file testfile is opened with the filehandle `FH`. If the file doesn't open properly, the `die` function is run with an error message.

Line 4: The entire contents of `testfile` are read into the array `@stuff`.

Line 7: The array `@stuff` is reversed—the first line becomes the last line, and so on—and the resulting list is traversed by the `foreach` statement. Each line of the reversed list is assigned to `$_` and the body of the `foreach` loop is executed.

Line 8: Each line (now in `$_`) is itself reversed—from left-to-right to right-to-left—and printed. The `scalar` function is needed because `print` expects a list; also, `reverse` used in a list context reverses a list, so nothing would happen to `$_`. The `scalar` function forces `reverse` into a scalar context, and it reverses `$_` character by character.

Probably, only small files should be reading in their entirety into array variables for manipulation. Reading a very large file into memory, although allowed, might cause Perl to use all the available memory on your system.

If you ever exceed Perl's memory by reading too large a file into memory, or do anything else to exceed your system's memory, Perl displays the following error message:

Out of memory!

and your program terminates. If this happens when you are reading an entire file into memory at once, you should probably consider processing the file one line at a time.

 PREV

< Day Day Up >

NEXT 

Writing

To write data to a file, you must first have a filehandle open for writing. Up till now, all `open` statements you have seen have opened the filehandle for reading only. The syntax for opening a file for writing is almost identical to that for reading:

```
open(filehandle, ">pathname")
```

```
open(filehandle, ">>pathname")
```

The first syntax line should look familiar, except for the `>` in front of the pathname. The `>` signifies to Perl that the file specified at `pathname` should be overwritten with new data, that any existing data should be discarded, and that `filehandle` is open for writing. In the second example, `>>` tells Perl to open the file for writing but, if the file already exists, not to discard the contents but to append the new data to the end of it. Check out these examples:

```
# Overwrite existing data, if any  
  
open(NEWFH, ">output.txt") || die "Opening output.txt: $!";  
  
# Simply append to whatever data may be there.  
  
open(APPFH, ">>logfile.txt") || die "Opening output.txt: $!";
```

By the Way

Until now, it's been almost impossible for your Perl programs to *harm* anything. Now that you know how to write to files, you must be very careful to write only to files that you mean to change. On systems in which the operating system files are vulnerable (Windows 95/98, Mac), you can damage your operating system by carelessly writing to files. Be very aware of what files you are writing to. Recovering the data in files accidentally opened with `>` is nearly impossible. Cleaning out data in files opened accidentally with `>>` also is difficult, so be careful.

When you're done with a filehandle opened for writing, closing the filehandle is especially important. Your operating system doesn't commit data to disk as you write it; it buffers the data up and writes it occasionally. The `close` function notifies the operating system that you are done writing and that the data should be moved to permanent storage on disk:

```
close(NEWFH);
```

```
close(APPFH);
```

If your program terminates with `exit`, or simply "falls off" the end, any written but not-yet-flushed information in filehandles is committed (written out to the file) as if you had closed the filehandles yourself.

After you've opened the filehandle for writing, actually putting data in the file is easy, and you're already familiar with the `print` function. Until now, you've been using `print` to simply display data to the screen. The `print` function is actually used for writing to any filehandle. The syntax for printing to filehandles is as follows:

```
print filehandle LIST
```

`filehandle` is the filehandle you want the data to be written to, and `LIST` is the list of things you want written.

In the `print` syntax, notice that no comma appears between the filehandle name and the list; this point is important. Inside the list, commas are used to separate items, as you've done until now. The lack of a comma between the filehandle and list signifies to Perl that the token following `print` is a filehandle and not the first element in the list. If you include this comma and have Perl's warnings turned on, Perl warns you with the following message: `No comma allowed after filehandle`.

Now consider the following:

```
open(LOGF, ">>logfile") || die "$!";  
  
if (! print LOGF "This entry was written at", scalar(localtime), "\n" ) {  
    warn "Unable to write to the log file: $!";  
}  
  
close(LOGF);
```

In this snippet, the file named `logfile` is opened for appending. The `print` statement writes a message to the `LOGF` filehandle. The return value from `print` is checked, and if false (meaning that the log entry could

not be printed), a warning is issued. The filehandle is then closed.

You can have multiple filehandles open for reading and writing at the same time, as this snippet demonstrates:

```
open(SOURCE, "sourcefile") || die "$!";  
  
open(DEST, ">destination") || die "$!";  
  
@contents=<SOURCE>; # Slurp in the source file.  
  
print DEST @contents; # Write it out to the destination  
  
close(DEST);  
  
close(SOURCE);
```

This snippet implements a simple file copy. As a matter of fact, you can do the read and the write at the same time and shorten the routine, just a little:

```
print DEST <SOURCE>;
```

Because `print` expects a list as an argument, `<SOURCE>` is evaluated in a list context. When the angle operator is evaluated in a list context, the entire file is read in and then printed to the filehandle `DEST`.

 PREV

< Day Day Up >

NEXT 

Free Files, Testing Files, and Binary Data

Files and file systems don't contain just the data that you put into files. On the one hand, filehandles sometimes represent more than a simple file. They may stand for such things as the keyboard, the monitor screen, network sockets, or mass storage devices such as tape drives.

Also, in addition to the data stored in the files, a file system contains what's called *metadata*—that is, data about the data in its files. Perl can be used to get metadata out of the file system; for example, to determine how big your file is, when it was last changed, who changed it, and information about what's in the file. With some operating systems, the file metadata can even determine whether the file is treated as a text file or a binary file.

Free Filehandles

Perl began life as a Unix utility, and sometimes pieces of its heritage show through—even on non-Unix platforms. When your Perl program is started, it receives three filehandles that are "free"—that is, they are opened automatically without you having to do any programming work to open them. They are **STDOUT** (standard output), **STDIN** (standard input), and **STDERR** (standard error). By default, they are connected to your terminal.

As you type, Perl can read your input from the **STDIN** filehandle:

```
$guess=<STDIN>;
```

When you want to display output, you use **print**. The reason why your output has appeared on the screen is that, by default, **print** uses the **STDOUT** filehandle:

```
print "Hello, World!\n";                                # is the same as...
print STDOUT "Hello, World!\n";
```

In Hour 12, you'll learn how to change **print**'s default filehandle.

STDERR is usually set to your terminal, like **STDOUT**, but it is used for displaying error messages. In Unix, error messages and normal output can be sent to different display devices, and writing error messages to the **STDERR** filehandle is traditional. The **die** and **warn** functions both write their messages to **STDERR**. If your operating system does not have a separate error-reporting filehandle—for example, Windows or MS-DOS—your **STDERR** output will go to the **STDOUT** device.

By the Way

Redirecting error and output messages in Unix is well beyond the scope of this book, and the technique varies depending on what shell you use. Any good book on using Unix should cover this topic thoroughly.

Text Files and Binary Files

Some operating systems—such as VMS, Atari ST, and notably Windows and MS-DOS—make the distinction between binary files (raw) and text files. This distinction causes problems because Perl can't really tell the difference—and you wouldn't want it to.

Text files are simply records that end in end-of-line characters, called *record separators*. Binary files, on the other hand, are collections of bits that need to be treated literally, such as images, executable programs, and data files.

When you're writing a text file, Perl translates the `\n` character sequence into the record separator that your operating system uses. In Unix, `\n` becomes an ASCII 10 (LF); on a Macintosh, ASCII 13 (CR); and on DOS and Windows systems, it becomes the sequence ASCII 13 and ASCII 10 (CRLF). When you're writing text, this behavior is appropriate.

When you're writing binary data—GIF files, EXE files, MS Word documents, and so on—translation isn't what you want. Any time you really need to write binary data, and don't want Perl or the operating system to translate it for you, you must use the `binmode` function to mark the filehandle as binary. Use `binmode` after the filehandle is opened but before you do any input or output from it:

```
open(FH, ">camel.gif") || die "$!";
binmode(FH);           # The filehandle is now binary.
# Start of a valid GIF file...
print FH "GIF87a\056\001\045\015\000";
close(FH);
```

You have to use `binmode` on the filehandle only once, unless you close it and reopen it. Using `binmode` on systems that do not distinguish between binary files and text files (Unix, Macintosh) causes no harm.

File Test Operators

Before you open a file, sometimes it's nice to know whether the file exists, whether the file is really a directory, or whether opening the file will give a [permission denied](#) error. If you could examine the file's metadata, you could get answers to these questions. For these situations, Perl provides the *file test operators*. The file test operators all have the following syntax:

`-X filehandle`

`-X pathname`

Here, `X` is the particular test you want performed, and `filehandle` is the filehandle you want tested. You can also test a `pathname` without having an open filehandle. Table 5.1 lists some of the operators.

Table 5.1. Short List of File Test Operators

Operator	Example	Result
<code>-r</code>	<code>-r 'file'</code>	Returns true if ' <code>file</code> ' is readable
<code>-w</code>	<code>-w \$a</code>	Returns true if the filename contained in <code>\$a</code> is writeable
<code>-e</code>	<code>-e 'myfile'</code>	Returns true if ' <code>myfile</code> ' exists
<code>-z</code>	<code>-z 'data'</code>	Returns true if ' <code>data</code> ' exists but is empty
<code>-s</code>	<code>-s 'data'</code>	Returns size of ' <code>data</code> ' in bytes if it exists
<code>-f</code>	<code>-f 'novel.txt'</code>	Returns true if ' <code>novel.txt</code> ' is a regular file rather than a directory
<code>-d</code>	<code>-d '/tmp'</code>	Returns true if ' <code>/tmp</code> ' is a directory
<code>-T</code>	<code>-T 'unknown'</code>	Returns true if ' <code>unknown</code> ' appears to be a text file
<code>-B</code>	<code>-B 'unknown'</code>	Returns true if ' <code>unknown</code> ' appears to be a binary file
<code>-M</code>	<code>-M 'foo'</code>	Returns the age (in days) since the file ' <code>foo</code> ' was modified since this program began

You can view the full list of file test operators in the online documentation. Type `perldoc perlfunc` at a command prompt, and look in the section "Alphabetical List of Perl Functions."

The following snippet uses file test operators to verify that files don't already exist before they are overwritten and to determine the age of the file since it was last modified:

```
print "Save data to what file?";  
  
$filename=<STDIN>;  
  
chomp $filename;  
  
if (-s $filename) {  
  
    warn "$file contents will be overwritten!\n";  
  
    warn "$file was last updated ",  
  
        -M $filename, "days ago.\n";  
  
}  
  
}
```

 PREV

< Day Day Up >

NEXT 

Summary

This hour you learned about opening and closing filehandles in Perl: You open files by using the [open](#) functions and close them by using [close](#). When filehandles are open, they can be read with [<>](#) or [read](#) and written with [print](#). Also, you learned about some of the oddities of how your operating system handles files and how to deal with them by using [binmode](#).

Along the way, I hope that you also learned something about defensive programming.

Q&A

Q1: *My open statement keeps failing, and I'm not sure why. What's wrong?*

A1: First, check the syntax of the `open` statement. Make sure you're opening the right filename. Print the name before the `open` if you need to be sure. If you intend to write to the file, make sure you put a `>` in front of the filename; you need to. Most importantly, did you check the exit status of `open` by using `open() || die "$!"`; syntax? The `die` message might be very important in helping you find your mistake.

Q2: *I'm writing to the file, but nothing seems to go into it. Where's my output going?*

A2: Are you sure that the filehandle opened properly? If you used the wrong filename, your data could be going to the wrong file. A common mistake is to open a file for writing by using backslashes in the pathname and enclosing the pathname in double quotation marks, as shown here:

```
open(FH, ">c:\temp\notes.txt") || die "$!" ; #WRONG!
```

This line creates a file called `c:(tab)emp(newline)otes.txt`—probably not what you had in mind. Also, make sure that your `open` function succeeded. If you write to a filehandle that hasn't been properly opened, Perl discards the output silently un

Workshop

Quiz

1: To open a file named `data` for writing, you should use which of the following:

a.

```
open(FH, "data", write);
```

b.

```
open(FH, "data"); and simply print to FH
```

c.

```
open(FH, ">data") || die "Cannot open data:  
$!";
```

2: `(-M $file > 1 and -s $file)` is true if

a.

`$file` has been modified more than one day ago and has data.

b.

That expression cannot be true.

c.

`$file` is writeable and has no data.

Answers

A1: c. Choice a is false because this is not how `open` opens a file for writing; b is false because it opens the filehandle only for reading. Choice c is true because it does what is asked—and uses good form by checking for errors.

A2: a. `-M` returns the number of days old a file is (>1 is more than 1 day), and `-s` returns true if the file has data in it.

Activities

- Modify the Hangman program from Hour 4, "Stacking Lists and Arrays," to take the list of possible words from a data file.



< Day Day Up >



Hour 6. Pattern Matching

What You'll Learn in This Hour:

- How to construct simple regular expressions
- How to use regular expressions to match patterns
- How to edit strings using regular expressions

In the preceding hour, you learned about reading data from files. With that information, combined with your knowledge of scalars, arrays, and operators, you're prepared to manipulate that data to do what you want—almost. The data in the file might not have an easy-to-use format that can be broken up with a simple `split` on spaces. You might have lines that contain data you're not interested in and want to edit out.

What you need is the ability to recognize patterns in the input stream, to pick and choose data based on those patterns, and possibly to edit the data to a more usable form. One of Perl's tools for performing these tasks is regular expressions. Throughout this text, **regular expression** and **pattern** will be used almost interchangeably.

Regular expressions are almost a language unto themselves. Regular expressions are a formal method of describing patterns to match. In this hour, you'll learn just a little bit of this pattern-matching language.

By the Way

The online documentation has a much deeper (but more terse) description of the full regular expression language used by Perl. You can look at the `perlre` documentation included with Perl. The topic is so deep that an entire book has been published on regular expressions. This book, highly recommended by the Perl community, is called *Mastering Regular Expressions* by Jeffrey E. F. Friedl (Sebastopol: O'Reilly, 2002). It is about regular expressions in general, but close attention is paid to Perl.

Regular expressions are used in other programming languages as well, including TCL, JavaScript, and Python. Many of the Unix operating system utilities also use regular expressions. Perl happens to have a very rich set of expressions—very similar to those used in other systems, but more powerful—and learning them will help you not just in Perl, but in other languages as well.

Simple Patterns

In Perl, patterns are enclosed inside a pattern match operator, which is sometimes represented as `m//`. A simple pattern might appear as follows:

```
m/Simon/
```

The preceding pattern matches the letters *S-i-m-o-n* in sequence. But where is it looking for *Simon*? Previously, you learned that the Perl variable `$_` is frequently used when Perl needs a default value. Pattern matches occur against `$_` unless you tell Perl otherwise (which you'll learn about later). So the preceding pattern looks for *S-i-m-o-n* in the scalar variable `$_`.

If the pattern specified by `m//` is found anywhere in the variable `$_`, the match operator returns true. Thus, the normal place to see pattern matches is in a conditional expression, as shown here:

```
if (m/Piglet/) {  
    # the pattern "Piglet" is in $_  
}
```

Inside the pattern, every character matches itself unless it is a **metacharacter**. Most of the "normal" characters match themselves: *A* to *Z*, *a* to *z*, and digits. Metacharacters are characters that change the behavior of the pattern match. The list of metacharacters is as follows:

```
^ $ ( ) \ | @ [ { ? . + *
```

You'll shortly read about what the metacharacters do. If your pattern contains a metacharacter that you want to match for its literal value, simply precede the metacharacter with a backslash, as shown here:

```
m/I won \$10 at the fair/;    # The $ is treated as a literal dollar sign.
```

Earlier, you read that the pattern match operator is *usually* represented by `m//`. In reality, you can replace the

slashes with any other character you want, such as the commas in the following example:

```
if ($m,Waldo,) { print "Found Waldo.\n"; }
```

The slash or other character that marks the beginning and end of the pattern is called the **delimiter**. Often you replace the delimiter when the pattern contains slashes (/) and the end of the pattern could be confused with the slashes inside the pattern. If you stick with slashes to delimit that pattern, the enclosed slashes need to have backslashes in front of them, as shown here:

```
if ($m/>\usr/>\local/>\bin/>\hangman/) { print "Found the hangman game!" }
```

By changing the delimiter, you could write the preceding example more legibly as follows:

```
if ($m:/usr/local/bin/hangman:) { print "Found the hangman game!" }
```

If the delimiters around the pattern are slashes, you also can write the pattern match without the \$m. This way, you also can write \$m/Cheetos/ as /Cheetos/. Normally, unless you need to use delimiters other than slashes (//), you write pattern matches with just slashes and no \$m.

Variables can also be used in a regular expression. If a scalar variable is seen in a regular expression, Perl first evaluates the scalar and interpolates it, just as in a double-quoted string (recall Hour 2, "Perl's Building Blocks: Numbers and Strings"); then it examines the regular expression. This capability allows you to build regular expressions dynamically. The regular expression in the following if statement is based on user input:

```
$pat=<STDIN>; chomp $pat;  
  
$_="The phrase that pays";  
  
if (/ $pat /) { #Look for the user's pattern  
  
    print "\$_ contains the pattern $pat\n";  
  
}
```

By the Way

Regular expressions in the manual pages and in other documentation are sometimes called *REs* or *regexps*. For clarity, I'll continue to refer to them as regular expressions throughout this book.

Rules of the Game

As you begin to write regular expressions in Perl, you should know that a few rules are involved in the way Perl interprets them. There are not many rules, though, and most of them make sense after you think about them. They are as follows:

1.

Normally, pattern matches start at the left of the target string and work their way to the right.

2.

Pattern matches return true (in whatever context) if and only if the entire pattern can be used to match the target string.

3.

The first possible match (the leftmost) in the target string is matched first. Regular expressions don't leave behind one good match to go looking for another further along. However...

4.

The largest possible first match is taken. Your regular expressions might find a match immediately and then try to stretch that match as far as possible. Regular expressions are *greedy*, meaning they try to match as much as possible.

The Metacharacters

In all the examples that follow, the portions of text matched by patterns are represented with underlines. Remember that the entire target string is said to match even if just a portion of it matches the regular expression. The underline marks are to help demonstrate exactly what part of the target they match.

It's important to read through the following sections; however, don't worry if the information doesn't immediately make sense; it will shortly. The application of these metacharacters will be demonstrated in a bit.

A Simple Metacharacter

The first of the metacharacters is the dot (.). Inside a regular expression, the dot matches any single character except a newline character. For example, in the pattern `/p.t/`, the `.` matches any single character. This pattern would match pot, pat, pit, carpet, python, and pup_tent. The `.` requires that one character be there between the p and the t, but no more. Thus, the pattern would not match *apt* (no character at all between p and t) or *expect* (too many characters between p and t).

The Unprintables

Earlier you read that, to include a metacharacter inside a regular expression, you have to precede the character with a backslash, as shown here, to make it lose its meta-ness:

```
/\^\$/; # A literal caret and dollar sign
```

When preceded by a backslash, normal characters *become* metacharacters. As you saw in [Hour 2](#), some characters take on special meaning in (double-quoted) string literals when they are preceded by a backslash; almost all those same characters represent the same values in regular expressions, as shown in [Table 6.1](#).

Table 6.1. Special Characters

Character Matches

\n	A newline character
\r	A carriage return
\t	A tab
\f	A formfeed

Quantifiers

Until now, all the characters in patterns, whether text characters or metacharacters, have had a one-to-one relationship with characters in the target string they were trying to match. For example, in `/Simon/`, `S` matches an `S`, `i` matches an `i`, `m` matches an `m`, and so on. A **quantifier** is a kind of metacharacter that tells the regular expression how many consecutive occurrences of something to match. A quantifier can be placed after any single character or a group of characters (you'll learn more details on that topic momentarily).

The simplest quantifier is the `+` metacharacter. The `+` causes the preceding character to match at least once, or as many times as it can and still have a matching expression. Thus, `/do+g/` would:

Match These But Not These Why Not

hound <u>dog</u>	badge	The required <code>o</code> is missing.
hot <u>dog</u>	doofus	The <code>g</code> is missing.
<u>doogie</u> howser	Doogie	<code>D</code> is not the same as <code>d</code> .
<u>doooooog</u> doog	pagoda	The <code>d</code> , <code>o</code> , and <code>g</code> do not appear in order.

The `*` metacharacter is similar to the `+` metacharacter, but it causes the preceding character to be matched zero or more times. In other words, the `/t*/` pattern means to match as many `t`'s as possible, but if none exist, that's okay. Thus, `/car*t/` would:

Match These But Not These Why Not

<u>carted</u>	carrot	The <code>o</code> intrudes into the pattern, but the <code>ed</code> follows the pattern.
<u>cat</u>	carl	The <code>t</code> in the pattern isn't optional, but the <code>r</code> is.
<u>carrrt</u>	caart	The <code>a</code> in the pattern can't be repeated, but the <code>r</code> can.

One step down from the `*` metacharacter is `?`. The `?` metacharacter causes the preceding character to be matched either zero times or once (but no more). So the pattern `/c?ola/` causes a `c` to be matched if it's available; otherwise, that's okay. Then it is followed by `o`, `l`, and `a`; essentially, this pattern matches any string with `ola` in it, and if `ola` is preceded by a `c`, that string is matched as well.

The difference between the `?` and `*` metacharacters is that `/c?ola/` would match `cola` and `ola`, but not `ccola`. The extra `c` requires *two* matches. The pattern `/c*ola/` would match `cola`, `ola`, and `ccola` because the `c` can be repeated as many times as necessary, not just zero or one time.

If matching zero, one, or many occurrences of a pattern isn't specific enough for you, Perl allows you to match exactly as many occurrences as you need by using braces, `{ }`. The quantifier with braces has the following format:

`pat{n,m}`

Here, *n* is the minimum number of matches, *m* is the maximum number of matches, and `pat` is the character or group of characters you're trying to quantify. You can omit either *n* or *m*, but not both. Consider the following examples:

`/x{5,10}/` `x` occurs at least 5 times, but no more than 10.

`/x{9,}/` `x` occurs at least 9 times, possibly more.

`/x{0,4}/` `x` occurs up to 4 times, possibly not at all.

`/x{8}/` `x` must occur exactly 8 times.

A common idiom in regular expressions is `.*`. You can use it to match *anything*—usually anything between two other things that you're interested in. For example, `/first.*last/` attempts to match the word *first*, followed by anything, and then the word *last*. Observe how `/first.*last/` matches the following strings:

first then last

The good players get picked **first, the bad last**.

The first shall be last, and the last shall be first.

Look at the match in the third line carefully. The match starts on the word *first* as expected. The match then matches the word *last*, but it doesn't consider itself done. It continues searching until it finds the second (and final) occurrence of the word *last*. Here, the `*` follows the fourth rule listed in the section "[Rules of the Game](#)": It matches the largest possible string, while still completing the match. Often, matching the largest string is not what you want, so Perl offers another solution called **minimal matching**, which is documented further in the perlre manual page.

Character Classes

Another common practice in regular expressions is to ask for a match of "any of these characters." If you're trying to match numbers, it would be nice to be able to write a pattern that matches "any digit 0–9"; if you're searching a list of names and want to match *Van Beethoven* and *van Beethoven*, a pattern that matches "either v or V" would be helpful.

Perl's regular expressions have such a tool; it's called a **character class**. To write a character class, you enclose the characters it contains in square brackets, `[]`. Characters in a character class are treated as a single character during the match. Inside a character class, you can specify ranges of characters (where ranges make sense) by putting a dash between the upper and lower bounds. The following are some examples:

Character Class Explanation

[abcde]	Match any of <i>a</i> , <i>b</i> , <i>c</i> , <i>d</i> , or <i>e</i>
[a-e]	Same as above; match any of <i>a</i> , <i>b</i> , <i>c</i> , <i>d</i> , or <i>e</i>
[ls]Gg[rs]	Match an uppercase <i>G</i> or lowercase <i>g</i>
[0-9]	Match a digit
[0-9]+	Match one or more digits in sequence
[A-Za-z]{5}	Match any group of five alphabetic characters
[* !@#\$%& ()]	Match any of these punctuation marks

The last example is interesting because the characters in that class are usually metacharacters. Inside a character class, most metacharacters lose their "meta-ness"; in other words, they behave like any other ordinary character. Thus, the ***** really represents a literal *****.

If a caret (**^**) occurs as the first character of a character class, the character class is negated. That is, the character class matches any single character that is *not* in the class, as in this example:

```
/[^A-Z]/; # Matches non-uppercase-alphabetic characters.
```

Because **]**, **^**, and **-** are special in a character class, some rules apply about trying to match those characters literally in a character class. To match a literal **^** in a character class, you must make sure it does not occur first in the class. To match a literal **]**, you either need to put it first in the class or put a backslash in front of it (for example, `/[abc\]]/`). To put a literal hyphen (**-**) in a character class, you can simply put it first in the class or put a backslash in front of it.

Perl contains shortcuts for certain commonly used character classes. They are represented by a backslash and a nonmetacharacter, as shown in [Table 6.2](#).

Table 6.2. Special Character Classes

Pattern Matches

\w	A word character; same as [a-zA-Z0-9_]
\W	A nonword character (the inverse of \w)
\d	A digit; same as [0-9]

\D	A nondigit
\s	A whitespace character; same as [\t\f\r\n]
\S	A nonwhitespace character

The following are some examples:

```
/\d{5}/;      # Matches 5 digits
/\s\w+\s/;    # Matches a group of word-characters surrounded by whitespace
```

Be careful, though. The last example here doesn't necessarily match a word; it can also match an underscore surrounded by spaces. Also, not all words are matched by the last pattern; they need to have whitespace around them, and words such as "don't" wouldn't be matched because of the apostrophe. You'll learn better patterns for word matching later in this hour.

Grouping and Alternation

Sometimes in a regular expression, you might want to know whether any of a set of patterns is found. For example, does this string contain *dogs* or *cats*? The regular-expression solution to this problem is called **alternation**. Alternation happens in a regular expression when possible matches are separated with a | character, as in this example:

```
if (/dogs|cats/) {
    print "\$_ contains a pet\n";
}
```

Alternation can be fun, but it also can be tedious when you want to match lots of similar things. For example, if you want to match the words *frog*, *bog*, *log*, *flog*, or *clog*, you could try the expression /frog|bog|log|flog|clog/ except that it's horribly repetitive. What you really want is to alternate on just the first part of the string like this:

```
/fr|b|l|f1|clog/;      # Doesn't QUITE work.
```

The preceding example doesn't quite work because Perl has no way of knowing that the alternations are one thing you want to match and *og* is another.

To solve this problem, you can use Perl's regular expressions to group parts of the pattern with parentheses, (), as shown here:

```
/(\fr|b|l|f1|cl)og/;
```

You can nest parentheses to have groups within groups. For example, you could write the preceding expression as `/(fr|b|(f|c)?l)og/` as well.

In a list context, the match operator returns a list of the portions of the expression matched that were in parentheses. Each parenthesized value is a return value to the list, or 1 if the pattern contains no parentheses. Check out this example:

```
$_= "apple is red";  
($fruit, $color)=/(.* )\sis\s(.*)/;
```

In this snippet, the pattern matches anything (as a group), and then whitespace, the word *is*, more whitespace, and then anything (also as a group). The two grouped expressions are returned to the list on the left side and assigned to `$fruit` and `$color`.

Anchors

The last two metacharacters (I bet you thought they'd never end) are the anchors. You use anchors to tell the regular expression engine exactly where you want to look for the pattern—at the beginning of a string or at the end.

The first of these anchors is the caret (^). The caret at the beginning of a regular expression causes the expression to match only at the beginning of a line. For example, `/^video/` matches the word *video* only if it occurs at the beginning of a line.

Its counterpart is the dollar sign (\$). The dollar sign at the end of a regular expression causes the pattern to match only at the end of the line. For example, `/earth$/` matches *earth*, but only at the end of a line.

Patterns	What They Do
<code>/^Help/</code>	Matches only lines that begin with <i>Help</i> .
<code>/^Frankly.*darn\$/</code>	Matches lines that begin with <i>Frankly</i> and end in <i>darn</i> . Everything in between is matched as well.

/^hysteria\$/

Matches lines that contain only the word *hysteria*.

/^\\$/

Matches the beginning of a line, followed immediately by the end of the line. That is, it matches only blank lines.

 PREV

< Day Day Up >

NEXT 

Substitution

Just finding patterns in strings and lines of input isn't enough; sometimes you need to modify the data as well. One way—but certainly not the only way—is to use the substitution operator `s///`. The syntax is as follows:

```
s/searchpattern/replacement/;
```

The substitution operator searches `$_` by default for `searchpattern` and replaces the entire matched regular expression with `replacement`. The operator returns the number of matches or substitutions performed, or 0 if no matches were made. The following is an example:

```
$_= "Our house is in the middle of our street".  
  
s/middle/end/;          # Is now: Our house is in the end of our street  
  
s/in/at/;              # Is now: Our house is at the end of our street.  
  
if (s/apartment/condo/) {  
  
    # This code isn't reached, see explanation.  
  
}
```

Here, the substitutions happen as you would expect. The word *middle* is changed to *end*, and *in* is changed to *at*. The `if` statement, however, fails, because the word *apartment* does not appear in `$_` and therefore can't be substituted.

The substitution operator can also use delimiters other than slashes (`/`), just as the match operator can. Simply put whatever delimiter you want immediately after the `s`, as shown here:

```
s#street#avenue#;
```

Exercise: Cleaning Up Input Data

The "blind" substitutions in the preceding example—where substitutions are made, but the exit status isn't checked—are common when you're trying to *cook data*. Cooking data is taking data from a user or a file that is not formatted exactly the way you would like and reformatting it. Listing 6.1 shows a routine to convert your weight on the earth to your weight on the moon, which demonstrates data manipulation.

Using your text editor, type the program in Listing 6.2 and save it as `Moon`. Of course, don't type the line numbers. Be sure to make the program executable according to the instructions you learned in Hour 1, "Getting Started with Perl."

When you're done, try running the program by typing the following at a command line:

Moon

or, if you cannot make the program executable,

perl -w Moon

Some sample output is shown in Listing 6.1.

Listing 6.1. Sample Output from Moon

```
1: $ perl Moon
2: Your weight: 150lbs
3: Your weight on the moon: 25.00005 lbs
4: $ perl Moon
5: Your weight: 90 kg
6: Your weight on the moon: 30.00066 lbs
```

Listing 6.2. Your Moon Weight

```

1:  #!/usr/bin/perl -w
2:
3:  print "Your weight:";
4:  $_=<STDIN>;
5:  chomp;
6:  s/^[\s]+//;    # Remove leading spaces, if any.
7:  if (m/(lbs?|kgs?|kilograms?|pounds?)/i) {
8:      if (s/[\s]*(kgs?|kilograms?).*/);
9:          $_*=2.2;
10:     } else {
11:         s/[\s]*(lbs?|pounds?).*/;
12:     }
13: }
14: print "Your weight on the moon: ", $_*.16667, " lbs\n";

```

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Lines 3–5: These lines prompt the user for his weight, assign the input to `$_`, and `chomp` off the newline character. Remember that `chomp` changes `$_` if no other variable is specified.

Line 6: The pattern `/^\s+ /` matches whitespace at the beginning of the line. No replacement string is listed, so the portion of `$_` matching the pattern is simply removed.

Line 7: If a unit of measurement is found in the user's input, this `if` block removes the unit and converts it, if applicable.

Lines 8–9: The pattern `/[\s]*(kgs?|kilograms?)/i` matches whitespace and then either *kg* or *kilogram* (each with an optional *s* on the end). This means that if the input contains `kg` or `kg` (with no space), it is removed. If the pattern is found and removed, what is left over in `$_` is multiplied by 2.2—in other words, converted to pounds.

Line 11: Otherwise, *lbs* or *pounds* is removed from `$_` (along with optional leading whitespace).

Line 14: The weight in `$_`—converted to pounds already—is multiplied by 1/6 and printed.

 PREV

< Day Day Up >

NEXT 

Pattern Matching Odds and Ends

Now that you can match patterns against `$_` and you know the basics of substitution, you're ready for more functionality. To be really effective with regular expressions, you need to match against variables other than `$_`, be able to do sophisticated substitutions, and work with Perl's functions that are geared toward—but not exclusive to—regular expressions.

Working with Other Variables

In Listing 6.2, the weight gathered from the user is stored in `$_` and manipulated with substitution operators and matching operators. This listing does have a problem, however: `$_` isn't exactly the best variable name to store "weight" in. It's not very intuitive for starters, and `$_` might get altered when you least expect it.

Watch Out!

In general, storing anything in `$_` for long is playing with fire; eventually, you will get burned. Many of Perl's operators use `$_` as a default argument, and some of them modify `$_` as well. `$_` is Perl's general-purpose variable, and trying to keep a value in `$_` for very long (especially after what you learn in Hour 8, "Functions") will cause bugs eventually.

Using a variable called `$weight` would have been better in Listing 6.2. To use the match operator and substitution operator against variables other than `$_`, you must *bind* them to the variable. You do so by using the **binding operator**, `=~`, as shown here:

```
$weight="185 lbs";  
  
$weight=~s/ lbs//;      # Do substitution against $weight
```

The `=~` operator doesn't make assignments; it merely takes the operator on the right and causes it to act on the variable to the left. The entire expression has the same value as it would if `$_` were used, as you can see in this example:

```
$poem="One fish, two fish, red fish";  
  
$n=$poem=~m/fish/;      # $n is true, if $poem has fish
```

Modifiers and Multiple Matching

Until now, all the regular expressions you've seen have been **case sensitive**. That is, upperand lowercase characters are distinct in a pattern match. To match words and not care about whether they're in upperor lowercase would require something like this:

```
/[Mm][Aa][Cc][Bb][Ee][Tt][Hh]/;
```

This example doesn't just look silly; it's error-prone because it would be really easy to mistype an upper-/lowercase pair. The substitution operator (`s///`) and the match operator (`m//`) can match regular expressions regardless of case if followed with the letter `i`:

```
/macbeth/i;
```

The preceding example matches *Macbeth* in uppercase, lowercase, or mixed case (*MaCbEtH*).

Another modifier for matches and substitutions is the global-match modifier, `g`. The regular expression (or substitution) is done not just once, but repeatedly through the entire string, each match (or substitution) taking place starting immediately after the first one.

The `g` modifier (and other modifiers) can be combined by simply specifying all of them after the match or substitution operator. For example, `gi` matches all occurrences of the pattern in the string, whether uppercase or lowercase.

In a list context, the global-match modifier causes the match to return a list of all the portions of the regular expression that are in parentheses:

```
$_= "One fish, two frog, red fred, blue foul";  
@F=m/\W(f\w\w\w)/g;
```

The pattern matches a nonword character, and then the letter *f*, followed by four word characters. The *f* and the four word characters form a group, marked by parentheses. After the expression is evaluated, the array variable `@F` will contain four elements: *fish*, *frog*, *fred*, and *foul*.

In a scalar context, the `g` modifier causes the match to iterate through the string, returning true for each match and false when no more matches are made. Now consider the following:

```

$letters=0;

$phrase="What's my line?";

while($phrase=~/\w/g) {

    $letters++;

}

```

The preceding snippet uses the match operator (`//`) with a `g` modifier in a scalar context (which is provided by the condition of `while`). The pattern matches a word character. The `while` loop continues (and `$letters` gets incremented) until the match returns false. When the snippet is all done, `$letters` will be 11.

By the Way

You'll find much more efficient ways of counting characters than this presented in [Hour 9, "More Functions and Operators."](#)

Backreferences

When you use parentheses in regular expressions, Perl remembers the portion of the target string matched by each parenthesized expression. These matched portions are saved in special variables named `$1` (for the first set of parentheses), `$2` (for the second), `$3`, `$4`, and so on, as follows:

`/(\d{3}) - (\d{3}) - (\d{4}) /`

 ↓ ↓ ↓
 \$1 \$2 \$3

The pattern shown matches well-formed U.S./Canadian telephone numbers—for example, 800-555-1212—and remembers each portion in `$1`, `$2`, and `$3`. The values are assigned for each set of parentheses found, from left to right. If there are nested and overlapping parentheses, the captures are numbered from left to right for each opening parenthesis. These variables can be used after the following expression:

```

if (/(\d{3})-(\d{3})-(\d{4})/) {
    print "The area code is $1";
}

```

Or they can be used as part of the replacement text in a substitution, as follows:

```
s/(\d{3})-(\d{3})-(\d{4})/Area code $1 Phone $2-$3/;
```

Be careful, however; the variables `$1`, `$2`, and `$3` are reset every time a pattern match is successfully performed (regardless of whether it uses parentheses), and the variables are *set if and only if the pattern match succeeds completely*. Based on this information, consider the following example:

```
m/(\d{3})-(\d{3})-(\d{4})/;  
print "The area code is $1"; # Bad idea. Assumes the match succeeded.
```

In this snippet, `$1` was used without making sure the pattern match worked. This will probably cause trouble if the match ever fails.

A New Function: `grep`

A common operation in Perl is to search arrays for patterns—for example, if you've read a file into an array and need to know which lines contain a particular word. Perl has one function in particular that you can use in this situation; it's called `grep`. The syntax for `grep` is as follows:

```
grep expression, list
```

```
grep block list
```

The `grep` function iterates through each element in `list` and then executes the `expression` or `block`. Within the `expression` or `block`, `$_` is set to each element of the list being evaluated. If the expression returns true, the element is returned by `grep`. Consider this example:

```
@dogs=qw(greyhound bloodhound terrier mutt chihuahua);  
@hounds=grep /hound/, @dogs;
```

In the preceding example, each element of `@dogs` is assigned, in turn, to `$_`. The expression `/hound/` is then tested against `$_`. Each of the elements that returns true—that is, each name that contains *hound*—goes into a list that is returned by `grep` and stored in `@hounds`.

You need to remember two points here. First is that `$_` within the expression refers to the actual value in the list, not a copy of it. Modifying `$_` changes the original element in the list:

```
@hounds=grep s/hound/hounds/, @dogs;
```

After running this example, `@hounds` contains *greyhounds* and *bloodhounds*, with an *s* on the end. The original array `@dogs` is also modified—by way of changing `$_`—and it now contains *greyhounds*, *bloodhounds*, *terrier*, *mutt*, and *chihuahua*.

The other point to remember—which Perl programmers forget sometimes—is that `grep` isn't *necessarily* used with a pattern match or substitution operator; it can be used with any operator or function. The following example collects just the names of dogs longer than eight characters:

```
@longdogs=grep length($_)>8, @dogs;
```

By the Way

The `grep` function gets its name from a Unix command by the same name that is used for searching for patterns in files. The Unix `grep` command is so useful in Unix (and hence, Perl) that in the culture it has become a verb: "to grep." "To grep through a book" means to flip through the pages looking for a pattern.

A related function, `map`, has an identical syntax to `grep`, except that the return value from the expression (or block) is returned from `map`—not the value of `$_`. You use the `map` function to produce a second array based on the first. The following is an example:

```
@words= map { split ' ', $_ } @input;
```

In this example, each element of the array `@input` (passed to the block as `$_`) is split apart on spaces, producing a list of words; this list is added to the list that the `map` function returns. After every consecutive

line of @input has been split apart, the accumulated words are stored in @words.

 PREV

< Day Day Up >

NEXT 

Summary

In this hour, you learned what regular expressions are, how they're constructed, and how they're used in Perl. Regular expressions are made up of normal characters and metacharacters. The normal characters all stand for themselves (usually), and the metacharacters alter the meanings of the normal characters (or each other). These regular expressions can be used to test for the presence of patterns or to substitute one pattern for another.

Q&A

Q1: *The pattern `/\W(\w)+\W/` doesn't seem to match all the words on the line, just the ones in the middle. Why?*

A1: You're looking for word characters surrounded by nonword characters. The first word of the line—assuming it starts at the beginning of the line—doesn't have a nonword character in front of it. It doesn't have a character in front of it at all.

Q2: *What's the difference between `m//` and `//`? I don't get it.*

A2: There's almost no difference at all. The only difference is that if you decide to specify a pattern delimiter other than `/`, you can do so only if you precede the pattern with an `m`—for example, `m!pattern!`.

Q3: *I'm trying to verify that the user typed a number, but `/\d*/` doesn't seem to work. It always returns true!*

A3: It returns true because a pattern using only the `*` quantifier always succeeds. It might match zero occurrences of `\d`, or it might match 2 or 100 or 1,000. Using `/\d+/` ensures that you have at least one digit.

Workshop

If you've started to figure out the pattern to regular expressions, try this quiz to see what you've learned.

Quiz

- 1:** If you have lines formatted "`x=y`", what expression would swap the left and right sides of the expression?

a.

`s/(.+)=(.+) /$2=$1/ ;`

b.

`s/(*)=(*) /$2=$1/ ;`

c.

`s/(.*)=(.*) /$2$1/ ;`

- 2:** After this code, what's the value in `$2`?

```
$foo="Star Wars: The Phantom Menace";  
$foo=~/star\s((Wars): The Phantom Menace)/;
```

a.

`$2` is not set after the pattern match because the match fails.

b.

`Wars`

c.

`Wars: The Phantom Menace`

3: What does the pattern `m/^[-+]?[0-9]+(\.[0-9]*)?\$/` match?

a.

Dates in the format 04-03-1969

b.

Well-formed numbers such as 45, 15.3, -0.61

c.

Addition-looking patterns: 4+12 or 89+2

Answers

A1: a. Choice c doesn't include the `=` symbol in the replacement string and it wasn't captured in `$1` or `$2` because the `=` occurred outside of the parentheses. Choice b is invalid; a character must appear in front of the `*`'s. Choice a does the job nicely.

A2: a. The match fails because `star` is not capitalized, and the match doesn't have the non-case-sensitive modifier `i`. For this reason, you should always test whether the match succeeds before using the values `$1`, `$2`, and so on. (If the pattern match had used the `i` modifier or `star` had been capitalized, choice b would have been the correct response.)

A3: b. The pattern reads, at the beginning of the line match, an optional `+` or `-`, followed by one or more digits, followed (optionally) by a decimal and possibly more digits at the end of the line. The pattern matches simple, well-formed numbers.

Activities

- See whether you can produce a pattern to match a standard time format. All the following should be acceptable: `12:00am`, `5:00pm`, `8:30AM`. These should probably not be accepted: `3:00`, `2:60am`, `99:00am`, `3:0pm`.
- Write a short program that does the following:
 - 1.

Opens a file

2.

Reads all the lines into an array

3.

Extracts all the words from each line

4.

Finds all words that have at least four consecutive consonants, or nonvowels (such as the words "thoughts" or "yardstick")

 PREV

< Day Day Up >

NEXT 

Hour 7. Hashes

What You'll Learn in This Hour:

- Create a hash
- Insert and remove elements from a hash
- Use hashes to manipulate arrays

The hash is the third basic data type in Perl. The first data type you learned about was the scalar, which is a simple data type designed to hold one thing (*any* one thing—of any size—but still *one* thing). Next came the array, which is a collection of scalars. An array can hold as many scalars as you like, but searching for the scalar value you need in an array usually involves sequential access of the array until you find the scalar you need.

Hashes are another kind of collective data type. Like arrays, hashes contain a number of scalars. The difference between arrays and hashes is that hashes access their scalar data by name, not by using a numeric subscript as arrays do. Hash elements have two parts: a *key* and a *value*. The key identifies each element of the hash, and the value is the data associated with that key. This relationship is called a *key-value pair*.

Many applications lend themselves to this type of data structure naturally. For example, if you wanted to store information on licensed drivers in a state, you might use the drivers' license numbers as logical keys to store the license information; these numbers are unique (per driver). The data associated with each number would be the driver's information (license type, address, age, and so on). Each driver's license would represent an element in the hash, with the number and information as the key-value pair. Other data structures that have a hash nature are inventory part numbers, hospital patient records, telephone billing records, disk file systems, music CD-ROM collections, Rolodex information, Library of Congress numbers and ISBNs (International Standard Book Numbers), and countless others.

A hash in Perl can contain as many elements as you like, or at least as many as the available memory on your system will allow. Hashes are resized as elements are added and deleted from the hash. Access to individual elements in the hash is extremely fast and does not degrade significantly as the hash gets larger. As a result, Perl is comfortable (and quick) with your hash whether it has 10 elements or 100,000 elements. The keys to the hash can be as long as you want (they're just scalars), and the data portions of the hash can also be as large as you want.

Historically, hashes were called *associative arrays* in Perl and other languages. This term is a long-winded way of indicating that keys were associated with a value. Because Perl programmers aren't much for long-windedness, associative arrays are now simply called hashes.

Hash variables are indicated in Perl by the percent sign (%). They do not share names with arrays and scalars. You can have, for example, a hash named %a, an array named @a, and a scalar named \$a. Each of these names refers to one variable unrelated to the others.

Filling Your Hash

Individual hash elements are created by assigning values to them, much as with array elements. For example, you can create individual hash elements, as in the following:

```
$Authors{ 'Dune' }='Frank Herbert';
```

In this example, you assign to an element in the hash `%Authors`. The key for this element is the word `Dune`, and the data is the name `Frank Herbert`. This assignment creates a relationship in the hash between `Dune` and `Frank Herbert`. The value associated with the key, `$Authors{ 'Dune' }`, can be treated like any other scalar; it can be passed to functions, modified by operators, printed, or reassigned. When you're changing a hash element, always remember that you're modifying the value stored in the hash element, not the hash itself.

Why does the example use `$Authors{ }` instead of `%Authors{ }`? Like arrays, when hashes are represented as a whole, they have their own marker in front of the variable name (%). When you access an individual element of a hash, you are accessing a scalar value, so you precede the variable name with a dollar sign (\$) indicating a single value is being referenced, and you use the braces around the key to indicate that you mean the value in the hash associated with that key, not an element in an array or an unrelated scalar with the same name as the hash. To Perl, `$Authors{ 'Dune' }` represents a single scalar value—in this case, `Frank Herbert`.

A hash with just one key isn't particularly useful. To put several values into a hash, you could use a series of assignments, as shown in the following:

```
$food{ 'apple' } = 'fruit';
$food{ 'pear' } = 'fruit';
$food{ 'carrot' } = 'vegetable';
```

To make this operation shorter, you can initialize the hash with a list. The list should consist of pairings of keys and values, as shown here:

```
%food = ( 'apple', 'fruit', 'pear', 'fruit', 'carrot', 'vegetable' );
```

This example looks similar to array initializations discussed in Hour 4, "Stacking Building Blocks: Lists and Arrays." In fact, as you'll learn later in this hour, hashes can be treated as a special kind of array in many contexts.

When you're initializing a hash, keeping track of which items are keys and which items are values in a large list can be confusing. Perl has a special operator called a comma-arrow operator, =>. Using the => operator and taking advantage of the fact that Perl ignores whitespace, you can write hash initializations like the

following:

```
%food = ( 'apple' => 'fruit',
          'pear'    => 'fruit',
          'carrot'  => 'vegetable',
        );
```

Perl programmers, holding laziness as a virtue, have two additional shortcuts for hash initializations. The left side of the `=>` operator is expected to be a simple string and does not need to be quoted. Also, a single-word hash key inside the curly braces is automatically quoted. So the initializations shown previously become the following:

```
$Books{Dune} = 'Frank Herbert';
%food =( apple => 'fruit',  pear => 'fruit',  carrot => 'vegetable' );
```

By the Way

The comma-arrow operator is called that because it acts like a comma (when it is separating list items) and it looks like an arrow.

 PREV

< Day Day Up >

NEXT 

Getting Data Out of a Hash

As you have seen, to retrieve a single value from a hash, simply use a \$, the name of the hash, and (in curly braces) the key whose value you want to retrieve. Now consider this example:

```
%Movies = ( 'The Shining' => 'Kubrick', 'Ten Commandments' => 'DeMille', Goonies =>
    'Spielberg');

print $Movies{'The Shining'};
```

These lines print the value associated with the key `The Shining` in the hash `%Movies`. This example would print `Kubrick`.

Sometimes examining all the elements of a hash is useful. If all the keys of the hash are known, you can access them individually by key as shown previously. Most of the time, though, accessing each key by name isn't convenient. Some of the key names might not be known, or the keys may be too numerous to enumerate individually.

You can use the `keys` function to retrieve all the keys of a hash returned as a list. That list can then be examined to find all the elements of the hash. The keys of the hash aren't stored in any particular order internally, and the `keys` function doesn't return them in any particular order. To print all the movies in this hash, for example, you could use the following:

```
foreach $film (keys %Movies) {
    print "$film\n";
}
```

Here, `$film` takes on the value of each element of the list returned by `keys %Movies`. If you want to print all the directors' names in addition to the movie titles, you can enter the following:

```
foreach $film (keys %Movies) {
    print "$film was directed by $Movies{$film}.\n";
}
```

This snippet might print the following output:

Ten Commandments was directed by DeMille.

The Shining was directed by Kubrick.

Goonies was directed by Spielberg.

Because \$film contains the value of a hash key, \$Movies{\$film} retrieves the value of the element of the hash represented by that key. You can print both of them to see the key-value relationship in the hash. (Remember that your output might appear in a different order, because the keys are returned in no particular order by the `keys` function.)

Perl also provides the `values` function to retrieve all the values stored in a hash. Retrieving the values alone usually isn't useful, because you can't tell which key is associated with which value. The values of a hash are returned in the same order as the `keys` function would return the keys. Now consider the following example:

```
@Directors=values %Movies;  
  
@Films=keys %Movies;
```

Here, each subscript of `@Directors` and `@Films` contains a reference to the same key-value pair from `%Movies`. The name of the director contained in `$Directors[0]` corresponds to the name of the movie stored in `$Films[0]`, and so on.

Sometimes you need to retrieve individual elements from the hash by value instead of by key. The best method of retrieving by value is to *invert the hash*. This means that you make a new hash in which all the keys of the original hash become values, and all their values in the original hash become keys in the new hash. The following is an example:

```
%Movies = ( 'The Shining' => 'Kubrick', 'Ten Commandments' => 'DeMille', Goonies =>  
    'Spielberg');  
  
%ByDirector = reverse %Movies;
```

What's this? When you use the `reverse` function on a hash, Perl unwinds the hash into a flat list—perhaps something like this:

```
('The Shining', 'Kubrick', 'Ten Commandments', 'DeMille', 'Goonies', 'Spielberg')
```

Perl then reverses the order of the elements in the list, and you get the following output:

```
('Spielberg', 'Goonies', 'DeMille', 'Ten Commandments', 'Kubrick', 'The Shining')
```

Notice that now all the key-value pairs are switched around (values are now first). When you assign this list to `%ByDirector`, the resulting hash is identical to the original, except that all the keys are now values and all the values are now keys. Beware, however, if your hash has duplicate values for some reason. If the values (which are to become keys) are not unique, you end up with a hash with fewer elements than you started with. As the duplicate values collide in the new hash, old keys are replaced with the new ones.

 PREV

< Day Day Up >

NEXT 

Lists and Hashes

When I discussed how to initialize a hash, I hinted that hashes and arrays are somehow related. Whenever a hash is used in a list context, Perl unwinds the hash back into a flat list of keys and values. This list can be assigned to arrays, like any other list:

```
%Movies = ( 'The Shining' => 'Kubrick', 'Ten Commandments' => 'DeMille', Goonies =>
  'Spielberg');

@Data = %Movies;
```

At this point, `@Data` is an array containing six elements. (The even elements—zero included—are directors' names, and the odd elements are movie titles.) You can perform any normal array operation on `@Data`, and then you can reassign the array to `%Movies`, as shown here:

```
%Movies = @Data;
```

By the Way

Perl stores hash keys in a seemingly random order, useful only to Perl. Perl makes no effort to remember the order in which the keys were placed into the hash, and it doesn't put them into any particular sequence when retrieving the keys. Getting them in order for display requires that you sort them (see the "[Useful Things to Do with a Hash](#)" section, later in this hour) or somehow remember the order they were inserted in (see the "[Q&A](#)" section at the end of this hour).

Arrays and hashes are similar in other respects. To copy a hash, you can simply assign the hash to another hash, like this:

```
%New_Hash = %Old_Hash;
```

When you put `%Old_Hash` on the right side of a hash initialization, where Perl would normally expect a list or array, Perl unwinds the hash into a list. This list is then used to initialize `%New_Hash`. In the same way, you can combine and manipulate hashes in ways similar to lists, as you can see here:

```
%Both = (%First, %Second);

%Additional = (%Both, key1 => 'value1', key2 => 'value2');
```

The first line combines two hashes, `%First` and `%Second`, into a third hash, `%Both`. Something to remember about this example is that if `%First` has keys that also appear in `%Second`, the second occurrence of the key-value pair replaces the first in `%Both`. In the second example, `%Both` is represented as a list of key-value pairs in the parentheses. Two additional key-value pairs are also in the parentheses. The entire list is then used to initialize `%Additional`.

 PREV

< Day Day Up >

NEXT 

Hash Odds and Ends

Some operations on hashes aren't obvious if you're new to Perl. Because of the special nature of hashes, a couple of common operations require functions that aren't necessary for scalars and arrays.

Testing for Keys in a Hash

To test to see whether a key exists in a hash, for example, you might be tempted to try the following syntax:

```
if ( $Hash{keyval} ) {           # WRONG, in this case  
:  
}
```

This example doesn't work, for a few reasons. First, this snippet doesn't test to see whether `keyval` is a key in a hash; it actually tests the value associated with the key `keyval` in the hash.

Does it work to test whether the key is defined, as in the following?

```
if ( defined $Hash{keyval} ) {     # WRONG, again in this case  
:  
}
```

Again, this example doesn't quite work. This snippet still tests the data associated with the key `keyval`, and not whether the key exists. `undef` is a perfectly valid value to associate with a hash key, as in the following:

```
$Hash{keyval} = undef;
```

The preceding test of `defined` returns false because it doesn't test for the existence of a key in a hash; it tests the data associated with the key. So what's the right way? Perl has a special function just for this purpose; it's called `exists`. The `exists` function, shown here, tests for the presence of the hash key in the hash and returns true if it's there or false otherwise:

```
if ( exists $Hash{keyval} ) {      # RIGHT!
:
}
```

Removing Keys from a Hash

The other operation that isn't obvious is removing a key from a hash. As you saw earlier, simply setting the hash element to `undef` doesn't work. To remove a single hash key, you can use the `delete` function, as follows:

```
delete $Hash{keyval};
```

To remove all the keys and values from a hash, simply reinitialize the hash to an empty list like this:

```
%Hash = ();
```

 PREV

< Day Day Up >

NEXT 

Useful Things to Do with a Hash

Hashes are often used in Perl for more reasons than to store records by keys for later retrieval. The advantages to using a hash are fast individual access to keys and the fact that all keys in a hash are unique. These properties lend themselves to some useful data manipulations. Not surprisingly, because arrays and hashes are so similar, many of the interesting things you can do with hashes are array manipulations.

Determining Frequency Distributions

In [Hour 6](#), "Pattern Matching," you learned how to take a line of text and split it into words. Examine the following snippet of code:

```
while ( <> ) {  
  
    while ( /(\w[\w-]*)/g ) {      # Iterate over words, setting $1 to each.  
  
        $Words{$1}++;  
  
    }  
  
}
```

The first line reads the standard input one line at a time, setting `$_` to each line.

The next `while()` loop then iterates over each word in `$_`. Recall from [Hour 6](#) that using the pattern-matching operator `(//)` in a scalar context with the `g` modifier returns each pattern match until no more are left. The pattern being looked for is a word character `\w`, followed by zero or more word characters or dashes `[\w-]*`. In this case, you use parentheses to remember the string matched in the special variable `$1`.

The next line, although short, is where the snippet gets interesting. `$1` is set, in turn, to each word matched by the pattern on the second line. That word is used as the key to the hash `%Words`. The first time the word is seen, the key does not already exist in the hash, so Perl returns a value of `undef` for that key-value pair. By incrementing it, Perl sets the value to 1, creating the pair. The second time a word is seen, the key (that word) already exists in the hash `%Words`, and it is incremented from 1 to 2. This process continues until no input is left.

When you're finished, the hash `%Words` contains a frequency distribution of the words read in. To look at the frequency distribution, you can use the following code:

```
foreach ( keys %Words ) {
```

```
print "$_ $Words{$_}\n";
}
```

Finding Unique Elements in Arrays

The technique shown in the preceding code is also useful for finding which elements in an array occur only once. Suppose you have already extracted all the words from the input into an array instead of a hash, and you have made no particular effort to make sure that a word wasn't already in the list before putting it in again. In this case, you would have a list with a lot of duplicated words.

If your input text were the opening lines of *One Fish, Two Fish*, the list would look something like the following:

```
@fishwords=('one', 'fish', 'two', 'fish', 'red', 'fish', 'blue', 'fish');
```

If you are given this list of words (in `@fishwords`), and you need only the unique elements of the list, a hash works nicely for this purpose, as shown in Listing 7.1.

Listing 7.1. Finding Unique Elements in an Array

```
1:  %seen = ();
2:  foreach (@fishwords) {
3:    $seen{$_} = 1;
4:  }
5:  @uniquewords = keys %seen;
```

Line 1: This line initializes a temporary

Exercise: Creating a Simple Customer Database with Perl

When you call a customer service center and finally make it through the Touch-Tone menus, the first thing the human being on the other end of the phone asks you for is your telephone number. Well, almost every time. Sometimes the customer service representative wants your customer number or even your Social Security number. What this person is after is something that uniquely identifies you to the computer he or she is using. These numbers serve as keys for retrieving information about you in a database. Sounds like Perl's hashes, doesn't it?

For this exercise, you're going to search a customer database. This program assumes that the database already exists, and it doesn't provide any way to update that database—yet. Here, you're going to allow the user to search on one of two different fields.

To begin this exercise, you need some data. Fire up your text editor, key in the text in [Listing 7.4](#) (or something similar), and save it as `customers.txt`. Don't worry about the number of spaces between the columns or aligning them, as long as you leave at least one space between each column.

Listing 7.4. Sample Data for the Customer Program

Smith, John (248)-555-9430 jsmith@aol.com

Hunter, Apryl (810)-555-3029 april@showers.org

Stewart, Pat (405)-555-8710 pats@starfleet.co.uk

Ching, Iris (305)-555-0919 iching@zen.org

Doe, John (212)-555-0912 jdoe@morgue.com

Jones, Tom (312)-555-3321 tj2342@aol.com

Smith, John (607)-555-0023 smith@pocahontas.com

Crosby, Dave (405)-555-1516 cros@csny.org

Johns, Pam (313)-555-6790 pj@sleepy.com

Jeter, Linda (810)-555-8761 netless@earthlink.net

Garland, Judy (305)-555-1231 ozgal@rainbow.com

In the same directory, key in the short program in [Listing 7.5](#) and save it as `Customer`. As usual, don't type the line numbers, and, if you can, be sure to make the program executable according to the instructions you

learned in [Hour 1](#), "Introduction to the Perl Language."

When you're done, try running the program by typing the following at a command line:

Customer

or, if you cannot make the program executable on your system,

perl -w Customer

Listing 7.5. Complete Listing of the Customer Program

```
1:  #!/usr/bin/perl -w
2:
3:  open(PH, "customers.txt") or die "Cannot open customers.txt: $!\n";
4:  while(<PH>) {
5:      chomp;
6:      ($number, $email) = ( split(/\s+/, $_) )[1,2];
7:      $Phone{$number} = $_;
8:      $Email{$email} = $_;
9:  }
10: close(PH);
11:
12: print "Type 'q' to exit\n";
13: while (1) {
14:     print "\nNumber? ";
15:     $number = <STDIN>;    chomp($number);
```

```
16:     $address = "";  
17:     if (! $number ) {  
18:         print "E-Mail? ";  
19:         $address = <STDIN>; chomp($address);  
20:     }  
21:  
22:     next if (! $number and ! $address);  
23:     last if ($number eq 'q' or $address eq 'q');  
24:  
25:     if ( $number and exists $Phone{$number} ) {  
26:         print "Customer: $Phone{$number}\n";  
27:         next;  
28:     }  
29:  
30:     if ($address and exists $Email{$address} ) {  
31:         print "Customer: $Email{$address}\n";  
32:         next;  
33:     }  
34:     print "Customer record not found.\n";  
35:     next;  
36: }  
36: print "\nAll done.\n";
```

[Listing 7.6](#) shows the output of the Customer program.

Listing 7.6. Sample Output from Customer

Type 'q' to exit

Number? <return>

E-Mail? cros@csny.org

Customer: Crosby, Dave (405)-555-1516 cros@csny.org

Number? (305)-555-0919

Customer: Ching,Iris (305)-555-0919 iching@zen.org

Number? q

All done.

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The `customers.txt` file is opened for reading on the filehandle `PH`. Of course, errors are checked for and reported.

Lines 4–5: The `PH` filehandle is read, each line being assigned to `$_. $_` is chomped to remove the trailing newline character.

Line 6: The line (in `$_`) is split on whitespace (`\s+`). Surrounding the `split` statement is a set of parentheses, and brackets follow them. Because you're interested in only the phone number and email address from each line, you take a slice of the return values from the `split`. The two values are assigned to `$number` and `$email`.

Lines 7–8: `%Email` is used to store the customer record, the key being the email address. `%Phone` is used to store the customer record also, but keyed by phone number.

Line 10: This line closes the filehandle.

Line 13: This `while` loop encloses the portion of code that needs to be repeated. The statement `while(1)` is a Perl idiom that means "loop forever." To exit this loop, a `last` statement will eventually be used.

Lines 14–15: The phone number is read, and the newline character is removed.

Lines 17–20: If no phone number is available, these lines prompt for an email address.

Lines 22–23: If nothing was entered, this line repeats the loop. If either response was a `q`, then the loop is exited.

Lines 25–28: If a number was entered, and it is valid, line 26 prints the customer record. Control is passed back to the top of the block with a `next` statement.

Lines 30–33: If an address was entered, and it's valid, the customer record is printed. Control is passed back to the top of the block with `next`.

Lines 34–35: Either an address or a phone number was entered, and it was found not to be valid. These lines print an appropriate message and repeat the block with `next`.

This example demonstrates a few Perl features. Hashes are used for quick lookups of data based on a key. Because Perl implements hashes very efficiently, response time for a query should not become inefficient even if this program has thousands or tens of thousands of records in the hashes. Also, this program serves as a demonstration of program flow control using a `simpleblock` instead of other control structures (`while`, `do`, `until`, and so on).



< Day Day Up >



Summary

Hashes provide the Perl programmer with many useful tools. Beyond simple record storage and retrieval, hashes provide mechanisms for doing useful transformations and analysis on data. The formulas for array manipulation, record storage, and retrieval will benefit you greatly. In the hours to come, hashes will provide the gateway to your learning about such topics as handling DBM (Data Base Manager) files, working with complex data structures, and interacting with your system's environment.

Q&A

Q1: *Can I store more than one thing in a hash—for instance, if I need to store several things (a list) by one key?*

A1: Yes. Two basic methods are involved. The first (and most cumbersome) is to format the value portion of the hash element into something recognizable, such as a comma-separated list. Whenever you store the hash element, you assemble the list into a scalar by using `join`, and whenever you retrieve a value from a hash, you split the scalar back into a list by using `split`. This method is cumbersome and error-prone.

The other method is to use a *reference*. References allow you to create hashes of arrays, hashes of hashes, and other complex data types. Using references to create complex structures is easy when you get the hang of it. This topic will be covered in Hour 13, "References and Structures."

Q2: *How can I keep my keys in the order in which I assigned them to the hash?*

A2: Again, you can keep them in order in a couple of ways. The first method, which isn't easy, is to keep track of the insertion order yourself. One technique is to use an array that mirrors the hash. As you put new elements into the hash, you use `push` to put the same key into an array. When you need to see the insertion order, simply use the array instead of the `keys` function. This method is complicated and likely to be buggy.

The far superior method is to use the module `Tie::IxHash`. This module causes the `keys` function to return the hash keys in insertion order, just as you wanted. You can find an explanation of how to use modules in Hour 14, "Using Modules."

Q3: *Can you suggest a convenient way to write my hash into a file?*

A3: Yes. Modules such as `Data::Dumper` or `Storable` can reformat data types such as hashes and arrays into easily storable scalar values, which can be written to text files. These modules also have functions that take those formatted scalars and re-create the original structure you stored.

In Hour 15, "Finding Permanence," you'll see that an even easier way to write a hash into a file is to use DBM files. DBM files allow you to tie your hash to a disk file. When you alter the hash, the disk file changes. The disk file causes your hash to be preserved as long as the file is intact.

Workshop

Quiz

1: Why would `name` have been an inappropriate search key for the Customer program you wrote in this hour, especially if the customer list were much, much longer?

a.

Combinations of last name and first name would have been more than Perl's hashes would allow.

b.

People's names are not unique keys.

c.

No one would ever want to search a customer database by name.

2: What's the difference between an associative array and a hash?

a.

There is no difference.

b.

Associative arrays are used on more formal data sets, such as billing records.

c.

Hashes aren't really associative arrays in Perl, so they have a different name.

3: What kinds of data are best suited for hashes?

a.

Simple lists of items

b.

Potatoes and corned beef

c.

Lists of key-value pairs

Answers

A1: b. Perl's hashes are virtually unlimited in size, and users will (inevitably) ask for a search by name. But people's names are inappropriate, because they're not unique. The phone book has many, many duplicate names such as John Smith and Robert Jones.

A2: a. Hashes and associative arrays are identical. The only difference is that *hash* is much easier to say and to spell.

A3: c. Choice c is the correct answer, although a nice corned-beef hash can really round out breakfast.

Activities

- Modify the Customer program to allow a search by name. Because you cannot use `name` as a hash key, you need to search through the values in the hash.
- Modify the Customer program to allow searching by a partial key (for example, part of a phone number or part of an email address). You can use regular expressions to search for the patterns. Simply keep in mind that you could find multiple results and you should return each one.

Hour 8. Functions

What You'll Learn in This Hour:

- How to define your own functions and call them
- How to pass values into functions and return values
- How to write programs with `use strict` to enforce structure

Almost all computer languages support *functions*. A function is a grouping of code statements that can be called by name to do some work and then return some value. You've been using functions throughout this book; for example, you've used `print`, `chomp`, `sort`, `open`, `close`, `split`, and so on. These functions are built into Perl.

Perl also allows you to write your own functions. In Perl, user-defined functions are called *subroutines*, or *subs*. Like Perl's built-in functions, user-defined functions can take arguments and return values to the caller.

Perl also supports the concept of *scope*. Scope determines the set of variables that a given statement in the program can access as it is executed. Because of Perl's scope features, you can write functions that can behave autonomously from the rest of your program. Well-written functions can be reused in other programs.

Creating and Calling Subroutines

You create user-defined subroutines in Perl by using the following syntax:

```
sub subroutine_name {  
    statement1;  
    :  
    statementx;  
}
```

Subroutine names in Perl follow the same naming conventions as scalars, arrays, and hashes outlined in [Hour 2](#), "Perl's Building Blocks: Numbers and Strings." Subroutine names can have the same names as existing variables. However, you should avoid creating subroutine names that are the same as those of Perl's built-in functions and operators. Creating two subroutines with the same name in Perl causes Perl to emit a warning if warnings are enabled; otherwise, the second definition causes the first to be forgotten.

Here is an example of a subroutine:

```
sub countdown {  
    for ($i=10; $i>=0; $i--) {  
        print "$i -";  
    }  
}
```

When the program needs to use the code in the subroutine, you can use either of the following syntax lines to *call* (or *invoke*) the subroutine:

```
&countdown();
```

or

```
countdown( );
```

You can use the second syntax (without the `&`) if the subroutine has been declared in the code already, although the `&countdown()` syntax is acceptable anywhere, and some consider it more readable. In this book, I'll declare subroutines ahead of where they're first called, and I'll use the form without the `&`, although either is acceptable.

When the subroutine is invoked, Perl remembers where it was, executes the subroutine's code, and then returns to the remembered place in the program when the subroutine is complete, as in the following example:

```
print "T-minus: ";
countdown();
print "Blastoff!\n";
```

prints

```
T-minus: 10 -9 -8 -7 -6 -5 -4 -3 -2 -1 -0 -Blastoff!
```

Perl subroutines can be called anywhere within your program, including within other subroutines, as shown here:

```
sub world {
    print "World!";
}

sub hello {
    print "Hello, ";
    world();
}
```

```
}
```

```
hello();
```

Returning Values from Subroutines

A subroutine isn't just for grouping code together by a convenient name. A subroutine, like Perl's functions, operators, and expressions, also has a value, called the subroutine's *return value*. The return value of a subroutine is the value of the last expression evaluated in the subroutine or a value explicitly returned by the `return` statement.

The return value of the subroutine, computed when the subroutine is called, can then be used in whatever code called the subroutine. Consider the following example:

```
sub two_by_four {      # A silly subroutine
    2 * 4;
}

print 8*two_by_four();
```

In the preceding snippet, for Perl to evaluate the expression `8*two_by_four()`, the subroutine `two_by_four()` is run and the value `8` returned. The expression `8*8` is then evaluated, and `64` is printed.

Values can also be explicitly returned from a subroutine with the `return` statement. Use a `return` statement when your program needs to return before the end of the subroutine or when you want to be explicit about what value is being returned—instead of just "falling off" the end of the subroutine and using the last expression's value. The following snippet uses both methods:

```
sub x_greaterthan100 {

    # Relies on the value of $x being set elsewhere
    return(1) if ( $x > 100 );
    0;
}
```

`$x = 70;`

```

if (x_greaterthan100()) {
    print "$x is greater than 100\n";
}

```

Subroutines can return arrays and hashes as well as scalars, as shown here:

```

sub shift_to_uppercase {
    @words = qw( cia fbi un nato unicef );
    foreach (@words) {
        $_ = uc($_);
    }
    return (@words);
}

@acronyms=shift_to_uppercase();

```

Arguments

All the preceding examples of subroutines have one thing in common: They act on either data that was hard-coded (`2*4`) or variables that just happened to have the right data in them (`$x` for `x_greaterthan100()`). This limitation creates a problem because functions that rely on hard-coded data or expect variables to be set outside the function aren't really portable. It would be nice to call a function and say, "Take *this* data and do something with it," and then call it later and say, "Take this *other* data and do something with it." The result of calling a function, then, can change depending on what is passed to it.

These values given to functions to change their results are called *arguments*, and you've been using them throughout this book. Perl's built-in functions (`grep`, `sort`, `reverse`, `print`, and so on) take arguments, and now yours can too. To pass a subroutine arguments, you can use any of the following:

```

subname(arg1, arg2, arg3);

subname arg1, arg2, arg3;

&subname(arg1, arg2, arg3);

```

You can use the second form—without the parentheses—only if Perl has encountered the subroutine definition already.

In a subroutine, any arguments that were passed to it are accessible through the Perl special variable `@_`. This snippet demonstrates passing arguments—the three string literals—and printing them with a function:

```
sub printargs {  
    print join( ', ', @_ );  
}  
  
printargs( 'market', 'home', 'roast beef' );
```

To access individual arguments that were passed in, as shown in the following example, you use an index on the array `@_` as you would any other array. Just remember that `$_[0]`—an element of `@_`—has nothing to do with the scalar variable `$_`.

```
sub print_third_argument {  
    print $_[2];  
}
```

Working with variable names like `$_[3]` isn't exactly a "clear" style of programming. Functions that take multiple arguments often begin by giving names to those arguments; this practice makes it clearer what they do. To see what I mean, check out the following example:

```
sub display_box_score {  
    ($num_hits, $num_at_bats) = @_;  
  
    print "For $num_at_bats trips to the plate, "  
  
    print "he's hitting ", $num_hits / $num_at_bats, "\n";  
}
```

```
display_box_score(50, 210);
```

In the preceding subroutine, the array `@_` is copied to the list (`$num_hits, $num_at_bats`). The first element of `@_`—`$_[0]`—becomes `$num_hits`, and the second `$num_at_bats`. The variable names are simply used here for readability.

By the Way

The variable `@_` actually contains aliases to the original arguments that were passed to the subroutine. Modifying `@_` (or any element of `@_`) modifies the original variables in the argument list. Doing so unexpectedly is considered bad form; your function shouldn't interfere with arguments from callers to the function unless the users of the function are expecting this.

Passing Arrays and Hashes

Arguments to subroutines don't have to be scalars. You can pass arrays and hashes to subroutines, but doing so requires some thought. You pass an array—or hash—to a subroutine the same way you pass a scalar:

```
@sorted_items=sort_numerically(@items);
```

In the subroutine, the entire array `@items` is referenced through `@_`:

```
sub sort_numerically {  
    print "Sorting...";  
  
    return( sort { $a <=> $b } @_ );  
}
```

You do hit a small snag when passing hashes and arrays to subroutines: Passing two or more hashes (or arrays) to a subroutine doesn't usually do what you want. Examine this snippet of code:

```
sub display_arrays {  
    (@a, @b) = @_;
```

```

print "The first array: @a\n";
print "The second array: @b\n";
}

display_arrays(@first, @second);

```

The two arrays `@first` and `@second` are put into a list together, and the elements are put into `@_` during the subroutine call. The end of the `@first`'s elements are indistinguishable from the beginning of `@second`'s elements in `@_`; it's just one large flat list. Inside the subroutine, the assignment `(@a, @b) = @_` takes all the elements in `@_` and assigns them to `@a`. The array `@b` gets no elements. (The reason for this was explained in [Hour 4](#), "Stacking Building Blocks: Lists and Arrays.")

One or more scalars can be passed along with a *single* array or hash, as long as the scalars are passed first in the argument list and you know how many you have. That way, the hash or array encompasses all the values beyond the last scalar, as in this example:

```

sub lots_of_args {
    ($first, $second, $third, %hash) = @_;
    # rest of subroutine...
}

lots_of_args($foo, $bar, $baz, %myhash);

```

If you must pass multiple arrays and hashes into a subroutine (and be able to distinguish them later), you must use *references*. You'll learn how to pass references in [Hour 13](#), "References and Structures."

Scope

In the introduction for this hour, you learned that subroutines are used to take pieces of code and bundle them up and give them a name. You then can use this name to execute the code whenever you need it. You can also write the code within the subroutine so that it functions autonomously. That is, you can make it run and produce its return value using nothing but its arguments, the language's built-in functions and operators, and literal expressions within it. In that case, you can reuse the subroutine in other programs, because it no longer relies on the context it's being called in; it simply takes its arguments and data defined internally and produces a return value. The subroutine becomes a black box—stuff goes in, stuff comes out—and outside it you don't care what happens on the inside. Such a subroutine is called a *pure function*.

Now examine the following two snippets:

```
# One fairly good way to write this function
```

```
sub moonweight {
```

```
    ($weight) = @_;
```

```
    return($weight / 6);
```

```
}
```

```
print moonweight(150);
```

```
# A poor way to write this function.
```

```
sub moonweight {
```

```
    return( $weight/6 );
```

```
}
```

```
$weight = 150;
```

```
print moonweight;
```

In the long run, the first implementation shown here is a better one. It doesn't rely on any external variables—those outside the function—being set. It takes its argument, which it copies to `$weight`, and then does its calculations. The second implementation cannot easily be reused in another program; you would have to ensure that `$weight` is set properly and is not being used for some other value. If it were being used for

something else, you would have to edit the `moonweight()` subroutine to use a different variable. That's not very efficient.

So, the first example is a better subroutine, but it's still missing something. The variable `$weight` could conflict with a variable named `$weight` somewhere else in the program.

Perl permits you to reuse variable names over and over and for different purposes within a large program. Perl's variables, by default, are visible in the main body of your program and in the subroutines; these kinds of variables are called *global variables*.

What you need to do in the `moonweight` example is to make the variable `$weight` a *private variable* in the subroutine. To do so, you use the `my` operator:

```
sub moonweight {  
    my $weight;  
    ($weight)=@_;  
    return($weight/6);  
}
```

Inside `moonweight()`, `$weight` is now private. Neither the main body nor any other subroutines in the program can access the value of `$weight`. Any other variables with the name `$weight` are completely separate from the `moonweight()` subroutine's `$weight`, so they are not affected by the assignment to `$weight` inside `moonweight()`. This subroutine is now completely self-contained.

The portion of the program in which the variable is visible is called the variable's *scope*.

You can use the `my` operator to declare scalar, array, and hash variables private to a subroutine. Filehandles, subroutines, and Perl's special variables—`$!`, `$_`, and `@_`, for example—cannot be marked as private to a subroutine. You can declare multiple variables private if you use parentheses with `my`:

```
my($larry, @curly, %moe);
```

Variables that are private to subroutines are stored in a completely different manner than global variables are. Global variables and private variables can have the same names, but they have nothing to do with each other, as shown here:

```
sub myfunc {  
    my $x;  
  
    $x=20;      # This is a private $x  
  
    print "$x\n";  
  
}  
  
$x=10;          # This is a global $x  
  
print "$x\n";  
  
myfunc();  
  
print "$x\n";
```

The preceding snippet prints 10, 20, and then 10. The \$x inside the `myfunc()` subroutine is a completely different \$x from the one outside. (Is it possible for a subroutine to use *both* its private \$x and the global \$x? Yes, but the answer is somewhat complicated and beyond this introductory book on Perl.)

Most of the time, Perl subroutines start by assigning `@_` to a list of variable names and then declaring the list private to the subroutine:

```
sub player_stats {  
    my($num_at_bats, $num_hits, $num_walks)=@_;  
  
    # Rest of function...  
  
}
```

This technique makes for a programmer-friendly subroutine: The variables are all private to the subroutine, so they cannot affect or be affected by other subroutines or the main body of the program. When a subroutine ends, any private variables are destroyed.

Other Places for `my`

You can also declare variables to have a scope that's even smaller than a subroutine. The `my` operator actually declares variables to be private to the enclosing block, which might or might not be a subroutine block. For example, in this snippet, the private variable `$Y` (the one declared with `my`) is visible only within the block:

```
$y=20;  
{  
    my $y=500;  
  
    print "The value of \$y is $y\n";    # Will print 500  
}  
  
print "$y\n";                      # Will print 20.
```

The declaration can even take place within one of the control structures, such as `for`, `foreach`, `while`, or `if`. Essentially, anywhere you have a block, a variable can be scoped so that it's visible only within that block, as in this example:

```
while($testval) {  
  
    my $stuff;      # Visible only within the while() loop.  
  
    :  
  
}  
  
foreach(@t) {  
  
    my %hash;      # Visible only within the foreach loop.  
  
}
```

In the preceding snippet, the `my` variables—`$stuff` and `%hash`—are created anew each time through the loops.

Perl versions 5.004 and newer allow the iterator in `for` and `foreach` loops and the test conditions in `while` and `if` structures to be declared as private to the block:

```
foreach my $element (@array) {  
  
    # $element is only visible in the foreach()
```

}

```
while(my $line=<STDIN>) {  
    # $line is visible only in the while()  
}
```

Again, when an enclosing block ends, any variables private to that block—and their values—are destroyed.

 PREV

< Day Day Up >

NEXT 

Exercise: Statistics

Now that you've learned about subroutines, you should begin to see the benefits of encapsulating code in self-contained subroutines. They provide code that can be easily reused. In this exercise, three subs provide some analysis on groups of numbers.

Just to refresh your memory from school, the *mean*, also called an *arithmetic mean* or *average*, of a set of numbers is simply the sum of all the numbers in the set divided by how many numbers there are in the set. The *median* is the number that would be in the middle if you were to sort the set numerically; with an even number of elements, the median is the average of the two numbers that would be in the middle. The *standard deviation* gives an idea of how "bunched" the numbers are around the mean. A high standard deviation that means the numbers are widely distributed; a small one means they're bunched tightly around the average. In many sets of numbers commonly found in nature, the mean, plus or minus the standard deviation, represents about 68 percent of the set of the numbers; plus or minus two standard deviations, 95 percent of the set of numbers. Here is a program that finds the mean, median, and standard deviation for a set of numbers that the user types at the keyboard.

Using your text editor, type the program from [Listing 8.1](#) and save it as `Stats`. Again, do not type line numbers, and if possible, be sure to make the program executable according to the instructions you learned in [Hour 1](#), "Introduction to the Perl Language."

When you're all done, try running the program by typing the following at a command line:

`stats`

or, if your system cannot make the program executable,

`perl stats`

Listing 8.1. Complete Listing for Stats Program

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  sub mean {
```

```
5:     my(@data) = @_;
6:     my $sum;
7:     foreach(@data) {
8:         $sum += $_;
9:     }
10:    return($sum / @data);
11: }
12: sub median {
13:     my(@data)=sort { $a <=> $b } @_;
14:     if (scalar(@data) % 2) {
15:         return($data[@data / 2]);
16:     } else {
17:         my($upper, $lower);
18:         $lower=$data[@data / 2];
19:         $upper=$data[@data / 2 - 1];
20:         return(mean($lower, $upper));
21:     }
22: }
23: sub std_dev {
24:     my(@data)=@_;
25:     my($sq_dev_sum, $avg)=(0,0);
26:
27:     $avg = mean(@data);
28:     foreach my $elem (@data) {
29:         $sq_dev_sum += ($avg - $elem) **2;
```

```

30:      }
31:      return(sqrt($sq_dev_sum / ( @data - 1 )));
32:  }
33:  my($data, @dataset);
34:  print "Please enter data, separated by commas: ";
35:  $data = <STDIN>;  chomp $data;
36:  @dataset = split(/[\s,]+/, $data);
37:
38:  print "Median: ", median(@dataset), "\n";
39:  print "Mean: ", mean(@dataset), "\n";
40:  print "Standard Dev.: ", std_dev(@dataset), "\n";

```

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The `use strict` directive means that all variables must be declared with `my` and that bare words must be quoted.

Lines 4–11: The `mean()` function works by using a `foreach` loop to add up all the numbers in `$sum` and then divides by the number of numbers.

Lines 12–21: The `median()` function works in two ways. With an odd number of elements, it simply picks the middle element by taking the length of the array and dividing it by two and then using the integer portion. With an even number of elements, it does the same but instead takes the two middle numbers. Those numbers—in `$upper` and `$lower`—are then averaged with the `mean()` function and returned as the median.

Lines 23–32: The `std_dev()` function is simple, but mostly just math. In short, each element in `@data` is subtracted from the mean and squared. The result is then accumulated in `$sq_dev_sum`. To find the standard deviation, the sum of the squared differences is divided by the number of elements minus 1, and then the square root is taken.

Lines 33–35: The variables needed in the main body of the program are declared as lexicals (with `my`), and the user is prompted for `$data`. The variable `$data` is then split into the

array @dataset using the pattern `/[\s,]+/`. This pattern splits the line on commas and spaces. Extra spaces and commas are ignored.

Lines 38–40: The output is produced. Notice that this isn't the only place the functions `mean()`, `median()`, and `std_dev()` are called. They also call each other: `std_dev()` and `median()` both use `mean()`, which is a good example of code reuse!

[Listing 8.2](#) shows a sample of the statistics program's output.

Listing 8.2. Sample Output from Stats

```
Please enter data, separated by commas: 14.5,6,8,9,10,34
```

```
Median: 9.5
```

```
Mean: 13.583333333333
```

```
Standard Dev.: 10.3943093405318
```

 PREV

< Day Day Up >

NEXT 

Function Footnotes

Now that you know about scope, there are some things that can be done effectively only with scope. One is recursive subroutines, and the other is the `use strict` Perl statement that enables a stricter Perl, possibly preventing you from making mistakes.

Declaring Variables `local`

Perl version 4 didn't have private variables. Instead, Perl 4 had variables that were "almost private." This concept of almost-private variables is still around in Perl 5. You declare these variables by using the `local` operator, as in this example:

```
sub myfunc {  
    local($foo)= 56;  
    # rest of function...  
}
```

In the preceding snippet, `$foo` is declared to be local to the `myfunc()` subroutine. A variable that has been declared with `local` acts almost identically to one declared with `my`: It can be scoped to a subroutine, block, or `eval`, and its value is destroyed upon leaving the subroutine or block.

The difference is that a variable declared `local` can be seen within its scoped block and within any subroutines called from that block. [Table 8.1](#) shows a side-by-side comparison.

Table 8.1. Comparing `my` and `local`

```

sub mess_with_foo {    sub mess_with_foo {
    $foo=0;                      $foo=0;
}

sub myfunc {               sub myfunc {
    my $foo=20;                  local $foo=20;
    mess_with_foo();            mess_with_foo();
    print $foo;                 print $foo;
}

myfunc();                  myfunc();

```

The two snippets shown in Table 8.1 are identical except for the declaration of `$foo` in `myfunc()`. On the left, it's declared with `my`; on the right, with `local`.

When the code on the left runs, `$foo` is created private to `myfunc()`. When `mess_with_foo()` is called, the `$foo` that's changed inside `mess_with_foo()` is the global variable `$foo`. When control returns to `myfunc()`, the value of 20 is printed, because the `$foo` in `myfunc()` was never changed.

When the code on the right runs, `$foo` is created and declared local to `myfunc()`. When `mess_with_foo()` is called, `$foo` is set to 0. That `$foo` is the same `$foo` from `myfunc()`; the "privateness" is passed down to the called subroutine. Upon returning to `myfunc()`, the value of 0 is printed.

By the Way

If you're picky about terminology, local variables are technically known as *dynamically scoped* variables, because their scope changes depending on what subroutines are called. Variables declared with `my` are called *lexically scoped* variables, because their scope can be determined by simply reading the code and noting what block they were declared in, and the scope does not change.

Whenever your program needs a variable that's private to a subroutine, you almost always want a variable declared with `my`.

Making a Stricter Perl

Perl is a permissive language. It doesn't try to get in your way; it allows you to just get your work done without complaining too much about what your code looks like. However, you can also tell Perl to be a bit more strict about your code. For example, Perl can help you avoid silly mistakes if you use the warning switch on the command line—or on the `#!` line. Perl warns you when you're using undefined variables, using a variable name only once, and so on.

In large software projects, and as your own programs get larger and larger, it's nice to have Perl help you keep yourself in line. In addition to using the `-w` switch, you can tell the Perl interpreter at compile time to turn on more warnings. You do so by using `use strict`:

```
use strict;

sub mysub {
    my $x;
    :
}

mysub();
```

The `use strict` statement is actually something called a *compiler directive*. It tells Perl to flag the following situations as runtime errors that point onward in the current block or file:

- Attempts to use variable names (other than special variables) that are not declared with `my`
- Attempts to use a bare word as a function name when the function definition hasn't been seen yet
- Other potential errors

The `use strict` directive, for now, helps you avoid the last two problems. Having Perl flag variables not declared with `my` prevents you from using a global variable when you actually intend to use a private variable; it's a way of helping you write more self-contained code and not rely on global variables.

The last trap that `use strict` catches is that of bare keywords. Examine this code:

```
$var=value;
```

In this case, do you mean for `value` to be interpreted as a subroutine call, as a string (but you forgot the

quotes), or as a variable (but you forgot the type identifier)? Perl's `use strict` directive would note that this code is ambiguous and disallow the syntax, unless the subroutine `value` was already declared before this statement was reached.

From this point on, I will include `use strict` in all the exercises and longer program listings in this book.

Recursion

You will probably run into a special class of subroutines sooner or later. These subroutines actually call themselves in order to do their work. They are called *recursive* subroutines.

Recursive subroutines are used wherever tasks can be broken down into smaller and smaller identical tasks. One example of a recursive task is searching a directory tree for a file. After searching the topmost directory, as subdirectories are found, those directories must be searched. And within those directories, if subdirectories are found, those directories must be searched. You should begin to see a pattern here.

Another recursive task is computing factorials, which are frequently used in statistics. For example, the number of ways the six letters *ABCDEF* can be arranged is 6-factorial, written $6!$ in mathematics. The factorial of an integer is the product of that integer and all the smaller integers down to 1. So the factorial of 6 is $6 \times 5 \times 4 \times 3 \times 2 \times 1$, or 720, and the factorial of 5 is $5 \times 4 \times 3 \times 2 \times 1$, or 120. Notice that $6!$ is equal to $6 \times 5!$. So one way to compute the factorial of 6 is to compute the factorial of 5 and multiply that by 6. To compute the factorial of 5, you compute the factorial of 4 and multiply that by 5, and so on. The factorial of 1, on the other hand, is just 1 by the definition. The factorial of 0 is also defined as 1, so that $1!$ equals $1 \times 0!$ regularly. These facts lead us to a recursive subroutine for computing factorials, which is shown in Listing 8.3.

Listing 8.3. Factorials

```
1:  sub factorial {  
2:      my ($num) = @_;  
3:      return(1) if ($num <= 1);  
4:      return($num * factorial($num - 1));  
5:  }  
6:  print factorial(6);
```

Line 2: The argument to the `factorial()` subroutine is copied to `$num`, which is declared private to the subroutine.

Line 3: Every recursive subroutine needs to have a *termination condition*. That is, there needs to be some argument value at which the subroutine no longer calls itself for an answer.

Otherwise the subroutine would try to call itself infinitely and never return an answer. For the `factorial()` subroutine, the termination condition is 1-factorial (or 0-factorial); as stated previously, those two values are both 1. The `factorial()` subroutine will execute the `return(1)` when called with 1 or 0 (`$num<=1`).

Line 4: Otherwise, if the argument is not 0 or 1, the factorial of the next smaller sequence must be computed. Observe the following:

If `$num` is: Line #4 evaluates as follows:

- 6 `return(6 * factorial(5))`
- 5 `return(5 * factorial(4))`
- 4 `return(4 * factorial(3))`
- 3 `return(3 * factorial(2))`
- 2 `return(2 * factorial(1))`
- 1 Line #4 is not reached; `factorial(1)` returns 1.

After the next-smaller factorials are computed (all the way down to 1), the subroutine starts returning values up through the chain of calls until finally 6-factorial can be computed.

Recursive subroutines aren't very common. Large ones are complicated to construct and difficult to debug. Any task that can be done through iteration (using `for`, `while`, `foreach`) can be done with recursion, and any recursive task can be accomplished iteratively. Recursion is usually reserved for a small set of tasks that lend themselves to it easily.

Summary

Perl supports user-definable functions, called subroutines or subs, that behave just like the built-in functions: They can take arguments, do work, and then return values to the caller if necessary. Perl subroutines can call other subroutines and even themselves. Perl also allows you to declare variables private to a subroutine (or any block of code) and create self-contained pieces of code that can then be reused.

Q&A

Q1: *Is there really any difference between using & and not using & when calling a subroutine?*

A1: None that should concern you at this point. A subtle difference does exist between `&foo` and `foo` when subroutine prototypes are used or when a subroutine is called without parentheses. These are topics well beyond the scope of this book, but for the curious, they're documented in the `perlsub` manual page.

Q2: *When I use `my($var)` in a program, Perl responds with a syntax error, next 2 tokens my(message.*

A2: You either have mistyped something or have Perl version 4 installed. Type `perl -v` at a command prompt. If Perl responds with version 4, you should upgrade immediately.

Q3: *How do I pass (or return) functions, filehandles, multiple arrays, or hashes to a subroutine?*

A3: For passing functions, multiple arrays, and hashes, you need to use a reference, which is covered in [Hour 13](#). For passing filehandles to and from a subroutine, you need to use either something called a *typeglob* or the `IO::Handle` module, both of which are beyond the scope of this book.

Q4: *A function I'm using returns many values, but I'm interested in only one. How do I skip some return values?*

A4: One method is to make a literal list out of the function by wrapping the entire function call in parentheses. When it's a list, you can use ordinary list slices to get pieces of the list. The following extracts just the year (current year minus 1900) from the built-in function `localtime`, which actually has nine return values:

```
print "It is now ", 1900+ (localtime)[5];
```

The other method is to assign the return value from the function into a list and simply assign the unwanted values to `undef` or a dummy variable:

```
(undef, undef, undef, undef, undef, $year_offset)=localtime;
```

 PREV

< Day Day Up >

NEXT 

Workshop

Quiz

Examine the following block of code:

```
sub bar {  
    ($a,$b)=@_;  
    $b=100;  
    $a=$a+1;  
}  
  
sub foo {  
    my($a)=67;  
    local($b)=@_;  
    bar($a, $b);  
}  
  
foo(5,10)
```

1: After you run `bar($a, $b)`, what is the value in `$b`?

a.

5

b.

100

c.

68

2: What is the return value from `foo()`?

a.

67

b.

68

c.

`undef`

3: Inside `foo()`, how is `$b` scoped?

a.

Lexically

b.

Dynamically

c.

Globally

Answers

A1: b. `$b` is declared with `local` in `foo()` so that every called subroutine shares the same value for `$b` (unless they later declare `$b` again with `local` or `my`). After calling `bar()`, where `$b` is modified, `$b` is set to 100.

A2: b. Surprised? The last statement in `foo()` is `bar($a, $b)`. `bar()` returns 68 because the value of `$a` is passed to `bar()`, and it's incremented. `foo()` returns the value of the last expression, which is 68.

A3: b. Variables declared with `local` are called dynamically scoped variables.

Activities

- Use the functions from the statistics exercise in this hour and the word-counting code from [Hour 7](#), "Hashes," to examine the length of the words in a document. Compute their mean, median, and standard deviation.
- Write a function to print part of the Fibonacci series. This series begins 0, 1, 1, 2, 3, 5, 8 and continues forever. The Fibonacci series is a recurring pattern in mathematics and nature. Each successive number is the sum of the previous two (except 0 and 1). These numbers can be computed iteratively or recursively.

 PREV

< Day Day Up >

NEXT 

Part II: Advanced Features

[HOUR 9](#) More Functions and Operators

[HOUR 10](#) Files and Directories

[HOUR 11](#) System Interaction

[HOUR 12](#) Using Perl's Command-Line Tools

[HOUR 13](#) References and Structures

[HOUR 14](#) Using Modules

[HOUR 15](#) Finding Permanence

[HOUR 16](#) The Perl Community

Hour 9. More Functions and Operators

What You'll Learn in This Hour:

- How to do simple string searches on scalars
- How to do character substitutions
- How to use the `print` function
- How to use arrays as stacks and queues

The Perl culture follows the tradition of "There Is More Than One Way To Do It," and in this hour, you will look more closely at that philosophy. You will learn a potpourri of new functions and operators.

For scalar searching and manipulation, you've been using regular expressions until this point. There is more than one way to do the job, so Perl provides a variety of functions for searching and editing scalars. In this hour, some of the other ways will be presented.

Also, you've viewed arrays as linear lists of items to be iterated through with `foreach` or joined together with `join` to make scalars. In this hour, you'll learn a whole new way of looking at arrays.

Finally, the plain vanilla `print` function will be revisited—and spiced up a bit. With a new and improved `print` function, you too can write nicely formatted reports suitable for presentation to others.

Searching Scalars

Regular expressions are nice for searching scalars for patterns, but sometimes they're overkill. In Perl, some overhead—but not much—is involved with assembling the pattern and then searching for the pattern within scalars. Also, you can easily make mistakes when writing regular expressions. Perl provides several functions for searching and extracting simple information from scalars.

Searching with `index`

If you merely want to find one string within another scalar, Perl provides the `index` function. The syntax for `index` is as follows:

```
index string, substring
```

```
index string, substring, start_position
```

The `index` function starts at the left of `string` and searches for `substring`. The `index` function returns the position at which `substring` is found, with 0 being the leftmost character. If the `substring` is not found, `index` returns -1. The string to be searched can be a string literal, a scalar, or any expression that returns a string value. The `substring` is not a regular expression; it's just another scalar.

The following are some examples. Remember that you can write Perl's functions and operators with or without parentheses enclosing the arguments.

```
index "Ring around the rosy", "around";      # Returns 5  
  
index("Pocket full of posies", "ket");        # Returns 3  
  
$a="Ashes, ashes, we all fall down";  
  
index($a, "she");                            # Returns 1  
  
index $a, "they";                            # Returns -1 (not found)  
  
@a=qw(oats      peas      beans);  
  
index join(" ", @a), "peas";                 # Returns 5
```

Optionally, you can give the `index` function a start position in the string to start searching, as shown in the

following example. To start searching at the left, you use the start position of -0.

```
$reindeer="dasher dancer prancer vixen";  
  
index($reindeer, "da");           # Returns 0  
  
index($reindeer, "da", 1);        # Returns 7
```

You also can use the `index` function with a starting position to "walk" through a string and find all the occurrences of a smaller string, as shown here:

```
$source="One fish, two fish, red fish, blue fish.;"  
  
$start=-1;  
  
# Use an increasing beginning index, $start, to find all fish  
  
while( ($start=index($source, "fish", $start)) != -1) {  
  
    print "Found a fish at $start\n";  
  
    $start++;  
  
}
```

The preceding example slides through `$source`, as shown here:

\$start begins at -1
↓
one fish, two fish, red fish, blue fish
↑
index finds "fish" at position 4, saves that in \$start
index is run again at new \$start position +1 (5)
↓
one fish, two fish, red fish, blue fish
↑
index finds another "fish" at position 14

Searching Backward with `rindex`

The function `rindex` works the same as `index`, except that the search starts on the right and works its way left. The syntax is as follows:

```
rindex string, substring
```

```
rindex string, substring, start_position
```

When the search is exhausted, the `rindex` function returns `-1`. The following are some examples:

```
$a="She loves you yeah, yeah, yeah.";  
  
rindex($a, "yeah");           # Returns 26.  
  
rindex($a, "yeah", 25);      # Returns 20
```

The walk-through loop used with `index` looks a little different searching backward with `rindex`. The `rindex` start position must start at (or after) the end of the string—`length($source)` in the following example—but still finishes when `-1` is returned. After each find, `$start` must be decremented by 1, rather than incremented as it was with `index`.

```
$source="One fish, two fish, red fish, blue fish.";  
  
$start=length($source);  
  
while( ($start = rindex($source, "fish", $start)) != -1) {  
  
    print "Found a fish at $start\n";  
  
    $start--'  
  
}
```

Picking Apart Scalars with `substr`

The `substr` function is often overlooked and easily forgotten, but it provides a general-purpose method for extracting information from scalars and editing scalars. The syntax of `substr` is as follows:

```
substr string, offset
```

```
substr string, offset, length
```

The `substr` function takes `string`, starting at position `offset`, and returns the rest of the string from `offset` to the end. If `length` is specified, then `length` characters are taken—or until the end of the string is found, whichever comes first—as shown in this example:

```
#Character positions in $a
#   0          10         20         30
$a="I do not like green eggs and ham.";

print substr($a, 25);      # prints "and ham."
print substr($a, 14, 5);   # prints "green"
```

If the `offset` specified is negative, `substr` starts counting from the right. For example, `substr($a, -5)` returns the last five characters of `$a`. If the `length` specified is negative, `substr` returns from its starting position to the end of the string, less `length` characters, as in this example:

```
print substr($a, 5, -10);    # prints "not like green egg"
```

In the preceding snippet, `substr` starts at position 5 and returns the rest of the string except the last 10 characters.

You can also use the `substr` function on the left side of an assignment expression. When used on the left, `substr` indicates what characters are to be replaced in a scalar. When `substr` is used on the left side of an assignment, the first argument must be an assignable value—such as a scalar variable—and not a string literal. The following is an example of editing a string with `substr`:

```
$a="countrymen, lend me your wallets";

# Replace first character of $a with "Romans, C"
substr($a, 0, 1)="Romans, C";
```

```
# Insert "Friends" at the beginning of $a  
  
substr($a, 0, 0)="Friends, "  
  
substr($a, -7, 7)="ears." # Replace last 7 characters.
```



< Day Day Up >



Transliteration, Not Substitution

For this next operator, the *transliteration* operator (sometimes called the *translation* operator), think back to how regular expression substitutions work: The substitution operator, which looks like `s/pattern/replacement/` and was discussed in Hour 6, works against the `$_` variable unless another scalar is specified with a binding operator `=~`. The transliteration operator works something like that, except that it doesn't use regular expressions and works completely differently. Still follow? The syntax for the transliteration operator is as follows:

```
tr/searchlist/replacementlist/
```

The transliteration operator—`TR///`—searches a string for the elements in *searchlist* and replaces them with the corresponding elements in *replacementlist*. By default, the transliteration operator searches and modifies the variable `$_`. To search and modify other variables, you use a binding operator as you would for regular expression matches, as shown here:

```
tr/ABC/XYZ/;          # In $_, replaces all A's with X's, B's with Y's, etc..  
$r=~tr/ABC/XYZ/;    # Does the same, but with $r
```

Logical groups of characters are accepted with dashes between them. For example, `A-Z` represents the capital letters `A` through `Z`, so that you don't have to write them all out, as in this example:

```
tr/A-Z/a-z/;          # Change all uppercase to lowercase  
tr/A-Za-z/a-zA-Z/;   # Invert upper and lowercase
```

If *replacementlist* is empty or identical to *searchlist*, the characters matched are counted by `tr///` and returned. The target string is not modified, as in the following example:

```
$eyes=$potato=~tr/i//;      # Count the i's in $potato, return to $eyes  
$nums=tr/0-9//;            # Count digits in $_, return to $nums
```

Finally, for historical reasons, `TR///` can also be written as `y///` with the same results, because `y` is a synonym for `tr`. The `TR///` operator (and hence, `y///`) also allows you to specify an alternate set of delimiters for `searchlist` and `replacementlist`. They can be any naturally paired set such as parentheses or any other character, as you can see here:

```
tr(a-z)(n-zA-m);          # Rotate all characters 13 to the left in $_
y[ ,.-_-][;:=|];          # Switch around some punctuation
```

By the Way

The `TR///` operator actually has additional functionality, but it isn't used often. To read about all the other tasks `TR///` can perform, look at the online documentation in the perlop section.



< Day Day Up >



A Better Way to print

The `print` function is a very simple output function; it provides almost no formatting capability. To have fine-grained control over the appearance of expression values, such as left and right alignment, number of digits of decimal precision, and fixed-width output, you must use the Perl `printf` function instead.

Formatted Printing with `printf`

The `printf` function was borrowed (almost verbatim) from the C programming language, but other languages have a similar function, such as BASIC's `print using` function. The syntax for `printf` is as follows:

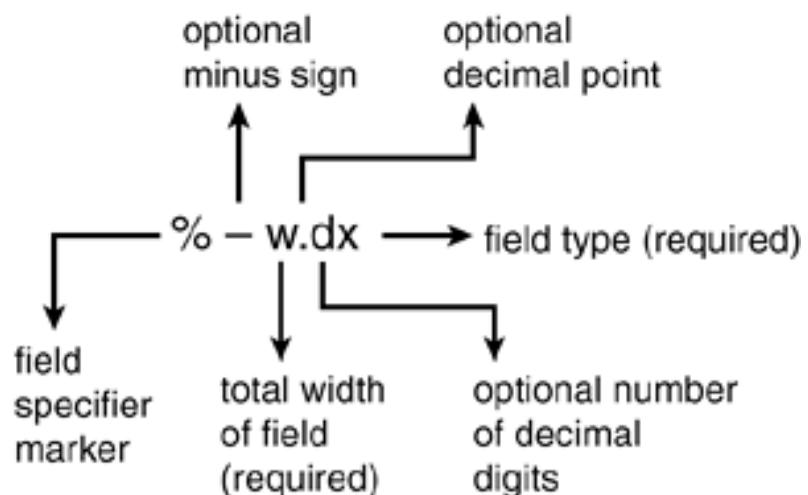
```
printf formatstring, list
```

```
printf filehandle formatstring, list
```

The `formatstring` is a string that describes the format of the output, which is described shortly. The `list` is a list of values that you want `printf` to display—somewhat similar to the arguments to the `print` function. Normally, `printf` displays its output to the `STDOUT` filehandle, but, as with `print`, if you specify a filehandle, `printf` uses that filehandle instead. Note that no comma is used between the `filehandle` name and the `formatstring`.

Specifying the Field Formats

The `formatstring` is a string literal (usually) or a scalar that describes what the output should look like. Every character in the `formatstring` is printed literally, except those character sequences beginning with a `%`. The `%` indicates the start of a *field specifier*. The format of a field specifier is shown here:



As seen in the figure, field specifiers have the form `%-w.dx`, where `w` is the total desired width of the field. `d`

is the number of positions to the right of the decimal point (for numbers) and the total allowable width of the field (for strings). `x` indicates what kind of data is to be printed. A minus sign (a hyphen) before the `w` specifier indicates that the field is to be left-justified within `w` characters; otherwise, it's right-justified. Only the `%` and `x` are mandatory. [Table 9.1](#) lists some of the different types of field specifiers (that is, values for the `x` position).

Table 9.1. Partial List of `printf` Field Specifier Types

Field Type Meaning

c	Character
s	String
d	Decimal integer; fraction is truncated
f	Floating-point number

A full list of field specifier types is included in the online manual. You can view it by typing `perldoc -f printf` at a command prompt.

The following are some examples of using `printf`:

```
printf("%20s", "Jack");           # Right-justify "Jack" in 20-characters
printf("%-20s", "Jill");          # Left-justify "Jill" in 20 characters
$amt=7.12;
printf("%6.2f", $amt);           # prints "      7.12"
$amt=7.127;
printf("%6.2f", $amt);           # prints "      7.13", extra digits rounded
printf("%c", 65);                # prints ASCII character for 65, "A"
$amt=9.4;
printf("%6.2f", $amt);           # prints "      9.40"
printf("%6d", $amt);             # prints "      9"
```

Each field specifier uses one item from the list, as shown in the following figure. For each item, there should be a field specifier; for every field specifier, there should be a list element.

```
print f ("Totals: %6.2f %15s %7.2f %6d", $a $b $c $d);
```

To print leading zeros in a number, simply put a 0 in front of the width in the field specifier, as follows:

```
printf("%06.2f", $amt); # prints "009.40"
```

What if you want a literal percent sign in your output? The sequence `%%` represents a percent sign:

```
$newprice = 1.449;  
  
$oldprice = 1.229;  
  
printf("The price went up by %3.1f%%!\n", 100*(($newprice/$oldprice) - 1));  
  
# prints "The price went up by 17.9%!"
```

Formatted Output to a String

The `sprintf` function is nearly identical to `printf`, except that instead of being printed, the formatted output is returned from `sprintf`—ready for you to assign to a scalar or use in another expression, as you can see here:

```
$weight=85;  
  
# Format result nicely to 2-decimals  
  
$moonweight = sprintf("%.2f", $weight / 6);  
  
print "You weigh $moonweight on the moon.";
```

Remember that `printf` and `sprintf` with the `%f` format specifier will round your results to the specified number of decimal places for you.

 PREV

< Day Day Up >

NEXT 

Exercise: A Formatted Report

A task that comes up inevitably when you're dealing with computers is formatting raw data into a report. Computer programs exchange data in formats that are difficult for humans to read, and a common task is taking that data and formatting it into a human-friendly report.

For this exercise, you're given a set of employee records that contain information about some mythical employees, including hourly wages, number of hours worked, names, and employee numbers. The exercise takes that data and reformats it into a nice report.

You can easily modify this same kind of program to print other kinds of reports. The data for the exercise is contained within an array initialized at the beginning of the program. In a real report, the data would probably come from a file on disk. Modifying this exercise to use an external file is left as an exercise for later.

Using your text editor, type the program from [Listing 9.1](#) and save it as `Employee`. Do not type the line numbers. Make the program executable according to the instructions you learned in [Hour 1](#), "Getting Started with Perl."

When you're all done, try running the program by typing the following at a command line:

`Employee`

or, if you cannot make the program executable,

`perl Employee`

Listing 9.1. Complete Listing of Employee Program

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:
5:  my @employees = (
```

```
6:      'Smith,Bob,123101,9.35,40' ,
7:      'Franklin,Alice,132912,10.15,35' ,
8:      'Wojohowicz,Ted,198131,6.50,39' ,
9:      'Ng,Wendy,141512,9.50,40' ,
10:     'Cliburn,Stan,131211,11.25,40' ,
11:   );
12:
13: sub print_emp {
14:   my($last,$first,$emp,$hourly,$time)=
15:     split(',',$_[0]);
16:   my $fullname;
17:   $fullname = sprintf("%s %s", $first, $last);
18:   printf("%6d %-20s %6.2f %3d %7.2f\n",
19:         $emp, $fullname, $hourly, $time,
20:         ($hourly * $time) + .005 );
21: }
22:
23: @employees = sort {
24:   my ($L1, $F1)=split(',',$a);
25:   my ($L2, $F2)=split(',',$b);
26:   return( $L1 cmp $L2  # Compare last names
27:          ||  # If they're the same...
28:          $F1 cmp $F2  # Compare first
29:        );
30: } @employees;
```

31:

```
32: foreach (@employees) {  
33:     print_emp($_);  
34: }
```

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The `use strict` directive means that all variables must be declared with `my` and that bare words must be quoted.

Lines 5–11: The list of employees is assigned to `@employees`. Each element in the array consists of a last name, first name, employee number, hourly wage, and number of hours worked.

Lines 23–30: The `@employees` array is sorted by last name and first name.

Line 24: The first element to be sorted (`$a`) is split apart into fields. The last name is assigned to `$L1` and the first name to `$F1`. Both of these are declared private to the sort block with `my`.

Line 25: The same is done with another element, `$b`. The names are assigned to `$L2` and `$F2`.

Line 26–29: The names are compared alphabetically using something similar to the sort shown in [Listing 4.1](#) in [Hour 4](#), "Stacking Building Blocks: Lists and Arrays."

Lines 32–24: The sorted list in `@employees` is passed, one element at a time, to `print_emp()`.

Lines 13–21: The `print_emp()` function prints the employee records nicely formatted.

Lines 14–15: The passed-in record—in `$_[0]`—is split into fields and assigned to variables—`$last`, `$first`, and so on—which are all private to this subroutine. Remember, `$_[0]` represents the first argument passed to a function -- the first element of `@_`.

Line 17: The last name and first name are combined into a single field so that the two fields can be padded to a certain width and justified together.

Lines 18–20: The record is printed. `$hours` and `$time` are multiplied to give the total amount earned. The amount .005 is added to the total so that, when the product is truncated to two digits, it's properly rounded.

[Listing 9.2](#) shows a sample of the Employee program's output.

Listing 9.2. Output from the Employee Program

131211 Stan Cliburn 11.25 40 450.00

132912 Alice Franklin 10.15 35 355.25

141512 Wendy Ng 9.50 40 380.00

123101 Bob Smith 9.35 40 374.00

198131 Ted Wojohowicz 6.50 39 253.50

 PREV

< Day Day Up >

NEXT 

New Ways with Arrays

In the following section, we will discuss a set of array functions that make several powerful data structures possible in Perl.

A List as a Stack

Until now, lists (and arrays) have been presented as linear arrangements of data with indexes that make it possible to access each element, as shown here:

0	1	2	3	4	5
Apple	Peach	Pear	Plum	Mango	Guava

Use your imagination for a moment and imagine the array elements piled up vertically with the element having the highest index at the top, as shown here:

Guava
Mango
Plum
Pear
Peach
Apple

Notice that there are no index numbers in this figure. The reason is that the functions we will discuss here are concerned with only the top and the bottom of the array, not with anything in the middle. In computer terminology, this list is called a *stack*. Stacks are useful for accumulating tasks that you want processed in order. A good example is a game of Klondike solitaire (like the Solitaire game often installed with Windows). Each of the seven piles of cards represents a stack. Initially, cards are placed onto the stack face down. As they are needed, they are turned over and removed from the top of the stack, or additional cards can be placed on top of the newly turned-over cards, but cards cannot be removed or inserted in the middle of the stack or at the bottom.

Stacks in Perl are (usually) implemented with arrays. To place items on top of a stack, you use the `push` function. To remove from the top of the stack, you use the `pop` function.

A stack implemented as a Perl array can also be modified from the bottom, like dealing from the bottom of the deck of cards. The `shift` function removes elements from the bottom, and `unshift` adds elements to the bottom of the stack. All four operations are shown here:



Summary

In this hour, you learned that searching for strings within other strings doesn't have to involve regular expressions; you can perform simple searches by using `index` and `rindex`. You also can make simple substitutions by using the `TR///` operator. The `substr` function can be used for both extracting data from strings and editing them. You can create nicely formatted output in Perl by using the `printf` and `sprintf` statements. Additionally, you saw arrays used as stacks of items instead of flat lists, and you learned how to manipulate those stacks.

Q&A

Q1: Are `substr`, `index`, and `rindex` really necessary? Why do they exist when regular expressions can be used for most of those operations?

A1: First, regular expressions for simple string searches are generally slower than `index` and `rindex`. Second, writing substitution expressions for fixed character positions with regular expressions can be messy; `substr` is a much more elegant solution sometimes. Third, Perl is a rich language. Use what you like; you have plenty of choices.

Q2: What happens with `substr` (or `index` or `rindex`) if I specify an index that's beyond the end of the scalar?

A2: One nice thing about computers is that they're consistent and have limitless amounts of patience. With questions like "What happens if I...?" sometimes just trying it is easiest! What's the worst that can happen?

In this case, accessing a portion of a scalar that doesn't exist might cause a `use of undefined value` error if you have warnings enabled, as you should. For example, if you use `$a="Foo"; substr($a, 5);`, the `substr` function returns `undef`.

Workshop

Quiz

1: Given the following code, what's left in @A afterward:

```
@A=qw(oats peas beans);  
  
shift @A;  
  
push @A, "barley";  
  
pop;
```

a.

oats peas beans

b.

beans barley

c.

peas beans barley

2: What does `printf("%18.3f", $a)` do?

a.

It prints a floating-point number that is 18 characters wide: 15 to the left and 3 to the right of the decimal point.

b.

It prints a floating-point number 18 characters to the left of the decimal, a decimal point, and then 3 more digits.

c.

It prints a floating-point number that is 18 characters wide, 14 to the left and 3 to the right of the decimal point.

- 3:** If `TR/a-z/A-Z/` is run against a string, will `tr/A-Z/a-z/` restore the string to its original form?

a.

Yes, of course

b.

Probably not

Answers

A1: c. The `shift` removed `oats`, and the `push` added `barley` to the end. The final `pop` was a ruse: No array was specified, so it popped something from the array `@_` or from `@ARGV`, which causes nothing to happen to `@A`.

A2: c. If you guessed a, you didn't count the decimal point, which occupies a position in the total ($18=14+1+3$).

A3: b. The string `"Rosebud"` transformed by `tr/a-z/A-Z/` becomes `"ROSEBUD"`. TRYing to transform it back with `tr/A-Z/a-z/` yields `"rosebud"`—not the original string.

Activities

- Rewrite the Hangman game from [Hour 4](#) to use scalars instead of arrays. You can use `substr` to manipulate individual characters in the scalars.
- Modify [Listing 9.2](#) to read from a file instead of getting data from an array. Open the file, read the data into an array, and continue as normal. You will have to create the file on disk, of course.

Hour 10. Files and Directories

What You'll Learn in This Hour

- How to get a directory listing
- How to create and remove files
- How to create and remove directories
- How to get information about files

Files in your operating system provide a convenient set of storage concepts for data. The OS enables a name to be given to the data (a filename) and provides an organizational structure, called a **file system**, so that you can find the data later. Your computer's file system then organizes files into groups called **directories**—sometimes called **folders**. These directories can store files or other directories.

This nesting of directories inside directories provides a treelike structure to the file system on your computer. Each file is part of a directory, and each directory is part of a parent directory. In addition to providing an organizational structure for your files, the operating system also stores data about the file: when the file was last read, when it was last modified, who created it, the current size of the file, and so on—called metadata (see [Hour 5](#), "Working with Files"). This organization is true of almost all modern computer operating systems.

In the case of the Macintosh (pre-Mac OS X), this structure still holds true, except that the top-level directory is called a Volume, and the subdirectories area is called Folders.

Perl allows you to access this structure, modify the organization, and examine the information about the files. The functions that Perl uses for these tasks are all derived from the Unix operating system, but they work just fine under whatever operating system Perl happens to be running on. Perl's file system manipulation functions are portable, meaning that if you use Perl's functions to manipulate your files and query them, you should have no problems running your code under any operating system Perl supports, providing that the directories are structured similarly.

Getting a Directory Listing

The first step in obtaining directory information from your system is to create a **directory handle**. A directory handle is something like a filehandle, except that instead of a file's contents, you read the contents of a directory through the directory handle. To open a directory handle, you use the `opendir` function:

```
opendir dirhandle, directory
```

In this syntax, `dirhandle` is the directory handle you want to open and `directory` is the name of the directory you want to read. If the directory handle cannot be opened—because you don't have permission to read the directory, the directory doesn't exist, or because of some other reason—the `opendir` function returns false. Directory handle variable names should be constructed similarly to filehandles—using the rules for variable names outlined in Hour 2, "Perl's Building Blocks: Numbers and Strings"—and, like filehandles, they should be all uppercase to avoid conflicts with Perl's keywords. The following is an example:

```
opendir(TEMPDIR, '/tmp') || die "Cannot open /tmp: $!" ;
```

All the examples in this hour use forward slashes (/) in the Unix style because it is less confusing than the backslashes (\) used by Windows and MSDOS and works just as well with those operating systems as with Unix.

Now that the directory handle is open, you use the `readdir` function to read it:

```
readdir dirhandle;
```

In a scalar context, `readdir` returns the next entry in the directory, or `undef` if none are left. In a list context, `readdir` returns all the (remaining) directory entries. The names returned by `readdir` include files, directories, and (for Unix) special files; they are returned in no particular order. The directory entries `.` and `..` (representing the current directory and its parent directory) are also returned by `readdir`. The directory entries returned by `readdir` do not include the pathname as part of the name returned.

When you're done with the directory handle, you should close it by using `closedir`:

```
closedir dirhandle;
```

The following example shows how to read a directory:

```
opendir(TEMP, '/tmp') || die "Cannot open /tmp: $!";  
@FILES=readdir TEMP;  
closedir(TEMP);
```

In this preceding snippet, the entire directory is read into `@FILES`. Most of the time, however, you're not interested in the `.` and `..` files. To read the filehandle and eliminate those files, you can enter the following:

```
@FILES=grep( !/^\.\\.?$/ , readdir TEMP );
```

The regular expression `(/^\\.\\.?$/)` matches a leading literal dot (or two) that is also at the end of the line, and `grep` eliminates them. To get all the files with a particular extension, you use the following:

```
@FILES=grep( /\.txt$/i, readdir TEMP);
```

The filenames returned by `readdir` do not contain the pathname used by `opendir`. Thus, the following example will probably not work:

```
opendir(TD, "/tmp") || die "Cannot open /tmp: $!";  
while($file = readdir TD) {  
    # The following is WRONG  
    open(FILEH, $file) || die "Cannot open $file: $!\n";  
    # Process the file here...  
}  
closedir(TD);
```

Unless you happen to be working in the `/tmp` directory when you run this code, the `open(FILEH, $file)` statement will fail. For example, if the file `myfile.txt` exists in `/tmp`, `readdir` returns `myfile.txt`. When you open `myfile.txt`, you actually need to open `/tmp/myfile.txt` using the full pathname. The corrected code is as follows:

```
opendir(TD, "/tmp") || die "Cannot open /tmp: $!" ;  
  
while($file=readdir TD) {  
  
    # Right!  
  
    open(FILEH, "/tmp/$file") || die "Cannot open $file: $!\n" ;  
  
    # Process the file here...  
  
}  
  
closedir(TD);
```

Globbing

The other method of reading the names of files in a directory is called **globbing**. If you're familiar with the command prompt in MS-DOS, you know that the command `dir *.txt` prints a directory listing of all the files that end in `.txt`. In Unix, the globbing (sometimes called wildcard matching) is done by the shell, but `ls *.txt` has nearly the same result: The files whose names end in `.txt` are listed.

Perl has an operator for doing just this job; it's called `glob`. The syntax for `glob` is

`glob pattern`

where `pattern` is the filename pattern you want to match. The `pattern` can contain directory names and portions of filenames. In addition, the `pattern` can contain any of the special characters listed in [Table 10.1](#). In a list context, `glob` returns all the files (and directories) that match the pattern. In a scalar context, the files are returned one at a time each time `glob` is queried.

Table 10.1. Globbing Patterns

Character	Matches	Example
?	Single character	f?d matches <i>fud</i> , <i>fid</i> , <i>fdd</i> , and so on
*	Any number of characters	f*d matches <i>fd</i> , <i>fdd</i> , <i>food</i> , <i>filled</i> , and so on
[<i>chars</i>]	Matches any of <i>chars</i> ; this feature is not supported in MacPerl	f[ou]d matches <i>fod</i> and <i>fud</i> but not <i>fad</i>
{ <i>a</i> , <i>b</i> ,...}	Matches either of the strings <i>a</i> or <i>b</i> ; this feature not supported in MacPerl	f*.{txt,doc} matches files that begin with <i>f</i> and that end in either <i>.txt</i> or <i>.doc</i>

Watch Out!

Glob patterns are not the same as regular expressions.

Watch Out!

Unix fans, please note: Perl's `glob` operator uses C shell-style file globbing, as opposed to Bourne (or Korn) shell file globbing. This is true on any Unix system in which Perl is installed, regardless of whatever shell you personally use. The Bourne shell globbing and Korn shell globbing are different from the C shells. They are very similar in some respects—* and ? behave the same—but quite different in others. Beware.

Now check these examples of globbing:

```
# All of the .h files in /usr/include
my @hfiles=glob('/usr/include/*.h');

# Text or document files that contain 1999
my @curfiles=glob('*1999*.{txt,doc}')

# Printing a numbered list of filenames
$count=1;

while( $name=glob('*') ) {
    print "$count. $name\n";
```

```
$count++;
```

```
}
```

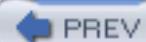
An important difference between `glob` and `opendir/readdir/closedir` is that `glob` returns the pathname used in the pattern, whereas the `opendir/readdir/closedir` functions do not. For example, `glob('/usr/include/* .h')` returns '`/usr/include`' as part of any matches; `readdir` does not.

So which should you use? It's completely up to you. However, using the `opendir/readdir/closedir` functions tends to be a much more flexible solution and will be used in most of the examples throughout this book.

Perl offers an alternative way to write pattern globs. Simply placing the pattern inside the angle operator (`<>`) makes the angle operator behave like `glob`:

```
@cfiles = <*.c>; # All files ending in .c
```

The syntax that uses the angle operator for globbing is older and can be confusing. In this book, I will continue to use the `glob` operator instead for clarity.

 PREV

< Day Day Up >

NEXT 

Exercise: The Unix grep

As you get further along in this book, the exercises will present you with more and more useful tools. This exercise presents a stripped-down version of the Unix `grep` utility. The Unix `grep`—not to be confused with Perl's `grep` function, introduced in [Hour 6](#), "Pattern Matching"—searches files for patterns. This exercise presents a utility that will prompt for a directory name and a pattern. Every file in that directory will be searched for that pattern, and lines matching that pattern will be printed.

In future exercises, this utility will be modified to search subdirectories (see [Hour 15](#), "Finding Permanence") and to take command-line arguments (see [Hour 12](#), "Using Perl's Command-Line Tools"). Stay tuned for details.

Using your text editor, type the program from [Listing 10.1](#) and save it as `mygrep`. If possible, be sure to make the program executable according to the instructions you learned in [Hour 1](#), "Getting Started with Perl." Also, make sure that you don't rename this file to `grep` on a Unix system because it could be mistaken for the real `grep` utility.

When you're all done, try running the program by typing the following at a command line:

```
perl -w mygrep
```

or, if your system enables you to make the file executable,

```
mygrep
```

Listing 10.1. Complete Listing for `mygrep`

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:
5:  print "Directory to search: ";
6:  my $dir=<STDIN>; chomp $dir;
```

```

7:     print "Pattern to look for: ";
8:     my $pat=<STDIN>; chomp $pat;
9:
10:    my($file);
11:
12:    opendir(DH, $dir) || die "Cannot open $dir: $!";
13:    while ($file=readdir DH) {
14:        next if (-d "$dir/$file");
15:        if (! open(F, "$dir/$file")) {
16:            warn "Cannot search $file: $!";
17:            next;
18:        }
19:        while(<F>) {
20:            if (/^$pat/) {
21:                print "$file: $_";
22:            }
23:        }
24:        close(F);
25:    }
26:    closedir(DH);

```

Line 1: This line contains the path to the interpreter (you can change it so that it's appropriate to your system) and the `-w` switch. Always have warnings enabled!

Line 3: The `use strict` directive means that all variables must be declared with `my` and that bare words must be quoted.

Lines 5–8: `$dir`, the directory to be searched, and `$pat`, the pattern to search for, are retrieved from STDIN. The newlines at the end of each are removed.

Line 10: `$file` is declared as private to satisfy `use strict`. `$file` is used later in this program.

Line 12: The directory `$dir` is opened; an error message is printed if this operation fails.

Line 13: The entries are retrieved from the directory one at a time and stored in `$file`.

Line 14: Any directory entry that's really a directory itself (`-d`) is rejected. Notice that the pathname checked is `$dir/$file`. This path must be checked because `$file` doesn't necessarily exist in the current directory; it exists in `$dir`. So the full pathname to the file is `$dir/$file`.

Lines 15–18: The file is opened, again using the full pathname `$dir/$file`, and rejected if it does not open.

Lines 19–23: The file is searched, line by line, for a line that contains `$pat`. A matching line is printed.

[Listing 10.2](#) shows a sample of the `mygrep` program's output.

Listing 10.2. Output from `mygrep`

```
Directory to search: /home/clintp
```

```
Pattern to look for: printer
```

```
mailbox: lot of luck re-inking Epson printer ribbons with
```

```
config.pl: # the following allows the user to pick a printer for
```

Directories

Thus far in this hour, I've been sort of handwaving over the topic of directory structure. Full pathnames are sometimes needed to open files, and the `readdir` function can read directories. But actually navigating directories, adding or removing them, and cleaning them out takes a little bit more Perl.

Navigating Directories

When you run software, your operating system keeps track of what directory you're in when you run the software. When you log in to a Unix machine and run a software package, you are usually placed in your home directory. If you type the operating system command `pwd`, the shell shows you what directory you are in. If you use MS-DOS or Windows and open a **command prompt**, the prompt reflects what directory you are in at the time—for example, `C:\WINDOWS`. Alternatively, you can type the operating system command `cd` at the MS-DOS prompt, and MS-DOS tells you what directory you are in. The directory that you're currently using is called your **current directory** or your **current working directory**.

Watch Out!

If you're using a program editor or an integrated editor/debugger and running your Perl programs directly from there, your "current directory" might not be what you think. It might be the directory that the Perl program is in, the directory your editor is in, or any other directory—it depends on your editor. Use the `cwd` function below from inside your Perl program to determine exactly where your current directory is.

If you do not specify a full pathname when you try to open a file—for example, `open(FH, "file") || die`—Perl will attempt to open the file in your current working directory. To change your current directory, you can use the Perl `chdir` function, as follows:

```
chdir newdir;
```

The `chdir` function changes the current working directory to `newdir`. If the `newdir` directory does not exist, or you don't have permission to access `newdir`, `chdir` returns false. The directory change from `chdir` is temporary; as soon as your Perl program ends, you return to the directory that you were working in before you ran the Perl program.

Running the `chdir` function without a directory as an argument causes `chdir` to change to your home directory. On a Unix system, the home directory is usually the directory that you were placed in when you logged in. On a Windows 95, Windows NT, or MS-DOS machine, `chdir` takes you to the directory indicated in the `HOME` environment variable. If `HOME` isn't set, `chdir` doesn't change your current directory.

at all.

Perl doesn't have a built-in function for figuring out what your current directory is—because of the way some operating systems are written, it's not easy to tell. To find the current directory, you must use two statements together. Somewhere in your program—preferably near the beginning—you must use the statement `use Cwd` and then, when you want to retrieve the current directory, use the `cwd` function:

```
use Cwd;  
  
print "Your current directory is: ", cwd, "\n";  
  
chdir '/tmp' or warn "Directory /tmp not accessible: $!";  
  
print "You are now in: ", cwd, "\n";
```

You have to execute the `use Cwd` statement only once; afterward, you can use the `cwd` function as often as necessary.

By the Way

The statement `use Cwd` actually instructs Perl to load a module called `Cwd` that extends the Perl language with new functions, such as `cwd`. If the preceding snippet returns an error saying `Can't locate Cwd.pm in @INC`, or you don't understand modules completely, don't worry about it right now. Modules will be fully discussed in [Hour 14](#), "Using Modules."

Creating and Removing Directories

To create a new directory, you can use the Perl `mkdir` function. The `mkdir` function's syntax is as follows:

```
mkdir newdir, permissions;
```

The `mkdir` function returns true if the directory `newdir` can be created. Otherwise, it returns false and sets `$!` to the reason that `mkdir` failed. The `permissions` are really important only on Unix implementations of Perl, but they must be present on all versions. For the following example, use the value `0755`; this value will be explained in the section "[Unix Stuff](#)" later in this hour. For MS-DOS and Windows users, just use the value `0755`; it's good enough and will spare you a long explanation.

```
print "Directory to create?";

my $newdir=<STDIN>;

chomp $newdir;

mkdir( $newdir, 0755 ) || die "Failed to create $newdir: $!";
```

To remove a directory, you use the `rmdir` function. The syntax for `rmdir` is as follows:

```
rmdir pathname;
```

The `rmdir` function returns true if the directory *pathname* can be removed. If *pathname* cannot be removed, `rmdir` returns false and sets `$!` to the reason that `rmdir` failed, as shown here:

```
print "Directory to be removed?";

my $baddir=<STDIN>;

chomp $baddir;

rmdir($baddir) || die "Failed to remove $baddir: $!";
```

The `rmdir` function removes only directories that are completely empty. This means that before a directory can be removed, all the files and subdirectories must be removed first.

Removing Files

To remove files from a directory, you use the `unlink` function:

```
unlink list_of_files;
```

The `unlink` function removes all the files in the *list_of_files* and returns the number of files removed. If *list_of_files* is omitted, the filename named in `$_` is removed. Consider these examples:

```
unlink <*.bat>;  
  
$erased=unlink 'old.exe', 'a.out', 'personal.txt';  
  
unlink @badfiles;  
  
unlink;      # Removes the filename in $_
```

To check whether the list of files was removed, you must compare the number of files you tried to remove with the number of files removed, as in this example:

```
my @files=<*.txt>;  
  
my $erased=unlink @files;  
  
# Compare actual erased number, to original number  
  
if ($erased != @files) {  
  
    print "Files failed to erase: ",  
          join( ', ', <*.txt> ), "\n";  
  
}
```

In the preceding snippet, the number of files actually erased by `unlink` is stored in `$erased`. After the `unlink`, `$erased` is compared to the number of elements in `@files`: They should be the same. If they're not, an error message is printed showing the "leftover" files.

Watch Out!

Files erased with `unlink` are irrevocably erased. There is no recovery, and they are not moved to a "trash" or "recycle bin." Be careful with `unlink`!

Renaming Files

Renaming files or directories in Perl is simple; you use the `rename` function, as follows:

```
rename oldname, newname;
```

The `rename` function takes the file named `oldname`, tries to change its name to `newname`, and returns true if the rename is successful. If the `oldname` and `newname` are directories, those directories are renamed. If the rename is unsuccessful, `rename` returns false and sets `$!` to the reason why, as shown here:

```
if (! rename "myfile.txt", "archive.txt") {  
    warn "Could not rename myfile.txt: $!";  
}
```

The `rename` function also moves the file from one directory to another if you specify full pathnames instead of just filenames, as in this example:

```
rename "myfile.txt", "/tmp/myfile.txt";      # Effectively it moves the file.
```

If the file `newname` already exists, it is destroyed.

By the Way

The `rename` function will not move files from one directory to another if they're on different file systems.

Unix Stuff

This portion of the hour is primarily for the users of Perl who are working under the Unix operating system. If you don't use Perl on a Unix system, you can safely skip this section and not miss anything important; you can read it if you're curious about Unix, however.

Unix users should know that Perl has a deep Unix heritage, and some Perl functions come right from Unix commands and operating system functions. Most of these functions you will not use. Some functions—like `unlink`—come from Unix but have meanings that are not really tied to Unix at all. Every operating system deletes files, and Perl ensures that `unlink` does the right thing on every operating system. Perl makes a great effort to ensure that concepts—such as file I/O—that should be portable between operating systems actually are portable, and it hides all the compatibility issues from you if it can.

The fact that many Unix functions and commands are embedded in the Perl language—which has been ported to many non-Unix operating systems—is a tribute to the desire of Unix developers and administrators to take a little bit of the Unix toolkit with them wherever they go.

By the Way

As the next section header suggests, this description is not to be taken as a full explanation of Unix file system permissions and how to manipulate them. For a full explanation, consult your operating system's documentation or any good book on Unix, such as *Sams Teach Yourself Unix in 24 Hours*.

A Crash Course in File Permissions

In Hour 1, to make a Perl program run as though it were a regular command, you were given the command `chmod 755 scriptname` without any real explanation of what it meant. First of all, `chmod` is a command in Unix that sets the permissions on files. Next, the `755` is a description of the permissions being given to the file `scriptname`. Each of the three digits represents one set of permissions, given to the owner of the file, the group that the file belongs to, and any nonowner and nongroup—called other—user. In this case, the owner has permission `7`, whereas the group and others have permission `5`, as shown in Figure 10.1.

Figure 10.1. Breakdown of Unix permission bits.

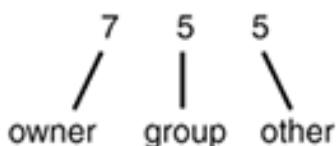


Table 10.2 lists the possible values of the digit for each possible combination of permissions.

Table 10.2. File Permissions

Permission Allows That Permission Group

7	Reading, writing, and executing the file
6	Reading and writing the file
5	Reading and executing the file
4	Reading the file only
3	Writing and executing the file
2	Writing the file only
1	Executing the file only
0	No access to the file

To set permissions on a file in Perl, you use the built-in function called `chmod`:

```
chmod mode, list_of_files;
```

The `chmod` function changes the permission on all the files in `list_of_files` and returns the number of files whose permissions were changed. The `mode` must be a four-digit number beginning with 0 (because it's a literal octal number, as mentioned in Hour 2 and later in this hour) followed by the digits that you want to indicate permission. The following are some examples of `chmod`. Note that R stands for read permission, W for write permission, and X for execute permission:

```
chmod 0755, 'file.pl';      # Grants RWX to owner, RX to group and non-owners  
chmod 0644, 'mydata.txt';  # Grants RW to owner, and R to group and non-owners  
chmod 0777, 'script.pl';    # Grants RWX to everyone (usually, not very smart)  
chmod 0000, 'cia.dat';     # Nobody can do anything with this file
```

Earlier this hour, you learned about the `mkdir` function. The first argument of `mkdir` is a file permission—the same kind that `chmod` uses:

```
mkdir "/usr/tmp", 0777;  # Publicly readable/writable/executable  
mkdir "myfiles", 0700;   # A very private directory
```

By the Way

The permissions on a UNIX file are often called its **mode**. So `chmod` is simply short for "change mode."



< Day Day Up >



Everything You Ever Wanted to Know About THAT File

To find out, in excruciating detail, everything you might want to know about a file, you can use Perl's `stat` function. The `stat` function originated in Unix, and the return values differ slightly between Unix and non-Unix systems. The syntax for `stat` is as follows:

```
stat filehandle;
```

```
stat filename;
```

The `stat` function can retrieve information either about an open filehandle or about a particular file. Under any operating system, `stat` returns a 13-element list describing the attributes of the file. The actual values in the list differ slightly depending on which operating system you're running because some operating systems include features that others do not implement. Table 10.3 shows what each element in `stat`'s return value stands for.

Table 10.3. Return Values from `stat`

Index	Name	Unix	Windows
0	<code>dev</code>	Device number	Drive number (C: is usually 2, D: is usually 3, and so on)
1	<code>ino</code>	Inode number	Zero, always
2	<code>mode</code>	File's mode (permissions)	Not relevant
3	<code>nlink</code>	Number of links	Usually 0; Windows NT; file system might allow links
4	<code>uid</code>	User ID (UID) of the owner	Zero, always
5	<code>gid</code>	Group ID (GID) of the owner	Zero, always
6	<code>rdev</code>	Special file info	Drive number (again)
7	<code>size</code>	Size of file in bytes	Size of file in bytes
8	<code>atime</code>	Time of last access	Time of last access
9	<code>mtime</code>	Time of last modification	Time of last modification
10	<code>ctime</code>	Inode change time	File creation time
11	<code>blksz</code>	Disk block size	Zero, always

12	<code>blocks</code>	Number of blocks in file	Zero, always
----	---------------------	--------------------------	--------------

Many of the values in Table 10.3 you will probably never use, but they are presented for completeness. For the more obscure values—especially in the Unix return values—you might want to consult your operating system's reference manual.

The following is an example of using `stat` on a file:

```
@stuff=stat "myfile";
```

Normally, the returned values from `stat` are copied into an assignable list of scalars for clarity:

```
($dev, $ino, $mode, $nlink, $uid, $gid, $rdev, $size,
$atime, $mtime, $ctime, $blksize, $blocks)=stat("myfile");
```

To print the permissions for a file in the three-digit form described in the section "[A Crash Course in File Permissions](#)," you use this bit of code, where `@stuff` contains the permissions:

```
printf "%04o\n", $mode&0777;
```

The preceding snippet contains elements you might not understand. That's okay; some of it has not been presented to you yet. The permissions as they're retrieved by `stat`—in `$mode` in this case—contain lots of "extra" information. The `&0777` strips out just the portion of the information you're interested in here. Finally, `%o` is a `printf` format that prints numbers in octal—the 0–7 form in which Unix expects permissions to be formatted.

By the Way

Octal is a base-8 representation of numbers. It's used in Unix for largely historical reasons, but it came along with Perl anyway. If you're still lost, don't panic. Just use the `printf` shown previously if you ever need to display a file's permissions. It doesn't happen that often; don't worry if you don't quite get that part.

The three time stamps mentioned in [Table 10.3](#)—access, modification, and change (or create) time—are stored in a peculiar format. The time stamps are stored as the number of seconds from midnight, January 1, 1970, Greenwich mean time. To print them in a usable format, you use the `localtime` function, as follows:

```
print scalar localtime($mtime);
```

This function prints the modification time of the file in a format similar to `Sat Jul 3 23:35:11 EDT 1999`. The access time is the time that the file was last read (or opened for reading). The modification time is the time when the file was last written to. Under Unix, the "change" time is the time when the information about the file—the ownership, number of links, permissions, and so on—was changed; it is not the creation time of the file but often happens to be by coincidence. Under Microsoft Windows, the `ctime` field actually stores the time the file was created.

Sometimes you might be interested in retrieving just one value from the list returned by `stat`. To do so, you can wrap the entire `stat` function in parentheses and use subscripts to slice out the values you want from it:

```
print "The file has", (stat("file"))[7], " bytes of data";
```

 PREV

< Day Day Up >

NEXT 

Exercise: Renaming Files En Masse

This exercise will provide another (small) tool for your toolkit. This utility allows you to rename files given a directory name, a pattern to look for, and a pattern to change it to. For example, if a directory contains the filenames `Chapter_01.rtf`, `Chapter_02.rtf`, `Chapter_04.rtf`, and so on, you could rename all the files to `Hour_01.rtf`, `Hour_02.rtf`, `Hour_04.rtf`, and so on. This task normally isn't easy at a command prompt and just silly when you're using a GUI-based file browser.

Using your text editor, type the program from [Listing 10.3](#) and save it as `Renamer`. If you can, be sure to make the program executable according to the instructions you learned in [Hour 1](#).

When you're done, try running the program by typing the following at a command line:

```
perl -w Renamer
```

or, if you could make the program executable,

```
Renamer
```

[Listing 10.4](#) shows some sample output from this program.

Listing 10.3. Complete Listing for `Renamer`

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:
5:  my($dir, $oldpat, $newpat);
6:  print "Directory: ";
7:  chomp($dir=<STDIN>);
8:  print "Old pattern: ";
```

```

9:     chomp($oldpat=<STDIN>);
10:    print "New pattern: ";
11:    chomp($newpat=<STDIN>);
12:
13:    opendir(DH, $dir) || die "Cannot open $dir: $!";
14:    my @files=readdir DH;
15:    close(DH);
16:    my $oldname;
17:    foreach(@files) {
18:        $oldname=$_;
19:        s/$oldpat/$newpat/;
20:        next if (-e "$dir/$_");
21:        if (! rename "$dir/$oldname", "$dir/$_" ) {
22:            warn "Could not rename $oldname to $_: $!";
23:        } else {
24:            print "File $oldname renamed to $_\n";
25:        }
26:    }

```

Lines 13–15: The entries in the directory indicated by `$dir` are read into `@files`.

Lines 17–19: Each file from `@files` is assigned to `$_` and that name is saved in `$oldname`. The original filename in `$_` is then changed to the new name on line 19.

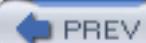
Line 20: Before renaming the file, this line makes sure the target filename doesn't already exist. Otherwise, the program could rename a file into an existing name—destroying the original data.

Lines 21–25: The file is renamed, and warnings are printed if the rename fails. Notice that the

original directory name needs to be appended onto the filenames—for example, `$dir/` `$oldname`—because `@files` doesn't contain full pathnames, which you need for the rename.

Listing 10.4. Sample Output from *Renamer*

```
1: Directory: /tmp  
2: Old Pattern: Chapter  
3: New Pattern: Hour  
4: File Chapter_02.rtf renamed to Hour_02.rtf  
5: File Chapter_10.rtf renamed to Hour_10.rtf
```

 PREV

< Day Day Up >

NEXT 

Summary

This hour you learned how to create, remove, and rename directory entries by using the `mkdir`, `rm`, and `rename` functions in Perl. Also, you learned how to query the file system with `stat` to find out information about files, not just what's in them. During the course of this hour, the two exercises provided small but useful tools that you can use to make yourself more productive.

Q&A

- Q1:** I'm having problems with the following program. No files are read in, even though they're in the directory!

```
opendir(DIRHANDLE, "/mydir") || die;  
  
@files=<DIRHANDLE> ;  
  
closedir(DIRHANDLE);
```

- A1:** Whoops! The problem is in the second line. `DIRHANDLE` is a directory handle, not a filehandle. You cannot read directory handles with the angle operators (`<>`). The correct way to read the directory is `@files=readdir DIRHANDLE`.
- Q2:** Why doesn't `glob("*.*)` match all the files in a directory?
- A2:** Because `"*.+"` matches only filenames with a dot in them. To match all the files in a directory, use `glob("*)`. The `glob` function's patterns are designed to be portable across many operating systems and thus do not behave the same as the `*.*` pattern in DOS.
- Q3:** I modified the `mygrep` exercise to search subdirectories by using `opendir` and some more loops, but it seems to have bugs.
- A3:** In short: Don't Do That. Descending a directory tree is an old problem, it's not particularly easy, it's been solved many times before, and you don't need to do it for yourself. (Doing all that work on a problem that's already been solved is called "reinventing the wheel.") If you're doing it only for fun, that's wonderful, but don't spend too much time on it. Hold on until Hour 15, where you'll learn to use the `File::Find` module instead. It's a lot simpler to use—and, more importantly, it's debugged.
- Q4:** The program in Listing 10.3 gives an error if I try to change `*.bat` to `*.tmp`. Why?

A4: The program wasn't expecting you to type `* .bat` as a pattern to look for. `* .bat` isn't valid to have in a regular expression—`*` must follow some other character. If you had entered `*\.bat`, the program would have accepted the input just fine—although it might not have worked as expected because filenames hardly ever have a literal `*` in them.

To fix this error, either give the program the input it expects (simple strings) or change line 19 to read `s/\Q$oldpat/$newpat/` so that "special characters" are disabled in the regular expression patterns.

 PREV

< Day Day Up >

NEXT 

Workshop

Quiz

1: To print the last modification time of the file `foofile`, use

a.

```
print glob("foofile");
```

b.

```
print (stat("foofile"))[9];
```

c.

```
print scalar localtime (stat("foofile"))[9];
```

2: The `unlink` function returns

a.

The number of files actually deleted

b.

True or false, depending on success

c.

The number of file deletions that were attempted

Answers

A1: b or c. Choice b prints the time expressed as the number of seconds since 1970—not very useful. Choice c prints the time as a nicely formatted string.

- A2:** a. However, choice b is also somewhat true. If no files can be deleted, `unlink` returns 0, which is false.

Activities

- As a programming exercise *only*, try to write a program that lists all the files in a directory, in its subdirectories, and so on.

 PREV

< Day Day Up >

NEXT 

Hour 11. System Interaction

What You'll Learn in This Hour:

- The system() function
- Capturing output
- Portability

All the Perl you've learned up until now has been fairly self-contained: If you wanted something done, you had to do it yourself—sorting data, creating directory listings, divining configuration information, and so on. The problem is that it's a lot of work, and you're repeating work that might have been done elsewhere.

One of the current buzzwords being touted about Perl is that it's an excellent *glue* language. What this means is that Perl can use other programs that are installed with your operating system as components, connecting them together to create larger programs. It can start your OS utilities and use them to gather information, interact with you, and then shut them down.

Perl can glue together these smaller utilities to produce a much larger and more useful utility. This capability has the added benefit of allowing you to quickly write code that might otherwise have taken a long time to write and debug. You should use any leverage at your advantage to write code quickly and accurately. Gluing system utilities together can offer a large advantage.

By the Way

The examples in this hour—for the most part—have two versions, one for Windows and MS-DOS systems and the other for Unix systems. Where only one example is provided, you'll find a note in the text about what to change for the other kind of system—and usually that change is minor.

The `system()` Function

The simplest way to run a command outside Perl is to use the `system` function. The `system` function pauses the Perl program, runs the external command, and then continues running your Perl program. The syntax for `system` is

```
system command;
```

where *command* is the command that you want run. The return value from `system` is 0 if everything goes okay or nonzero if a problem occurs. Notice that this result is *backward* from the normal Perl values of true and false.

The following is an example of `system` under Unix:

```
system("ls -lF");           # Print a file listing  
  
# print system's documentation  
  
if ( system("perldoc -f system") ) {  
  
    print "Your documentation isn't installed correctly!\n";  
  
}
```

This next example shows `system` under MS-DOS/Windows:

```
system("dir /w");           # Print a file listing  
  
# print system's documentation  
  
if ( system("perldoc -f system") ) {  
  
    print "Your documentation isn't installed correctly!\n";  
  
}
```

For the most part, `system` works the same under both architectures. The point to remember is that the commands are fundamentally different from one OS to another. To get a file listing in MS-DOS, you use the `dir` command; to get a file listing under Unix, you use the `ls` command. In very rare instances—`perldoc`, for example—the commands are similar under Unix and MS-DOS.

When the `system` function is running the external command, the output of that command is displayed on your screen as though your Perl program were printing it. If that external command needs input, the input comes from your terminal the same way your Perl programs read input from the terminal. The command run by `system` inherits the `STDIN` and `STDOUT` file descriptors, so the external command does I/O from the exact same place as your Perl program. Invoking fully interactive programs from `system` is possible.

Consider the following example under Unix:

```
$file="myfile.txt";  
  
system("vi $file");
```

Now consider this example under Windows/MS-DOS:

```
$file="myfile.txt";  
  
system("edit $file");
```

Each of the preceding examples runs an editor on `myfile.txt`—`vi` for Unix and `edit` for DOS. The editor runs full screen, of course, and all the normal editor commands work. When the editor exits, control returns to Perl.

You can use the `system` function to run any program, not just console-mode programs (text). Under Unix, this example runs a graphical clock:

```
system("xclock -update 1");
```

Under MS-DOS/Windows, the following example invokes a graphical file editor:

```
system( "notepad.exe myfile.txt" );
```

 PREV

NEXT 

Capturing Output

The `system` function does have a small shortcoming: It doesn't offer any particularly good way to capture the command's output and bring it into Perl for analysis. To do so in a roundabout way, you could use this workaround:

```
# 'ls' and 'dir' used for example only. opendir/readdir  
# would be more efficient in most cases.  
  
system("dir > outfile");      # Use "ls" instead of "dir" for Unix  
  
open(OF, "outfile") || die "Cannot open output: $!";  
  
@data=<OF>;  
  
close(OF);-
```

In the preceding snippet, the command run by `system` has its output redirected to a file called `outfile`. The file is then opened and read into an array. The array `@data` now contains the output of the `dir` command.

This method is messy and not too clever. Not surprisingly, Perl has another way of dealing with this situation: *backticks*, also called *backquotes*. Any command that is surrounded by backticks (` `) is run by Perl as an external command—as though through `system`—and the output is captured and returned as the return value from the backticks. Consider this example using backticks:

```
$directory=`dir`;      # Unix users, use ls instead of dir
```

In the preceding snippet, the `dir` command is run, and the output is captured in `$directory`.

Inside the backticks, all normal shell processing is observed: `>` does redirection, `|` does piping, and under Unix, `&` starts tasks in the background. Keep in mind, though, that commands that have been run in the background or that have had their output redirected with `>` have no output to capture.

In a scalar context, backticks return the output of the command as a single string. If the command output contains many lines of text, those lines all appear in the string, separated by record separator characters ("`\n`"). In a list context, the output is assigned to the list, with record separators at the end of each line.

Now consider this example:

```
@dir=`dir`; # Use 'ls' for Unix users  
  
foreach(@dir) {  
  
    # Process each line individually.  
  
}
```

In the preceding snippet, the output in `@dir` is processed in the `foreach` loop, one line at a time.

Perl has another way of representing backticks—that is, to use the notation `qx{ }`. The command you want to execute goes between the braces (`{ }`), as in this example:

```
$perldoc=qx{perldoc perl};
```

By using the `qx` operator, you can avoid the trouble of having to put backslashes in front of backticks when they appear as part of the command, as shown here:

```
$complex=`sort \`grep -l 'conf' *\``; # Somewhat messy
```

You can rewrite the preceding snippet as fo

 PREV

< Day Day Up >

NEXT 

Pipes

Pipes in Unix and MS-DOS/Windows are used for connecting different processes together so that the output of one process becomes the input of the next process. Consider the following set of commands, which would almost work in Unix (if you change the `dir` to `ls`) or MS-DOS:

```
dir > outfile
sort outfile > newfile
more newfile
```

The output of `dir` is collected in `outfile`. Then `sort` is used to sort `outfile`, and the output from `sort` is stored in `newfile`. Then `more` shows the contents of `newfile` a screenful at a time.

Pipes allow you to perform that same sequence, but without `outfile` and `newfile`, like this:

```
dir | sort | more
```

The output of `dir` is given to `sort`, which then sorts the data. The output from `sort` is then given to `more` for page-at-a-time display. No redirection (`>`) or temporary files are needed.

This kind of command line is called a *pipeline*, and the vertical line between the commands is called a *pipe*. Unix relies heavily on pipes to connect its small utilities. MS-DOS and Windows support pipes but have far fewer command-line utilities that work with them.

Perl programs can participate in a pipeline in different ways. First of all, if you have a Perl program that accepts standard input, transforms it, and sends it to standard output, then you can write a command line to insert that Perl program into a pipeline, as in the following example:

```
dir /B | sort | perl Totaler | more
```

In the preceding pipeline, `Totaler` could be a Perl program you write to print the sum of the directory listing and maybe some statistics, along with the directory listing itself. If you're using Unix, change the `dir /B` to `ls -1`, and the pipeline works as expected. [Listing 11.1](#) contains the `Totaler` program.

Listing 11.1. Complete Listing for `Totaler`

```
1:  #!/usr/bin/perl
2:
3:  use strict;
4:  my($dirs,$sizes,$total);
5:
```

```

6:  while(<STDIN>) {
7:      chomp;
8:      $total++;
9:      if (-d $_) {
10:          $dirs++;
11:          print "$_\n";
12:          next;
13:      }
14:      $sizes+=(stat($_))[7];
15:      print "$_\n";
16:  }
17:  print "$total files, $dirs directories\n";
18:  print "Average file size: ", $sizes/($total-$dirs), "\n";

```

Line 6: Each line of input is read from `STDIN` and assigned to `$_`. On a pipeline, a program's `STDIN` is connected to the previous program's `STDOUT`. So, in the example given, `STDIN` is being fed by `dir /B` through `sort`.

Lines 9?3: If a directory is encountered, its number is totaled separately in `$dirs`, the directory name is printed, and the loop is started again.

Lines 14?5: Otherwise, the sizes of the files are accumulated in `$sizes`, and the filenames are printed.

Lines 17?8: The average size of the file is printed, along with the total number of files and directories.

The other way Perl can participate in a pipeline is to treat a pipeline like a file that can either be read from or written to. This is done with the `open` function in Perl, as shown here:

```
# Replace "dir /B" with "ls -1" for Unix
open(RHANDLE, "dir /B| sort |") || die "Cannot open pipe for reading: $!";
```

In the preceding snippet of code, the `open` function opens a pipeline for reading from `dir/B | sort`. The fact that Perl is reading from this pipeline is indicated by having the final pipe (`|`) on the right. When the `open` function is run, Perl starts the `dir /B | sort` commands. When the filehandle `RHANDLE` is read, the output from `sort` is read into the Perl program.

Now consider this example:

```
open(WHANDLE, " | more") || die "Cannot open pipe for writing: $!";
```

This `open` function opens a pipeline for writing to the `more` command. The pipe symbol on the left means

that Perl is going to write to the pipe. All printing to the `WHANDLE` filehandle is buffered by `more` and displayed a page at a time. Writing the function like this might be a good way to get your program's output displayed a page at a time.

When you are done with a filehandle that has been opened to a program—like `RHANDLE` and `WHANDLE`—it is very important that you close the filehandle properly. The reason is that the programs started by `open` must be properly shut down, and using `close` on the filehandle ensures that. Failing to close the filehandle when you're done with it could result in the programs continuing to run even after your Perl program has ended.

When closing a filehandle that's been opened on a pipe, the `close` function indicates whether the pipeline was successful. Therefore, you should be careful to check the return value of `close` like this:

```
close(WHANDLE) || warn "pipe to more failed: $!" ;
```

By the Way

The reason that the `open` function might not tell you whether the pipeline was successfully started has to do with Unix's design. When Perl constructs the pipeline and starts it, it's not sure that the pipeline will actually work; if the pipeline is assembled properly and starts, it's assumed that it will finish properly. When the last program in the pipe completes, it should return a successful exit status. The `close` function can read that status to tell whether everything went all right; otherwise, an error results.

First Lesson in Portability

Portability: It's one of the things that Perl is good at. Whether your Perl code is run on a VMS machine, Unix, a Macintosh, or under MS-DOS, there's a very high probability that the Perl code you write will work seamlessly under any architecture that Perl supports. When you need to interact with the underlying operating system, such as when you're doing file I/O, Perl tries to hide all the nitty-gritty details for you so that your code will just work.

By the Way

Some of the reasons that Perl is so portable are discussed at great length in [Hour 16](#), "The Perl Community."

Sometimes, though, there's a limit to what Perl can hide from you.

Throughout this hour, the examples have read "do this for Windows and MS-DOS, do this for Unix," and either one has worked, depending on the architecture you're using. Trying to accommodate both Windows and Unix with your own programs means that you'll have to create two versions of each program—one for Windows and one for Unix. Developing two versions creates further problems when your program is successful and moved to an even more alien architecture like Mac OS 9.

It's not at all uncommon for a program to be written for one architecture—such as Windows NT—to find itself being run under a different architecture—such as Unix. Because Perl runs on so many different architectures, many people assume that running a Perl program under Windows NT is the same as running it under Unix. Web servers and other applications move frequently between architectures; it's just good business sense to keep your software portable.

Creating a different version of your program for each architecture that would work under every possible situation is time consuming, wasteful, and unproductive. By following a few rules, you can create programs that will work everywhere, or at least try to work everywhere, and that are easy to fix if they do not.

The following are some general rules for writing "go anywhere" code:

1.

Always have warnings turned on, and use the `use strict` directive. This way, you can make sure that your code will probably run with various versions of Perl and that there aren't any obvious bugs.

2.

Always check the return value from system requests—for example, use `open || die`, never just `open`. Checking the value will help you find errors when moving your application from one server to

another—not just between architectures.

3.

Print good, descriptive error messages.

4.

Use Perl's built-in functions to perform tasks you might otherwise do by using `system` or backticks (`` ``).

5.

Wrap system-dependent operations (file I/O, terminal I/O, process control, and so on) in subroutines, and check to make sure those operations are supported on the current operating system.

The first two points you should already be familiar with. All throughout this book, the examples have checked the exit status of critical functions, and since [Hour 8](#), "Functions," all larger examples have demonstrated `use strict` and warnings.

Point 3, having good error messages, cannot be overlooked. Of the following messages, which is the most helpful?

(no message, or wrong output)

Died at line 15.

Cannot open Foofile.txt: No such file or directory

Cannot open Foofile.txt: No such file or directory at myscript.pl line 24

Obviously, the last one is the most helpful. After you've installed the program, and a problem arises months (or years) later, the last message indicates what program failed ("myscript.pl"), what it wanted ("Foofile.txt"), why it failed ("No such file..."), and where it failed ("line 24"). This information will help you fix the problems quickly. A little bit of time spent writing a good, descriptive error message always pays off.

Point 4 simply means that you should use Perl whenever possible. To retrieve a directory listing, it's tempting just to use `$dir=`dir``; but this will fail if the program is ever moved to a non-Windows system. A good solution would be to use `<*>`. A better solution would be to use the `opendir/readdir/closedir` functions whenever possible. These solutions will work no matter where your program is moved.

Telling the Difference: An Example

The last two points for writing "go anywhere" code—wrapping system-dependent things in subroutines and checking the architecture of the machine that the program is running on—bear a little more explanation and a demonstration.

As you sit typing your Perl programs, you should remember that one day your Perl program might be used on a machine other than the one you're currently using. You may develop the next Amazon.com Web site, and it may move from your PC to a large Windows NT server to a cluster of Sun Enterprise 10000 Unix servers. Or you may simply have personal CGI programs and change Web providers only to find that your new provider has a different kind of server. These situations happen all of the time and need consideration.

So how does your program know the difference between Windows NT and Unix? Simple. The Perl special variable `$^O`—that's a dollar sign, a caret, and the capital letter *O*—contains the architecture name that the program is running under. For example, under Windows and MS-DOS, it contains the string `MSWin32`. Under Unix, it contains the type of Unix you're running—`linux`, `aix`, `solaris`, and so on.

The following are some tasks that depend on what operating system you are running:

- Finding out anything about the system's configuration
- Working with the disk and directory structure
- Using system services (such as email)

For this example, you're going to examine a piece of code to find the amount of available disk space on a system. This exercise might be useful if someone wants to upload a file to a server and you need to find out whether the file would fit first. A code snippet to find the free disk space in the current drive of a Windows system might look like the following:

```
# The last line of 'dir' reports something like:  
  
#      10 dir(s)    67,502,080 bytes free  
  
# Or on Win98, "MB" instead of "bytes"  
  
my(@dir,$free);  
  
@dir= `dir`;  
  
$free=$dir[$#dir];  
  
$free=~s/.*( [\d,]+ ) \w+ free/\$1/;  
  
$free=~s/,//g;
```

The preceding snippet takes the last line of the directory listing—in `@dir`—and uses regular expressions to remove everything but the size—the digits and commas preceding `bytes free`. Finally, the commas are removed so that `$free` contains just the raw free disk space. This approach works well for Windows

systems. For Unix—Linux, in particular—the following snippet works:

```
# Last lines of df -k . reports something like this:  
  
# Filesystem      1024-blocks  Used   Available Capacity Mounted on  
  
# /dev/hda1        938485       709863  180139    80%      /  
  
# And the 4th field is the number of free 1024K disk blocks  
  
# This format may be particular to Linux.  
  
my(@dir, $free);  
  
@dir=`df -k .`;  
  
$free=(split(/\s+/, $dir[$#dir]))[3];  
  
$free=$free*1024;
```

Notice the differences between this snippet and the previous one. The utility under Windows to find the disk space is `dir`; under Unix, it is `df -k ..`. The last line of output of `df -k ..` is split apart, and the fourth field is placed in `$field`. Output of `df` varies slightly between Unix systems; usually, the number of fields reported is different, or they're in a different order. Your Perl code can easily compensate by simply picking a different field.

So now you have two completely different routines to find free disk space. You can combine them and have the appropriate one run on each architecture, as follows:

```
if ( $^O eq 'MSWin32' ) {  
  
    # The last line of 'dir' reports something like:  
  
    #      10 dir(s)    67,502,080 bytes free  
  
    my(@dir,$free);  
  
    @dir=`dir`;  
  
    $free=$dir[$#dir];  
  
    $free=~s/.*( [\d, ]+) \w+ free/$1/;  
  
    $free=~s/,//g;
```

```

} elsif ($^O eq 'linux') {

    # Last line of df -k . reports something like this:

    # /dev/hda1      938485 709863 180139  80%  /

    # And the 4th field is the number of free 1024K disk blocks

    my(@dir, $free);

    @dir=`df -k .`;

    $free=(split(/\s+/, $dir[$#dir]))[3];

    $free=$free*1024;

} else {

    warn "Cannot determine free space on this machine\n";

}

```

The sample program has now been expanded to include both the MS-DOS/Windows version and the Linux version. If it's run under any other kind of machine, a warning is printed.

The routine is almost finished. What you need to do now is isolate this routine in a subroutine so that the variables needed can be declared private and the final product can be cut and pasted into any program and used whenever needed.

```

# Computes free space in current directory

sub freespace {

    my(@dir, $free);

    if ( $^O eq 'MSWin32' ) {

        # The last line of 'dir' reports something like:

        #      10 dir(s)   67,502,080 bytes free

        @dir=`dir`;

        $free=$dir[$#dir];

        $free=~s/.*(\d,)+ bytes free/\$1/;

```

```

$free=~s/,//g;

} elsif ($^O eq 'linux') {

    # Last line of df -k . reports something like this:

    # /dev/hda1      938485 709863 180139  80%  /

    # And the 4th field is the number of free 1024K disk blocks

    @dir=`df -k .`;

    $free=(split(/\s+/, $dir[$#dir]))[3];

    $free=$free*1024;

} else {

    $free=0; # A default value

    warn "Cannot determine free space on this machine\n";

}

return $free;

}

```

Now whenever your programs need to find the amount of free disk space, you can simply call the `freespace()` subroutine, and the answer is returned. If you try running this subroutine on another architecture that's not listed, an error message is printed. However, adding another Unix-like OS to the function wouldn't be difficult; you can just add another `elsif` clause.

Summary

This hour you learned about how to leverage your system's utilities to do work for you. The `system` function allows you to run a system utility (or a pipeline). The backticks (`` ``) run a system utility and then capture the output. The captured output can then be stored in a variable for use by Perl. The `open` function can open not only files but programs as well. The programs can be written to with `print` or read from with the angle operator (`<>`). Finally, you were shown some techniques for using these utilities on many different kinds of systems without having to write different programs for each.

Q&A

Q1: How do I open a pipe both to and from a command? For example, `open(P, "| cmd |")` doesn't seem to work.

A1: This task can actually be quite complicated because reading and writing from the same process can cause *deadlock*. At this point, your program is expecting `cmd` to print something and is waiting for data with `<P>`. Meanwhile, because of some snafu, `cmd` is actually waiting for your program to print something with `print P "...."`. In fact, if you have warnings enabled, Perl will inform you with this message: `Can't do bidirectional pipe.`

If you're prepared for this kind of problem, the `IPC::Open2` module will allow you to open a bidirectional pipe. Modules will be discussed in Hour 14, "Using Modules."

Q2: The code `$a=system("cmd")` didn't capture the output of `cmd` in `$a` as I expected. Why not?

A2: You're confusing `system` with backticks (`` ``). The `system` function doesn't capture `cmd`'s output. What you probably wanted was `$a=`cmd``.

Q3: When I run external programs under Unix with backticks (`` ``), the error messages are not captured. Why not?

A3: Because all Unix programs—including Perl—have two output file descriptors: `STDOUT` and `STDERR`. The `STDOUT` file descriptor is for normal program output. The `STDERR` file descriptor is reserved for error messages. Backticks—and `open` with pipes—capture only `STDOUT`. The short answer is to use the shell to redirect `STDOUT` into `STDERR` and then run your command as follows:

```
$a=`cmd 2>&1`;      # run "cmd", capturing output and errors
```

The Perl FAQ has a lengthy explanation of this and many other techniques for capturing a command's errors. Type `perldoc perlfaq8` for the appropriate section of the FAQ.

Workshop

Quiz

1: To have data generated by your program display a page at a time, use

a.

```
perl myprog.pl | more
```

b.

```
open(M, " | more") || die; print M "data...data...
data....\n";
```

c.

```
open(M, ">more") || die; print M "data..data...data...
\n";
```

2: Which value of \$foo is used in this statement: \$r=`dir \$foo`?

a.

The shell's value for \$foo.

b.

Perl's value of \$foo is substituted, and then the dir is run.

3: Which of the following tasks changes depending on what system you have?

a.

Finding the amount of free disk space

b.

Getting a directory listing

c.

Deleting a directory

Answers

A1: Either a or b. In the case of a, all output of myprog.pl is fed to `more`. In the case of b, any data written to the `M` file descriptor is fed to `more` for paging.

A2: b. To protect `$foo` from Perl's expansion, you use `qx`dir $foo``.

A3: a only. Choice b can be done with `glob, <*>` or `opendir` and `readdir`, and c is done with `rmdir`.

Activities

- Use the statistics functions in Hour 8, "Functions," to display even more statistics about the sizes of files in Listing 11.1.
- If you have a Unix system, add your particular flavor of Unix to the `freespace()` function. Use the Linux example as a starting point.

Hour 12. Using Perl's Command-Line Tools

What You'll Learn in This Hour:

- How to use Perl's debugger
- How to use command-line switches to write small Perl programs

Until now, the Perl interpreter has been a fairly simple program to use. You type a program into a file and then invoke the Perl interpreter to run your program. The Perl interpreter is much more flexible than this, however.

Built into the interpreter is a debugger. The debugger allows you to play your Perl program like a videotape. You can rewind your program to the beginning, play it slow, play it fast, and freeze-frame it to examine the innards closely. The debugger is an often underutilized tool for finding problems in Perl programs.

Perl can also run programs that are not typed into files. For example, you can run small programs directly from your system's command prompt.

What Is the Debugger?

The Perl debugger is a built-in feature of the Perl interpreter. It allows you to take any Perl program and step through that program statement by statement. Along the way, you can examine variables, change them, let the program run for a while longer, interrupt the program, or start it over again.

From your program's perspective, nothing is different. Input still comes from the keyboard, and output still goes to the screen. The program doesn't know when it's stopped or when it's running. In fact, you can examine your program's workings without disturbing them at all.

Starting the Debugger

If you are a Macintosh user running Perl, to use the debugger you need only select Debugger from the Script menu, and a Debugger window opens with a prompt for you.

Under any other operating system, to start the Perl debugger you must be at your operating system's command prompt. For DOS and Windows users, this means the standard MS-DOS `C:\>` prompt. For Unix users, this means the prompt that you were presented with when you logged in (usually `%` or `$`).

All the examples in this section use the Employee program from [Listing 9.1](#) in [Hour 9](#), "More Functions and Operators." You might find it handy to put a bookmark on that page and flip back and forth for reference.

To start the debugger at the prompt—an MS-DOS prompt is used for the example—type this line:

```
C:\> perl -d Employee
```

The `-d` switch to `perl` causes Perl to start up in debugging mode; the program to be debugged is also indicated on the command line. Some messages giving version information are then displayed, as shown here:

```
Loading DB routines from perl5db.pl version 1.0401
```

```
Emacs support available.
```

```
Enter h or `h h' for help.
```

```
main:::(Employee:5):      my @employees=(  
main:::(Employee:6):          'Smith,Bob,123101,9.35,40',
```

```
main::(Employee:7):      'Franklin,Alice,132912,10.15,35',
main::(Employee:8):      'Wojohowicz,Ted,198131,6.50,39',
main::(Employee:9):      'Ng,Wendy,141512,9.50,40',
main::(Employee:10):     'Cliburn,Stan,131211,11.25,40',
main::(Employee:11):    );
DB<1>__
```

The debugger first displays the version number (`1.0401`—yours will vary) and the help prompt. Next, the first executable line of the program is displayed. Because the first statement is actually seven lines long—starting with `my @employees=()` and ending with `;`—all seven lines of the statement are shown along with a description that shows what file they came from (`Employee`) and what line or lines of the file they were found on (5 through 11).

Last, you see the debugger prompt `DB<1>`. The `1` signifies that the debugger is waiting for its first command. The cursor waits at the debugger prompt for your command.

At this point, your Perl program is not running. The program is actually paused just before the first instruction `-my @employees=()`, and so on. Whenever the debugger shows you a statement from your program, it is the statement *about to be executed*, not the last statement run.

The debugger is now ready for your commands.

Basic Debugger Commands

The first—and most important—command you can give the debugger is the help command. If you type `h` at the debugger prompt, all the available debugger commands are printed. You also can use some variations: `h h` prints a summary of commands and syntax, and `h cmd` prints help for a specific command.

The list of help commands is probably longer than your screen will allow, and the first few commands will probably scroll right off. To make any debugger command's output display a screen at a time, put a `|` in front of the command. So, to see the help a screen at a time, the command is `|h`.

The most-used feature of the debugger is the capability to run Perl code an instruction at a time. So, continuing with the previous example, to go to the next statement in your Perl program, you use the debugger's `n` command:

```
main::(Employee:5):      my @employees=
main::(Employee:6):      'Smith,Bob,123101,9,35,40',
main::(Employee:7):      'Franklin,Alice,132912,10,15,35',
main::(Employee:8):      'Wojohowicz,Ted,198131,6,50,39',
main::(Employee:9):      'Ng,Wendy,141512,9,50,40',
main::(Employee:10):     'Cliburn,Stan,131211,11,25,40',
main::(Employee:11):    );
DB<1> n                  ← Command for "next" instruction
main::(Employee:24):     my ($L1, $F1)=split(',',$a):
DB<1>
```

After you type the `n` command, Perl executes the statement from lines 5 to 11 of the Employee program. The debugger then prints, but does not yet run, the next statement to be executed—`my($L1, $F1)=split(',',$a);`—and displays another prompt.

At this point in the execution, `@employees` is initialized to the five names, salaries, and so on. To view them, you can simply print them:

```
DB<1> print @employees
```

```
Smith,Bob,123101,9.35,40Franklin,Alice,132912,10.15,35Wojehowicz,Ted,198131,6.50  
,39Ng,Wendy,141512,9.50,40Cliburn,Stan,131211,11.25,40
```

```
DB<2> __
```

In fact, any Perl statement can be run at a debugger prompt. Notice that the array elements from `@employees` are all run together. You can enter the following to print them nicely:

```
DB<2> print join("\n", @employees)
```

```
Smith,Bob,123101,9.35,40
```

```
Franklin,Alice,132912,10.15,35
```

```
Wojehowicz,Ted,198131,6.50,39
```

```
Ng,Wendy,141512,9.50,40
```

```
Cliburn,Stan,131211,11.25,40
```

```
DB<3> __
```

To continue stepping through the program, just keep typing `n`, as shown here:

```
DB<3> n
```

```
main::(Employee:23):      @employees=sort {
```

```
DB<3> n
```

```
main::(Employee:25):      my ($L2, $F2)=split(',',$b);
```

```
DB<3> n
```

```

main::(Employee:26):           return( $L1 cmp $L2
main::(Employee:27):           || )
main::(Employee:28):           $F1 cmp $F2
main::(Employee:29):           );
DB<3> n
main::(Employee:23): @employees=sort {
DB<3>

```

Apparently, the debugger moves backward through the program here; line 23 is about to be executed again. The Perl `sort` statement is actually a kind of loop, and the debugger steps through each statement in the block of the `sort`. If you keep typing `n`, the debugger keeps looping until the sort is finished—which might take awhile.

To repeat the previous command, you can also just press the Enter key at a debugger prompt.

Breakpoints

Instead of stepping through the program an instruction at a time, you also can have the debugger continue to run your Perl program until a certain statement is reached and then stop. These places to stop are called *breakpoints*.

To set a breakpoint, you must first select a place in the program to stop. The `l` command lists the next 10 lines of the program. Typing `l` again lists the next 10 lines, and so on. To list the program starting at a particular line, type `l lineno`, where `lineno` is the line number of the program. You can also specify a range of lines by using the command `l start-end`.

In the listing, a `==>` marks the current line about to be executed by the debugger, as you can see here:

```

DB<3> l 11
23:==> @employees=sort {
24:           my ($L1, $F1)=split(',',$a);
25:           my ($L2, $F2)=split(',',$b);
26:           return( $L1 cmp $L2
27:           || )
28:           $F1 cmp $F2
29:       );

```

```
30: } @employees;  
31:  
32: foreach(@employees) {  
33:     print_emp($_);  
  
DB<3> __
```

In this case, line 33 is a good place for a breakpoint: It's after the `sort` statement, and it's the first statement inside the main loop of the program. You can set a breakpoint anywhere in a Perl program as long as the breakpoint is a valid Perl statement; you cannot break on a brace (line 30), punctuation (line 29), a blank line (line 31), or a line containing only a comment.

To set the breakpoint, use the `b breakpoint` command, where *breakpoint* can be a line number or a subroutine name. To set the breakpoint at line 33, for example, you enter the following:

```
DB<3> b 33  
  
DB<4>
```

The other command you need to know with breakpoints is the continue command, `c`. The `c` command instructs the debugger to let the Perl program run until the next breakpoint is reached or the end of the program is reached—whichever comes first:

```
DB<5> c  
  
main::(Employee:33):          print_emp($_);  
  
DB<6>
```

In this case, the debugger stops the Perl program at line 33 before the `print_emp` function is called, as expected. The breakpoint is still set, so typing another `c` causes the program to continue, run the `print_emp()` function, and stop again on line 33:

```
DB<8> l  
  
Employee:  
33:     print_emp($_);
```

To examine the breakpoints you have set in your program, you can use the `L` command like this:

```
DB<7> C
131211 Stan Cliburn 11.25 40 450.00
main::(Employee : 33): print_emp($_);
          ^ output of print_emp()
          ^ next line to be run
DB<8>
```

This example shows that the debugger has one breakpoint, in the file Employee at line 33.

To get rid of a breakpoint in the program, you use the `d` command in the same way you set the breakpoint—`d line` or `d subname`:

```
DB<9> d 33
```

```
DB<10>
```

Other Debugger Commands

If you want to examine the workings of the `print_emp()` function, you can do so in a few different ways. First, restart the program by using the `R` command:

```
DB<11> R
```

```
Warning: some settings and command-line options may be lost!
```

```
Loading DB routines from perl5db.pl version 1.0401
```

```
Emacs support available.
```

```
Enter h or `h h' for help.
```

```
main::(Employee:5):    my @employees=
main::(Employee:6):        'Smith,Bob,123101,9.35,40',
```

```
main::(Employee:7):      'Franklin,Alice,132912,10.15,35',  
main::(Employee:8):      'Wojohowicz,Ted,198131,6.50,39',  
main::(Employee:9):      'Ng,Wendy,141512,9.50,40',  
main::(Employee:10):     'Cliburn,Stan,131211,11.25,40',  
main::(Employee:11):    );
```

```
DB<12> b 33
```

The `R` command resets the Perl program to the beginning and prepares to execute it again. Any breakpoints you have set remain set, and any variables in the Perl program are reset. In the preceding example, a breakpoint is set at line 33. Then you can continue the program with the following:

```
DB<13> c
```

```
main::(Employee:33):      print_emp($_);  
DB<14> __
```

Executing the `n` command would execute the next instruction:

```
DB<14> n
```

```
131211 Stan Cliburn      11.25  40  450.00
```

```
main::(Employee:32):      foreach(@employees) {
```

```
DB<15> n
```

```
main::(Employee:33):      print_emp($_);
```

```
DB<16> __
```

Moving through the program this way doesn't allow you to examine what's *in* `print_emp()`. To step into `print_emp()`, instead of the `n` command, you should use `s`, the step command. The `s` command works just like `n`, but instead of simply executing subroutines and moving to the next instruction, `s` stops at the first instruction *inside* the subroutine, as you can see here:

```
main::(Employee:33):      print_emp($_);
```

DB<16> s

```
main::print_emp(Employee:14):      my($last,$first,$emp,$hourly,$time)=  
main::print_emp(Employee:15):          split(',',$_[0]);
```

DB<17>__

Here, the first statement of



Exercise: Finding the Bug

This exercise shows you how to use the debugger. The program in [Listing 12.1](#) has a problem—actually two. It's supposed to print these messages:

20 glasses of Lemonade on the wall

19 glasses of Lemonade on the wall

:

1 glass of Lemonade on the wall

0 glasses of Lemonade on the wall

But it does not. Your task is to type the program in [Listing 12.1](#) and try to find the bugs. Neither of the bugs is a syntax problem—Perl's warnings are not triggered, and `use strict` prints no messages—but the debugger should make the bugs fairly obvious to find.

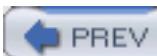
After you type the program, run Perl with the debugger to try to find the errors. Remember to print relevant variables and expressions occasionally and step through the function calls one at a time.

Listing 12.1. Buggy

```
1:  #!/usr/bin/perl -w
2:  # This program contains TWO errors
3:  use strict;
4:
5:  sub message {
6:      my($quant)=@_;
7:      my $mess;
8:      $mess="$quant glasses of Lemonade on the wall\n";
9:      if ($quant eq 1) {
```

```
10:         $mess=s/glasses/glass/;  
11:     }  
12:     print $mess;  
13: }  
14:  
15: foreach(20..0) {  
16:     &message($_);  
17: }
```

The solution is presented in the "[Quiz](#)" section at the end of this hour.

 PREV

< Day Day Up >

NEXT 

Other Command-Line Stuff

The debugger isn't the only feature of the Perl interpreter that can be activated by command-line switches. In fact, many useful Perl programs can be written just at the command prompt.

By the Way

Macintosh users should run these command-line examples by selecting 1-liners from the Script menu and then typing the command in the dialog box.

One-Liners

The key to such programs is the `-e` switch given to Perl on the command line. Following the `-e` can be any Perl statements at all, as in this example:

```
C:\> perl -e "print 'Hello, world';"
```

```
Hello, world
```

You can use multiple `-e` statements to insert multiple statements or separate them with semicolons, as shown here:

```
C:\> perl -e "print 'Hello, world';" -e "print 'Howzit goin?'"
```

```
Hello, worldHowzit goin?
```

A word of caution: Most command interpreters have rules about quotation marks. The Windows/MS-DOS command interpreter—command.com or the NT command shell—allows you to use double quotes to group words—such as `print` and `Hello, World` in the preceding example—but you cannot easily put double quotes inside other double quotes. Nor can you easily put `>`, `<`, `|`, or `^` inside the double quotes. Consult your operating system's manual for a complete list of MS-DOS/Windows quoting rules.

Under Unix, generally as long as the quotes are balanced—every open quote has a close quote—and embedded quotation marks have a `\` in front of them, you should be okay:

```
$ perl -e 'print "Hello, World\n";' -e 'print "Howzit goin?\n"'
```

The preceding example works under most Unix shells—`csh`, `ksh`, `bash`, and so on—and prints the message with the correct newlines. For a complete list of your shell's quoting rules, consult your shell's online manual page.

One frequent—and very useful—use of the `-e` switch is to combine it with `-d` and drop into Perl's debugger directly—without having a program to debug:

```
C:\> perl -d -e 1
```

```
Loading DB routines from perl5db.pl version 1
```

```
Emacs support available.
```

```
Enter h or `h h' for help.
```

```
main:::(-e:1):    1
```

```
DB<1> __
```

The debugger is now waiting for your commands. This particular incantation is useful for testing Perl syntax without your actually having to write a whole program, test, debug, edit, test, debug, and so on. Simply try your statements in the debugger until they work. The `1` on the command line is a minimalist Perl program; it's an expression that evaluates to—and returns—1.

Other Switches

The `-c` switch to the Perl interpreter causes Perl to examine your code for syntax problems—but not actually run the program:

```
C:\> perl -c Example.pl
```

```
Example.pl syntax OK
```

If a syntax error occurs, Perl may print a message like this:

```
C:\> perl -c Example.pl
```

```
syntax error at Example.pl line 5, near "print"
```

```
Example.pl had compilation errors
```

Combined with `-w`, the `-c` switch compiles your program and shows any warning messages that Perl thinks is appropriate.

When you're asking a more knowledgeable Perl user or your system administrator about debugging code, you are often requested to supply the version of the interpreter being used. The version of the Perl language that's in use primarily is Perl 5. The interpreter itself has a version that you can gather by using the `-v` switch on the command line, as shown here:

```
C:\> perl -v
```

```
This is perl, version 5.004_02
```

```
Copyright 1987-1997, Larry Wall
```

```
Perl may be copied only under the terms of either the Artistic License or the  
GNU General Public License, which may be found in the Perl 5.0 source kit.
```

In the preceding example, the version of the Perl interpreter is 5.004_02. To get even more detailed information about how the interpreter was built, when it was built, and so on, you can run the interpreter with the `-V` switch, like this:

```
C:\> perl -V
```

```
Summary of my perl5 (5.0 patchlevel 4 subversion 02) configuration:
```

```
Platform:
```

```
osname=MSWin32, osvers=4.0, archname=MSWin32
```

```
:
```

```
Compiler:
```

```
cc='bcc32', optimize='-O', gccversion=
:
:
```

Characteristics of this binary (from libperl):

Compile-time options: DEBUGGING

Built under MSWin32

Compiled at Aug 9 1997 21:42:37

@INC: C:\PERL\lib\site C:\PERL\lib

c:\perl\lib c:\perl\lib\site c:\perl\lib\site .

This output might be useful if you're trying to debug a problem with the interpreter itself—perhaps a problem with the installation. At the end, note the values for `@INC`. This particular installation of Perl expects to find its modules in these directories. (Modules will be discussed in Hour 14, "Using Modules.") After Perl is installed, it cannot simply be moved from one directory to another. The interpreter itself has a built-in idea of where to find its modules, and moving it causes Perl to look in the wrong place for its modules.

Empty Angle Brackets and More One-Liners

The angle operator (`<>`) discussed so far have had a couple of functions:

1.

With a filehandle, the angle operators allow you to read the filehandle—for example, `<STDIN>`.

2.

With a pattern, the angle operators return a list of files matching that pattern, called a glob—for example, `<*.bat>`.

The angle operators have yet another function: A set of angle operators with nothing between them reads all the contents of all the files on the command line or, if no filenames are given, from the standard input. Sometimes an empty angle operator is called a *diamond operator* (the name comes from its shape). For example, examine this small Perl program:

```
#!/usr/bin/perl -w
```

```
while(<>) {
```

```
print $_;
```

```
}
```

If you save the preceding program as `Example.pl`, then running the program with the command line

```
C:\> perl -w Example.pl file1 file2 file3
```

would cause the `<>` to read the contents of `file1` one line at a time, then `file2`, and then `file3`. If no files are specified, the angle operators read from the `STDIN` filehandle. This behavior mimics Unix utilities—such as `sed`, `awk`, and so on—that read input from files if given on the command line, and otherwise use standard input.

By the Way

After the Perl interpreter switches (`-w`, `-c`, `-d`, `-e`, and so on) are stripped away, the arguments to the Perl program are stored in an array called `@ARGV`. For example, for the previous snippet's arguments, `$ARGV[0]` would contain `file1`, `$ARGV[1]` would contain `file2`, and so on.

The `-n` switch to Perl wraps up any following statements on the command line into this small program:

```
LINE:
```

```
while(<>) {  
    ... # Command-line Perl statements here.  
}
```

So, to create a one-liner to remove leading spaces from input, you could write the following:

```
C:\> perl -n -e 's/^s+//g; print $_[0];' file1
```

The preceding command actually runs a Perl program that looks like this:

```
LINE:
```

```
while(<>) {  
    s/^\\s++//g;  
  
    print $_;  
}
```

In the preceding snippet, the file named file1 is opened and assigned to `$_` in the `while` loop, one line at a time. The line is edited with `s/^\\s++//g` and then printed.

The `-p` switch is identical to `-n` except that lines are printed automatically after your statements are executed. So, rewriting the previous example would yield the following line:

```
C:\> perl -p -e 's/^\\s++//g' file1
```

When you're editing a file with a Perl one-liner, you must be careful not to open the file for reading while trying to write to it at the same time, as in this example:

```
C:\>perl -p -e 's/\\r//g' dosfile > dosfile
```

The preceding snippet tries to remove carriage returns from a file called dosfile. The problem is that dosfile is overwritten by `> dosfile` before the Perl command is even processed. The correct way to edit a file is to redirect the output to another file and rename it to the original name, like this:

```
C:\>perl -p -e 's/\\r//g' dosfile > tempfile
```

```
C:\>rename tempfile dosfile
```

By the Way

Writing short "one-liners" is considered a pastime for some Perl enthusiasts; the more convoluted and capable the program, the better. *The Perl Journal*, a quarterly magazine for Perl, scatters one-liners throughout each issue.

 PREV

< Day Day Up >

NEXT 

Summary

In this hour, you learned how to use the debugger effectively to find problems in your Perl programs. You learned that the angle operators (`<>`) have yet another mode that allows Perl to process all the files on the command line. Also, you learned how to use the `-n` and `-p` switches to the Perl interpreter to write small, one-line Perl programs.

Q&A

Q1: *I'd really prefer a graphical debugger for Perl. Does such a thing exist?*

A1: Yes, there are a few. If you're using Perl under Windows, ActiveState has a very nice graphical debugger.

Q2: *What's the `main::` stuff that the debugger keeps printing?*

A2: It has to do with Perl's package naming conventions. Some of this information will be covered in the next hour, so don't worry too much about it for now.

Q3: *Does Perl have other command-line switches?*

A3: Yes, several. You can view the full list in the online manual. To access this information, type `perldoc perlrun` at a command prompt.

Workshop

Quiz

- 1:** What were the bugs in the Listing 12.1?
- 2:** If no files are given on the command line, reading `<>` returns
 - a.
`undef`
 - b.
lines from standard input
 - c.
True
- 3:** The Perl debugger can print Perl statements as they execute. This is called trace mode. How do you put the debugger into trace mode? (Hint: You need to look at the debugger's help message for this answer.)
 - a.
Use the `T` command, for trace
 - b.
Use the `t` command, for trace

Answers

A1: First, in line 15 the range `(20..0)` is not valid. The range operator—`..`—does not count down, only up. This line should be changed to a `for ($_=20; $_[> -1; $_--)` loop, `reverse(0..20)` or something similar. Second, at line 10 the `$mess=s/glasses/glass/` looks like a substitution on `$mess`, but it's not. The substitution is actually getting performed on `$_` because the assignment (`=`) should actually be a bind (`=~`).

A2: b. If no filenames are given, `<>` begins reading STDIN.

A3: b. The `t` command prints all your program's statements as they execute. The `T` command prints a *stack trace*, which is a listing of what function is currently being executed, the function that called that function, the function that called that function, and so on.

 PREV

< Day Day Up >

NEXT 

Hour 13. References and Structures

What You'll Learn in This Hour:

- The basics of references
- Common structures of references
- A brief example to help all this make sense

If Perl is your first programming language, you'll find this hour especially interesting. In most programming languages, one piece of data can be used as a reference to some other piece of data. Sometimes the piece of data used that way is called a *pointer* (Pascal or C); sometimes the technique is called *indirect referencing* (assembly language); some languages have no explicit concept of pointers at all (BASIC or Java). If you've never used references, pointers, or indirect referencing before, you may need to read parts of this chapter more than once; it can be confusing.

In Perl these pointers, called *references*, are used for many purposes. During this hour you will learn how to use references to call functions with multiple arguments of complex types and how to construct complex data types such as lists of lists.

A reference is much like a card in an old-fashioned library card catalog. Each index card in the catalog refers to a book on the library shelves. The card indicates the kind of book it was (fiction, nonfiction, reference) and its location. Card catalogs can have several references for the same book, in different categories, and even refer to other cards in the catalog ("See Also").

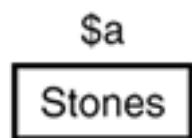
In Perl, references work much the same way by pointing to individual pieces of data. The references know the kind of data they point to (scalar, array, or hash) and its location. References can be copied without changing anything about the original data, and multiple references can be made to the same piece of data. References can, in fact, refer to other references.

Reference Basics

You create and assign a normal scalar variable by using an assignment operator, as follows:

```
$a="Stones"; # A normal scalar
```

After this snippet, a scalar variable called `$a` is created, and it contains the string "Stones". Now, somewhere within the computer's memory, there is a place labeled `$a` that contains that string, as illustrated here:



If you were to assign the scalar `$a` to `$b`—like `$b=$a;`—you would wind up with two copies of the data, with two different names, as shown here:

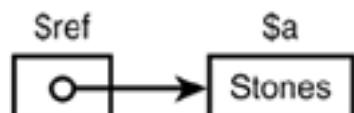


Having two copies might be acceptable if you want separate, independent copies of the data. However, if you want both `$a` and `$b` to refer to the *same* piece of data—not a copy—, you must create a *reference*. Just as the library catalog card for a book does not contain a copy of the book's text, a reference does not contain any real data; it is simply a pointer to a piece of data. The reference is usually stored in another scalar variable.

To create a reference to any given variable, you must put a backslash in front of the variable name with its type identifier. For example, to create a reference to `$a` called `$ref`, you would assign the reference to `$ref` as follows:

```
$ref=\$a; # Create a reference to $a
```

This assignment creates a situation like the following:



`$ref` doesn't contain any data for itself; it simply is a reference to `$a`. The variable `$a` isn't changed; it can

still be assigned to (`$a = "Foo"`) or displayed (`print $a`) as normal.

The variable `$ref` now contains a reference to `$a`. You cannot simply manipulate `$ref` as you would `$a`, because it doesn't contain a normal scalar value. In fact, printing `$ref` would display something similar to `SCALAR(0x0000)`. To get to the value inside `$a` through `$ref`, you must *dereference* `$ref`. Think of dereferencing as following the arrow in the preceding block diagram. To print the value of `$a` through the reference `$ref`, you would need to use an extra `$` like this:

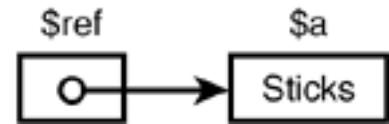
```
print $$ref;
```

In the preceding snippet, `$ref` contains, of course, the reference. The extra `$` tells Perl that the reference in `$ref` refers to a scalar value. The scalar value that `$ref` refers to is fetched and printed.

You can also modify the original value through the reference—something you can't do with a copy. The following code modifies the original value in `$a`:

```
$$ref = "Sticks";      # De-references $$ref
```

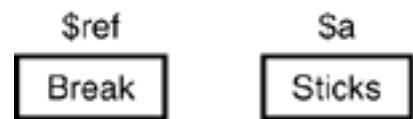
This modification creates something like the following:



If you had used `$ref` instead of `$$ref`

```
$ref = "Break";
```

then the reference stored in `$ref` would have been destroyed and replaced with a real value, as shown here:



After the preceding snippet, `$ref` no longer contains a reference; it is just another scalar. You can assign such a reference as you would any other scalar value:

```

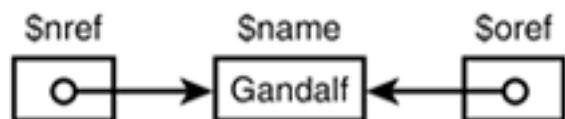
$name="Gandalf";

$nref=\$name;           # Has a reference to $name

$oref=$nref;           # Has a copy of the reference to $name

```

You get this result



After the preceding snippet, `$oref` and `$$nref` both can be used to get to the value "Gandalf". You can also store a reference to a reference, as follows:

```

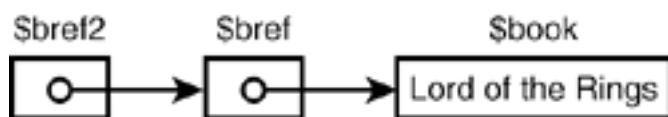
$book="Lord of the Rings";

$bref=\$book;      # A reference to $book

$bref2=\$bref;    # A reference to $bref (not to $book!)

```

In this case, the chain of references looks like the following:



To print the book title given `$bref`, you would use `$$bref`. To print the book title given `$bref2`, you would use `$$bref2`, with an extra dollar sign, requiring one more level of dereferencing to get to the original value.

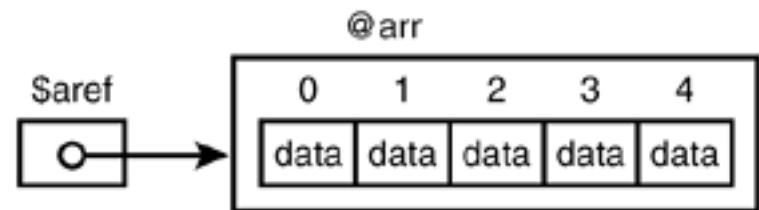
References to Arrays

References can also be created to arrays and hashes. Such references are created in the same way a reference is created to a scalar—by using a backslash:

```
$aref=\@arr;
```

Now the scalar variable `$aref` contains a reference to the entire array `@arr`. Visually, it might resemble the

following:



To access portions of `@arr` using the reference `$aref`, you would use one of the following examples:

`$$aref[0]` The first element of `@arr`

`@$aref[2,3]` A slice of `@arr`

`@$aref` The whole of array `@arr`

For clarity, you can use braces to separate the reference from the portions dealing with the array itself, as shown here:

`$$aref[0]` is the same as `$$aref[0]`

`@$aref[2,3]` is the same as `@$aref[2,3]`

`@$aref` is the same as `@$aref`

For example, to print all the elements of `@arr` using the array reference `$aref`, you can use this code:

```
foreach $element (@{$aref}) {  
    print $element;  
}
```

References to Hashes

To create a refer

Recipes for Structures

The following sections present some of the more common arrangements of lists and hashes into structures.

Example: A List of Lists

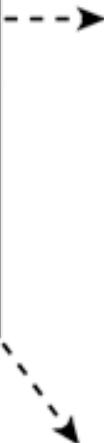
Lists of lists in Perl are often used to represent a common structure called a *two-dimensional array*. That is, a normal array is a linear list of values, as illustrated here:

[0]	[1]	[2]	[3]
VAL1	VAL2	VAL3	VAL4



A two-dimensional array is like a grid of values in which each element is addressed like a point with coordinates along two axes. The first part of the index indicates a row number (starting with 0), and the second is the column, as you can see here:

[0] [0] data	[0] [1] data	[0] [2] data
[1] [0] data	[1] [1] data	[1] [2] data
[2] [0] data	[2] [1] data	[2] [2] data



Perl doesn't really support a true two-dimensional array. However, Perl does allow you to emulate the two-dimensional array using an array of array references.

To create an array of arrays, use this literal representation:

```
@list_of_lists=(  
    [ qw( Mustang Bronco Ranger ) ],  
    [ qw( Cavalier Suburban Buick ) ],  
    [ qw( LeBaron Ram ) ],  
) ;
```

Look carefully at the preceding snippet. A regular list—`@list_of_lists`—is being created, but it consists of references to other lists. To access individual elements of the innermost lists—cells in the two-dimensional array—you can use the following code:

```
$list_of_lists[0][1];      # Bronco.  1st row, 2nd entry  
$list_of_lists[1][2];      # Buick.  2nd row, 3rd entry
```

To find the number of elements in the outer list, do so as you would for any other array—use the `$#` notation or use the array name in a scalar context:

```
$#list_of_lists;          # Last element of @list_of_lists: 2  
scalar(@list_of_lists);   # Number of rows in @list_of_lists: 3
```

Finding the number of elements in one of the inner lists is a little trickier. The syntax `$list_of_lists[1]` returns the reference in the second row of `@list_of_lists`. Printing it displays something like `ARRAY(0x00000)`. To treat an element of `@list_of_lists` as if it were an array, put an `@` sign in front of it, like this:

```
scalar(@{$list_of_lists[2]});    # Number of elements in the 3rd row  
${$list_of_lists[1]};           # Index of last element in the 2nd row
```

To traverse every element in the list of lists, you can use this code:

```
foreach my $outer (@list_of_lists) {  
    foreach my $inner (@{$outer}) {  
        print "$inner ";  
    }  
    print "\n";  
}
```

You can add to the structure like this:

```
push(@list_of_lists, [ qw( Mercedes BMW Lexus ) ]); # A new row  
push(@{$list_of_lists[0]}, qw( Taurus ) ); # A new element to one list
```

Other Structures

In the preceding section, you learned how to create a basic Perl structure, a lists of lists, by using references and arrays. Actually, an unlimited number of variations of arrays, scalars, and hashes can be combined to create more and more complex data structures, such as the following:

- Lists of hashes
- Hashes of lists
- Hashes of hashes
- Hashes that contain lists, which contain hashes, and so on

There isn't enough room in this book to describe all of these structures. However, the online documentation that comes with each Perl installation contains a document titled the *Perl Data Structures Cookbook*. It is a very detailed but understandable description of each of these structures and many others. With each structure, the *Perl Data Structures Cookbook* details

- Declaring your structure (literal representation)
- Filling your structure
- Adding elements
- Accessing elements
- Traversing the entire structure

To view the "Perl Data Structures Cookbook," at a command prompt, type `perldoc perldsc`.

Debugging with References

When debugging programs with references, it's not uncommon for new programmers to get confused about which references point to which kinds of structures. Also, the syntax can be confusing until you get used to it. For that reason, Perl provides some facilities to help you figure out what's going on.

First, you can simply print the reference. Perl displays what the reference points to. For example, the line

```
print $mystery_reference;
```

might display

```
ARRAY(0x1231920)
```

This result means that the variable `$mystery_reference` is a reference to an array. Other possibilities include references to scalars ([SCALAR](#)), hashes ([HASH](#)), or subroutines ([CODE](#)). To print the array referred to by `$mystery_reference`, you can treat it like an array, as shown he

 PREV

< Day Day Up >

NEXT 

Exercise: Another Game, Maze

After learning so many new and strange concepts—references and structures—you deserve a little diversion. The following exercise demonstrates a structure and a few references by allowing you to play a simple game.

In the manner of classic games such as Adventure, Zork, and Hunt the Wumpus, you will be placed in a maze from which you must find your way out. The maze is featureless, consisting of only rooms with at least one doorway. The doorways lead to adjacent rooms to the north, south, east, or west. The object is to find the secret room. You'll soon learn, however, that only a couple of paths actually lead to that room; there are many dead-end paths.

To get started, type [Listing 13.1](#) and save it as `Maze`. Running the program gives output similar to [Listing 13.2](#).

Listing 13.1. The Complete Listing for `Maze`

```
1:  #!/usr/bin/perl -w
2:  use strict;
3:
4:  my @maze=(
5:      [qw( e     swe we ws   ) ],
6:      [qw( se new sw ns   ) ],
7:      [qw( ns   -       ns n   ) ],
8:      [qw( ne w     ne w   ) ],
9: );
10: my %direction=( n=> [ -1, 0], s=> [1, 0],
11:                 e=> [ 0, 1], w=> [0, -1]);
12:
13: my %full=( e => 'East', n => 'North', w=>'West', s=>'South');
14: my($curr_x, $curr_y, $x, $y)=(0,0,3,3);
15: my $move;
```

```
16:  
17: sub disp_location {  
18:     my($cx, $cy)=@_;  
19:     print "You may move ";  
20:     while($maze[$cx][$cy]=~/([nsew])/g) {  
21:         print "$full{$1} ";  
22:     }  
23:     print "($maze[$cx][$cy])\n";  
24: }  
25: sub move_to {  
26:     my($new, $xref, $yref)=@_;  
27:  
28:     $new=substr(lc($new),0,1);  
29:     if ($maze[$$xref][$$yref]!~/[$new]/) {  
30:         print "Invalid direction, $new.\n";  
31:         return;  
32:     }  
33:     $$xref += $direction{$new}[0];  
34:     $$yref += $direction{$new}[1];  
35: }  
36:  
37: until ( $curr_x == $x and $curr_y == $y ) {  
38:     disp_location($curr_x, $curr_y);  
39:     print "Which way? ";  
40:     $move=<STDIN>; chomp $move;
```

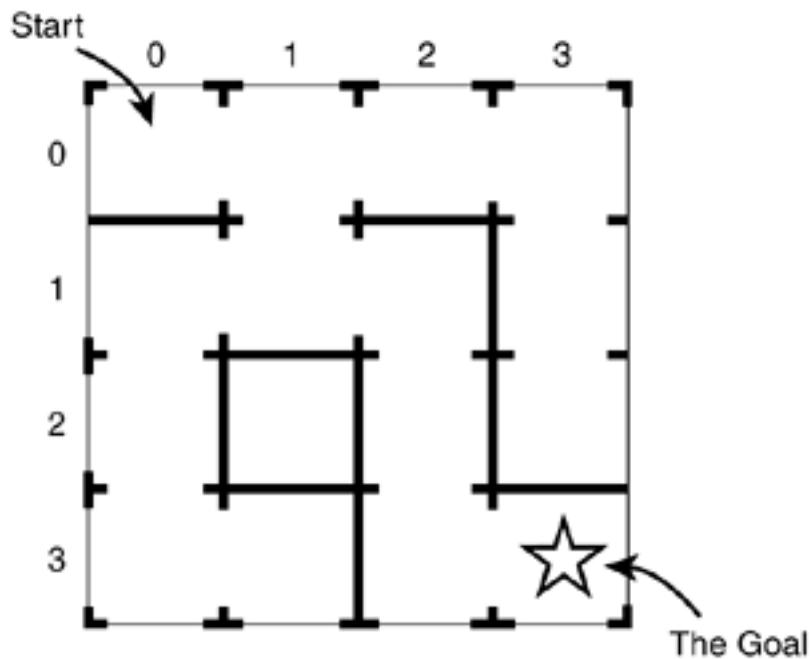
```

41:     exit if ($move =~ /q/);
42:     move_to($move, \$curr_x, \$curr_y);
43: }
44:
45: print "You made it through the maze!\n";

```

Lines 1–2: These lines are a normal beginning to a Perl program. Warnings are enabled with `-w`, and `use strict` is used to catch mistakes and bad programming practices.

Lines 4–9: The structure that describes the maze, `@maze`, is defined. The maze shown is a 4x4 grid, represented by a list of lists. Each element of the list describes which doors are available in any room in the maze, so if you redesign this maze, make sure you leave a path out! The current maze looks like this:



One room is inaccessible (2,1) and is designated by a `-` in the structure; actually, any string not matching `n`, `s`, `e`, or `w` would work as well.

Lines 10–11: As the player moves north, south, and so on, the current position in the maze needs to be changed. The hash `%direction` is used to calculate the player's new position, given the old position and direction. Moving "north" changes the player's x coordinate by `-1` (going up), and the y coordinate remains unchanged. Moving "east" leaves the player's x coordinate where it is but increases the y coordinate by `1`. You'll see how this is used in lines 33–34.

Lines 13–15: The variables used in the program are declared with `my` to make `use strict` happy. The player's current position—stored in `$curr_x` and `$curr_y`—is set to 0,0. The final destination—`$x` and `$y`—is set to 3,3.

Line 17: Given an *x,y* coordinate pair in the grid, this function displays the directions the player is allowed to move in each room.

Line 20: The letters `n`, `s`, `e`, and `w` are picked out of the room description in `$maze[$cx][$cy]` one letter at a time. The appropriate description for the `nsew` directions is displayed from the hash `%full`. This hash is used only for translating the short name (`n`) into the long name (`North`) for display.

Line 25: This function takes a direction (stored in `$new`) and references to the player's current coordinates.

Line 28: The direction is shifted to lowercase with `lc`, and `substr` takes just the first letter and assigns it to `$new`. Thus, `East` becomes `e`, `West` becomes `w`, and `s` remains `s`.

Line 29: The current room—`$maze[$$xref][$$yref]`—is searched for the direction given (`n`, `s`, `e`, or `w`). If it doesn't exist, it's not valid for this room, and a message is printed.

Lines 33–34: The player's *x* and *y* coordinates are changed. If the direction is `e`, then `$direction{e}` is a two-element array reference to (0, 1). The *x* coordinate would be increased by 0—`$direction{e}[0]`. The *y* coordinate would be increased by 1—`$direction{e}[1]`.

Line 37: The main body of the program starts here. The loop will continue until the player's *x* and *y* coordinates (`$curr_x`, `$curr_y`) match the hidden room's coordinates (`$x`, `$y`).

Line 38: The "map" of the current room is displayed.

Lines 39–41: The desired direction is read into `$move`, and the newline character is removed with `chomp`. If the player types anything beginning with a `q`, the game ends.

Line 42: The `move_to()` subroutine is called with the player's current desired move and references to the player's coordinates. The `move_to()` subroutine moves the player appropriately by adjusting `$curr_x` and `$curr_y`.

Listing 13.2. Sample Output from `Maze`

```
You may move East (e)
```

```
Which way? e
```

You may move South West East (swe)

Which way? e

:

Which way? e

You made it through the maze!

To change the maze to have a different layout, simply change the grid stored in `@maze`. The maze doesn't necessarily need to be square, nor does each room need to be mapped, nor does a valid path need to exist. Remember, though, not to provide a door that would lead off the edge of the maze; the program does not check the validity of the maze, although Perl does emit warnings if you construct an invalid maze. You can move the secret room by altering the values of `$x` and `$y`.



< Day Day Up >



Summary

During this hour you learned the basics of references. First, you learned how to take references to Perl's basic data types: scalars, arrays, and hashes. Then, using those references, you learned how to manipulate the original structures. You then learned that it was possible to take a reference to a hash or array that didn't have a variable name associated with it—anonymous storage. Finally, you learned how to use references to create complex data structures and where to find a cookbook of ready-made data structures.

Q&A

Q1: When I print a list of lists with `print "@LOL"`, it prints `ARRAY(0x101210),ARRAY(0x101400)`, and so on. Why?

A1: With a normal array, `print "@array"` would print the elements of the array with a space between them. The `print "@LOL"` is doing just that, printing the array elements in `@LOL`. To print the components of each of the arrays in `@LOL`, you must use the technique described in the "Example: List of Lists" section earlier in this hour.

Q2: I tried to take a reference to a list by using `$ref=\($a, $b, $c)` and ended up with `$ref` containing a reference to a scalar value instead of a list. Why?

A2: In Perl, `\($a, $b, $c)` is actually shorthand for `(\$a, \$b, \$c)`! What you wound up with is a reference to the last element in the parentheses, `$c`. To take a reference to an anonymous array, you should have used `$ref=[$a, $b, $c]` ; instead .

Workshop

Quiz

1: What does \$ref contain after the statement \$ref=\\"peanuts\";?

a.

Nothing. That syntax isn't valid.

b.

peanuts

c.

A reference to an anonymous scalar.

2: What does this structure create?

```
$a=[
```

```
{ name=> "Rose" , kids=> [ qw( Ted Bobby John ) ] } ,  
 { name=> "Marge" , kids=>[ qw( Maggie Lisa Bart ) ] } ,  
 ] ;
```

a.

A hash of hashes, which contains a list

b.

A list of hashes, which contains a list

c.

A list of lists, which contains another list

Answers

- A1:** c. A reference can be taken to any value, not just scalar, array, and hash variables. You can take references to a number also with `$ref=\100;`. If you answered a, it's a good idea to try new things with a short program or in the debugger to see what they do if you're unsure.
- A2:** b. This structure wasn't covered explicitly in this hour, but you should have been able to guess the answer. A list (the outer brackets) of hashes (the braces) contains a list (the data for `kids`).

Activities

- Modify the Maze game to take diagonal direction as well. You might want to use four new keys to indicate those directions (`ne`, `nw`, `se`, and `sw` would be difficult to program). Hint: The secret is modifying `@maze` to use the new symbols and `%direction` to indicate which way they should go—`[1, 1]`, `[-1, -1]`, and so on.
- Design a structure—even if just on paper—to describe a phone bill. The bill itself contains keys and data like a hash (name, phone, address), and some parts of the bill are lists (itemized calls). Each itemized call could also be considered a hash (destination, time).

 PREV

< Day Day Up >

NEXT 

Hour 14. Using Modules

What You'll Learn in This Hour:

- Learn how to use modules in your Perl programs
- Take a quick tour of some built-in modules
- Review the list of the core modules distributed with Perl

Perl, as you've noticed, is an extremely flexible language. Not only does it deal with files, text, math, algorithms, and other issues normally found in any computer language; much of Perl is dedicated to special-purpose functions. Regular expressions are a core part of the language and are very important to the way that Perl is used, although many languages do just fine without them. Likewise, Perl's handling of external programs (backticks, pipes, and `system` from Hour 11, "System Interaction") is also extensive, whereas many languages don't handle them at all.

A temptation in designing a language is to include anything and everything useful in the core of the language itself. Being so inclusive can, however, create a language that is very large and difficult to use. For example, some language designers feel that support for access to the World Wide Web should be included in the core of the language. That's a useful feature, but not necessarily one that everyone needs. If the Web (using present protocols) is less important 10 years from now than it is today, a decision will have to be made to remove that support and a lot of already-written software will be broken.

For this reason, Perl has taken a different track. Starting with Perl 5, the language itself can be extended through the use of *modules*. Modules are collections of routines that allow you to extend Perl's reach. You'll find modules that add Web browsing, graphics, Windows Object Linking and Embedding (OLE), databases, and almost anything imaginable to Perl. Remember, however, that Perl does not need modules to run. It's completely functional without them.

By using modules, you can gain access to a large library of working code to help you write your programs. The last third of this book, in fact, is dedicated to writing CGI (Common Gateway Interface) programs using Perl modules.

At the time of this writing, more than 4000 modules are available, and more than two dozen are distributed with Perl. These modules, for the most part, can be freely redistributed. You can use them in your own programs for almost any purpose. Consequently, many difficult problems that you may encounter have already been solved for you; all you need to do is install the correct module and use it appropriately.

A Gentle Introduction

The use directive allows you to use a module in your Perl program. For example, to include the `Cwd` module in your program, simply place the following somewhere in your code:

```
use Cwd;
```

It doesn't matter where in your code you place the `use Cwd`, although for clarity and maintainability, it should probably be placed near the beginning of the program.

You may recall that we used the following module in [Hour 10](#), "Files and Directories." At that time, however, you didn't know how it worked. When you run a program with `use Cwd`, what actually happens is this:

1.

The Perl interpreter opens your program and reads in all the code until the `use Cwd` statement is found.

2.

When your Perl interpreter was installed, it was informed which directory it should search for modules

(which is usually Perl's installation directory). That directory is searched for a module called `Cwd`.

This module is a file containing Perl code.

3.

Perl reads the module, and all the functions and variables that are required for the module to work are initialized.

4.

The Perl interpreter continues reading and compiling your program where it left off.

That's all there is to it. Once Perl has read your entire program, it's ready to run, using all the functionality that the module provides.

By the Way

You might notice that `use strict` looks a lot like `use Cwd`. Without confusing the issue too much, let's just say that the `use` statement is a general-purpose instruction to the Perl

interpreter to do something (this is called a *pragma*). In the case of `use strict`, it changes the behavior of the interpreter to be strict about references and bare words; there is no module called `s`



< Day Day Up >



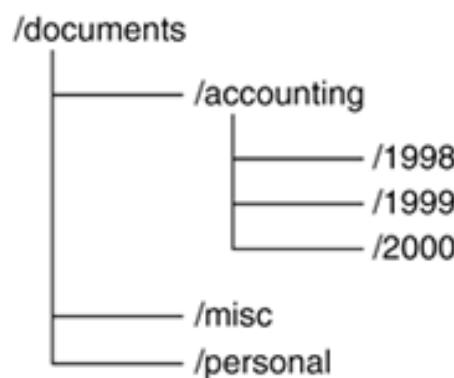
A Quick Tour

What follows is a quick tour of some of the modules already installed on your system as part of the core Perl distribution.

Exploring Files and Directories

In Hour 10, you learned how to open directories and read lists of filenames contained within those directories. At that time, when the subject of reading subdirectories came up, you were told not to be concerned. Now you'll learn how to traverse directories and subdirectories recursively.

You might write a common program to find a particular file without knowing exactly what directory it's in. For example, suppose you want to find a file named important.doc somewhere beneath the directory documents, as shown here:



This illustration shows a directory structure under a parent directory named documents. Finding a file somewhere under documents, using only `opendir/readdir/closedir`, is not easy. First, documents must be searched for the file. Then each of the directories under documents must be searched—accounting, misc, and personal—and then each directory under those directories, and so on.

This is an old problem, solved again and again by programmers over the last 30 years. Writing your own solution to it would be a waste of time. Naturally, the designers of Perl have included an easy solution: the `File::Find` module. To use the `File::Find` module in your program, simply enter the following in your program—preferably somewhere near the top:

```
use File::Find;
```

A new function called `find` then becomes available to you. The `find` function's syntax is as follows:

```
find subref, dirlist
```

The second argument to `find` is a list of directories to search. The first argument is new to you; it's a *subroutine*

reference. A subroutine reference is created just like a scalar or array reference: It's simply the subroutine name with a backslash in front of it. You *must* use the & in front of the subroutine name to take a reference to it. The subroutine named will be called for each file and directory found in *dirlist*.

The program to find the missing important.doc is shown in Listing 14.1.

Listing 14.1. Finding a File

```
1:  #!/usr/bin/perl -w
2:  use strict;
3:  use File::Find;
4:
5:  sub wanted {
6:      if ($_ eq "important.doc") {
7:          print $File::Find::name;
8:      }
9:  }
10: find \&wanted, '/documents';
```

Lines 1–2: These lines provide the usual start to your Perl programs. Warnings are enabled with `-w`, and `use strict` is used to catch errors.

Line 3: The `File::Find` module is brought into your program. It makes the `find` function available to you.

Line 5: This function is called for each file and directory under '`/documents`'. If you have 100 files and 12 directories, this routine will be called 112 times.

Line 6: When the `wanted()` function is called, `$File::Find::name` will contain the pathname to the current file being examined and `$_` will contain just the name. This line determines whether the filename is `important.doc`; if so, it prints the full pathname.

Line 10: The `find` function is called with a subroutine reference—`\&wanted`—and a directory. The `wanted()` function is called for each file and directory under '`/documents`'.

The function called by `find` will have the following variables available to it:

- `$File::Find::name`— The current pathname, directory, and filename.

- `$File::Find::dir`—The current directory name.
- `$_`—The current filename (without the directory). It's important that you don't change the value of `$_` in your function, or if you do, change it back when you're done.

[Listing 14.2](#) contains another `File::Find` example. This example removes all the files on the C: and D: drives with the extension .tmp; these temporary files accumulate and tend to clutter your hard disk. You can easily adapt this program to remove files from a Unix system or perform any kind of file maintenance.

Listing 14.2. Removing Temporary Files

```

1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:
5:  use File::Find;
6:
7:  sub wanted {
8:      # Check to see if the filename is not a directory
9:      if ( -f $File::Find::name ) {
10:          # Verify the filename ends in .tmp
11:          if ( $File::Find::name =~ /\.\tmp$/i ) {
12:              print "Removing $File::Find::name";
13:              unlink $File::Find::name;
14:          }
15:      }
16:  }
17:
18:  find(\&wanted, 'c:/', 'd:/');

```

Most of the program in [Listing 14.2](#) is similar to the one shown in [Listing 14.1](#), with the following differences.

Line 7: The filename passed in is tested to ensure it is a regular file. Remember, this subroutine gets called for both files and directories.

Lines 9–11: The filename is checked to see whether it contains .tmp at the end of the name. If so, the file is deleted with `unlink`.

Copying Files

Another common task—copying files—can be done the hard way in Perl:

1. Open the source file for reading.
2. Open the destination file for writing.
3. Read the source file and write to the destination.
4. Close both files.

And, of course, after each step, you must make sure that no errors occurred and that each write was successful. That's the hard way. An easier way is to take advantage of Perl's `File::Copy` module, which does the file copying for you. The following is an example of this module:

```
use File::Copy;  
  
copy("sourcefile", "destination") || warn "Could not copy files: $!";
```

The preceding snippet copies the contents of `sourcefile` to `destination`. The `copy` function returns 1 on success or 0 if a problem occurred, and it sets the variable `$!` to the proper error.

The `File::Copy` module also provides a `move` function, which moves a file from one directory to another. If the file can be moved by simply renaming it, it's renamed; this is usually the case when the source file and the destination file are on the same file system or disk. If simply renaming the file isn't possible, first the file is copied to the destination filename, and then the original file is removed. Consider this example:

```
use File::Copy;  
  
if (not move("important.doc", "d:/archives/documents/important.doc")) {  
    warn "important.doc could not be moved: $!";  
  
    unlink "d:/archives/documents/important.doc";  
}
```

In the preceding snippet, the file `important.doc` is moved from its current directory to the target directory `d:/archives/documents`. If the `move` function fails, it's possible that a partial target file exists. The `unlink` function removes the partially copied target file if the `move` fails.

Is Anybody Out There?

Perl's modules aren't limited to just manipulating your files and directories. For example, you can use the `Net::Ping` module to determine whether your system can communicate properly on a network.

`Net::Ping` gets its name from a Unix utility called `ping`, which got its name from the "ping" sound that submarines use in echo location. The `ping` utility sends a packet to another system on the network. If that system is up and running, it sends a reply, and the `ping` command reports success. `Net::Ping`, shown here, works the same way:

```
use Net::Ping;

if ( pingecho("www.yahoo.com", 15) ) {
    print "Yahoo is on the network.";
} else {
    print "Yahoo is unreachable.";
}
```

In the preceding snippet, the `Net::Ping` module provides a function called `pingecho`. The `pingecho` function takes two arguments, the first is a host to find—www.yahoo.com in this case. The second argument indicates how long `pingecho` should wait for a response in seconds.

By the Way

Due to the nature of Perl on Microsoft Windows platforms, the `Net::Ping` module does not currently work. `Net::Ping` relies on the `alarm` function, which does not work under Windows. ActiveState, a major developer of Perl for Windows, has announced plans to implement many missing functions for Windows and may incorporate those changes into Perl eventually.

Once Again, in English?

The `English` module allows some of Perl's obscure special variables to be addressed by more verbose names, as in this example:

```
use English;

while(<>) {
    print $ARG;
```

}

In the preceding snippet, `while(<>)` normally reads a line of input from `STDIN` and assigns it to `$_`. It still does. With `use English`, however, the variable `$_` is also known as `$ARG`. A partial list of special variables and their English equivalents is shown here:

Special Variable English Name

<code>\$_</code>	<code>\$ARG</code>
<code>@_</code>	<code>@ARG</code>
<code>\$!</code>	<code>\$OS_ERROR</code>
<code>\$^O</code>	<code>\$OSNAME</code>
<code>\$0</code>	<code>\$PROGRAM_NAME</code>

You can find a full list of special variables and the English equivalents in the `English` module's online documentation.

More Diagnostics

The Perl module `diagnostics` can help you find bugs in your program. As you were typing in the examples from this book, the Perl interpreter may have emitted an error message that you didn't quite understand. For example, the short program

```
#!/usr/bin/perl -w

use strict;

print "For help, send mail to help@support.org\n";
```

causes Perl to emit this warning message:

```
In string, @support now must be written as \@support at line 4
```

```
Global symbol "@support" requires explicit package name at line 4
```

The `diagnostics` module causes Perl to emit a wordy explanation of its errors and warnings. You can change the

sample program to include the `diagnostics` module like this:

```
#!/usr/bin/perl -w

use strict;

use diagnostics;

print "For help, send mail to help@support.com\n";
```

The revised program causes a wordier diagnostic message to print:

```
In string, @support now must be written as \@support at line 4
Global symbol "@support" requires explicit package name at ./diag.pl line 5 (#1)

(F) You've said "use strict vars", which indicates that all variables
must either be lexically scoped (using "my"), or explicitly qualified to
say which package the global variable is in (using "::").
```

If you think for a moment about the two messages, it's apparent that they're related. The first message is obvious. Perl wants your email message to be written as `help@support.com`. The second message, now that it's been explained, is a bit more obvious: Because `use strict` was in effect, the `@support` variable should have been declared with `my`. Of course, `@support` isn't a variable; it was supposed to be part of an e-mail address, but Perl was misled by the special character `@`.

The letter in front of the message indicates what kind of error you have. (`W`) indicates a warning, (`D`) indicates you've used a deprecated syntax, (`S`) is a severe warning, and (`F`) is a fatal error. For all message types except (`F`), your Perl program will continue to run.

Perl has 60 pages of descriptions to its error messages. If you're having problems figuring out Perl's terse error messages, `use diagnostics` can sometimes help.

Did you Know?

The full list of error messages and diagnostic information is available for you to brows

Full List of Standard Modules

A lengthy description of the full list of modules included with Perl is well beyond the scope of this book. The following is a listing of the modules in the standard Perl distribution with a brief description. If you're curious about what the module does and how it works, use `perldoc` to view the documentation for a specific module.

Module Name	Description
AutoLoader	Allows Perl to compile functions only when needed.
AutoSplit	Splits modules for autoloading.
Benchmark	Allows repetitive timing of Perl functions for speed benchmarking.
CGI	Allows easy access to the Common Gateway Interface for Web programming, covered in Hours 17 to 24 .
CPAN	Provides access to the archive of Perl's modules, for installing new modules.
Carp	Generates error messages.
DirHandle	Provides an object interface to directory handles.
Env	Maps the operating system's environment into variables.
Exporter	Allows you to write your own modules.
ExtUtils::*	Allow you to write your own modules or install modules.
File::*	Offer more file-manipulation operations, such as <code>File::Copy</code> .
File::Spec::*	Allows OS-independent operations on filenames.
FileCache	Opens more files than the OS normally allows.
FindBin	Finds the name of the current executing program.
Getopt::*	Allow you to process command-line options in your programs.
I18N::Collate	Allows locale-specific sorting.
IPC::*	Provide Interprocess Communications; two- and three-ended pipes, for example.
Math::*	Allow you to use extended math libraries with arbitrary-precision floating-point, integer, and complex numbers.
Net::*	Allow you to get information on network hosts. For example, <code>Net::hostent</code> translates IP addresses—such as <code>204.71.200.68</code> —into host names—such as <code>www.yahoo.com</code> .

<code>Pod::*</code>	Provide access to Perl's Plain Old Documentation formatting routines.
<code>Symbol</code>	Allows you to view or manipulate Perl's own symbol table.
<code>Sys::Hostname</code>	Obtains your system's IP hostname.
<code>Sys::Syslog</code>	Allows writing to a Unix system's error log.
<code>Term::*</code>	Provide terminal-controlling functions interface for cursor positioning, screen cleaning, and so on.
<code>Text::Abbrev</code>	Builds abbreviation tables.
<code>Text::ParseWords</code>	Allows you to parse text to search for words.
<code>Text::Soundex</code>	Categorizes words based on pronunciation, using the <code>Soundex</code> method.
<code>Tie::*</code>	Connect Perl's variables to functions so that you can implement your own arrays and hashes.
<code>Time::*</code>	Allow you to parse and manipulate time. For example, you can convert times such as "Sat Jul 24 16:21:38 EDT 1999" back into the number of seconds since January 1, 1970.
<code>constant</code>	Allows you to define constant values.
<code>integer</code>	Causes Perl to do its math in integers instead of floating-point numbers (sometimes).
<code>locale</code>	Causes locale-based string comparison (international character string comparison).

Where Do You Go from Here?

If you want to get a feel for the kinds of modules available to you—for free—use a Web browser and head to <http://www.cpan.org>. The modules are organized (roughly) by category.

Some modules require a C compiler and a minimal development environment for installation. They may not be available on a Windows machine. ActiveState's Perl installation contains a utility called `PPM`, which can be used to browse and install prebuilt modules.

The Appendix contains step-by-step instructions for installing modules on Unix and Windows machines. These instructions will explain how to use the `CPAN` module (for Unix) and ActiveState's `PPM` utility to install new modules.

Summary

In this hour, you learned how modules are used to extend the Perl language to perform many other tasks. This general-purpose method of adding new functionality to Perl will be used throughout the rest of the book. Also, a few common modules were presented, and the full listing of modules included with the standard distribution was reviewed.

Q&A

Q1: In the `File::Find` module, what are the colons (`::`) in the variable names? Like `$File::Find::dir`?

A1: Perl modules can create alternate areas for variable names, called *namespaces*, so that the module's global variable names and your global variable names don't get mixed up. So a global variable in the `Cwd` module would be known as `$Cwd::x`. Most of your global variables actually have the full name of `$main::x` for `$x`. But for now, that's not important.

Q2: I have a Windows 95/98/NT machine, and the module I want to use is not available through ActiveState's `PPM`. How can I install it?

A2: Unfortunately, most of CPAN's modules assume that you have a full Unix-like development environment to compile and install modules; this kind of environment is not easy to put together on a Windows machine. If you're very handy with a C compiler, you can download a development environment and build your own module, but doing so is not easy. Some modules are written entirely in Perl, and it may just be a matter of opening the distribution and putting the files in the correct places.

Q3: I have an older Perl program with `require` instead of `use`. What's `require`?

A3: The `require` statement is similar to `use`, and because Perl 4 did not have the `use` keyword, it used `require`. (Follow that?) The `require` statement causes the interpreter to find a library file and include it in your program—similar to `use`. However, the key difference is that a `require` happens whenever the `require` statement is run (at runtime), whereas the `use` directive happens when your program is first loaded (at compile time).

Workshop

Quiz

1: If you want to use the `cwd` function twice in a program, how many times do you have to use `Cwd ;`?

a.

Once.

b.

Once for each instance of `cwd`, so two times.

c.

None. `cwd` is a built-in function.

2: What module provides an alias for the `$_` variable?

a.

`LongVars`

b.

`English`

c.

`$_` does not have an alias.

Answers

A1: a. After a module has been pulled into your program with `use`, all its functions are available to use throughout the rest of your program.

A2: b. `use English` causes `$_` to be known by the name `$ARG`.

Activities

1.

Flip to [Appendix A](#) and try using the instructions there for installing the module `Bundle::LWP` from CPAN. You'll need a module in that bundle for the examples in [Hour 24](#).



< Day Day Up >



Hour 15. Finding Permanence

What You'll Learn in This Hour:

- Create a DBM file and store data in it
- [Treat regular text files as a database](#)
- Read and write to random positions in files
- Lock files for concurrent access

Writing Perl programs that find data in files or interact with the user is all well and good. But what happens when the program ends? The results are gone, you've got nothing to show for your effort, and you're left with this empty feeling—as though nothing had happened at all.

You may already be thinking of ways in which you could solve this problem in Perl based on what you have learned in previous hours, such as files, lists, and hashes. That would, of course, be the Hard Way.

Storing information for later retrieval the Easy Way is what databases are all about. A database is a set of one or more files in a particular format that specialized software can access for storing data and retrieving it later. A properly designed database can be used by many kinds of programs for queries, reports, and data entry. To design a database, you need to think carefully about what kind of data you want to store and how you want to store it. You also need to consider how the data will be accessed: by one person at a time or by many users simultaneously.

DBM Files

One of the simplest ways for your Perl program to remember data in an organized way is through the use of DBM files. A DBM file is a file that's been connected to a Perl hash. To read and write a DBM file, you simply manipulate a hash, just as you've been doing since [Hour 7](#), "Hashes."

To tie your hash to a DBM file, use the Perl function `dbmopen`, as shown here:

```
dbmopen(hash, filename, mode)
```

The `dbmopen` function connects the *hash* to a DBM file. The *filename* you supply actually creates one or two different files on your hard disk with names that are a variation on *filename*: *filename.pag*, *filename.dir*, or *filename.db*. Perl uses these files to store the hash. These files are not text files, and you should *not* attempt to use an editor on them. Also, if one of the files is empty, or seems very large relative to the amount of data in it, don't worry about it; that's normal.

The *mode* is the permissions on the two DBM files that Perl creates. Recall that permissions were discussed in the section "[A Crash Course in File Permissions](#)" in [Hour 10](#), "Files and Directories." For Unix, use a sensible set of permissions; they control who can access your DBM file. For example, `0666` allows everyone read and write access to your DBM files. Mode `0644` allows you to read and write the files, but others to read them only. For Windows, you can simply use `0666`, because you don't have any file system permissions to worry about.

The `dbmopen` function returns true the hash has been successfully tied to the DBM file; it returns false otherwise. Now consider this example:

```
dbmopen(%hash, "dbmfile", 0644) || die "Cannot open DBM dbmfile: $!" ;
```

After the preceding statement is executed, the hash `%hash` is connected to a DBM file called `dbmfile`. Perl creates a pair of files on your disk, called `dbmfile.pag` and `dbmfile.dir`, to keep the hash in. If you assign a value to the hash, as shown here, Perl updates the DBM files with that information:

```
$hash{feline}="cat" ;  
  
$hash{canine}="dog" ;
```

Fetching information causes Perl to retrieve the key and data from the DBM file, as follows:

```
print $hash{canine};
```

To disconnect the hash from the DBM file, use `dbmclose` with the hash name like this:

```
dbmclose(%hash);
```

After you've disconnected from the DBM file, the items that were stored in it—the keys `feline` and `canine` with the values associated with them—are still in the DBM file, which is the whole point of a DBM file. The key-value pairs stored in a hash tied to a DBM file are retained between invocations of a Perl program.

Functions you would normally perform with a hash can be performed with a hash that's tied to a DBM file. The hash functions `keys`, `values`, and `delete` work normally (`exists` does not work). You can empty the hash—and the DBM file—by assigning the hash to an empty list like this: `%hash=()`. You can even initialize the hash by assigning it to a list after it's tied to the DBM file with `dbmopen`.

Important Points to Know

The following are just a few small points you need to know about tying hashes to DBM files:

- The length of your keys and data is now limited. Normal hashes have (virtually) unlimited key length and value length. Hashes that are tied to a DBM file have a limited key and value length, typically around 1,024 characters for both the key and the value combined—this is a limitation of DBM files. The total number of keys and values that can be stored has not changed; it is limited only by your file system.
- The values in the hash before `dbmopen` are lost, so it's best just to use a fresh hash. Now consider this example:

```
%h=();  
  
$h{dromedary}="camel";  
  
dbmopen(%h, "database", 0644) || die "Cannot open: $!";  
  
print $h{dromedary}; # Likely will print nothing at all  
  
dbmclose(%h);
```

In the preceding snippet, the value in the hash `%h`—the key of `dromedary`—is lost when `dbmopen` is executed.

- The values in the hash while it's tied to the DBM file disappear after `dbmclose` is performed:

```
dbmopen(%h, "database", 0644) || die "Cannot open: $!" ;  
  
$h{bovine} = "cow" ;  
  
dbmclose(%h) ;  
  
print $h{bovine};      # Likely will print nothing
```

The values in the hash *while* it is connected to the DBM file are retained in the DBM file. Afterward, the hash itself is empty.

Walking Through DBM-Tied Hashes

Consider for a moment a hash that's been tied to a DBM file. For example, if you write a Perl program to keep appointments, phone numbers, and other information in that hash, after a while the hash can become large. Because the hash's values are retained between runs of your Perl program, they never quite go away—unless you delete them on purpose.

If your DBM file (called `records`) has collected a lot of information, this snippet might have some problems:

```
dbmopen(%recs, "records", 0644) || die "Cannot open records: $!" ;  
  
foreach my $key (keys %recs) {  
  
    print " $key = $recs{$key}\n" ;  
  
}  
  
dbmclose(%recs) ;
```

There's no problem with the code. The hash `%recs` is first tied to a DBM file, and then the keys from the hash are extracted with `keys %recs`. The list of keys is iterated over with `foreach my $key`, and each key and its value are printed.

If the list of keys in `%recs` is large, however, the statement `keys %recs` could take some time to execute. Perl has another function that allows you to iterate over a hash one key at a time; it's called `each`. The syntax for `each` is as follows:

```
($key, $val)=each(%hash);
```

The `each` function returns a two-element list: a key and a value from thehash. Each successive call to `each` returns the next key-value pair from the hash. When the keys are exhausted, `each` returns an empty list. A better way to iterate over a potentially large hash would be to use the following example:

```
dbmopen(%recs, "records", 0644) || die "Cannot open records: $!";  
while( ($key, $value)=each %recs) {  
    print " $key = $value\n";  
}  
dbmclose(%recs);
```

By the Way

You don't necessarily have to use `each` with a hash tied to a DBM file; you can use it with any hash at all.

Exercise: A Free-Form Memo Pad

Now that you have an easy way to keep data on the disk, it's time to put it to good use. This exercise presents a free-form memo pad that keeps information based on keys and allows you to search and retrieve that information with simple queries.

A sample session with the program (called `memopad`) is shown in Listing 15.1.

To query the `memopad` program, simply type the name of a topic, followed by a question mark. To program in a new fact, type a phrase in the form "X is Y", where X is the topic and Y is the information to associate with that topic. You can search the database for similarities by typing "`like pattern?`" where *pattern* is a regular expression to search for in the topics. All topics matching that expression will be printed. To exit the program, type `quit` at the prompt.

Listing 15.1. Sample Session with `memopad`

Your question: `perl?`

I don't know about "perl"

Your question: `perl is a programming language`

Ok, I'll remember "perl" as "a programming language"

Your question: `perl's homepage is at http://www.perl.org`

Ok, I'll remember "perl's homepage" as "at http://www.perl.org"

Your question: `perl?`

`perl` is a programming language

Your question: `like perl?`

`perl` is like perl

`perl's homepage` is like perl

Your question: `quit`

All the information given to the `memopad` program will be remembered each time the program is run, because the data is stored in a hash tied to a DBM file. The code for `memopad` is shown in Listing 15.2.

Listing 15.2. Full Listing of `memopad`

```
1:  #!/usr/bin/perl -w
2:  use strict;
3:
4:  my(%answers, $subject, $info, $pattern);
5:
6:  dbmopen(%answers, "answers", 0666) || die "Cannot open answer DBM: $!";
7:  while(1) {
8:      print "Your question ('quit' to quit): ";
9:      chomp($_=lc(<STDIN>));
10:     last if (/^quit$/);
11:     if (/like\s+(.* )\?/) {
12:         $pattern=$1;
13:         while( ($subject,$info)=each(%answers) ) {
14:             if ($subject=~/$pattern/) {
15:                 print "$subject is like $pattern\n";
16:             }
17:         }
18:     } elsif (/(.*)\?/) {
19:         $subject=$1;
20:         if ($answers{$subject}) {
21:             print "$subject is $answers{$subject}\n";
22:         } else {
23:             print qq{I don't know about "$subject"\n};
24:         }
25:     } elsif (/(.*)\sis\s(.*)/) {
```

```
26:     $subject=$1;
27:     $info=$2;
28:     $answers{$subject}=$info;
29:     print qq{Ok, I'll remember "$subject" as "$info"\n};
30: } else {
31:     print "I'm sorry, I don't understand.\n";
32: }
33: }
34: dbmclose(%answers);
```

Lines 1–2: These lines provide the usual start to Perl programs: a `#!` line with `-w`, which means that warnings are enabled. Also, a `use strict` directive keeps you from making silly mistakes.

Line 6: The hash `%answers` is tied to the DBM file `answers` with `dbmopen`. Two files will be created on disk—`answers.pag` and `answers.dir`.

Line 7: `while(1)` executes the loop forever. Somewhere in the loop is a `last` or `exit` to get out of the loop.

Line 9: This line looks confusing because several things are happening at once. `lc` shifts its scalar argument to lowercase. Because `<STDIN>` is being used in a scalar context, one line of input from `STDIN` is read in and subsequently shifted to lowercase, and the result is assigned to `$_`. `chomp` is used to remove the trailing newline character.

Line 10: If the input line consists of only the word `quit`, the `while` loop is exited.

Line 11: If the input line (now in `$_`) matches the word `like` and then some text and then a `?`, the text is saved in `$1` by the parentheses in the pattern match...

Line 12: and the string in the pattern match from line 11 is saved into `$pattern`.

Lines 13–17: The hash `%answers` is searched key by key for a key that matches the string in `$pattern`. As each key is found, it is printed.

Line 18: (This line is a continuation of the `if` started on line 11.) Otherwise, if the input line ends in a question mark, everything up to—but not including—the question mark is remembered

in \$1 with parentheses.

Line 19: The pattern in \$1 is saved into \$subject.

Lines 20–24: If the key \$subject is in the hash %answers, then the key and the associated data are printed. Otherwise, the program responds with I don't know.

Lines 25–27: (This line is a continuation of the if started on line 11.) Otherwise, if the input line is in the form X is Y, the first part (X) is remembered in \$subject and the last part in \$info.

Line 28: The information in \$info is stored in the hash %answers as \$subject.

Line 34: The DBM files are disconnected from %answers.

 PREV

< Day Day Up >

NEXT 

Text Files as Databases

Often, databases are small, simple arrangements: a list of users on a small system, local hosts on a small network, a list of favorite Web sites, or a personal address file. These are all simple forms of databases, and for simple databases, normal text files will often do. But before using text files as databases, you need to consider some pros and cons.

The good news: Using a text file as a database has a few distinct advantages over using more complicated alternatives such as DBM files or large databases such as Oracle or Sybase. Some of these advantages are:

- Text file databases are portable. They can be moved between vastly different kinds of systems without too much trouble.
- Text file databases can be edited with a text editor and printed to paper without any special tools.
- Text file databases are simple to construct initially.
- Text file databases can be imported into other programs—spreadsheets, word processors, or other databases—with hassles. Almost any program that allows you to import data allows you to import text.

Now, as you would expect, there is some bad news. To understand the bad news fully, consider how text files are usually constructed. Text file databases are traditionally arranged so that each line in the text file is a record and columns within each line are fields. To your system, however, a text file is simply a stream of characters. So a text file database that looks like

Bob 555-1212

Maury 555-0912

Paul 555-0012

Ann-Marie 555-1190

is actually stored as a continuous stream of characters like

Bob[space]555-1212[newline]Maury[space]555-0912[newline]Paul[space]...

where [space] represents a space character and [newline] represents a record separator (newline character, "\n") for your operating system, as discussed in [Hour 5](#). The characters for each record and each field are all packed together in one long stream of characters; the nice column-row display is simply the human-readable way that editors, printers, and Perl represent the data.

Keeping that structure in mind, consider the disadvantages of a text file database:

- Text files cannot be inserted into; they can only be overwritten—partially or completely. Inserting new data anywhere, except at the end of the file, involves copying all the data following the newly inserted data further down in the file.

New data: Susan 555-6613 to be inserted after "Bob"

Bob[space]555-1212[newline]Maury[space]555-0912[newline]Paul[space]...

All this data must be copied →

Bob[space]555-1212[newline]Susan[space]555-6613[newline]Maury[space]...

Copying data within a file is error-prone and slow.

- The reverse is also true: Removing data from the middle of a text file is difficult. All the data following the removed portion must be copied into the gap. For example, you would remove Maury from the original text database like this:

By removing this

Bob[space]555-1212[newline]Maury[space]555-0912[newline]Paul[space]...

← All this data must be copied

Bob[space]555-1212[newline]Paul[space]555-0012[newline]Ann-Marie[space]...

- To find a particular record in a text file database, you must search the file sequentially—normally from the top down. Unlike a DBM file, in which finding a record is as easy as looking for it in a hash, each line of a text file must be examined to see whether it's the correct record. This process is slow, and it gets progressively slower the larger the database gets.

Inserting into or Removing from a Text File

Text file databases aren't completely hopeless. With a small text file database, you can easily insert or delete from the database if you treat the text file like an array.

For example, if the database

Bob 555-1212

Maury 555-0912

Paul 555-0012

Ann-Marie 555-1190

were saved into a file called phone.txt, a short Perl program could read the database into an array like this:

```
#!/usr/bin/perl -w
```

```

use strict;

sub readdata {
    open(PH, "phone.txt") || die "Cannot open phone.txt: $!";
    my(@DATA) = <PH>;
    chomp @DATA;
    close(PH);
    return(@DATA);
}

```

Here, the `readdata()` function reads `phone.txt` and puts the data into `@DATA`—without the newline characters—and returns the array. If you add another function, `writedata()`, as follows, the database can be read and written:

```

sub writedata {
    my(@DATA) = @_;      # Accept new contents
    open(PH, ">phone.txt") || die "Cannot open phone.txt: $!";
    foreach(@DATA) {
        print PH "$_\n";
    }
    close(PH);
}

```

Now, to insert records into the database, simply read the data with `readdata()` into an array; use `push`, `unshift`, or `splice` to insert a record into the array; and then write the array out again with `writedata()` like this:

```
@PHONELIST=readdata();      # Put all of the records in @PHONELIST  
  
push(@PHONELIST, "April 555-1314");  
  
writedata(@PHONELIST);    # Write them out again.
```

To remove text from the text file database, use `splice`, `pop`, or `shift` on the array `@PHONELIST` before writing it back out. You can also manually edit the array with a loop, such as with `grep`:

```
@PHONELIST=readdata();      # Read all records into @PHONELIST  
  
# Remove everyone named "Ann" (or Annie, Annette, etc...)  
  
@PHONELIST=grep( ! /Ann/, @PHONELIST);  
  
writedata(@PHONELIST);
```

In the preceding snippet, the records are copied into `@PHONELIST` from `readda()`. The `grep` iterates over `@PHONELIST`, testing each element to see whether it does not match `Ann`; those that do not match are assigned to `@PHONELIST` again. The `@PHONELIST` array is then given back to `writedata()` for writing.

 PREV

< Day Day Up >

NEXT 

Random File Access

If you're adventurous, you can do random reads and writes within a file, as previously mentioned. The following sections briefly cover some tools you'll need to do random reads and writes; they are not covered more extensively because you shouldn't need them often.

Opening Files for Read and Write

Until now, you've looked at three methods for opening files. Files could be opened for reading, they could be opened for writing, or they could be opened for appending. Files can also be opened for both reading and writing at the same time. [Table 15.1](#) lists the various modes available for opening files.

Table 15.1. Summary of Open Modes

open Command	Reading?	Writing?	Append?	Creates If Does Not Exist?	Truncates Any Existing Data?
open(F, "<file")	Yes	No	No	No	No
<i>or</i>					
open(F, "file")					
open(F, ">file")	No	Yes	No	Yes	Yes
open(F, ">>file")	No	Yes	Yes	Yes	No
open(F, "+<file")	Yes	Yes	No	No	No
open(F, "+>file")	Yes	Yes	No	Yes	Yes
open(F, "+>>file")	Yes	Yes	Yes	Yes	No

Notice the following:

- Modes that specify "Append" can be tricky. On some systems such as Unix, data written to the file is always written at the end of the file, regardless of where the read pointer is. (You'll read more about that in a moment.)
- You should almost never use `+>`. The contents of the file are erased as soon as it's opened.

Moving Around in a Read/Write File

When a file is opened, the operating system keeps track of where in the file you happen to be. This pointer is

called the *read pointer*. For example, when a file is first opened for reading, the read pointer is at the beginning of the file, as shown here:



After you've read through the whole file, the read pointer is at the end of the file, as shown here:



To reposition the pointer to any spot within the file, you must use the `seek` function. The `seek` function takes two arguments: The first is an open filehandle, and the second is the offset in the file that you want to seek. The last argument is what the offset is relative to: 0, the beginning of the file; 1, the current position in the file; or 2, the end of the file. The following are some examples of seeking within a file:

```
# open existing file for reading and writing  
  
open(F, "+<file.txt") || die "file.txt error: $!";  
  
seek(F, 0, 2);                      # Seek to the end of the file  
  
print F "On the end";                 # Appended to the end of the file  
  
seek(F, 0, 0);                      # Seek back to the beginning of the file  
  
print F "This is at the beginning";
```

The `tell` function returns the position of the current read pointer in a file. For example, after running the preceding snippet, `tell(F)` would return 24—the length of "This is at the beginning"—because the read pointer is sitting immediately after that text.

By the Way

This section barely touched upon the `seek`, `tell`, and `open` commands. For more information on these commands, consult your online documentation. The `seek`, `tell`, and `open` functions are documented in the `perlfunc` manual, which you can access by typing `perldoc perlfunc` at the command prompt. Additionally, a longer explanation of `open` is in a manual section called `perlopentut`; for that documentation, type `perldoc perlopentut` at a command prompt.

 PREV

< Day Day Up >

NEXT 

Locking

Imagine that you've written a wonderful Perl program and that the whole world wants to use it. If you're on a Unix or Windows NT machine, or even on a Windows 95 or 98 machine, more than one person might be running your program *at the same time*. Or, your program may be put onto a Web server, and it's run so frequently that instances of your program overlap.

Now suppose that your program uses a database for its work, such as the text file database just described—but this discussion applies to any kind of database. Look at the following code, which uses functions described in the preceding section:

```
chomp($newrecord=<STDIN>);           # Get a new record from the user  
  
@PHONEList=readdata();                 # Read data into @PHONEList  
  
push(@PHONEList, $newrecord);          # Put the record into the array  
  
writedata(@PHONEList);                # Write out the array
```

Looks harmless, doesn't it? But if two people run your program at nearly the same time and try adding different records, it's not harmless at all; it's quite buggy. In the following diagram, this particular set of Perl statements is run at nearly the same time, on the same system, by two different people (Person 2 is working slightly behind Person 1). Watch carefully.

Time	Sequence	Person 1	Person 2
1		\$newrecord = "David 555-1212";	
2		@PHONEList = readdata();	\$newrecord = "Joy 555-6611";
3		push (@PHONEList, \$newrecord);	@PHONEList = readdata();
4		writedata (@PHONEList);	push (@PHONEList, \$newrecord);
5			writedata (@PHONEList);

From Person 1's perspective, the data is read in at step 2, and the new record ("David") is added to @PHONEList in step 3 and written at step 4.

From Person 2's perspective, the data is read in at step 3, the new record ("Joy") is added to @PHONEList in step 4, and @PHONEList is written out in step 5.

Here's the bug: The data read in by Person 2 in step 3 does not contain the record "David." That record

hasn't been written yet by Person 1. So Person 2 adds "`Joy`" to the array `@PHONE1`, which does not contain "`David`." At the same time, Person 1 writes a copy of `@PHONE1` to the database—which does contain "`David`."

When Person 2's instance of the program finally makes it to step 5, it overwrites the data written by Person 1. The database winds up with "`Joy`", but not "`David`"—clearly a bug.

Watch Out!

The problem is actually *worse* than you just learned; the preceding explanation is vastly oversimplified. The additional headaches come from the fact that the `writedata()` functions seemingly open and write the data all in one burst—but they don't. Multiprocessing operating systems can actually stop a program in the middle of writing data, go run another program for a moment, and resume it later—milliseconds later, but later nonetheless. Both programs can be writing to the same file, with different data, at the same time. This can cause your data file to become corrupted or even erased.

The kind of problem we have just seen has a formal name; it's called a *race condition*. That is, what gets stored in the file depends one which user "wins" or "loses" the race. Race conditions are difficult to debug in programs, because they come and go depending on how many instances of the program are running at the same time, and because the bugs associated with race conditions aren't always obvious.

Allowing multiple programs to update the same data at the same time is a tricky proposition, but it can be handled using a mechanism called *locks*. File locks are used to prevent multiple instances of a program from altering a file at the same time.

Locking files poses several problems, but the foremost among them is that different operating systems and different kinds of file systems require different types of locking mechanisms. The next sections describe how to lock files to prevent this kind of disaster.

Locking with Unix and Windows

To lock files under Unix and Windows, you can use Perl's `flock` function. The `flock` function provides an *advisory* locking mechanism. This means that any programs you write that need to access the file must also use `flock` to make sure that no one else is writing to the file at the same time. However, other programs can still modify the file; this is why the mechanism is called *advisory locking*, not *mandatory locking*.

You're familiar with one kind of advisory lock already: a traffic stop light. The signal is there to prevent multiple vehicles from entering the same part of the intersection at the same time. But the signal works only if everyone obeys it. The same holds true for file locks. Every program that can potentially access a file at the same time must use `flock` to prevent an accident. An advisory lock doesn't keep other processes from accessing the data; it prevents other processes only from obtaining a lock.

The `flock` function takes two arguments—a filehandle and a lock type—as you can see in the following syntax:

```
use Fcntl qw(:flock);

flock(FILEHANDLE, lock_type);
```

The `flock` function returns true if the lock is successful; it returns false otherwise. Sometimes calling `flock` causes your program to pause and wait for other locks to clear (more about this in a moment). The `use Fcntl qw(:flock)` allows you to use symbolic names for the `lock_type` instead of harder-to-remember numbers.

There are two kinds of locks: shared and exclusive. Normally, you get a shared lock when you want to read a file and an exclusive lock when you want to write to it. If a process has an exclusive lock on a file, then that's the only lock there is—no other process can have locks at all. But many processes can have shared locks at the same time as long as there are no exclusive locks. In such cases, it's safe to have many processes reading the file at the same time as long as nobody is writing.

Some possible values for `lock_type` are as follows:

- `LOCK_SH`— This value requests a shared lock on the file. If another process has an exclusive lock on the file, then `flock` pauses until the exclusive lock is cleared before taking out the shared lock on the file.
- `LOCK_EX`— This value requests an exclusive lock on a file opened for writing. If other processes have a lock (either shared or exclusive), then `flock` pauses until those locks are cleared.
- `LOCK_UN`— This value releases a lock, but it is rarely needed; simply closing the file writes out any unwritten data and releases the lock. Releasing the lock on a file that is still open can cause data corruption.

Did you Know?

Locks taken with `flock` are released when you close the file or when your program exits—even if it exits with an error.

Taking out a lock on a file that you're also trying to read or write can be tricky. Problems arise because opening a filehandle and locking the file is at least a two-step process: You have to have the file open before you can lock it. If you open a file with `open(FH, ">filename")` and then get a lock with `flock`, you've modified the file (truncated with `>`) before you obtain the lock. This could potentially modify the file (by truncating it) while some other process has a lock on it.

Solving this problem involves something called a *semaphore file*. A semaphore file is just a sacrificial file

whose contents aren't important; whoever holds a lock on that file can proceed.

To use a semaphore file, all you need is a filename that can be used as a semaphore and a couple of functions to lock and unlock the semaphore file, as shown in [Listing 15.3](#). This is not a complete program; it is meant to be included as a part of other programs.

Listing 15.3. General-Purpose Locking Functions

```
1:  use Fcntl qw(:flock);  
2:  # Any file name will do for semaphore.  
3:  my $semaphore_file="/tmp/sample.sem";  
4:  
5:  # Function to lock (waits indefinitely)  
6:  sub get_lock {  
7:      open(SEM, ">$semaphore_file")  
8:          || die "Cannot create semaphore: $!";  
9:      flock(SEM, LOCK_EX) || die "Lock failed: $!";  
10: }  
11:  
12: # Function to unlock  
13: sub release_lock {  
14:     close(SEM);  
15: }
```

These locking functions can surround any code that you do not want to be run concurrently, even if that code has nothing to do with reading or writing files. For example, this snippet—even if run by several processes at the same time—allows only one process at a time to print a message:

```
get_lock();      # waits for a lock.  
  
print "Hello, World!\n";
```

```
release_lock(); # Let someone else print now...
```

The `get_lock()` and `release_lock()` functions we have just seen will be used throughout the rest of this book for locking files when locks are needed.

Watch Out!

Waiting for user input (or any other potentially slow events) while holding a lock is not a good idea. All other programs that need that lock will stop and wait for the lock to be released. You therefore should obtain your lock, do your locking-sensitive code, and then release your lock.

Reading and Writing with a Lock

Now is a good time to show the text file database `readdata()` and `writedata()` functions being used with file locking. To do this, all you'll need is a semaphore file and the `get_lock()` and `release_lock()` subroutines from the previous section.

The first part of [Listing 15.4](#) is the locking code from the previous section.

Listing 15.4. Demonstration of Text File I/O with Locking

```
1: #!/usr/bin/perl -w
2:
3: use strict;
4:
5: use Fcntl qw(:flock);
6:
7: # Function to lock (waits indefinitely)
8: sub get_lock {
9:     open(SEM, ">$semaphore_file")
10:    || die "Cannot create semaphore: $!";
11:    flock(SEM, LOCK_EX) || die "Lock failed: $!";
```

```
12:  }
13:
14: # Function to unlock
15: sub release_lock {
16:     close(SEM);
17: }
18:
19: sub readdata {
20:     open(PH, "phone.txt") || die "Cannot open phone.txt $!";
21:     my(@DATA) = <PH>;
22:     chomp(@DATA);
23:     close(PH);
24:     return(@DATA);
25: }
26: sub writedata {
27:     my(@DATA) = @_;
28:     open(PH, ">phone.txt") || die "Cannot open phone.txt $!";
29:     foreach(@DATA) {
30:         print PH "$_\n";
31:     }
32:     close(PH);    # Releases the lock, too
33: }
34: my @PHONEL;
35:
36: get_lock();
```

```
37: @PHONE1=readdata( );
38: push(@PHONE1, "Calvin 555-1012");
39: writedata(@PHONE1);
40: release_lock()
```

Most of Listing 15.4 is code you've already seen. The functions `get_lock()`, `release_lock()`, `readdata()`, and `writedata()` were all outlined earlier in this hour.

The meat of this program starts at line 34. There, a lock is taken out with `get_lock()`. The file is then read into `@PHONE1` with `readdata()`, the data is manipulated, and then written back out to the same file with `writed`

 PREV

< Day Day Up >

NEXT 

Summary

In this hour, you learned a few ways to store data between invocations of your Perl programs. First you learned about DBM files and how they can be used to tie a hash to your hard disk. Next you learned about text files and how they can be used for simple database purposes. Finally, to prevent problems with concurrent access to files, you learned about how to lock files and keep data safe.

Q&A

Q1: *Can I store the structures from Hour 13, "References and Structures," in a DBM file or a text file?*

A1: The short answer is no, not easily. The longer answer is yes, but first you need to convert the "structure" over to a string that represents the data and the structure that contains it, and you then need to use that as a value in the DBM-tied hash. The module to do so is [Data::Dumper](#).

Q2: *How do I lock DBM files?*

A2: DBM files can be locked using the semaphore locking system shown previously. Simply use the `get_lock()` and `release_lock()` functions from [Listing 15.3](#) and put these around the DBM open and close functions:

```
get_lock();

dbmopen(%hash, "foo", 0644) || die "dbmopen: $!";

$hash{newkey} = "Value";

dbmclose(%hash);

release_lock();
```

Q3: *Can I somehow just check to see whether `flock` is going to pause without actually having it stop?*

A3: Yes you can. A value that can be passed to `flock` causes it not to pause. (It is called a nonblocking `flock`.) To find out whether a `flock` would pause, put a `|LOCK_NB` after the lock type like this:

```
use Fcntl qw(:flock);

# Attempt to get an exclusive lock, but don't wait for it.

if (not flock( LF, LOCK_EX|LOCK_NB )) {

    print "Could not get the lock: $!";

}
```

You can even wait for a lock for a while, and then print a message if you don't get it eventually:

```
use Fcntl qw(:flock);

$lock_attempts = 3;

while (not flock( LF, LOCK_EX|LOCK_NB )) {

    sleep 5;  # Wait 5 seconds

    $lock_attempts--;  # Count down chances...

    die "Could not get lock!" if (not $attempts);

}
```

Workshop

Quiz

1: Keys in hashes tied to DBM files can store keys with unlimited length.

a.

True

b.

False

2: Why is it difficult to insert data into normal files?

a.

The surrounding data has to be moved to accommodate it.

b.

Normal files cannot be opened for reading and writing at the same time.

c.

The file must be locked while the file is being edited.

3: Locking files is covered in which section of the FAQ?

Answers

A1: b. The total combined length of a key and value in a DBM file is 1,024 characters, by default.

A2: a. Data cannot easily be moved "up" and "down" within a file, so moving the surrounding data can be exceedingly difficult. Choice c is also true, but only when the file may be used by more than one program at a time.

Activities

- Write a small program to update a counter in a file so that each time the program is run, the counter increases by 1. Remember to use locking in case multiple instances of your program are run at the same time.

 PREV

< Day Day Up >

NEXT 

Hour 16. The Perl Community

What You'll Learn in This Hour:

- A bit about Perl's history
- What the CPAN is, and how you can use it
- Where to go for help from here

This hour, you get to take a break. So, sit down in a comfortable spot with something to nibble on, and learn a bit about Perl's history and culture.

You may have expected this hour to be an appendix or an introduction, but those sections of any book are often overlooked. To use Perl to its fullest potential, you must understand at least a little about the Perl community.

Knowing what makes the Perl community tick will help you understand what resources are available to you, why they're there, how they work, and why Perl is what it is. Many resources are out there to help you, and this hour will help you find them.

What's Perl All About, Anyway?

To get a sense of the culture of Perl, how it works, and what resources are available to you, it's necessary to learn what makes Perl tick.

A Brief History of Perl

In 1988, the Internet was a different place. First, it was a lot smaller and looked much different than it does today. The Internet, at the time, had roughly 60,000 machines on it. Today there are hundreds of millions.

At the time, the World Wide Web didn't exist. It wouldn't be conceived until 1991 at CERN in Europe, and the first graphical browser, Mosaic, wouldn't be invented until 1993.

Most of the traffic on the Internet was textual. Usenet news provided a messaging system so that interest groups could keep in touch. Email strongly resembled what it is today—that is, mostly textual. File transfers and remote logins rounded out the traffic on the Internet.

In January 1988, Larry Wall announced that he had just written a replacement for the `awk` and `sed` tools under Unix; he called it "Perl." The original manual for Perl described it as follows:

Perl is a interpreted language optimized for scanning arbitrary text files, extracting information from those text files, and printing reports based on that information. It's also a good language for many system management tasks. The language is intended to be practical (easy to use, efficient, complete) rather than beautiful (tiny, elegant, minimal). It combines (in the author's opinion, anyway) some of the best features of C, sed, awk, and sh, so people familiar with those languages should have little difficulty with it. (Language historians will also note some vestiges of csh, Pascal, and even BASIC-PLUS.) Expression syntax corresponds quite closely to C expression syntax. If you have a problem that would ordinarily use sed or awk or sh, but it exceeds their capabilities or must run a little faster, and you don't want to write the silly thing in C, then perl may be for you. There are also translators to turn your sed and awk scripts into perl scripts.

Perl version 2 was released in June 1988, and it looked very much like modern-day Perl: You would recognize and be able to use most of the features of Perl 2. It was, and is, a rich, fully functional programming language. As the description said, Perl's features at the time were geared toward text processing and system programming tasks.

For Perl, 1991 was a banner year. In January, the first edition of *Programming Perl* by Larry Wall and Randal Schwartz was published. This book was (and remains, in its later editions) the definitive reference book for the Perl language. On the pink cover was featured a camel, the official mascot of the Perl language. (It's not a handsome animal, but it's reliable, trustworthy, and very utilitarian.)

The first edition of *Programming Perl* coincided with the release of Perl version 4. This version of Perl was the first widely distributed version of Perl, and remnants of it may still be found today in far-flung corners of the Net, despite the fact that it was last patched in 1992. You probably shouldn't use it if you run across it.

In October 1994, Perl version 5 was released. Features such as private variables, references, modules, and objects—which you haven't seen yet—were introduced. In October 1996, the second edition of *Programming Perl* ("The Blue Camel") was released, documenting these new features.

Open Source

One of the reasons for Perl's success has to do with how the Perl language is developed and distributed. The Perl interpreter is a piece of Open Source software. **Open Source** is the new term given to an old concept among software programmers: freely redistributable software. This software can be given away for free, and the source code for the software can be viewed, fixed, and adapted by anyone who wishes to change it. Other software packages that follow this model are the Linux and FreeBSD operating systems, the Apache Web server, and the Firefox Open Source browser.

Using the Open Source model is a very effective way to develop software. Because the code is written by volunteers, unnecessary code tends not to be included in the package. Features that are deemed necessary are proposed and included. The software's quality is very good because everyone with an interest in the software has the right, and duty, to watch the development carefully to look for bugs. The more people who look at the code, the fewer bugs tend to survive.

By the Way

Eric S. Raymond has written an excellent series of essays on the Open Source model of development that discuss why it works so well, why it's economically beneficial, and how it developed. The first essay, "The Cathedral and the Bazaar," is a good introduction to how the Open Source development model works. The URL for these essays is listed in the "[Other Resources](#)" section of this hour.

The Perl interpreter is copyrighted by Larry Wall; he owns it, and it's his to do with as he pleases. However, like most software, Perl can be licensed for use. A software license describes how the software can be used and distributed; it's the fine print you find when opening store-bought software. Larry Wall offers two different licenses for you to choose from: the GNU General Public License and the Perl Artistic License. After reading both, you can choose to follow the terms of either agreement when redistributing Perl.

The text of both licenses is quite lengthy, but let me summarize briefly:

- You may redistribute the source code for the Perl interpreter, but you must duplicate the copyright notices.
- You may modify the original source code if your changes are clearly marked as your own and you either give away the changes or clearly indicate that it is not the standard version of Perl. You must provide the standard version as well.
- You can charge a reasonable fee for distributing Perl. You can charge a fee for support, but you cannot sell Perl itself. You can include Perl in other products you do sell.
- Programs that are written using Perl are not subject to this license.
- No warranties are made about Perl.

You shouldn't rely on a brief summary like this for legal purposes; I provided it only to give you a feel for the licenses themselves.

Watch Out!

Before you attempt to include Perl within another package, it is important that you read the licenses yourself and determine whether what you're doing follows either of the licenses. The Perl Artistic License is included with every Perl distribution as a file named *Artistic*. You can view the GNU General Public License at <http://www.gnu.org>.

The licenses allow Perl to be developed and enhanced in an open forum. In this way, all the source code for Perl is visible to whoever wants to read it and suggest changes. This approach encourages good programming and prevents getting locked into proprietary, hidden, and obscured software solutions.

The Development of Perl

The development of the Perl interpreter, the language, and the modules that go along with it in the standard distribution takes place on a mailing list where the developers of Perl propose changes, examine bug reports, and debate what changes should be made to the Perl source code.

Anyone is welcome to participate in this process; that's what Open Source is all about. However, to prevent chaos, the changes are filtered through a group of core developers who approve or reject the changes and maintain the core of Perl's development. Changes are evaluated on the basis of what's good for Perl, how useful the changes are, and whether anyone has attempted the changes successfully. Larry Wall, who oversees this process, acts in the role of a benevolent dictator, allowing the beneficial changes to occur and vetoing changes he considers harmful.

The current release of Perl is 5.8.6. This reflects a numbering scheme for releases of: *major.minor.patchlevel*. The **major** number is the version of the language (5), the **minor** number is the version of the interpreter (8), and the **patchlevel** (6) indicates the patch level of this interpreter's code.

The even-numbered *minor* releases are production quality, the odd ones (5.7, 5.9) are unstable development releases. The next significant production release of Perl will be called 5.10.

The Comprehensive Perl Archive Network (CPAN)

Perl offers additional modules to expand your development environment even further. These modules are contained in the **Comprehensive Perl Archive Network (CPAN)**.

What Is CPAN?

CPAN is a large collection of Perl modules, programs, and documentation, contributed to the Perl community by the volunteers who wrote them.

The list of modules available in CPAN is extensive. CPAN was formed in 1995, and at the time of this writing, contains more than 7,000 modules available for installation. The modules cover a wide range of programming problems. [Table 16.1](#) provides a short list to give you a taste of what's in CPAN.

Table 16.1. Modules in CPAN

Module Group	General Description
Tk	Graphical interface for Perl programs; specific toolkit modules are available for accessing specialty graphics libraries such as the Win32 API, Gtk, Gnome, Qt, or the X11 toolkit
Net::*	Networking modules: interfaces for Mail, Telnet, IRC, LDAP, and 40+ others
Math::*	More than 30 modules for such constructs as complex numbers, fast Fourier transforms, matrix manipulation, and so on
Date::*, Time::*	Modules for converting dates/times into and out of various formats and doing manipulations on them
Data::*, TRee::*	Modules for manipulating data structures such as linked-lists and B-trees
DBI::*	Generic interface to databases
DBD::*	Interface to commercial and free databases such as Oracle, Informix, Ingres, ODBC, Msqli, MySQL, Sybase, and many others
Term::*	Fine control over text-mode screens such as DOS Command window or Unix xterms
String::*, Text::*	Dozens of modules for parsing text and formatting text
CGI::*,	Modules for producing, serving, fetching, and parsing
URI::*,	Web pages

`HTML::*`,

`LWP::*`,

`WWW::*`

`GD, Graphics::*`, `Image::*` Modules for manipulating graphics and images

`Win32::*`, `Win32API::*` Modules for manipulating Microsoft Windows

The most important point to remember is that a module is already available for solving most problems at least partially. The solutions in CPAN have been coded, tested, and reviewed by many programmers for completeness and correctness.

All the modules in CPAN are copyrighted by their respective authors, so you should read the [README](#) file that comes with each module to see the terms under which the module can be used. Most often, the modules are distributed under the same terms as Perl itself: by the Artistic License or by the GNU General Public License.

CPAN is also the name of a standard module used to help add additional modules into your Perl installation. The CPAN module is documented in [Appendix A](#), "Installing Modules."

Why Do People Contribute?

Over the last half century of computer programming, programmers have solved the same problems over and over again. Searching, sorting, communicating, reading, writing—these problems have actually changed very little since the 1950s. Some books on computer programming theory and management are still applicable after 20 or 30 years.

Programmers are driven, ultimately, by solving interesting problems. Solving the same problems over and over again ("reinventing the wheel") is not always an interesting exercise and often results in inferior solutions. A typical frustration for programmers is to spend a long time and a great deal of effort to solve a complex problem, only to discover afterward that a simple and elegant solution was possible. This leads to programmers seeking ways to share code with one another. Sharing has the interesting side effect of creating better code because other programmers may notice problems in your code that you do not.

CPAN is an effort by the Perl community to save itself from unnecessary work. The modules are there to keep you from having the frustrating experience of reinventing someone else's wheel.

The quality in most of the CPAN modules is also very good because the modules, as well as Perl, are developed under the Open Source model of production. When you install a module on your system, you automatically have the source code for the module. You can examine the source yourself and—based on the license—use portions of the source for your own programs, modify the source, and even contact the author with suggested changes.

On the surface, the idea of CPAN seems very communal, but the actual reasons that authors contribute to CPAN are widely varied. Sometimes they contribute to help others with similar problems. Sometimes they

contribute for the sake of giving to a good cause. Sometimes they do so to help gain the respect and admiration of their peers—a powerful motivator. Whatever the reason, the end result is a very large body of work that you can use in your own programs.



< Day Day Up >



Your Next Steps

After reading though the first two-thirds of this book, you should have an understanding of the basics of Perl. You have by no means learned the whole language. On my bookshelf are at least half a dozen books on the Perl language with—discounting overlapping subjects—2,300 pages, and still some topics in Perl are not included.

A single resource can't cover *all* of Perl; however, the following sections provide recommendations for your next few steps.

The resources outlined here are presented roughly in the order in which you should seek them out. There are exceptions, but for the most part, following this order will solve your problems in the quickest manner possible.

Your First Step

When you have a Perl problem, trying to figure out what to do first is tough. You're frustrated, and if you've been working on the problem for a while, you're probably upset. Take a deep breath, don't panic, and convince yourself that everything will be all right. Believe it or not, this is an important first step. Most people will become blocked on a problem after a while, and all the frustration from being blocked will prevent clear thinking, so you might wind up making things worse.

Walk away, get something to drink, calm down, and relax. You will solve this problem.

Your Most Useful Tool

The most useful tool in your Perl toolbox is Perl itself. First, you need to determine what kind of problem you have. Usually, problems fall into two categories: faulty syntax or bad and incomplete logic.

If you have a problem with syntax, you can usually break it down into two smaller problems: Either you're using some bit of Perl incorrectly, or you've made a typo. Run your program and examine the error message; it usually indicates Perl's best guess as to which line is wrong. Look at the line for any of the following:

- Does Perl's error message indicate specifically where you should look? Look there! The Perl interpreter can be your best ally in finding bugs.
- Do all the open parentheses, brackets, and braces have matches somewhere?
- Did you check your spelling carefully? Check it again. You'd be surprised at how many bugs turn out to be spelling errors.
- Did you leave something out? A comma? A period?
- Do the lines immediately before the indicated line look okay?
- If you go back to the section in this book where that particular type of syntax was addressed, can you find examples that look similar to yours?
- If you copied the code from another source, did you check someplace else for a similar piece of code? Mistakes happen.

If your Perl program runs but simply doesn't produce the right results, you probably have a problem in your logic. Before you go tearing things apart, follow these steps:

1. Make sure that the `#!` line in your program contains a `-w`.
2. Make sure you have `use strict` somewhere near the top of your program.

Many apparent logic problems turn out to be simple mistakes that `-w` and `use strict` will catch. Use them, and if you still have problems, read on.

Debug Your Program

If you're certain that the syntax of your program is right, but it's just not doing the right thing, it's time for some elementary debugging.

The first, and probably most used, technique in debugging a program is to use the lowly `print` statement. Used carefully in your programs, it can provide some runtime diagnostics as to what's going on. Look how `print` operates in this example:

```
sub foo {  
  
    my($a1, $a2)=@_;  
  
    # Diagnostic added to see if everything's OK.  
  
    print STDERR "DEBUG: Made it to foo with $a1 $a2\n"  
  
}
```

Just remember that when your program is complete, you need to take out all the debugging `print` statements. I recommend putting some kind of string in them ("DEBUG") so that you can find them all later. Also, by printing to the `STDERR` filehandle, you can separate your normal output from the diagnostics. If you include the literal symbols `__LINE__` and `__FILE__` in your diagnostic message, Perl will print the name of the current line and file.

The other approach you should try is to use the Perl debugger. You can use the debugger on almost any Perl program. Watching your program run step-by-step can be very enlightening. Instructions for using the Perl debugger are found in Hour 12, "Using Perl's Command-Line Tools."

First, Help Yourself

If your syntax is all right and your logic seems to be okay, but you're just not getting the results you want, it

might be time to seek outside help. The first place you should look for answers is in the Perl documentation.

As indicated in [Hour 1](#), "Getting Started with Perl," every Perl installation comes with a full set of documentation. For the 5.8 release, more than 2,200 pages of documentation are included with the distribution. Every module, every function, and most aspects of the Perl language are covered in this documentation, as well as a large Frequently Asked Questions list and tutorials.

To get a list of the available documentation, type `perldoc perl` at a command prompt. Each of the manual sections is listed, as well as a general description of Perl.

The Frequently Asked Questions list contains a list of the most commonly asked questions that beginners—and experts—have about the Perl language. It's worth browsing at least once, to get an idea of what kinds of questions are there, even if you don't completely understand the answers yet.

If, for some reason, you do not have the Perl documentation installed on your system or `perldoc` does not present the documentation, you should first talk to your system's administrator to find the documentation. Having the documentation correctly installed is important because the online documentation matches the version of Perl you're running. Any other documentation might have differences.

If you cannot access the online documentation, you also can find the documentation at <http://www.perl.com>.

Learn from the Mistakes of Others

Usenet is a distributed messaging system that was developed in the early 1980s and spread to the fledgling Internet. Usenet is divided into tens of thousands of discussion groups, with topics ranging from meditation, gardening, computing, and science fiction to hockey and inline skating, along with regional groups for every region of the world. A few newsgroups are specific to Perl:

<code>comp.lang.perl.announce</code>	News about new Perl releases, modules, and information
<code>comp.lang.perl.moderated</code>	Low-traffic group with moderated discussions about Perl
<code>comp.lang.perl.misc</code>	High-traffic discussion group about anything related to Perl

To read Usenet news, you need a newsreader, and newsreaders are not hard to find. You can go to any software download site and grab a newsreader. You can also go to Web Archives such as groups.google.com that mirror the Usenet newsgroups in a Web format and require only that you have a Web browser to read news.

In these newsgroups, people post questions about problems that they're having with Perl, and other people answer the questions—all on a voluntary basis. Also, discussions about general-interest topics related to Perl take place here.

An observation that has held true through my entire programming career is as follows: *There are no original problems in computing*. Any problem that you are having, someone else has had before. The trick is finding who asked the question and what answer that person found. It's very likely that at least one person has asked a

remarkably similar question to yours in one of these newsgroups.

Google maintains an online history for much of Usenet. Using its search engine, with a few well-chosen keywords, you'll likely find the answers to your question.

For example, say that you want to know how to write a Perl program to fetch a Web page. Go to groups.google.com's Advanced Search screen, and fill in the screen with the following:

Keywords: **fetch web page**

Group: **comp.lang.perl.misc**

For this example, leave all the other fields blank. When the search returns—with thousands of matches—most of the matches will have to do with the topic you asked about. You should remember a few points about the articles that you read in Usenet:

- Not all the answers are correct. Anyone can ask a question, and anyone can answer it. Read a few responses, and decide for yourself which ones seem authoritative. Your mileage will vary.
- If you're unsure about whether an answer is correct, use the answer as a starting place to check the information yourself. Go read manual pages on the topic now that you know where to look.
- Google archives news for five or six years. Answers that may have been true five years ago may not be true now.

When All Else Fails, Ask

If you've checked the online documentation, your books, and Usenet history, and you still don't have an answer to your question, you might need to ask someone.

Asking for help should be your last resort, certainly not your first. Experts are a unique medium for answering questions. They can take your badly phrased question and sometimes come up with a genuinely good solution to your problem. However, unlike the other resources previously mentioned, people do not have an unlimited capacity for answering questions. They will get tired, they will have bad days, and they will especially get tired of answering the same questions over and over.

Although it's very likely that the person you're asking may know the answer, remember that you're borrowing that person's time and experience to get the answer to your question. It's your responsibility to do a reasonable amount of searching before asking someone else.

To ask a question in Usenet, you need to use a newsreader or one of the previously mentioned Web interfaces to Usenet News. As you're assembling your question, follow these guidelines:

1.

Before you do anything else, see whether the group has a Frequently Asked Questions (FAQ) list. The

Perl groups do; it was shipped with your Perl interpreter. For any other group, search Google for the group's FAQ before posting a message.

2.

Post the question to the correct newsgroup. A general Perl language question should be posted to the group [comp.lang.perl.misc](#). A CGI-specific programming question should probably be posted to [comp.infosystems.www.authoring.cgi](#). By reading the FAQ for the group, you'll know whether you're posting to the right place.

3.

Pick a good subject line for your post. Subject lines should describe your problem, avoid useless phrases such as "help me" or "newbie question," and be descriptive but concise.

4.

Make sure that the body includes the following:

a.

A description of what you're trying to do (maybe even an explanation why)

b.

A description of what you've tried so far

c.

A description of the errors you've encountered

For example, if you post error messages or references to your code, you should also include enough of your code so that the responders can tell what's going on. If you're trying to process data, include a few lines for an example.

The body should not include the following:

a.

Large code segments

b.

Postings of binaries such as .EXEs or uuencoded files

MIME attachments (instead, include your examples and code in the body of the text)

5.

Make sure you post with a valid email address, in case someone wants to reply but not publicly.

6.

Above all, be polite. You are asking for the kindness of strangers. No one is obligated to help you. Say "please" and "thank you," and avoid inflammatory remarks. Don't use gimmicks to try to get help—for example, "Help a poor little girl with her CGI program" or "I'll give you a free Web page if you...." These gimmicks are rude and demeaning.

After you post your article to the newsgroup, wait. Usenet news can take a few days to propagate to the entire world, and people don't always keep up and read every article. Be patient, and go on to the next problem while you wait. Whatever you do, do *not* post to Usenet again with your question too soon. Wait *at least* a couple of weeks before asking again. Rephrase your question, make sure your subject line is clear, and then try again.

The responses to your article can start immediately (within a few minutes) or could show up a month or more after you posted. As I said before, the quality of the responses will vary greatly. Some will be informative, and others will be wrong. Some respondents will be gracious and polite, and others will be exceedingly rude. It's considered good Netiquette to thank the responders and ignore any flames you've received.

Another Place to Look and Ask: PerlMonks

The PerlMonks web site (www.perlmonks.org) is another useful resource for Perl programmers. It's a fairly easy-to-navigate web message-board system that provides answers to Perl questions and discussion of Perl-related topics. Between 25 and 40 questions are answered daily in the "Seekers of Perl Wisdom" portion of the site and most questions are answered within 24 hours.

A wide range of people seek help at PerlMonks, from those just starting out with Perl to knowledgeable users who are looking for opinions on modules or articles before publication.

Like Usenet, those answering questions are all volunteers and have diverse backgrounds and experiences with Perl. As with any free resource, the quality and depth of answers can vary from response to response and from day to day.

In addition to the guidelines mentioned for Usenet, Tim Vroom who runs PerlMonks recommends a few more suggestions:

- Try to find the answer first using Google to search Usenet, the Perl FAQs, and PerlMonk's search feature.
- Realize that people are usually much more receptive to answering a question when it's clear you've put some effort into solving the problem yourself.
- Use the "Chatterbox" feature—a live chat with other visitors to PerlMonks—to ask smaller and less involved questions, or get clarification on how to use the site.

- Be sure to include any relevant code in your posting. Using <CODE></CODE> tags to surround your code will format it properly, and make it easy for others to download and test.

I've used this resource myself a few times to ask questions, and visited to answer questions on occasion, and I recommend it highly.



< Day Day Up >



Other Resources

If you want to learn more about Perl, Perl programming, and the Perl community, you might want to check out some additional resources:

- The *Perl Developer's Dictionary* by Clinton Pierce— This book serves as a reference once you've got the basics of the language down. It contains hundreds of code examples covering every topic in Perl.
- *The Perl Cookbook* by Tom Christiansen and Nathan Torkington— This comprehensive collection of various problems, examples, solutions, and commentary on hundreds of problems is written in a cookbook format. Each problem is posed, a solution is presented, and then examples and explanations about the solution are included.
- *The Perl Review*— This quarterly journal is a subscriber supported magazine about the language and community of Perl. It contains articles contributed by members of the Perl community—programmers who use Perl every day. It is supported through subscriptions, but some back issues are available for free download.

For further reading about these topics, check out the following:

Internet History: Hobbe's Internet Timeline

<http://www.isoc.org/zakon/Internet/History/HIT.html>

History of Perl: CPAST

<http://history.perl.org>

The *Perl Journal*

<http://www.tpj.com>

CPAN:

<http://www.perl.com/CPAN>

Online Documentation:

On your system, also <http://www.perl.com>

Eric S. Raymond's Open Source essays:

<http://www.catb.org/~esr/writings/>.

Summary

In this hour you learned a bit about the history of Perl and how the Open Source development model is used with Perl. You also learned about CPAN, why it exists, and who maintains it. Finally, you learned about what kinds of resources are available when you're having problems with your Perl programs.

Q&A

Q1: *If the Web was invented after Perl, why is Perl a CGI language?*

A1: Perl is a CGI language for the same reason that computers are used to play games: It's not what they were invented for—they're just good at it. The next hour presents a more detailed explanation of why Perl is a good CGI language.

Q2: *I posted to Usenet and got a rude, flaming response. What should I do?*

A2: First, does that flame contain some good advice? If so, take that and throw the rudeness away. Otherwise, just ignore it and go on. Life's too short to waste time in a flame war.

Q3: *Is there an easy way to search CPAN?*

A3: Yes! The Web page at <http://search.cpan.org> contains a general search function and allows you to browse recent changes to CPAN and look through the modules by category.

Workshop

Quiz

1: Questions about CGI programming in Perl should be sent to which Usenet group first?

a.

<comp.infosystems.www.authoring.cgi>

b.

<comp.lang.perl.misc>

2: If your system doesn't seem to have the documentation, what should you do?

a.

Bug the administrator to install it

b.

Post a message to <comp.lang.perl.misc>

c.

Try a secondary source for documentation, such as <http://www.perl.com>

Answers

A1: a. <comp.lang.perl.misc> might be the second place you should consider posting CGI questions, but not the first.

A2: a and c, and probably in that order. The documentation might already be installed, and the administrator can help you find it. If not, <www.perl.com> has a copy of a recent set of documentation.

 PREV

< Day Day Up >

NEXT 

Part III: Applying Perl

[HOUR 17](#) Writing Modules

[HOUR 18](#) Object Primer

[HOUR 19](#) Data Processing

[HOUR 20](#) Perl as a Glue Language

[HOUR 21](#) Introduction to CGI

[HOUR 22](#) Forms

[HOUR 23](#) Complex Forms

[HOUR 24](#) Manipulating HTTP and Cookies

Hour 17. Writing Modules

What You'll Learn in This Hour:

- Building a module
- Scoping
- The our Statement

In [Hour 8](#), "Functions," you learned about subroutines. Subroutines allow you to take a frequently used pieces of code and assign it a name to create a small self-contained "mini-program" with its own variables.

After you've written enough programs, you'll find yourself creating the same functions over and over again for tasks. If you write a lot of file-manipulation programs, you'll probably have code to lock files and unlock them in several places; programs with user interaction (prompts and input) will have functions to display prompts, wait for answers, and reject bad input; and programs that work with a particular database might have functions for reading, writing, and editing data.

If every program contained copies of the required functions, cut-and-pasted into the program, your programs would soon become unwieldy to maintain because

- Every enhancement or bug fix would require hunting down and changing each version of each function in every program that used it.
- If you weren't perfect in your cutting-and-pasting or updating, you've now got several versions of that function in your programs.
- The programs themselves would become huge

Instead programmers have settled on the concept of **function libraries**. Every programming language and system has them. Sometimes they're called libraries, modules, or archives. For example, with operating systems, Microsoft Windows systems have dynamically linked libraries (DLLs) and Linux systems have shared libraries. Perl's method of storing collections of functions is called a **module**.

Entire books have been written on creating modules and function libraries. In this hour, you're going to explore writing modules to accomplish a few simple goals, namely making your code easier to maintain, share, and debug.

Building a Module

To understand how a module is put together, let's walk through a simple example: a small statistics library. This will re-use functions from the end of [Hour 8](#), abbreviated for space. The sample module, which we're going to call `TYPStats`, is shown in [Listing 17.1](#).

List

Scoping Revisited

In Hour 5, "Working with Files," you learned about variables and variable scope. Specifically you learned that the `my` keyword declares that a variable is local only to the program, or to any curly braces that might be around it. Look at the example in Listing 17.3.

Listing 17.3. Small Scoping Example

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  my $max = 15;
5:
6:  sub read_max {
7:      my $counter = 0;
8:
9:      while(<DATA>) {
10:         my $line = $_;
11:         if ($counter++ < $max) {
12:             print $line
13:         }
14:     }
15: }
16: read_max();
```

The variable `$max` is visible anywhere within the program, `$counter` is visible only within the subroutine, and `$line` is visible only within the while loop.

In Hour 5, I said that with `strict` in place you had to declare all variables with the `my` declaration. That was misleading. The `strict` compiler directive will also allow you to use variables not declared with `my` as long as you use the fully qualified package name shown in Listing 17.4.

Listing 17.4. Small Program Using Package and Lexical Variables

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  $main::max = 15;
5:
6:  sub read_max {
7:      my $counter = 0;
8:
9:      while(<DATA>) {
10:         my $line = $_;
11:         if ($counter++ < $main::max) {
12:             print $line
13:         }
14:     }
15: }
16: read_max();
```

Notice that `$main::max` isn't declared with `my` anywhere. This program will compile and run without errors because `$max` was only used by its fully qualified package name. These variables are called **package variables**, whereas the variables declared with `my` are called **lexical variables**.

This becomes important when your program is spread over two or more files. If every variable were declared with `my`, there would be no way to directly affect variables in other modules. (Some programming languages might consider this a good thing, but Perl thinks this is too restrictive.)

By the Way

In Hour 14, "Using Modules," in Listing 14.1 you saw the `File::Find` module use these kinds of variables. In that example, you changed the variable `$File::Find::name` which contained the name of the file that was currently being examined.

 PREV

< Day Day Up >

NEXT 

Another statement: our

Within a package, if you want to use a global (package) variable without prefixing the package name, use the `our` statement. The `our` statement allows the package variable name to be used within the module without being fully qualified (as shown in Listing 17.5), and visible outside of the package if it's fully qualified.

Listing 17.5. The `our` Statement

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  our $max = 15;
5:
6:  sub read_max {
7:      my $counter = 0;
8:
9:      while(<DATA>) {
10:         my $line = $_;
11:         if ($counter++ < $our) {
12:             print $line
13:         }
14:     }
15: }
16: read_max();
```

Forcing Your Names on Others

There is one small difference between the `File::Find` module from Hour 14 and the statistics module

from Listing 17.1 (beyond the obvious): `File::Find` gave the program a `find` function that wasn't fully qualified. In other words, for convenience, the module author of `File::Find` decided that instead of forcing everyone to write

```
use File::Find;

File::Find::find \&wanted, 'documents';
```

users of the module could just write

```
use File::Find;

find \&wanted, 'documents';
```

How did `File::Find` make its `find` subroutine part of your program? `File::Find` uses a special module called Exporter to take subroutine names like `find` and inject them into the `main` package so your program can see them.

To make your statistics module do the same thing, just add a little code to the beginning as shown in Listing 17.6.

Listing 17.6. Statistics Program with Function Name Export

```
1:  #!/usr/bin/perl -w
2:  use strict;
3:  package TYPStats;
4:
5:  use Exporter;
6:  our @ISA=qw(Exporter);
7:  our @EXPORT=qw(mean median);
8:
9:  sub mean {
```

```

10:     my(@data) = @_;
11:     my $sum;
12:     $sum += $_ foreach(@data);
13:     return @data ? ($sum / @data) / 0;
14: }
15: sub median {
16:     my(@data)=sort { $a <=> $b } @_;
17:     if (scalar(@data) % 2) {
18:         return($data[@data / 2]);
19:     }
20:     return(mean($data[@data / 2],
21:                 $data[@data / 2 - 1]));
22: }
23: 1;

```

Everything is exactly as it was before, except for the following:

Line 5: You're using the Exporter module. The Exporter module requires that you declare a couple of package variables to make it work.

Line 6: The `@ISA` variable is needed by Exporter to do its work. Notice that `@ISA` is being declared explicitly to be in this package by using the `our` statement.

Line 7: The `@EXPORT` variable contains a list of subroutine names and variables that you want to allow to be imported into the calling program. In this case, you're exporting the subroutines `mean` and `median`.

Now you can call `mean` without any qualification:

```
print mean(3, 6, 10, 51, 3, 99);
```

Exporting your module's names into other people's packages usually isn't good form. You're causing new functions or variables to suddenly appear in someone's program. If they had written a function called `mean` and you're importing a `mean` function, which one gets run? Writing modules that stomp on other program's variables is called **namespace pollution**—avoid this whenever possible.

By the Way

The particular combination of statements used in Listing 17.6 (`package`, `use`, `ISA`, `EXPORT`, and so on) is put together differently by different module authors and documentation. Some will use the keyword `our` instead of using fully qualified names, some will use `require` for Exporter, and others will not. There is more than one way to do it.

It is possible to make names visible, but only if the `use` program wants them. Instead of using `@EXPORT`, use `@EXPORT_OK`. This offers the name to the calling program, but only if it wants the names. The example in the following snippet offers up two function names, `mean` and `median`.

```
package TYPStats;

use Exporter;

our @ISA=qw(Exporter);

our @EXPORT=qw();

our @EXPORT_OK=qw(mean median);
```

If you use the module like this:

```
use TYPStats;
```

the calling program will not have a `mean` subroutine name imported for it. It will still have to use the fully qualified name `TYPStats::mean`. If you use the module like this:

```
use TYPStats qw(mean);
```

you get a `mean` subroutine, but not `median`—that is still accessible by the fully qualified name, though. If you want both, use both:

```
use TYPStats qw(mean median);
```

Now `mean` and `median` have both been imported into your program to be used without the fully qualified name.

 PREV

< Day Day Up >

NEXT 

Example: A Module to Handle Common File Information Requests

For this example, you're going to create a module to present information about a file. All of the information presented by this module is readily available in Perl through a variety of interfaces. In this module you're going to create functions for retrieving this information with convenient function names.

The module is presented in [Listing 17.7](#).

Listing 17.7. File Information Module

```
1:  #!/usr/bin/perl -w
2:
3:  package TYPFileInfo;
4:  use strict;
5:
6:  use Exporter;
7:  our @ISA=qw(Exporter);
8:  our @EXPORT=qw(bytes lines name extension modified $FileInfoName);
9:
10: our $FileInfoName = "";
11:
12: sub bytes {      # Size of the file in bytes
13:     return -s $FileInfoName;
14: }
15: sub lines {      # Number of text lines in the file
16:     my $count = 0;
17:     open(FH, $FileInfoName) || die "Can't open $FileInfoName: $!";
18:     while(<FH>) {
```

```

19:         $count++;
20:     }
21:     close(FH);
22:     return $count;
23: }

24: sub name {      # Filename portion of the path
25:     if ($FileInfoName =~ m/([\w\.\.]+)$/) {
26:         return $1;
27:     }
28:     return $FileInfoName;
29: }

30: sub extension {# Extension portion of the filename
31:     if ($FileInfoName =~ m/\.(.*?)$/) {
32:         return $1;
33:     }
34:     return "";
35: }

36: sub modified {      # Last modified time, suitable for localtime()
37:     my @stats = stat($FileInfoName);
38:     return $stats[9];  # Modified time.
39: }

40:

41: 1;

```

Lines 3–8: This is the normal preamble for establishing a module. Five subroutine names will be imported into the calling program. In addition, one variable name will be imported into the

caller: \$FileInfoName.

Line 10: The variable \$FileInfoName is cleared. All of the functions will use this for the filename to operate on.

Line 13: The -s operator is used to find the size of the file and return it.

Lines 16–22: The entire file is read, and then the number of lines read is returned.

Lines 25–28: The regular expression matches word-characters or periods at the end of the filename and returns that portion. This is an overly simplified way of matching filenames and will not work for files containing punctuation or spaces. If no match is found, the entire filename is returned.

Lines 31–34: The regular expression matches word-characters following a period at the end of the file name. Again, this is overly simplified and will not match odd characters in extensions. If no match is found, the empty string is returned.

Lines 37–38: The stat function (see Hour 10, "Files and Directories") is used to retrieve information on the file. The ninth value in the returned list is the last modified time, and that value is returned.

Line 41: This is the normal "return true" line that modules require.

The program in Listing 17.8 will use the module from Listing 17.7.

Listing 17.8. Sample Usage of the File Information Module

```
1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  use TYPFileInfo;
5:
6:  $FileInfoName = "/temp/message.txt";
7:
8:  print "\nFilename: " . name();
9:  print "\nExtension: " . extension();
```

```
10: print "\nModified: " . localtime(modified());  
11: print "\nBytes: " . bytes();  
12: print "\nLines: " . lines();
```

Line 4: This line reads the module code, importing all of the default names (five subroutine names and one variable name).

Line 6: This sets the name of the file that the functions will work on. The name `$FileInfoName` has been imported by the module, so it's now equivalent to `$TYPFileInfo::FileInfoName`.

Lines 8–12: These simply call the functions in the `TYPFileInfo` module and display their values.

For this example, I set `$FileInfoName` to the name of a file that contained a mail message. The results looked something like this:

Filename:



< Day Day Up >



Summary

In this hour you learned how to collect your code into useful groupings called modules. These modules can be shared among multiple programs. This has the benefits of allowing you to use the same code in many places without cutting-and-pasting, making the code easier to maintain by having only one copy of shared code, and grouping useful subroutines together by function.

Q&A

Q1: *How do I create many modules, and group them together by name?*

A1: Perl groups module names by using `::` to separate each grouping from least specific to most specific. For example, under `File` there are modules called `File::Find` and `File::IO`—both are related to files, but perform different functions. To create your own grouping for a module named `MyModules::File::FileInfo`, you need to create a directory called `MyModules`, and under that a directory named `File`, and under that a file named `FileInfo.pm` where your module code will go.

Q2: *I want to put my modules someplace else, separate from the programs. How can I do this?*

A2: Perl uses an internal array called `@INC` to determine where to search for modules. To look for modules in another directory, there's a helper module called `lib` that will cause perl to look elsewhere for modules. To look for modules in `\home\mymodules` use this at the beginning of your program:

```
use lib '\home\mymodules';
```

For more information, type `perldoc lib`.

Q3: *Most of my module users will import large groups of subroutines. This will make my use line very long. Is there a remedy?*

A3: Yes. You can choose to collect a group of module subroutines or variable names together and export them as a group. Use the hash `%EXPORT_TAGS` instead of (or in addition to) `@EXPORT_OK`. The keys to the hashes are group names; the values are an array reference of the names themselves. The Exporter module manual page has more information.

Workshop

Quiz

- 1:** If there is no `package` declaration in a Perl program or module, what namespace will subroutines and package variables wind up in?

a.

`main`

b.

`perl`

c.

They're not in a package; they're lexical

- 2:** Which of the following is a valid way to end a module?

a.

`-1;`

b.

`"Hello, World!"`

c.

`return 1;`

d.

All of the above

e.

None of the above

3: If you export a subroutine name from a module into a program that already has a subroutine by that same name, which one runs when you call the subroutine? (*Try it and see.*)

a.

The subroutine from the module

b.

The subroutine from the main program

c.

Neither; the program won't compile.

d.

Neither; the program will compile, but it will crash when you call the ambiguous function.

Answers

A1: a. Generally, anything not declared with `my`, and not following a package declaration, is in the namespace `main`.

A2: d. All of the above. In Perl, a module must end in a statement that has a true value. The only false values in Perl are zero, `undef`, and the empty string. The traditional way is with `1;`.

A3: b. All `use` statements are processed during compile-time. Momentarily, the module's subroutine would have taken precedence. Later in the compilation it is replaced with the local version of the subroutine.

Activities

- Expand on the File Information module to include subroutines to return the type of file (text or binary), the contents of the file (as an array), and the directory name of the file if available. This module will be revisited in the next hour, so choose good extensions.

 PREV

< Day Day Up >

NEXT 

Hour 18. Object Primer

What You'll Learn in This Hour:

- What Object Oriented Programming is all about
- The basics of constructing a class and methods
- How to design a class in Perl

In the late 1980s a trend started towards **object-oriented programming** (sometimes abbreviated as OOP). In this type of programming, instead of viewing a programming task as just variables, functions, and sequential program flow, the view was changed to classes, methods, and properties. The concepts of polymorphism, inheritance, and encapsulation became built into languages such as Smalltalk, C++, and later Java.

Thousands of books and research papers have been written on object-oriented programming. At one point, it was thought to be the mythical "silver bullet" of programming that would allow code to be reused indefinitely by many programmers, and all of programming would be reduced to simply re-arranging prebuilt components. Although it didn't quite live up to this hype, it has proven itself indispensable as a tool.

By the Way

Before beginning this hour, you will certainly want to review [Hour 13](#), "References and Structures," and make sure that you have a good handle on references in Perl. This hour depends heavily on references.

Don't let all of the big words, books, and hype scare you. Enough object-oriented programming can easily be learned in an hour to make it a useful tool in your programming toolbox.

Classes, Properties, and Methods

Object-oriented programming comes with its own set of jargon. You will have to learn a little bit of it to keep up, but only a little.

The most basic concept in OOP is the **class**. A class is simply a generic description of something. In real life cars, apples, and cats are all types of classes. Not all cats are the same, but conceptually there's such a thing as a "generic cat": mammal, quadruped, carnivore, small ears, long tails, and they sleep a lot. OOP also has the concept of a *subclass*, which is a more refined concept of the generic class—like a housecat, *felis cattus*.

The next two terms go together: an **object** is an **instance** of the class. If I wanted a concrete, physical example of something from the housecat class my cat Kirby would suffice. There's only one of him, and he embodies all the properties and functionality of the cat class, but he's a real cat and not an ideal, imaginary cat.

Objects have two more terms associated with them: **methods** and **properties**. Object methods are things that the object does. Properties are things that describe the object. My housecat would have the methods sleep, pounce, and eat; his properties would be furry, gray, and wears a blue collar.

A Thought-class: Car

Now that you've got the terminology down let's design a simple [Car](#) class from scratch. This class won't be terribly practical, but it will exercise the OOP terminology you've learned.

You're not going to design your class to inherit from other classes (a [Car](#) might inherit properties and methods from a [MotorVehicle](#) class), and you're not expecting anyone to subclass your class, so you don't have to worry about inheritance.

The first thing to do is make a list of things that a car has. For simplicity's sake we'll make this an electric car that's always on. When you push the accelerator, it goes. So these are its properties:

- A catchy model name
- A Vehicle ID Number
- How fast it's moving

Then, there are some things that a car can do. These are its methods:

- Accelerate/Decelerate
- Brake

The properties and methods of the car represent things that the user can see and do. The car itself might have properties that it knows internally, which are none of the user's business. These are properties too, but they're private properties:

- Engine Speed (RPM)

Now that you've got the idea of a car (a class), how will it work? First you'll need to create a car to use. Each class has a special method called a constructor that creates an instance of the class:

- New Car

So how do you use your class? First you'd invoke the new car method:

"Create me a car!"

And like magic you'd have a car. You could then give commands to your car:

"Car, accelerate."

"Car, tell me the current speed."

"Car, brake."

What if you decided to have a race? You'd have to create two cars and give each individual car commands. Each car would track, on its own, internal information like engine speed, how fast it's moving, and the Vehicle ID number.

 PREV

< Day Day Up >

NEXT 

Example: Implementation of Car in Perl

[Listing 18.1](#) shows how a simple car could be represented in Perl. This listing should be saved as `Car.pm`.

To create a class in Perl, a constructor must **bless** a reference into the class. What this means is a reference (it doesn't matter what type, most people use hash references) is going to be created within the **new** subroutine. The **bless** function will take that reference and associate it with the class as an instance. The **new** subroutine should then return that reference.

Listing 18.1. First Part of the Car Class

```
1: #!/usr/bin/perl -w
2:
3: use strict;
4: package Car;
5:
6: my $model_name = "RoadPerl";
7: my $vin = 1;
8:
9: sub new {
10:     my $car = {
11:         vin => $vin++,
12:         rpm => 0,
13:     };
14:     bless $car;
15: }
```

Line 4: This establishes that all the package variables mentioned within, and all the subroutines, will be in the `Car` namespace. With a class, you don't need to export any

subroutine names, which differs from the modules you saw last hour.

Lines 6–7: These variables will be visible only within this class. Because any subroutine in the class can see them, and they're not inside of your blessed reference; these are called *class variables*.

Lines 10–13: A hash reference is created. Because it will be passed to every method called by this class, you've put some information in it that's unique to each car produced. The VIN, or vehicle identification number, starts at 1 and goes up for each car you produce, and the RPM (engine speed) is initialized to 0.

Line 14: The reference is blessed into the class. The bless function returns the reference, so it's usually the last statement in the constructor.

There's no rule that says that the constructor for `Car` has to be called new; it's just convention. In fact, your class can have as many constructors as it wants, or none at all. But it's a little pointless without a constructor.

[Listing 18.2](#) continues our Car Class by providing methods to get the speed of the car and the VIN. Methods that are simply returning properties of a class are called **accessor methods**.

Listing 18.2. Property Methods in the `car` Class

```
16: sub vin {  
17:     my($self)=@_;  
18:     return $self->{vin};  
19: }  
20: sub speed {  
21:     my($self)=@_;  
22:     return $self->{rpm}/100;  
23: }
```

Line 17: When the `vin` method is called (because it is a method) the first argument to the subroutine will be a copy of the original blessed reference from line 13. Normally, the first line of the method takes the argument list from `@_` and moves it into more meaningful variable names. By tradition, the object itself is called `$self`.

Line 18: With the original reference in hand, it's a simple matter of retrieving the VIN number

from the hash reference and returning it.

Lines 21–22: Again, `$self` is set to the original reference. This time you've taken the RPM speed of the engine, done a little math on it, and returned a speed for the vehicle.

The Car Class's next two methods, shown in Listing 18.3, control the car's internals and provide the Car's actions.

Listing 18.3. Methods in the `Car` Class

```
24: sub brake {  
25:     my($self)=@_;  
26:     $self->{rpm}=0;  
27:     print "Skreetch!\n";  
28: }  
29: sub accelerate {  
30:     my($self,$value)=@_;  
31:     $self->{rpm}=$self->{rpm} + $value;  
32:     if ($self->{rpm} > 10) {  
33:         die "Engine blew up in " . $self->{vin};  
34:     }  
35:     print "Vroom!\n";  
36: }
```

Lines 25–27: If the brake method is called on a car object, this subroutine will set the motor rpm back to zero and print a message. This is done by using the `$self` reference, and manipulating the RPM value in the hash.

Line 30: If the car wants to accelerate, it has to call this method with an additional value to indicate how much acceleration to give the car. This is stashed into `$value`.

Line 31: The value passed in is used to increment the RPM of the car.

Lines 32–35: If you accelerate too many times or too much, the engine will blow up, and the program will stop. Otherwise a satisfying message will print.

Finally, Listing 18.4 provides methods to get information about the entire Car Class. Because an object instance isn't required to use these, they're called **class methods**.

Listing 18.4. Remainder of the *Car* Class

```
37: sub model_name {  
38:     return $model_name;  
39: }  
  
40: sub models_sold {  
41:     return $vin-1;  
42: }
```

Lines 38 & 41: Notice that these methods don't use `$self`. Why? Because the value returned is the same for every instance of this class. This is information that the class knows about itself and about all the instances of the class.

You can call class methods in any of the following ways:

```
Car->model_name();    # Probably the preferred style  
Car::model_name();    # Using the package name directly, see Hour 17  
$car->model_name();  # Using an instance of Car works too.
```

Using the Car Class

Now you have a car class. How do you use it? Using the class, it would look like this:

```
my $newcar = new Car;
```

So now `$newcar` contains a blessed reference into the car class. What does this mean? If you use that

reference with the method operator `->`, functions within that class will be called, like so:

```
print $newcar->vin();
```

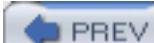
When Perl sees this it will call the `vin` subroutine inside of the `Car` class. As an added bonus it will pass a copy of that original reference as the first argument to the subroutine `vin`.

If you want your car to accelerate, you'll need to pass an additional value indicating how much:

```
print $newcar->accelerate(2);
```

Again, Perl finds the class that `$newcar` has been blessed into, finds the `accelerate` subroutine, and calls it. It will be called with two arguments: the reference (as all methods are) and the value `2`. You know that internally, it will be used to increase the car's RPM. But from our perspective, it will just print `Vroom!`.

[Listing 18.5](#) is a small program to create a sort-of race. Several cars will be created, each one will accelerate

 PREV

< Day Day Up >

NEXT 

Example: File Information Class

If you recall from last hour, you built a module to give information on files. After building the module, it had a series of defects. Each of those defects can easily be fixed by using a class instead of a regular module.

The object-oriented version of the file information module is presented in [Listing 18.6](#).

Listing 18.6. The *FileInfo* Module, as a Class.

```
1: #!/usr/bin/perl -w
2:
3: package TYPFileInfoOO;
4: use strict;
5:
6: sub new {
7:     my($class,$filename)=@_;
8:     if (! -e $filename) {
9:         die "$filename doesn't exist!";
10:    }
11:    bless { filename => $filename };
12: }
13: sub bytes { # Returns the size of the file in bytes
14:     my($self)=@_;
15:     return -s $self->{filename};
16: }
17: sub lines { # Returns the number of text lines in the file
18:     my($self)=@_;
19:     my $count = 0;
20:     open(FH, $self->{filename}) || die "Can't open $self->{filename}: $!";
21:     while(<FH>) {
22:         $count++;
}
```

```

23:      }
24:      close(FH);
25:      return $count;
26:  }
27: sub name {    # Returns just the filename portion of the path
28:     my($self)=@_;
29:     if ($self->{filename} =~ m/([\w\.\.]+)$/) {
30:         return $1;
31:     }
32:     return $self->{filename};
33: }
34: sub extension {        # Returns just the extension portion of the filename
35:     my($self)=@_;
36:     if ($self->{filename} =~ m/\.(.*?)$/ ) {
37:         return $1;
38:     }
39:     return "";
40: }
41: sub modified {        # Returns the last modified time, suitable for localtime()
42:     my($self)=@_;
43:     my @stats = stat($self->{filename});
44:     return $stats[9];  # Modified time.
45: }
46: 1;

```

This listing is virtually unchanged from 17.8. The primary difference surrounds the fact that the global variable `$FileInfoName` is now gone.

Line 7: This new constructor is a bit different than the last. In order to create a `TYPEFileInfoOO` object,

you must pass the filename in to the constructor. (See [Listing 18.7](#) for an example.) The `new` method is passed the class name as the first argument. You're really not interested in it, so just ignore it.

Lines 8–10: Because the filename is passed in during the constructor, and can't be changed anywhere else, you can check here to see if it exists. Contrast this to relying on `$FileInfoName` in the last hour.

Line 9: An anonymous hash reference is created, with one key-value pair to store the filename. This is blessed and returned (implicitly) from the subroutine.

Line 14: Because this is an OOP module, each of the subroutines will be called as methods and the first argument passed to the subroutine is the blessed reference created in the constructor—in this case you stash it in `$self`.

Line 15: To get to the original filename passed into the constructor, use `$self->{filename}`. This is used to get the size of the file.

The rest of the module continues the pattern of simply retrieving the filename from `$self->{filename}`. Remember that each instance of the class (each object) will have its own version of `$self`.

Using the File Information Class

[Listing 18.7](#) shows how the class would be used to write the same program you wrote in Hour 17, "Writing Modules."

Listing 18.7. A Reworking of Listing 17.9 with a Class

```
1: #!/usr/bin/perl -w
2:
3: use strict;
4: use TYPFileInfoOO;
5:
6: my $fileInfo = new TYPFileInfoOO("/temp/message.txt");
7:
8: print "\nFilename: " . $fileInfo->name();
9: print "\nExtension: " . $fileInfo->extension();
10: print "\nModified: " . localtime($fileInfo->modified());
11: print "\nBytes: " . $fileInfo->bytes();
12: print "\nLines: " . $fileInfo->lines();
```

Line 6: Notice that this constructor (the `new` part) takes a pathname as an argument. From now on, each

method call against the resulting object will use whatever path was passed in.

Lines 8–12: This time around, to call the functions, use the object that you've created with the appropriate method.

Virtually no change, right? Well, like the `Car` class, the power of this class doesn't come out until you have *multiple* objects in play simultaneously.

To demonstrate this, Listing 18.8 presents a program to sort a directory by file size.

Listing 18.8. Using Multiple `FileInfo` Objects at Once

```
1: #!/usr/bin/perl -w
2:
3: use TYPFileInfoOO;
4:
5: my @fileobjects;
6: foreach my $pathname ( glob("/temp/*.txt") ) {
7:     if (-d $pathname) {
8:         next;
9:     }
10:    push @fileobjects, new TYPFileInfoOO($pathname);
11: }
12:
13: foreach my $fileobj (
14:     sort { $a->bytes() <=> $b->bytes() }
15:     @fileobjects ) {
16:
17:     printf("%20s %9d %9d\n",
18:            $fileobj->name(), $fileobj->bytes(), $fileobj->lines());
19: }
```

Line 6: This uses the built-in `glob` function presented in Hour 10, "Files and Directories," to retrieve a list

of all the text files in a directory.

Lines 7–9: Directory names may also be returned by `glob`. This class isn't designed for directories, so skip them.

Line 10: Each pathname is used to construct a `TYPFileInfoOO` object. The resulting object is then accumulated in `@fileobjects`. If there are 100 text files, 100 independent objects will be created and thrown into `@fileobjects`.

Lines 13–15: The array `@fileobject` is sorted by size (the number of bytes in the file). Because you're sorting an array of objects, the sort subroutine uses `$a` and `$b` as objects, calling their `byte` method to get the size for sorting. The variable `$fileobj` is set to each object in the sorted list and the loop body is run.

Lines 17–18: `$fileobj` will take on the value of each object in the list. Because it's an object, simply call methods against the object to get the values.

The results (on my system) look something like this:

<code>chess.txt</code>	135	8
<code>sample.txt</code>	456	30
<code>outwords.txt</code>	1031	141
<code>message.txt</code>	1112	36
<code>parrot.txt</code>	1152	144
<code>outfile.txt</code>	3799	1
<code>NewSchedule.txt</code>	4119	57
<code>volleyball.txt</code>	5289	440
<code>log.txt</code>	15529	363

This program would have been much more difficult to write using the old non-OOP version of the `FileInfo` module.

Let's revisit the list of problems presented at the end of [Hour 17](#), and see if we've solved any of them.

1.

Problem: Namespace cluttering from importing names.

Solved: An OOP module doesn't need to import *any* names at all. Each object has its own methods all accessed from the object itself. Objects of different types can all share method names and there are no namespace collisions at all!

2.

Problem: Special-secret variable names to control behavior.

Solved: The constructor takes the pathname as an argument, and the dirty work is done inside of the class. The only names you have to know are method names to use with the object.

3.

Problem: Can't do concurrent operations with the module.

Solved: Because each object maintains its own internal state and structures, there's no real limit to how many independent objects you can create.

4.

Problem: With users accessing internals, the behavior becomes unpredictable.

Mostly Solved: No internal mechanisms are directly visible to the module user. Arguments get passed to the constructor and the methods, and the module can keep its internals private. In Perl, it's still possible to directly change the internals components of a class from outside of the class—but it's not recommended, and sometimes not easy.



Summary

In this hour you briefly explored object-oriented programming in Perl. Along the way you learned how to construct a simple class including constructors, methods, and properties, and you learned how to use that class.

Q&A

Q1: *Can I put functions in my class, and call them directly without an object?*

A1: Yes. In object-oriented terminology these are called **class methods**. To call a subroutine named `doit` (and retrieve its value) in a class named Widget you can do any of the following:

```
$value = Widget->doit();  
  
$value = Widget::doit();  
  
$value = doit Car;
```

The last one may look odd, but it's a piece of syntax sugar that means the same thing as the others.

Q2: *This hour missed a lot of object-oriented stuff. There was no inheritance! No polymorphism!*

A2: There's far, far too much to cover in an hour. Perl includes these manual pages dedicated to objects: perlboot, perltoot, perltootc, perlobj, and perlbot. For an entire book on Perl objects, I'd recommend *Object Oriented Perl* by Damian Conway.

Q3: *If my object has lots and lots of properties, how can I access them if I don't want to create a function for each one of them?*

Q4: If you call a method on an object, and that method doesn't exist in the class, Perl looks for a subroutine named `AUTOLOAD` and calls that instead. The first argument will be the object, the remaining arguments will be passed as well, and a special variable called `$AUTOLOAD` will be set to the `class::method` name you were trying to access. Here's a simple autoloader that could be used to implement a new method called `filename` in your `TYPEfileInfoOO` class.

```
our($AUTOLOAD); # This goes near the top of the class  
  
sub AUTOLOAD { # This can go anywhere  
  
    my($self, $value)=@_;
```

```
my( $property )=($AUTOLOAD=~m/::(.*)$/);  
die "No property $property" unless exists $self->{$property};  
  
if (defined $value) {  
  
    return $self->{$property}=$value;  
  
} else {  
  
    return $self->{$property}  
  
}  
}
```

If a property does not exist inside of `$self`, the `AUTOLOAD` subroutine dies. If it does, the property value is changed to `$value` if `$value` is defined; otherwise it's returned.

If you add more key/value pairs to the hash, this `AUTOLOAD` will give you access to them as properties.



< Day Day Up >



Workshop

Quiz

1: The first argument passed to an object method is always

a.

The class name

b.

The originally blessed reference

c.

Whatever the first argument was when you invoked the method

2: Can you have method names with the same names as built-in functions, or of already-defined subroutines?

a.

Yes, but it's difficult, as noted in [Hour 8](#), "Functions."

b.

Yes, and there are no concerns really.

c.

No, it's impossible.

3: Does the constructor subroutine in a class have to be called new?

a.

Yes

b.

No

Answers

A1: b. The reference that was blessed in the constructor is returned to each method. In this way the method gets a handle back to the object, and can use it to do its work.

A2: b. It's not really a problem because Perl can always tell that you meant the method—and not the built-in—because it will appear on the right-hand side of a -> operator.

A3: b. No. That's just the convention. If you use a different subroutine name, users of your class might get confused.

Activities

- In the Q&A section, question 3 explained the `AUTOLOAD` method. How do you think you could construct an `AUTOLOAD` method that would only be able to access *some* of the properties within the object and not others? For a hint, see the `perltoot` manual page.
- Sketch out a `DirectoryInfo` class. It would be a bit different than a `FileInfo` class. What methods would you need? What kinds of properties would you need to store?
- The `TYPEFileInfoOO` class recomputes the file size and line count each time the `byte` and `lines` methods are called. Is this a good idea? If you wanted to store those values so that subsequent calls would fetch the cached values instead of recomputing them each time, how would you do it?

Hour 19. Data Processing

What You'll Learn in This Hour:

- How to determine what kind of data you're dealing with
- Basic techniques for breaking down common data types in Perl
- How to parse table data and XML

If you're old enough, the phrase "data processing" probably conjures up images of magazine ads from the 1970s proclaiming that you could "make big money as a Data Processor," or minicomputers with whirling tape drives from the 1980s.

But that's not really what we're talking about in this hour. This hour is about taking data from its presented form and making it more useful. This takes a few different stages.

First, you have to examine the raw data to see if it can be cut apart, sliced up, stretched, and massaged into the final form that you need. Usually this step is obvious, but not to be forgotten. If you've got a CD collection, is it possible to assemble a (partial) discography for each band in the collection? Sure. Is it possible to take that and assemble a telephone directory for your company? No, because the raw data you need just isn't in there.

Next you need to pick your tools to read the data, pull it apart, and reassemble it. Our tool of choice, of course, is Perl.

Finally—and this is the part that requires some creativity—examine the data to determine *how* to pull it apart. Should you cut it into vertical slices (columns)? Horizontal slices (rows)? Make new tables and manipulate those? Do you have to glue two different sources of data together?

This hour will show you some basic techniques for pulling apart your data and reassembling it into a useful form.

By the Way

For further reading, an entire book has been written on the subject of using Perl to manipulate table data, XML data, and parse unstructured data: *Data Munging with Perl* by David Cross.

How to Look at Data

To give this chapter some context, I'm going to invent a fictional company, WidgetCo, and *you* are going to take on the role of programmer. WidgetCo is a small (10 employee) shop that sells retail replacement parts for antique bicycles. WidgetCo primarily sells through catalogs to bicycle enthusiasts and repair shops.

Your boss, Mr. Widget, has determined that his company will now be taking orders from the Internet. In fact, he would like you as Chief Programmer to streamline the process. Initially, orders will be received through email sent by someone else's web site. (The exact mechanism for taking such orders over the web will be covered in [Hours 21–24](#).) For now, your concern is dealing with incoming mail messages.

So what can you expect to see?

Unstructured Data

One kind of data representation is called *unstructured*. For example, the web site could simply have users type what they want in a text box and mail the results to you. What you'll likely see is something like the following mail message:

From: lisa@geeksalad.org

Subject: Order

I love your catalog! I have a 1962 Schwinn Astro Special that I'm trying to get working again. It used to be my older brother's bike, but he broke the seat clamp so I'll need one of those. The one on page [62](#) of the catalog in the right-hand corner should do nicely. Also, can you send me a half yard of the seat material on the next page? My VISA number is 0000-11-2222-3333 and expires in July 2006. Please ship it to 1000 Rose St in Monroe, MI 48089.

This contains all the relevant information, but there's no hope of having a computer program dissect all of this. A human is still going to have to read this message, and re-key the information into your ordering system. Unstructured data is not suitable for any kind of automation.

Table Data

Most likely, the web site will be designed to send data in a structured form. The most common representation is a table. Your email order might arrive and look like the following:

From: beth@geeksalad.org

Subject: Order

beth@geeksalad.org

404 Garden St.

Royal Oak

MI

48037

VISA 0000-11-2222-3333 06/06

12-31441 1

99-00129 1

This is more typical of the kind of data that would be sent through a web application (or any other kind of program). If you've ordered products on the web from an online retailer, you've probably received a confirmation email that looked very much like this. It's still human-readable, but it's structured in such a way that it can be easily processed by software.

Hierarchical Data

Hierarchical data refers to data that is structured, like table data, but contained in parent-child or container relationships.

The Table of Contents of this book is an example of hierarchical data. Each line in the Table of Contents isn't really meaningful unless you know its context. There are 24 entries called "[Summary](#)," but each Summary is only meaningful if you know in which chapter it appears. The Index of this book is also hierarchical. There are at least a half-dozen entries for the word "arrays," sometimes as a minor topic (data types/arrays, references/arrays, and so on) and once as a major topic. Exactly where it appears depends a great deal on what is being talked about.

You'll look at hierarchical data in a little while, specifically at a method of representing hierarchical data called XML.

Binary Data

The last category of data is *binary* data. It's usually quite unreadable by humans. Loading a binary file into your text editor results in gibberish.

Binary data is useful when interchanging data between programs where space and speed are concerns, and human readability is not. Normally to decode binary data, you have to be supplied with the data layout by the person who's also supplying the data. It's very structured—like the table data in [Figure 19.1](#)—it's just that the structure isn't apparent. JPEG images, ZIP files, and MP3 music files are all examples of binary data.

Figure 19.1. Marked-up incoming order

From: beth@geeksalad.org
Subject: Order

beth@geeksalad.org
404 Garden St.

Royal Oak
MI
48037

VISA 0000-11-2222-3333 06/06

12-31441 1
99-00129 1

email header

email address &
postal address of
customer

Credit card info

Order Information

 PREV

< Day Day Up >

NEXT 

Dealing with Table Data

To continue the example, for the moment have the recipient of the email simply cut and paste it into your Perl program. Now you'll need a Perl program to read the incoming email. [Figure 19.1](#) is the same as the previous example of table data, except it's marked up a bit to indicate what the Perl script is going to have to deal with.

The data occurs in four distinct blocks:

- The email header. You can discard this, as it will vary from email system to email system and isn't a reliable source for information.
- The email address and postal address of the customer. There are two lines for the street address, and in this example the second line is blank.
- A payment block. This seems to be three sections: a card type, card number, and card expiration.
- The order information. This consists of part numbers and a quantity.

Now that you've discovered the structure of the data (which you should confirm with the originator), you can write a Perl program to pull it apart.

Example: Email Order Taker

Initially, your email order program is going to read the pasted email and determine if the structure is valid and the order is complete. You'll add some more features later on. The program to read the email is presented in [Listing 19.1](#). Most of the program is a subroutine to create a structure to contain the order. You'll expand on it shortly to do something with the order.

This listing contains a feature you haven't seen before, the global variable `$/`. The global variable `$/` controls how Perl reads text files. If set to `undef`, the next read will continue to the end of the file. If set to the empty string (" ") the next read will read the next paragraph of data. That is, it will continue reading until the next empty line.

Listing 19.1. Email Reader

```
1: #!/usr/bin/perl -w
2:
3: use strict;
4:
5: sub read_message {
6:     {
7:         local $/= " ";
```

```
8:         my $headers = scalar <STDIN>;
9:     }
10:    my @body=<STDIN>;
11:    chomp @body;
12:    my($name, $addr1, $addr2, $city, $state, $zip)=splice(@body, 0, 6);
13:
14:    shift @body;      # Throw one blank line away
15:    $_=shift @body; # Grab the next line, should be CC Info
16:    my($cardname, $cardno, $exp);
17:    if (/^(\w+)\s(\S+)\s(\d+\/\d+)/) {
18:        $cardname=$1;
19:        $cardno=$2;
20:        $exp=$3;
21:    }
22:
23:    # Process the item lines
24:    my @items;
25:    foreach(@body) {
26:        if (! /(^[-\d]+)\s+(\d+)/) {
27:            next;
28:        }
29:        push @items, [ $1, $2 ];
30:    }
31:    return { name => $name,
32:             addr1 => $addr1, addr2 => $addr2,
```

```

33:           city  => $city,   state => $state,
34:           zip   => $zip,
35:           card  => $cardname, cardno => $cardno,
36:           expires => $exp,
37:           items => \@items
38:       };
39: }

```

Line 6: Setting `$/` to `" "` will cause the next read to slurp up an entire paragraph. By using the `local` keyword on this global value, at the end of the block `$/` will revert to its previous value.

Line 7: The first paragraph is read and discarded. This will eliminate the email headers.

Lines 10–11: The remainder of the input is read into `@body`, and the trailing `newline` characters are removed.

Line 12: The first six lines of `@body` are removed and assigned to meaningful variables.

Line 17: The value in `$_` is pattern-matched for a word, followed by something surrounded by spaces, and then a pair of numbers separated by slashes.

Lines 18–20: The matched values are assigned to meaningful variables.

Lines 26–27: Each remaining line in the body is matched to see if it's a series of dashes and digits followed by a space and then more digits. Invalid lines are thrown away.

Line 29: An array reference consisting of two elements—the item # and quantity—are pushed into `@items`. When you're done, `@items` will be an array of array references.

Lines 31–38: A hash reference is created containing all the elements of the order. The element `items` are a reference to the array `@items`.

To read the order, you'd call the subroutine with something like this:

```
my $order = read_message();
```

When that's done, use the hash reference to get to individual elements of the order:

```

print $order->{name};           # Name on the order

print $order->{items}->[0]->[0]; # First item's item number on the order

```

Now add a function to print an order. Print it to a file, and then launch an editor so that it can be printed. In this way, the order can follow the usual paper-trail through the office. Append Listing 19.2 onto Listing 19.1 to enable you to print the order.

Listing 19.2. Order Printer

```

1: sub print_order {
2:     my ($order)=@_;
3:     open(OF, ">orderfile.txt") || die "Can't open orderfile.txt: $!";
4:     print OF "Customer:\t$order->{name}\n";
5:     print OF "\t\t$order->{addr1}\n";
6:     print OF "\t\t$order->{addr2}\n";
7:     print OF "\t\t$order->{city}, $order->{state} $order->{zip}\n\n";
8:     print OF "$order->{card} $order->{cardno} $order->{expires}\n";
9:     print OF "\nItem #\tQuantity\n";
10:    foreach my $item (@{ $order->{items} }) {
11:        print OF "$item->[0]\t$item->[1]\n";
12:    }
13:    close(OF);
14:    system("notepad orderfile.txt");
15: }

```

There's no real rocket science in this program that you haven't seen before. An order file is created, filled with some data, and then closed. Afterwards, you make an external call to the operating system to open an editor. Presumably, the order can be printed from there.

By the Way

There is a cure for the bulky print statements that make up half of this function in [Listing 19.2](#). In [Hour 23](#), "Complex Forms," you'll read about here documents that allow for an easier presentation of this code.

Now to use these functions, you'd write the main body of the program like this:

```
my $order = read_message();
print_order($order);
```

When you start the program, it will wait for input. Cut and paste the order email message (or type it) into the waiting program and then press Control+D (Unix, Mac OS X) or Control+Z, Control+Z (Windows) to end the message.

This program could be modified easily to read input from a file containing the mail message, if you'd prefer.

Example: Verifier for the Email Order

For the next task, the boss has noted that the new order forms don't print the name of the item on them—just the item number, like so:

Item #	Quantity
12-31441	1
99-00129	1

He wants the item's description to print on the form to make filling the order easier. Fortunately for you, the vendor supplies a copy of the parts list as a table in a text file. (Had the vendor supplied the data in a different format such as XML or Excel, you'd soon be able to handle that as well; I discuss it later in this hour and in [Hour 20](#).) The vendor's parts list file looks something like [Listing 19.3](#).

Listing 19.3. Vendor Parts List

1231441	1' seat post clamp
3511221	brake cable set

```

3512314      brake lever assembly
3588123      brake calipers
6692818      1 1/8 carbon fork
6055232      fork oil
1600112      derailleur hanger
1619921      derailleur housing end caps
9900127      vinyl seat cover, red
9900129      vinyl seat cover, blue

```

For the rest of this example to work, [Listing 19.3](#) should be saved in a file called `partslist.txt`.

At this point, stop, step back, and take a good long look at the data here. Time spent now will make things easier later on. A few observations:

- The table is relatively straightforward: a part number, some number of spaces (or tabs), and then a description that goes to the end of the line.
- Your part numbers contain dashes; the vendor's do not. Your warehouse people like the dashes because they denote the type of item, and where to find them. So keep the dashes, but find a way to match the tables up somehow.
- Part numbers are usually unique. This case seems no different. Remember from [Hour 7, "Hashes,"](#) that tables that have a unique index are quickly and efficiently dealt with in hashes. This part table should wind up in a hash eventually.

[Listing 19.4](#) should be appended to [Listing 19.1](#) (and [Listing 19.2](#)) to make the final program.

Listing 19.4. Add Descriptions to Parts

```

1: sub add_descriptions {
2:     my($order) = @_;
3:
4:     my %parts;
5:     open(PL, "partslist.txt") || die "Can't read parts listing: $!";
6:     while(<PL>) {

```

```

7:         if (/^(\d\d)(\d+)\s+(.*)/) {
8:             $parts{"$1-$2"} = $3;
9:         }
10:    }
11:    close(PL);
12:
13:    foreach my $item (@{ $order->{items} }) {
14:        $item->[2] = $parts{$item->[0]};
15:    }
16: }

```

Lines 5–6: The parts table is opened and processed one line at a time.

Lines 7–8: Each line is matched. The first two digits are saved in `$1`, subsequent digits in `$2`, spaces are skipped, and then everything else is stored in `$3`. These values are then stored in the parts table `%parts`.

Line 13: Loop through each item on the order, assigning it to `$item`.

Line 14: Add a third element (`[2]`) to the item, which is the looked-up description from `%parts`.

In addition, Listing 19.2, line 11 should be altered to read the following:

```
print OF "$item->[0]\t$item->[1]\t$item->[2]\n";
```

So your description will print to the right of the quantity ordered.

To use your subroutines, only three function calls are needed:

```
my $order = read_message();
add_descriptions($order);
```

```
print_order($order);
```

Your order will be read, descriptions added, and then displayed in Notepad ready for printing.

 PREV

< Day Day Up >

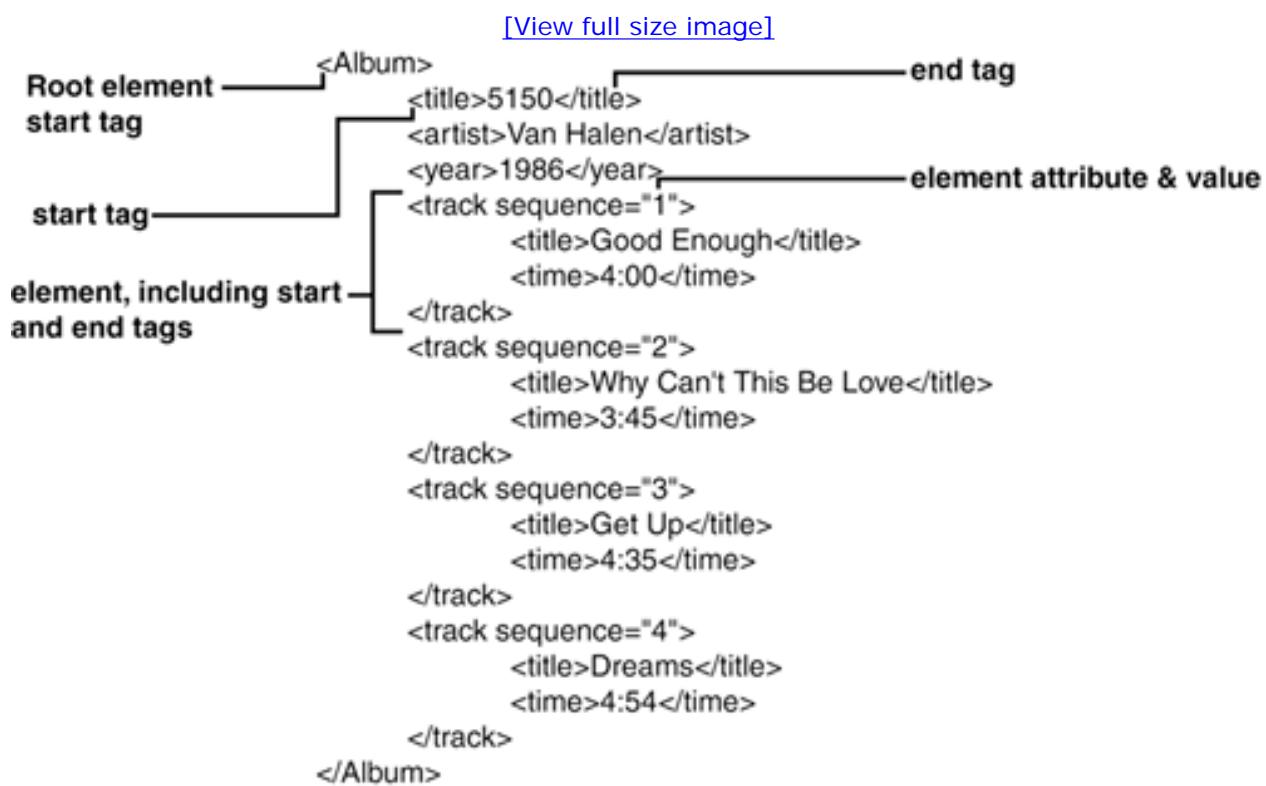
NEXT 

XML Data

Hierarchical data comes in many forms. As noted before, the index and table of contents of this book are largely hierarchical. Many methods of representing this kind of data structure have been developed over the years. The most common one in use now is called **eXtensible Markup Language** or **XML**.

XML is a format developed by the World Wide Web Consortium (W3C). It consists of a text **document**. Inside of that document are a series of **elements**, nested within each other. There is one element that encompasses the entire document called the **root element**. Each element is enclosed in a pair of **tags**—empty elements can be represented by a single start-and-end tag. Elements can have **attributes**, which are indicated in the opening tag. A sample XML document is shown in Figure 19.2.

Figure 19.2. A marked-up XML document.



By the Way

This is a *drastically* simplified description of XML, but it will suffice for your purposes. If you want all the details, visit the W3C's web site at <http://www.w3.org>.

XML documents are used for all kinds of data interchange. RSS feeds from blogs use XML as a data representation method, commercial databases can return results in XML format, Google can be queried in XML, and even Microsoft Office products can store data as XML documents. XML has become the lingua franca of data interchange since the late 1990s.

Because of this widespread use, there are *hundreds* of toolkits and methods for creating, reading, and modifying

XML data. You'll learn the most common two.

Reading XML Using Regular Expressions

XML documents, like any other kind of text, can be processed by regular expressions—almost. For example, take the document in [Figure 19.2](#). If it were saved in a disk file, you could open it and display the album title rather easily as shown in [Listing 19.5](#).

Listing 19.5. Reading the XML File, Almost

```
1: open(A, "album.xml") || die "Can't open album.xml: $!";  
2: while(<A>) {  
3:     if (/<title>(.*)</title>/) {  
4:         print $1;  
5:     }  
6: }  
7: close(A);
```

However, this doesn't quite work. Instead of seeing "5150," which is the name of the album, you'd see this:

```
5150Good EnoughWhy Can't This Be LoveGet UpDreams
```

This isn't quite what you wanted. The script picked up every instance of `<title>` and interpreted that as the album title. Could the script be fixed? Sure, you could exit after you've seen the first title, as shown in [Listing 19.6](#).

Listing 19.6. Still Flawed XML Reading

```
1: open(A, "album.xml") || die "Can't open album.xml: $!";  
2: while(<A>) {  
3:     if (/<title>(.*)</title>/) {  
4:         print $1;  
5:         exit;  
6:     }
```

```
7: }  
8: close(A);
```

This is closer, and would yield the correct results of "5150." On the surface, you've solved your XML parsing problem.

Except that you really haven't. Depending on the document, the order or elements or attributes might not be important. Whitespace might not be important. So your album XML file might be validly represented like [Listing 19.7](#).

Listing 19.7. The Album XML document, re-arranged.

```
<?xml version="1.0" encoding="UTF-8"?>  
  
<Album>  
  
<track sequence="1"><time>4:00</time>  
  
<title>Good Enough</title></track><track sequence="4"><title>Dreams</title>  
  
<time>4:54</time></track>  
  
<track sequence="2"><title>Why Can't This Be Love</title>  
  
<time>3:45</time></track>  
  
<track sequence="3"><title>Get Up</title>  
  
<time>4:35</time></track>  
  
<title>5150</title><artist>Van Halen</artist><year>1986</year>  
  
</Album>
```

Now your "working" XML reader generates this for an album title:

```
Good Enough</title></track><track sequence="4"><title>Dreams
```

Not even close. Why didn't this work? The regular expression `/<title>(. *)<\/title>/` matched everything between the `<title>` tags. In the case of [Listing 19.7](#) the first `<title>` tag shown belongs to the first tracks and the last `</title>` tag actually belongs to the track 4 disk on the same line of the input.

Could you fix this? In this specific case, yes. Listing 19.8 solves the problem in a different way.

Listing 19.8. Even More Convoluted XML Parse in Regular Expressions

```
1:  open(X, "album.xml") || die "Can't open album.xml: $!";
2:  {
3:      local $/=undef;
4:      $_=<X>;
5:  }
6:  close(X);
7:  s/<track.*</track>//gs;
8:  if (! m|<title>(.*)</title>|) {
9:      print "No match";
10:     exit;
11: }
12: print $1;
```

Line 3: You're using the `$/` like you did in the *Email Order Reader* in Listing 19.1, except by setting `$/` to `undef`, the next read will read the entire file at one time.

Line 4: The entire file's contents wind up in `$_`—newline characters and all.

Line 7: Removes all the track tags, and anything contained therein. This leaves the document looking something like this:

```
<?xml version="1.0" encoding="UTF-8"?>
<Album>
<title>5150</title><artist>Van Halen</artist><year>1986</year>
</Album>
```

Line 8: Now that all of the extra `<track>` elements have been removed, picking out the right

`<title>` element is easy—it's the only one left.

So to parse this XML document, you had to completely wreck the data. For a small document, or a task that's not going to be done very often, this inefficiency may be just fine. But it demonstrates that trying to handle XML using just regular expressions is an arms race that you'll eventually lose.

By the Way

Think you can come up with a general-purpose method for using regular expressions in XML? Think again. Consider that XML can include comments and the comment contents can look just like the surrounding XML but should be ignored; XML can also contain CDATA sections that contain arbitrary data that might or might not look like XML. Combine that with the order-of-elements problem you just explored, the fact that attribute tag order is never important, entity encodings, and that regular expressions can parse non-well-formed XML documents (which aren't really XML), and you'll realize that you just can't win.

Reading XML with `XML::Simple`

There are dozens of XML parsers for Perl, in two or three different styles. You're going to focus on `XML::Simple` because it's the easiest to use, and you don't need a thorough understanding of how parsers work to use it effectively.

To parse the `album.xml` document found in [Figure 19.2](#), use the program shown in [Listing 19.9](#). This program also contains some code to dump the resulting parse out for explanatory purposes.

Listing 19.9. Sample Using `XML::Simple` to Parse Your Album Document

```
1: use XML::Simple;
2: use Data::Dumper;
3:
4: my $ref = XMLin("./album.xml");
5:
6: print Dumper $ref;
```

The `XMLin` function takes a file and reads it, parsing the XML. It returns a reference to a structure that mirrors the XML file's structure and contents. If you recall from [Hour 13](#), "References and Structures," the `Data::Dumper` module can be used to dump a reference structure. In this case the result of the program would be the following:

```
$VAR1 = {
```

```

'track' => [
    {
        'title' => 'Good Enough',
        'sequence' => '1',
        'time' => '4:00'
    },
    {
        'title' => 'Why Can\'t This Be Love',
        'sequence' => '2',
        'time' => '3:45'
    },
    {
        'title' => 'Get Up',
        'sequence' => '3',
        'time' => '4:35'
    },
    {
        'title' => 'Dreams',
        'sequence' => '4',
        'time' => '4:54'
    }
],
'artist' => 'Van Halen',
'title' => '5150',
'year' => '1986'
];

```

Each top-level element in the XML structure became a key in a hash. To get to the album title, simply use `$ref->{title}`. Where more than one element would have the same key (`TRack`), `XML::Simple` produced an array of hashes. The first track's title would be referenced as `$ref->{track}->[0]->{title}`.

Example: Extending Your Ordering System for XML Input

As a final example this hour, suppose that your largest customer, BikeCo, would like to send you orders in bulk. Their standard data interchange format is an XML document that looks like [Listing 19.10](#).

Listing 19.10. Bulk Order Sample Document

```
<order>
```

```

<po number="1123110">

    <part number="35-88123" quantity="100"/>
    <part number="35-11221" quantity="34"/>
    <part number="35-12314" quantity="66"/>

</po>

<po number="0012231">

    <part number="16-00112" quantity="40"/>
    <part number="16-19921" quantity="40"/>

</po>

</order>
```

Looking at your sample, it appears as though there's a root element called `order` with one or more purchase orders in it. Each purchase order for this company contains parts and quantities. Because you *know* who the customer is—BikeCo—you won't need the address or credit card information in the file itself; you can generate that on your own. You'll save the PO number because it might be important for the order.

A function to read this bulk order XML file and create order documents as you did earlier this hour is shown in [Listing 19.11](#).

Listing 19.11. Reading Orders in Bulk Using XML

```

1:  use XML::Simple;
2:
3:  sub bulk_order {
4:      my($orderfile)= @_;
5:
6:      my $xml = XMLin($orderfile);
7:
8:      my @orders;
9:
10:     if (ref $xml->{po} ne "ARRAY") {
```

- 10: \$xml->{po} = [\$xml->{po}];

```

11:      }
12:      foreach my $po ( @{ $xml->{po} } ) {
13:
14:          my @items;
15:
16:          if (ref $po->{part} ne "ARRAY") {
17:
18:              $po->{part} = [ $po->{part} ];
19:
20:          }
21:
22:          push @orders,
23:
24:              { name => 'BikeCo',
25:
26:                  addr1 => '1 Hippodrome Ln', addr2 => '',
27:
28:                  city => 'Flint', state => 'MI',
29:
30:                  zip => 48506,
31:
32:                  card => 'PO #', cardno => $po->{number},
33:
34:                  expires => '',
35:
36:                  items => \@items
37:
38:          }
39:
40:      }
41:
42:      return @orders;
43:  }

```

Line 6: This function takes the name of the XML file full of orders as an argument. Using that filename, you're using `XMLin` to parse the XML.

Lines 9–10: You'd like to iterate over the array of Purchase Orders in the file, but what if there's only

1? Then `XML::Simple` won't create an array of POs. This line decides that if the `$xml->{po}` structure is not an array, it will coerce it into a single-element array.

Lines 15–16: The same trick needs to be applied to the items. If there's a PO with just one item, you create an array reference with exactly one item.

Lines 18–19: Each part line in the incoming XML is converted to an array reference with the first element being the part number and the second being the quantity.

Lines 22–30: This looks very much like the ending of [Listing 19.1](#). A hash reference is created that reflects the order. There are two big differences. First, much of the order information is hard-wired for this vendor. Secondly, the orders are pushed into an array instead of being returned directly.

Line 32: The array of orders is returned to the main body of the program.

To use this new function, you'd have to program something like [Listing 19.12](#).

Listing 19.12. Using the Bulk Order Function

```
1: my @orders = bulk_order("./bulkorder.xml");
2: foreach my $order (@orders) {
3:     add_descriptions($order);
4:     print_order($order);
5: }
```

Unlike the `read_order` function presented earlier, `bulk_order` returns an array of orders. The items in the array are selected one at a time, the descriptions are added to the items, and then they're printed.

Summary

In this hour, you learned to make some basic observations on data types that your Perl programs might receive. Based on the kind of data, you also learned how to select an appropriate parser, read the data, and do something useful with it.

Q&A

Q1: *`XML::Simple` isn't reading my XML correctly! I'm getting errors about well-formed documents. What does this mean?*

A1: XML documents should be well-formed. This means that there are matching opening and closing tags for every element, that they're correctly nesting, attributes are quoted, and that there are no stray XML characters inside of the element's data. If your XML document violates any of these rules, all XML parsers are supposed to stop. That's one big difference between reading XML with a regular expression and with a real parser. Fix your XML document and try again!

Q2: *`XML::Simple` isn't returning hashes where I'd expect it to, and I can't figure out the document structure.*

A2: Use `Data::Dumper` as shown in Listing 19.8 to figure out what the structure really looks like. `XML::Simple`'s handling of various elements (should it create a hash or an array here?) can be tweaked by using arguments to `XMLin`.

Q3: *I've been given a binary file to deal with in Perl. What are my options?*

A3: Binary data's not so bad to deal with in Perl, it's just unusual and unique enough that each case has to be dealt with individually. Functions that can be useful in decoding binary data are `unpack`, `substr`, and of course, regular expressions.

Workshop

Quiz

1: What kind of data does this appear to be?

01/30/2005, 05:02, PM, 2264, Listing19_1.pl

01/26/2005, 08:26, PM, 541, Listing19_2.pm

01/26/2005, 08:24, PM, 301, Listing19_3.pm

01/27/2005, 09:01, PM, 773, Listing19_4.pm

01/27/2005, 09:20, PM, 154, Listing19_5.pl

01/27/2005, 09:40, PM, 164, Listing19_6.pl

01/27/2005, 10:19, PM, 233, Listing19_7.pl

01/30/2005, 04:27, PM, 92, Listing19_8.pl

01/30/2005, 05:05, PM, 693, Listing19_9.pm

a.

Unstructured

b.

Table data

c.

Binary data

d.

Hierarchical data

2: What kind of data does this appear to be?

Othello, The Moor of Venice

Dramatis Personae

Act I

Scene 1: Venice, A street.

Scene 2: Another street.

Scene 3: Council chamber

Act II

Scene 1: Sea port in Cyprus

Scene 2: Street

Scene 3: Hall in the castle

a.

Unstructured

b.

Table data

c.

Binary data

d.

Hierarchical data

3: What kind of data does this appear to be?

There once was a man from Nantucket,
Who kept all of his cash in a bucket,
But his daughter, named Nan,
Ran away with a man,
And as for the bucket, Nantucket...

(Published in the *Princeton Tiger*, 1924)

a.

Unstructured

b.

Table data

c.

Binary data

d.

Hierarchical data

e.

Other

Answers

A1: (b) This is table data, specifically a format called **Comma Separated Value**, or **CSV**. It's a common data-interchange format, and many tools including Microsoft Excel deal with it quite well.

A2: (d) This is hierarchical data. The meaning of each line depends heavily on where it appears within the outline. Where the scenes appear changes the meaning of the document.

A3: (a) Unstructured. As far as this hour is concerned, this is unstructured data. To be fair though, this is an example of limerick, and limericks have a definite structure. Answering (e) Other will get you partial credit.

Activities

- The functions that were written in this chapter—`read_data`, `add_descriptions`, `print_order`, and `read_bulk`—are very well suited for object-oriented programming. Design an object-oriented module that:
 - Has one constructor that takes a mail message and returns an `Order` object.
 - Has a function that takes an XML filename and returns an array of `Order` objects.
 - Has `add_descriptions` and `print_order` as methods for the `Order` object.



< Day Day Up >



Hour 20. Perl as a Glue Language

What You'll Learn in This Hour:

- How to fetch web pages in Perl
- Practical applications of XML Parsing
- [How to create a PDF document](#)
- [How to create a Microsoft Excel spreadsheet](#)

Perl was born in the Unix world where there's a tool for every job imaginable. This was called the **toolkit** approach to software development. Every Unix utility was constructed to do one thing and do it well.

Perl breaks that model, a little. Perl was partly designed as a glue language. Its purpose was to connect different tools together in a useful way. To do that it has to be a very flexible multitasker.

In this hour you'll get to play with a smattering of technologies, using Perl to glue the whole mess together into some useful tools to add to your repertoire. I'm going to approach this as a series of small, independent tasks and hopefully you'll be able to glue them into larger programs as needed.

Weather Station

For your first Perl glue feat, you're going to create your own weather application. You might be familiar with those desktop applications that run on the desktop toolbar and display the current temperature. You're going to build your own (at least, a text version) in Perl.

Your requirements will be the following:

- Your data sources must be readily available, free (no money required), and freely redistributable (no license required).
- You want this to be easy for users to use. They'll supply just the ZIP code, and you'll figure out the rest.
- You'd like at least the current conditions (rainy, cloudy, sunny, and so on) and the temperature.

Not too difficult, eh? Now you have to construct how you're going to get from the ZIP code to a weather report:

- Hourly weather observations are freely available from NOAA (National Oceanic and Atmospheric Administration) at http://www.nws.noaa.gov/data/current_obs. They're arranged by observation station names, called METAR codes. These observation stations are frequently airports.
- A list of airports, locations, and their corresponding METAR codes is available from NOAA as well at http://www.nws.noaa.gov/data/current_obs/index.xml.
- The U.S. Census Bureau maintains a database that will roughly correlate a ZIP code to latitude and longitude. You can use the latitude and longitude to find the closest airport.

So reversing your steps gives you the procedure you need to write your program:

1. Use the ZIP code to find out the rough latitude and longitude of where that ZIP code is from the U.S. Census Bureau.
2. Take the list of airports and compute the distance between each airport and the ZIP code using a Great Circle formula—this is used to find the distance between two points on the surface of the Earth given latitude and longitude. Remember the closest airport and its METAR code.
3. Use that METAR code to look up the weather at NOAA.
4. Display the information you've collected.

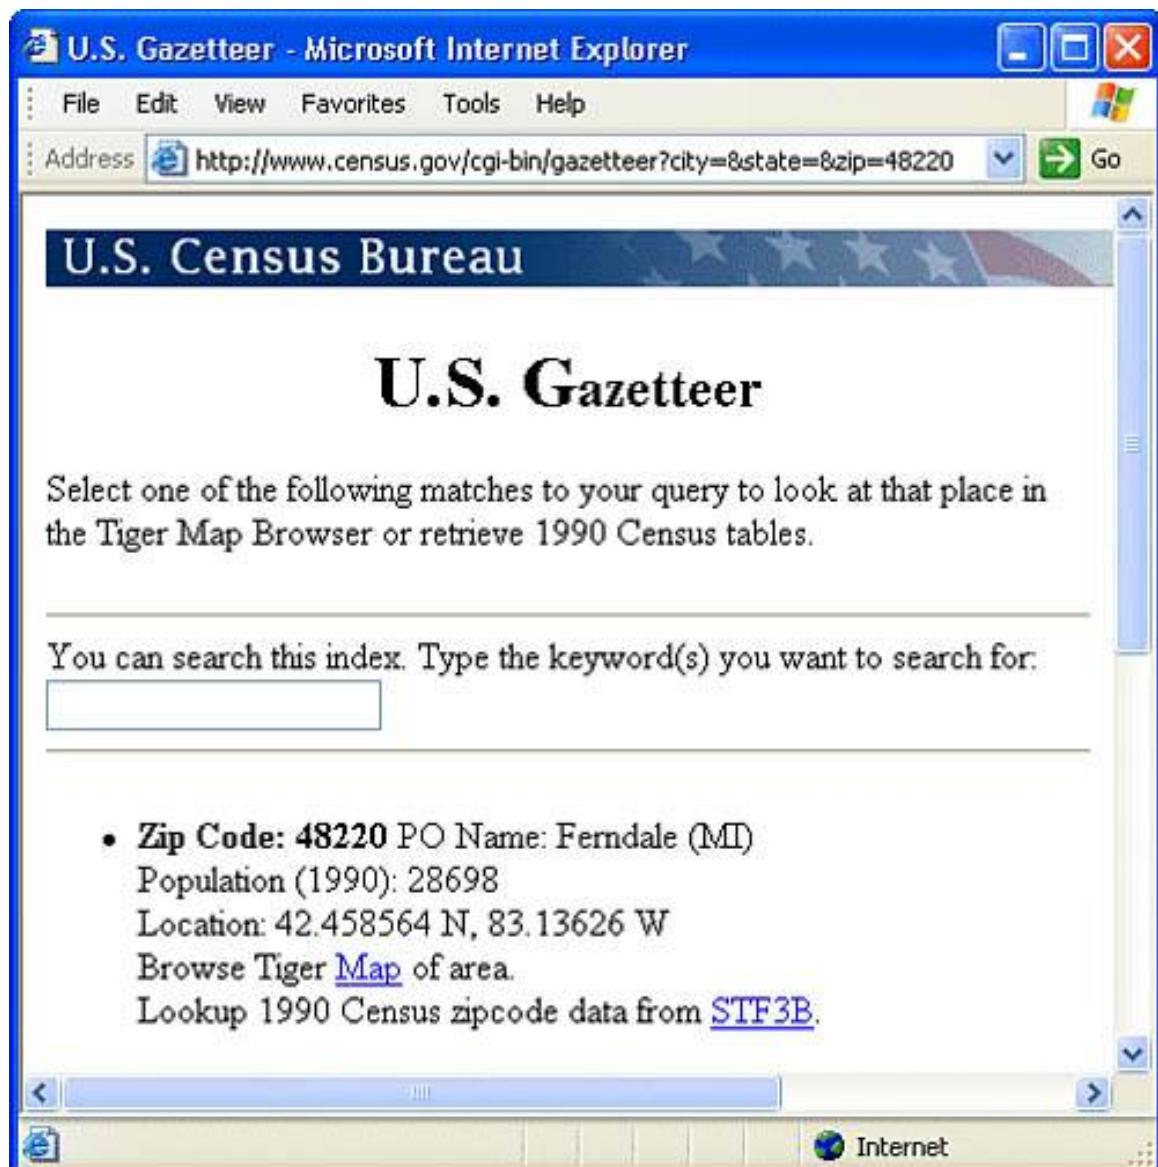
Let's do this program in three different parts. Of course, it should all be combined at the end into one large program.

Part 1: Finding Out Where You Are

The first part of your problem is going from a ZIP code to a latitude and longitude. The U.S. Census Bureau maintains the U.S. Gazetteer Place and ZIP code database. To query the database, go to <http://www.census.gov/cgi-bin/gazetteer> and enter a ZIP Code. If I use 48220, the resulting web page is shown in [Figure 20.1](#).

Figure 20.1. US Census Bureau's Gazetteer web interface.

[\[View full size image\]](#)



The first thing to notice is the URL in the browser's Address window. Notice how the ZIP code 48220 stands out in the URL? If you were to type in the URL

<http://www.census.gov/cgi-bin/gazetteer?city=&state=&zip=90210>

the web site returns information on Beverly Hills, CA.

The module `XML::Simple` has a function called `get` that works like this:

```
my $page =  
  
get("http://www.census.gov/cgi-bin/gazetteer?city=&state=&zip=48220");
```

The web page will be fetched and stored in the variable `$page`. What's `$page` going to look like? It will contain the HTML source for the web page. If you right-click in your browser and select View Source you can look at the source. It will look something like [Figure 20.2](#).

Figure 20.2. HTML output from a Gazetteer query.

```

<p>
Select one of the following matches to your query to look at the
place in the Tiger Map Browser or retrieve 1990 Census tables.<br>
<isindex>
<ul>
<li><strong>Zip Code: 48220</strong> PO Name: Ferndale (MI)<br>
Population (1990): 28698<br>
Location: 42.458564 N, 83.13626 W<br>
Browse Tiger <a href="http://tiger.census.gov/cgi-bin/mapbrowse">
Lookup 1990 Census zipcode data from <a href=http://venus.censu
<hr noshade>*Note: This dataset is derived from the Census GICS
does not contain unincorporated place names.

```

Because it's just text in \$page, you can use regular expressions to get the Location information you need. Listing 20.1 will retrieve a prompt for a ZIP code, retrieve the page, and get the location information.

Listing 20.1. Retrieving the ZIP Information

```

1:  #!/usr/bin/perl -w
2:
3:  # Part 1 Retrieve the Latitude & Longitude based on the ZIP Code
4:  #      Also, retrieve the modules you need for this Exercise
5:  use strict;
6:  use LWP::Simple;
7:  use XML::Simple;
8:  use Math::Trig;
9:
10: print "Weather for what ZIP code?";
11: my $zip = <>;
12: chomp $zip;
13:
14: my $zipdata = get("http://www.census.gov/cgibin/gazetteer?city=&state=&zip=$zip");
15:
16: if ($zipdata =~ /no matches found/) {
17:     die "Invalid zip code";
18: }

```

```
20: my ($ziplat, $ziplon) = $zipdata =~ /Location: (\d+\.\d+) N, (\d+\.\d+) W/;
21: my ($poName) = $zipdata =~ /PO Name: ([^<]+)/;
```

Lines 6–8: Pull in the modules you're going to need for this program. `LWP::Simple` will retrieve web pages, `XML::Simple` does your XML parsing, and `Math::Trig` contains the trig functions necessary to do the distance calculations in Part 2.

Line 14: Fetch the web page for the given ZIP code, and stuff the results into `$zipdata`.

Lines 16–18: This is a safety precaution in case the user types in a bad ZIP code.

Line 20: The HTML source for the web page is searched for the Location information. Figure 20.2 shows you what you're looking for. These latitude and longitude coordinates are in degrees and fractions of degrees.

Line 21: You're also pulling out the Post Office Name for the ZIP code. This is often the name of the city or neighborhood where the post office is located.

You now have the ZIP code's rough latitude and longitude information.

Part 2: Finding the Local Airport

Now that you know where the ZIP code is, you need to find the local airport. At the URL http://www.nws.noaa.gov/data/current_obs/index.xml there is an XML document that lists the NOAA observation stations and the latitude and longitude of each. A sample of the document is shown in Figure 20.3.

Figure 20.3. A Bit of the NOAA XML file showing the data structure.

```
<?xml version="1.0" encoding="UTF-8" ?>
- <wx_station_index>
  - <station>
    <station_id>KDTW</station_id>
    <state>MI</state>
    <station_name>Detroit Metropolitan Wayne County Airport</station_name>
    <latitude>42.13.53N</latitude>
    <longitude>083.19.51W</longitude>
    <html_url>http://weather.noaa.gov/weather/current/KDTW.html</html_url>
    <rss_url>http://weather.gov/data/current_obs/KDTW.rss</rss_url>
    <xml_url>http://weather.gov/data/current_obs/KDTW.xml</xml_url>
  </station>
  - <station>
    <station_id>KFNT</station_id>
    <state>MI</state>
    <station_name>Flint, Bishop International Airport</station_name>
    <latitude>42.58.00N</latitude>
    <longitude>083.44.58W</longitude>
    <html_url>http://weather.noaa.gov/weather/current/KFNT.html</html_url>
    <rss_url>http://weather.gov/data/current_obs/KFNT.rss</rss_url>
    <xml_url>http://weather.gov/data/current_obs/KFNT.xml</xml_url>
  </station>
  - <station>
```

This is a large file (more than 700KB), it doesn't change often, and you don't want to retrieve it every time you do a weather lookup. It should be saved to a file, and then the file should be referenced when a lookup is necessary. This example assumes it's stored in a file called `stations.xml`.

To find the closest airport, you're going to use this formula:

```
distance = 3963.1 * arccos( sin(Lat1) * sin(Lat2) +
    cos(Lat1) * cos (Lat2) *
    cos(Lon1-Lon2)))
```

This formula computes, in miles, the distance on the surface of the earth between two points given their latitude and longitude (expressed in radians).

If you apply this formula to every airport in the list and your ZIP code's latitude and longitude, and remember which airport is closest, you should get the most relevant METAR code.

The code in [Listing 20.2](#) takes the latitude and longitude gathered from Gazetteer and computes the distance to each airport listed in the NOAA database. The closest airport (METAR station) is recorded for later use in Part 3.

Listing 20.2. Finding the Nearest Airport

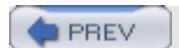
```
22:
23: # Part 2 Using the Latitude and Longitude
24: #         find the closest METAR station
25: my $metar = XMLin("./stations.xml");
26: my $bestDistance = 10000;
27: my $closestStation = undef;
28: my $rad = 57.29577951;
29:
30: my ($latmin, $lonmin);
31: foreach my $station ( @ { $metar->{station} } ) {
32:     my ($lat, $lon) = ($station->{latitude}, $station->{longitude});
33:     next if ($lat eq "NA");
34:     ($lat, $latmin) = $lat=~/(\d+)\.(\d+)/;
35:     ($lon, $lonmin) = $lon=~/(\d+)\.(\d+)/;
```

```

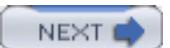
36:     $lat = (int($lat) + $latmin/60)/$rad;
37:     $lon = (int($lon) + $lonmin/60)/$rad;
38:
39:     my $greatcircle = 3963.1 * acos(
40:             sin($ziplat/$rad)*sin($lat) +
41:             cos($ziplat/$rad)*cos($lat) *
42:             cos(abs($ziplon/$rad - $lon)))
43:         );
44:
45:     if ($bestDistance > $greatcircle) {
46:         $bestDistance = $greatcircle;
47:         $closestStation = $station;
48:     }
49: }
```

Line 25: The `stations.xml` file is read in, and parsed into a reference. This will be used to traverse the structure and examine each station's information.

Lines 26–27: To compute the closest observation station, you need to keep an indicator to signify which is the closest airport yet seen. `$bestDistance` will remember how far away the most recent airport is, and `$closestStation` will remember th



< Day Day Up >



Presenting Data as PDF

The data on the weather report example is impressive, but the presentation isn't all that hot. And as you all know, presentation is very important.

For a printed document it's tough to beat PDF for presentation. PDF stands for Portable Document Format, and is a powerful and very common way of distributing electronic documents that require precise printing control. The most common tool for reading PDF documents is Adobe System's Acrobat Reader. This tool is freely available from <http://adobe.com>.

There is hardly enough room to go into all the details of creating PDF documents. The specification itself is over a thousand pages long! But a few short notes should be enough to get you started:

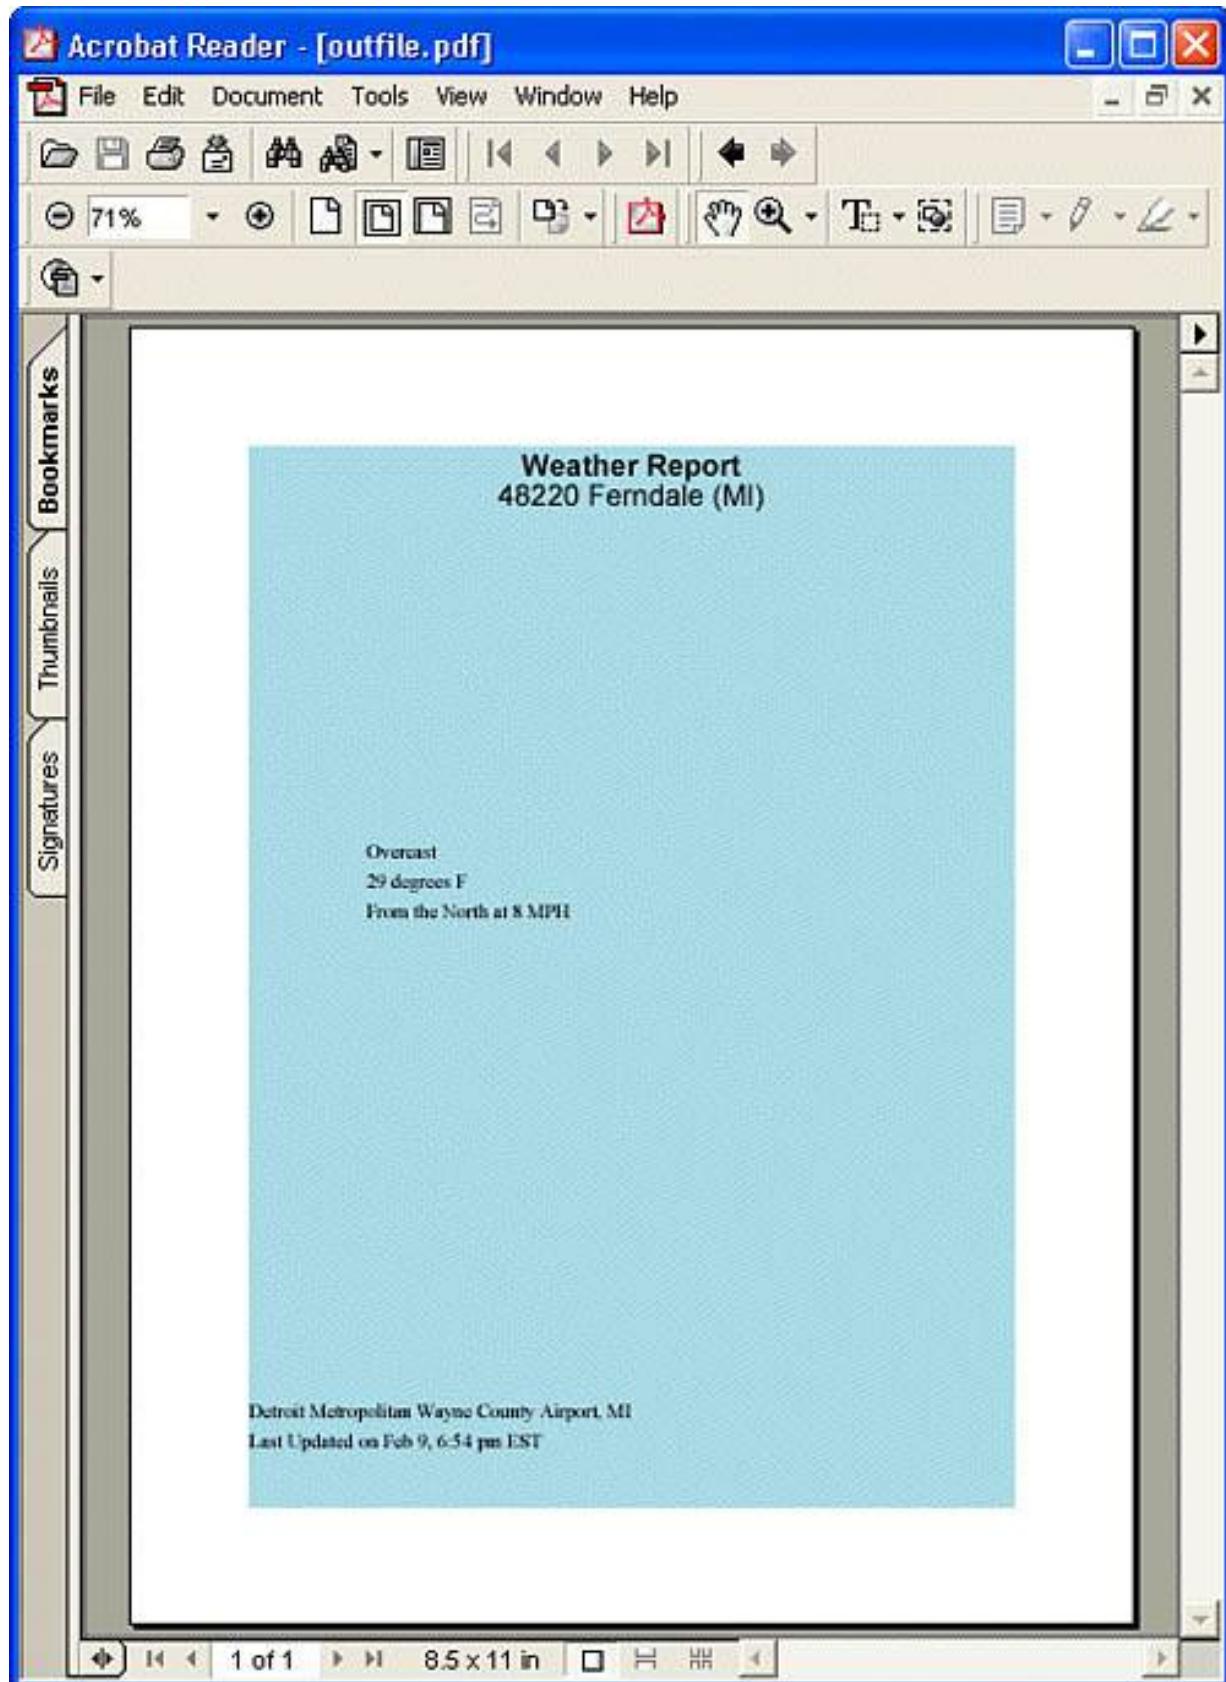
- PDF documents are object-oriented. PDF documents have properties (creator, pages, encoding schemes, and so on). Each page in the PDF document is an object.
- Page objects have properties like media size (paper size) and several properties to control clipping of text and graphics on the page (artbox, cropbox, and so on).
- Text and graphics are placed onto pages by specifying coordinates, pen types, pen colors, and font names.
- The coordinate system for PDF is upside down from what you're probably used to if you've done any graphics programming. On an 8 1/2"x11" sheet of paper, the 0,0 coordinate is at the lower left. The upper-right corner is 612,792.

There are several modules on CPAN for writing PDF files. The module you'll use is [PDF::API2](#). It's slightly more complicated than others (such as [PDF::Create](#)), but once mastered will provide a lot of flexibility.

Eventually, your weather report will look something like [Figure 20.6](#).

Figure 20.6. The weather report in all its glory.

[\[View full size image\]](#)



Example: Weather Report as PDF

To make this example work, you'll have to modify the Weather Report program in the prior example. Remove the last few lines from Listing 20.3 that look like this:

```
printf "The weather for Zip code %s %s is reported by\n", $zip, $poName;
```

```
printf "%s (%d miles away)\n%s\n", $w->{location},  
$bestDistance, $w->{observation_time};  
  
printf "%s and %s degrees Farenheit,\n", $w->{weather}, $w->{temp_f};  
  
printf "Winds are %s.\n", $w->{wind_string};
```

You can leave them in for debugging purposes, but it's not necessary. To the end of that program you'll add [Listings 20.4, 20.5, and 20.6](#).

The first part of your PDF output code appears in [Listing 20.4](#).

Listing 20.4. Setting Up the PDF Document

```
62: use PDF::API2;  
  
63: my $ipp = 1/72;           # Inches Per Point  
  
64:  
  
65: my $doc = new PDF::API2(-file => "outfile.pdf");  
  
66: my %font = (  
67:     Helvetica => {  
68:         Bold    => $doc->corefont('Helvetica-Bold', -encoding => 'latin1'),  
69:         Normal  => $doc->corefont('Helvetica', -encoding => 'latin1'),  
70:     },  
71:     Times   => {  
72:         Roman  => $doc->corefont('Times', -encoding => 'latin1'),  
73:     }  
74: );  
  
75:  
  
76:  
  
77: my $page = $doc->page();  
  
78: $page->mediabox(8.5/$ipp, 11/$ipp);
```

Line 62: This brings in the `PDF::API2` module that you'll need to produce the document.

Line 63: Because a PDF document is measured in points, but counting in points is a pain, you'll set `$ipp` to 1/72. There are 72 points per inch, so you can use inches everywhere and divide by this number to adjust to points.

Line 65: A `PDF::API2` document object is created and stored in `$doc`. When you add text and graphics to this document, it will eventually get written to a file called `output.pdf`.

Lines 66–74: You're using the object in `$doc` to help create a structure of fonts. The `corefont` method returns yet another object (a `PDF::API2::Font` object). You'll need these font objects for placing text on the page.

Line 77: To actually put content in the document, you need to create a `Page` object using `$doc->page()`. From now on, almost all calls will use the `$page`, and not the `$doc` object.

Line 79: The `mediabox()` method sets the size of the page that you need. In this case, it's 8 1/2x11 (adjusted to points).

Now that you've created the `PDF::API2::Page` object (`$page`), you can start putting things onto the page.

To put graphics on the page, you'll have to obtain a **graphics handle**, and for text a **text handle**.

Both handles are objects. For the graphics handle there are many methods for drawing shapes (ellipse, arc, box, and so on) and they're all detailed in the `PDF::API2::Gfx` manual page. For drawing shapes, simply set the color, indicate what kind of shape you want and where you want it, and then instruct the module to draw it as shown:

```
$ghandle = $page->gfx();    # Get the graphics handle  
  
$ghandle->fillcolor('lightblue');  
  
# 1-inch circle, nearly centered on the page  
  
$ghandle->circle(4/$ipp, 5/$ipp, 1/$ipp);  
  
$ghandle->fill();  # Fill it in.
```

To draw lines, set the stroke color (the color of the line), `move` your "pen" to the starting position, use the `line` method to draw lines from there, and then apply the color:

```
$ghandle = $page->gfx(); # Or, re use the previous one  
  
$ghandle->strokecolor('yellow');  
  
$ghandle->move(0, 0);  # Start at lower left-hand corner
```

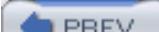
```
$ghandle->line(8.5/$ipp, 11/$ipp); # Upper right  
  
$ghandle->stroke(); # Actually draw the line
```

Text starts off with a text handle, but proceeds a bit differently. First, set the font using one of the font objects you created using the document; next, use the `TRanslate` method to move the "cursor" where you want it, set the direction (defaults to horizontal, left to right), and then use one of the text methods to set and align the text. The entire list of text methods is in the [PDF::API2::Text](#) manual page.

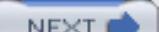
There are lots of examples of this in [Listing 20.5](#), where you continue your PDF weather report.

Listing 20.5. Writing the PDF Document

```
79: my $blue_box = $page->gfx;  
  
80: $blue_box->fillcolor( 'lightblue' );  
  
81: $blue_box->rect( 1/$ipp, 1/$ipp, 6.5/$ipp, 9/$ipp );  
  
82: $blue_box->fill;  
  
83:  
  
84: my $th = $page->text();  
  
85: $th->font( $font{'Helvetica'}->{'Bold'}, 18 );  
  
86: $th->fillcolor( 'black' );  
  
87: $th->translate( 4.25/$ipp, 9.75/$ipp );  
  
88: $th->text_center("Weather Report");  
  
89: $th->font( $font{'Helvetica'}->{'N
```

 PREV

< Day Day Up >

NEXT 

Reading and Writing Excel Spreadsheets

If PDF is the lingua franca of display formatting, then Microsoft Excel is the common tongue of spreadsheet data. Microsoft's ubiquitous Office suite means that nearly every office has some way of reading Microsoft Word documents and exchanging Microsoft Excel spreadsheets.

Most spreadsheet programs can accept a format called Comma Separated Values (CSV). This is simply a table where the values are separated by commas. An example of a CSV table is shown below.

CSV table of Academy Award winners.

Year	Best Picture	Gross	Budget	Best Actor	Best Actress	Best Director
1965	My Fair Lady	72000000	17000000	Rex Harrison	Julie Andrews	► George Cukor
1966	The Sound of Music	163214286	8200000	Lee Marvin	Julie Christie	► Robert Wise
1964	Lawrence of Arabia	70000000	15000000	Gregory Peck	Anne Bancroft	► David Lean
1960	Ben Hur	90000000	15000000	Charlton Heston	Simone Signoret	► William Wyler
1958	The Bridge on the River Kwai	40000000	3000000	Alec Guinness	► Joanne Woodward	► David Lean
1971	Patton	61700000	12000000	George C. Scott	Glenda Jackson	► Franklin Schaffner

You could just hand this table as-is to Excel, but the results are not impressive, as shown in [Figure 20.7](#).

Figure 20.7. Your CSV table imported into Excel.

[\[View full size image\]](#)

A	B	C	D	E	F	G	H
1	Year	Best Picture	Gross	Budget	Best Actor	Best Actress	Best Director
2	1965	My Fair Lady	72000000	17000000	Rex Harrison	Julie Andrews	George Cukor
3	1966	The Sound of Music	1.63E+08	8200000	Lee Marvin	Julie Christie	Robert Wise
4	1964	Lawrence of Arabia	70000000	15000000	Gregory Peck	Anne Bancroft	David Lean
5	1960	Ben Hur	90000000	15000000	Charlton Heston	Simone Signoret	William Wyler
6	1958	The Bridge at San Domingo	40000000	3000000	Alec Guiness	Joanne Woodward	David Lean
7	1971	Patton	61700000	12000000	George C. Scott	Glenda Jackson	Franklin J. Schaffner
8							
9							

What's wrong? For starters, the columns are smashed together, the movie grosses are not shown as monetary values, and there are no discernable column headers. It looks primitive.

Using Perl to Create a Spreadsheet

Let's make a spiffier spreadsheet with Perl. To follow this example, you'll have to install the module `Spreadsheet::WriteExcel`. You do not need Microsoft Excel to use the module.

Your spreadsheet should

- Have column headers highlighted in bold.
- Show monetary values as money.
- Resize the column widths so they display properly.

[Listing 20.6](#) contains a program to do just this.

Listing 20.6. Convert CSV to a Native Excel Spreadsheet

```

1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  use Spreadsheet::WriteExcel;
5:
6:  my $workbook = new Spreadsheet::WriteExcel("movies.xls");
7:  my $sheet = $workbook->add_worksheet();
8:
9:  my $moneyFormat = $workbook->add_format();
10: $moneyFormat->set_num_format('$.##0');
11: my $boldFormat = $workbook->add_format();

```

```
12: $boldFormat->set_bold();
13:
14: my $row=0;
15: my @longest=(0,0,0,0,0,0,0);
16: open(MOV, "<movies.csv") || die "Can't open movies.csv: $!";
17: while(<MOV>) {
18:     chomp;
19:     my @fields = split(/,/, $_);
20:     for(my $col = 0; $col < @fields; $col++) {
21:         if ($row == 0) {
22:             $sheet->write($row, $col, $fields[$col], $boldFormat);
23:         } else {
24:             if ($col == 2 or $col == 3) {
25:                 $sheet->write($row, $col,
26:                               $fields[$col], $moneyFormat);
27:             } else {
28:                 $sheet->write($row, $col, $fields[$col]);
29:             }
30:         }
31:
32:         if ($longest[$col] < length($fields[$col])) {
33:             $longest[$col] = length($fields[$col]);
34:         }
35:     }
36:     $row++;
37: }
38: $longest[2]+=3;
39: $longest[3]+=3;
40: for(my $i = 0; $i < @longest; $i++) {
41:     $sheet->set_column($i,$i,$longest[$i]);
42: }
```

43:

```
44: close(MOV);  
45: $workbook->close();
```

Line 6: The first step to creating a spreadsheet is calling the `Spreadsheet::WriteExcel` constructor. This will return a workbook object.

Line 7: A new sheet is created using the existing workbook. You can create as many sheets as needed, but this example needs only one sheet.

Lines 9–10: You're going to need three cell formats for this spreadsheet: normal, bold, and money. These two lines create a new format for displaying money.

Lines 11–12: This creates a bold format for cells. You'll apply these formats as needed.

Line 14: To write into the spreadsheet, you'll need a row and column. `$row` stores the current row you're writing.

Line 15: To resize the columns properly, you need to keep track of the longest field seen in each column. This array is initialized to zeroes, and as you process each field you'll edit the value to the new maximum length.

Line 17: Read a line from the CSV store into `$_`.

Line 19: The fields are split apart by commas and stored into `@fields`.

Line 20: This loop processes each column value (stored in `@fields`) one at a time. I'm using a `for (; ;)` loop instead of `foreach(array)` because you'll need the column number for formatting.

Lines 21–22: If you're writing the first row of the sheet, write the value into the sheet using the formatting in `$boldFormat`.

Lines 24–28: Otherwise, write using the format in `$moneyFormat` (if columns #2 or #3) or no formatting at all.

Lines 32–34: This stores the length of the longest field for each column in `@longest`.

Lines 38–39: The monetary column lengths are increased a bit. Add the dollar sign, and a couple of commas need to be accounted for.

Lines 40–42: Each column width is adjusted to be the length of the longest value in that column.

Line 45: The workbook is closed. This step is important; otherwise you might wind up with a corrupt Excel file.

The resulting pretty spreadsheet is shown in [Figure 20.8](#).

Figure 20.8. Nicely formatted Excel spreadsheet.

[\[View full size image\]](#)

The screenshot shows a Microsoft Excel window with the title bar 'Microsoft Excel - movies.xls'. The spreadsheet has a header row 'Year' and contains data for seven years (1965-1971). The columns represent categories: Year, Best Picture, Gross, Budget, Best Actor, Best Actress, and Best Director. The data includes film titles like 'My Fair Lady', 'The Sound of Music', and 'Lawrence of Arabia', along with their respective awards.

Year	Best Picture	Gross	Budget	Best Actor	Best Actress	Best Director
1965	My Fair Lady	\$72,000,000	\$17,000,000	Rex Harrison	Julie Andrews	George Cukor
1966	The Sound of Music	\$163,214,286	\$8,200,000	Lee Marvin	Julie Christie	Robert Wise
1964	Lawrence of Arabia	\$70,000,000	\$15,000,000	Gregory Peck	Anne Bancroft	David Lean
1960	Ben Hur	\$90,000,000	\$15,000,000	Charlton Heston	Simone Signoret	William Wyler
1958	The Bridge on the River Kwai	\$40,000,000	\$3,000,000	Alec Guinness	Joanne Woodward	David Lean
1971	Patton	\$61,700,000	\$12,000,000	George C. Scott	Glenda Jackson	Franklin Schaffner

Reading the Spreadsheet

Now that you've created a nice Excel spreadsheet, how would you take changes and put them back into the original CSV file?

Reading the Excel spreadsheet is even easier than writing it, especially because you really don't care about formatting. The `Spreadsheet::ParseExcel` module enables you to read the spreadsheet as a series of rows and columns, either with or without formatting. You do not need Microsoft Excel to use this module.

[Listing 20.7](#) shows an example to read the `movies.xls` file that you created with [Listing 20.6](#) (see [Figure 20.8](#)). This example only *prints* the spreadsheet as CSV. If you want to write it back into the original file, feel free to make the necessary modifications.

Listing 20.7. Reading an Excel Spreadsheet

```

1:  #!/usr/bin/perl -w
2:
3:  use strict;
4:  use Spreadsheet::ParseExcel;
5:
6:  my $excelObj = new Spreadsheet::ParseExcel;
7:  my $workbook = $excelObj->Parse('movies.xls');
8:  my $sheet = $workbook->{Worksheet}[0];
9:
10: my $row = $sheet->{MinRow};
11: while($row <= $sheet->{MaxRow}) {
12:     my $col = $sheet->{MinCol};
13:     my @rowval = ();
14:     while($col < $sheet->{MaxCol}) {
15:         my $cell = $sheet->{Cells}[$row][$col];
16:         push(@rowval, $cell->{Val});

```

```
17:           $col++;
18:       }
19:       print join( ',', @rowval), "\n";
20:       $row++;
21: }
```

Line 6: This creates a parser object that you'll use to read the spreadsheet.

PREV

NEXT

Summary

In this hour you used Perl to glue together several different technologies. In the span of an hour you glued together XML, web services, PDF, and Excel to produce some impressive documents. These documents are suitable for sharing with others and use data that would otherwise be just out of reach for you.

Q&A

Q1: *Reading and writing Excel documents is great. Is there a way to read and write Microsoft Word documents?*

A1: Not directly. First, you can use **Object Linking and Embedding (OLE)** to remotely control a running copy of Microsoft Word to load a document and retrieve the information you need. Of course, this only works on Windows systems with Word installed, and you'll need the [OLE::Win32](#) module. The module documentation has examples of how to do this.

The other approach is to use a common document format between Word and your programs. The most common is Rich Text Format. Microsoft Word can load and save RTF documents and there are Perl modules for reading and writing RTF documents.

Q2: *How can I put images in my PDF documents?*

A2: The document object has methods for creating images ([image_jpeg](#), [image_gif](#), and so on) and the graphics handle has a method for placing that image on a page ([gfx->image](#)). You can also rotate, skew, and create many kinds of effects for images and text. See the [PDF::API2](#), [PDF::API2::Text](#), and [PDF::API2::Gfx](#) manual pages for details.

Q3: *The web site I'm trying to retrieve information from requires login information to be sent. Can I use [LWP::Simple](#) for that?*

A3: Probably not. As your web requests get more and more complex and include cookies, encryption, and HTTP header manipulation, you're going to need the more complete—but complex—set of LWP modules ([LWP::UserAgent](#), [HTTP::Cookies](#), [HTTP::Request](#), [WWW::Mechanize](#), and so on). If you're going to do this regularly, I recommend the book *Perl & LWP* by Sean Burke.

Workshop

Quiz

1: XML documents can be parsed with regular expressions

a.

True

b.

False

2: PDF document coordinates are specified in

a.

Millimeters

b.

Points

c.

Inches

d.

Pixels

3: The top-level "component" in an Excel file `.xls` is

a.

Workbook

b.

Worksheet

c.

Grid

d.

None of the above

4: To compute the Great Circle distance in [Listing 20.2](#) you needed what module?

a.

`LWP::Simple`

b.

`XML::Simple`

c.

`Math::Trig`

d.

None. The `acos`, `sin`, and `cos` functions are built into Perl.

Answers

A1: a. But not really. As explained in Hour 19, "Data Processing," it's possible to extract specific pieces of information from an XML document with regular expressions. However, as a general-purpose parsing tool it doesn't work well.

A2: a. Points (72 points per inch).

A3: a. The hierarchy goes: a workbook has worksheets; worksheets have rows, columns, and cells.

A4: c. `Math::Trig`. If you answered (d), you were close. Perl has `sin` and `cos` as built-in functions, but not `acos` (arc cosine).

Activities

- Surf the net and find some icons (gifs, jpegs) for various weather phenomenon (rain, clouds, sun, and so on). Place these on the PDF weather report in an appropriate location.
- Almanac information for your area can be obtained by consulting the U.S. Naval Observatory web site. If you submit the following URL (all on one line):

`http://aa.usno.navy.mil/cgi-bin/aa_rstablew.pl?`

`yy0=1&yy1=45&yy2=0&xx0=-1&xx1=83&xx2=0&zz0=-1&zz1=0`

`&Compute%20Table=1`

a table of sunrise and sunset for your location will be generated. The fields are:

`yy0 1 = Latitude North, -1 = Latitude South`

`yy1 = Degrees Latitude`

`yy2 = Minutes Latitude`

`xx0 1 = Longitude East, -1 = Longitude West`

`xx1 = Degrees Longitude`

`xx2 = Minutes Longitude`

`zz0 1 = East of Greenwich, -1 = West of Greenwich`

`zz1 = Number of hours east or west of Greenwich`

Because you've computed latitude and longitude, you're using ZIP codes so it's reasonable to assume the addresses are all in the United States (west and north); the only real problem is computing how many hours west of Greenwich the ZIP code is. This is left as an exercise for the reader.

 PREV

< Day Day Up >

NEXT 

Hour 21. Introduction to CGI

What You'll Learn in This Hour:

- The basics of how the Web works
- What you need to know *before* writing CGI
- How to write your first CGI program

There's no question in anyone's mind that the explosion in popularity of the Internet is mostly due to the World Wide Web. Since the introduction of the first graphical web browser in 1993, the Internet has expanded at a phenomenal rate—going from the number of hosts doubling every 20 months around 1993 to doubling every 12 months currently. The growth of private networks—intranets—has increased even more rapidly.

The content of the web has become more sophisticated since 1993, and the users of the web expect each web page to do more than simply show **static** (unchanging) web pages. A successful web site requires dynamic web pages—web pages that provide up-to-date information. Complex pages with rapidly changing content would be almost impossible to keep current if not for **Common Gateway Interface (CGI)**.

By the Way

The next four hours require you to have some knowledge of Hypertext Markup Language (HTML). If you're unfamiliar with HTML, don't despair. It's really not hard to learn, and you won't need much to complete this book.

HTML is a markup language commonly used for constructing web pages. HTML consists of plain text with formatting codes embedded in the text to indicate how a web browser should display the text. For example, the text `HTML is <I>not</I> hard to learn.` is normal text except for the `<I></I>` markers. They are called **tags** and describe the formatting used to display the text. In this case, the word *not* should be displayed in italic by a web browser, if it can. (Remember, not all browsers are graphical.)

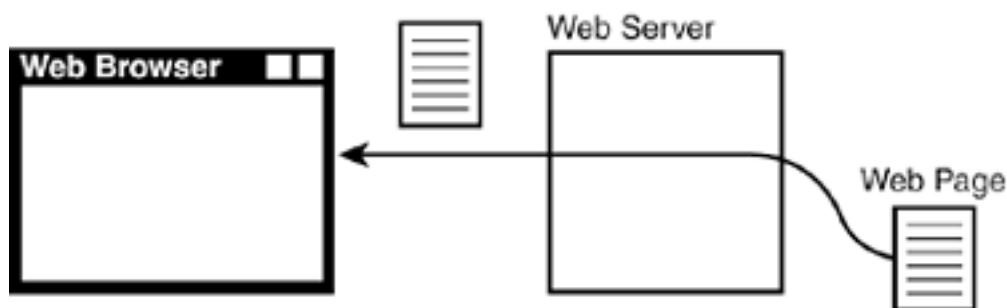
A full lesson in HTML is well beyond the scope of this book. It's not difficult; there's just a lot of material to cover. The specification for HTML is maintained by the World Wide Web Consortium (W3C) at <http://www.w3c.org>, and you can find some nice tutorials there as well. One good book on HTML is *Sams Teach Yourself HTML in 24 Hours*.

Browsing the Web

The Web, as you know it, is simply the interaction between two different systems trying to exchange data. The system that is trying to fetch a web page is known as the **client**. The client system usually runs a program called a **web browser**, such as Safari, Firefox, Internet Explorer, Opera, and so on. This is the extent of the Web that you're used to using every day. The web browser provides navigation buttons and bookmarks and is responsible for drawing web pages on your screen.

On the other end of the Web is a system known as the **web server**. This system takes the client's request for a page, retrieves the page from a local disk, and sends it to the client—your web browser. This interaction is shown in [Figure 21.1](#).

Figure 21.1. Web browser fetching a page.



Fetching a Static Web Page

A client requests a web page by examining a Uniform Resource Locator (URL) to determine the protocol, server, and request to make on that server. A typical URL might look like the following:

<http://www.google.com:80/more.html>

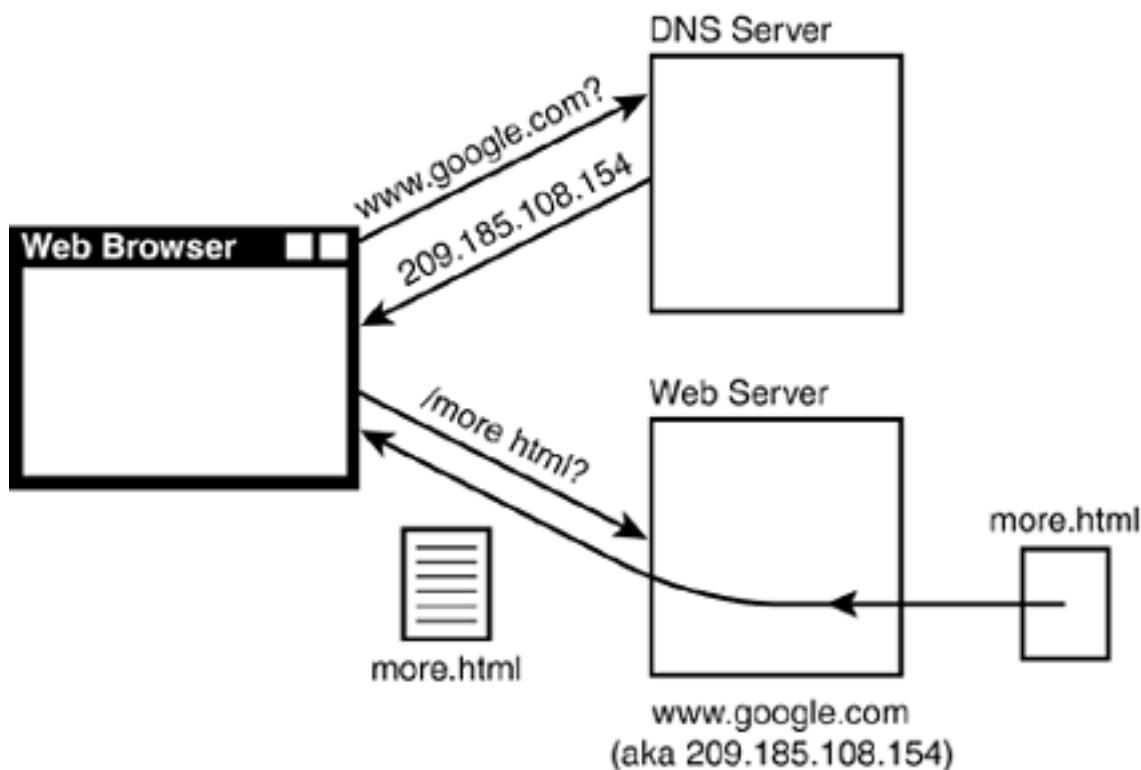
The parts of the URL can be broken down like this:

- **http**— This part is the protocol. HTTP, or Hypertext Transfer Protocol, is a protocol used for transferring web pages. You may also have seen File Transfer Protocol (ftp) or secure HTTP (https).
- **www.google.com**— This part is the name of the server—also called a **host name**—that contains the document you want. Sometimes, instead of a host name, you might see an IP address, usually written as four numbers separated by dots: 209.185.108.147. These addresses tend to be less reliable than the names, though.
- **:80**— This part is a port number that determines on which port your client and the server will connect with each other. This portion is usually optional; the protocol used determines what port will be used. http usually means "use port 80."
- **more.html**— This is the request being made on the server. Usually it's a document you want to retrieve. Sometimes it's written as a pathname, such as /archives/foo.html, or it has other characters trailing at the end, such as (?&), but essentially it is what the client needs the server to retrieve.

The client then follows these steps for http (see [Figure 21.2](#)):

1. The hostname (www.google.com) is converted to an IP address.
2. A connection is established with the server at www.google.com using the IP address and the port number.
3. The server is asked for the page more.html. The client waits for a response.
4. The server sends the response—in this case, the contents of more.html—and drops the connection to the server.
5. The client renders the response on the screen.

Figure 21.2. Requesting a page.

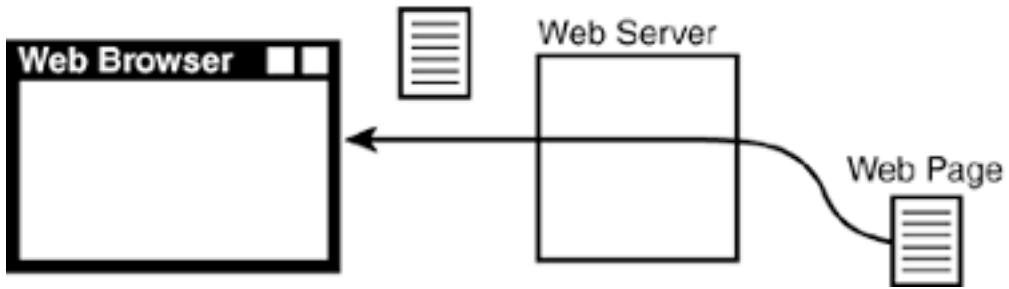


The nitty-gritty of the "conversation" between the client and the server is covered in depth in [Hour 24, "Manipulating HTTP and Cookies."](#)

Dynamic Web Content—The CGI

During a normal web page fetch, the server simply locates the document requested, retrieves it from its disk storage, and sends it to the client, as illustrated in [Figure 21.3](#).

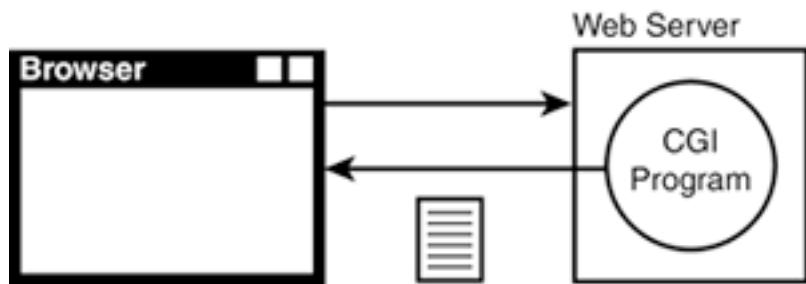
Figure 21.3. Static web page fetch.



The server in [Figure 21.3](#) doesn't process the data at all; it simply examines the request and passes the requested data back to the client.

One method to create dynamic content on the Web is through the use of CGI programs. CGI is an agreed-upon method that web servers use to run programs on the server to generate web content. When a URL indicates to a server that a CGI program should be run to generate the content, the server starts the program, the program generates the content, and the server passes the content back to the client, as illustrated in [Figure 21.4](#).

Figure 21.4. CGI script-generated web page.



Each time the client requests a page that's really a CGI program, the following occurs:

1.

The server starts a new instance of the CGI program.

2.

The CGI program generates a page, or another response, using whatever information it needs.

3.

The page is sent back to the client.

4.

The CGI program exits.

The CGI program can be any kind of program. It can be a Perl script, which is what you'll learn about here. It can also be programmed in C, the Unix shell, Pascal, Lisp, TCL, or nearly any other programming language. The fact that many CGI programs are written in Perl is a happy coincidence. Perl happens to be very well-suited to writing programs that deal with text, and the output of CGI programs is often text.

The output of CGI programs can be almost anything, however. It can be images, HTML-formatted text, Zip files, streaming video, or any other kinds of content you might find on the Web. For the most part, the CGI programs you'll be writing will generate HTML-formatted text.

By the Way

CGI is not a language; it has nothing to do specifically with Perl, it has nothing to do with HTML, and it has very little to do with HTTP. It's simply an agreed-upon interface between web servers and programs run on their behalf. This informal interface wasn't codified until October 2004 in RFC 3875. You can read about this at <http://www.ietf.org>. You'll pick up bits and pieces of these details over the next four hours.



< Day Day Up >



Don't Skip This Section

You're almost ready to write a CGI program. Almost. You need to check a few things first; otherwise, your first experience writing CGI programs will not be a happy one. It's much easier to check things before you write, rather than while you're debugging.

To use CGI, you need a web server. A common problem that new CGI programmers have is that they try writing CGI programs without having a web server properly installed. For obtaining a web server, you have a couple options: You can rent space on a commercial web server, or you can run your own. The decision is up to you, based on what you would like to pay, your bandwidth requirements, and how technically adept you are.

To obtain a commercial web server, surf the Web and find one. These commercial servers are commonly known as "web hosting" companies, and the rates they charge and the features offered can vary wildly. If you're planning on writing Perl CGI programs, you should be sure to check that Perl version 5 is supported as a CGI language. A very few web hosting companies do not support Perl as a CGI programming language or don't allow CGI at all. Avoid those companies; you have plenty of others to choose from.

You should also make sure that the web hosting company allows you to use your own scripts. Some companies claim that they allow Perl CGI but then require that you use their programs—sometimes for a fee. You should avoid those companies as well.

Still other companies will charge a fee to "review" your CGI programs before you're allowed to use them. If you choose one of these companies, you should set up your own server for testing because this "reviewing" can be expensive.

Running your own personal web server isn't too difficult, if you have some technical skills and are willing to read through the instructions. First, you need to pick a web server. If you're running Windows, you can choose from dozens of free or nearly free web servers to install. Make sure they support Perl as a CGI scripting language. A few commercial web servers are also available for Windows, notably Microsoft's Internet Information Server (IIS).

If you have a Unix machine, a few commercial web servers are available. Contact your Unix vendor for a list.

The most popular web server on the Internet, Apache, is completely free. The Apache web server is easy to install—if you have a C compiler—and fairly easy to operate if you're comfortable editing configuration files. Apache is even available precompiled for the Microsoft Windows platform. For information on Apache, go to <http://www.apache.org>.

If you run your own web server, make sure that the web server is working properly and can serve static web pages before you attempt to write CGI programs. If your web server does not serve static web pages, it's very unlikely that CGI programs will work.

You should also check the configuration for the web server to make sure that you have CGI scripts properly enabled. Not having this feature enabled is the cause of most beginners' frustrations with CGI programming.

The Checklist

Whether you run your own web server or rent space from a commercial web host, take the time to complete the following checklist—really! Write down this information; it will save you time and money.

 PREV

< Day Day Up >

NEXT 

Your First CGI Program

After all the fanfare, warnings, checklists, and runaround, you're ready to type in your first CGI program. It is shown in [Listing 21.1](#).

Type and save this program as `hello`. If, in the checklist, you are required to use a certain extension for CGI programs—as noted in the "Extension to be used for CGI programs" item—use that extension. Thus, if you're required to name your CGI programs with a `.cgi` extension, save this script as `hello.cgi`. If you're required to use `.pl`, save this script as `hello.pl`.

You *did* fill out the checklist, didn't you?

Listing 21.1. Your First CGI Program

```
1:  #!/usr/bin/perl -w
2:  use CGI qw(:standard);
3:  use strict;
4:
5:  print header;
6:  print "<b>Hello, World!</b>" ;
```

Line 1: This line is the standard `#!` line. You *must* substitute the path from the checklist item "Location of Perl on the web server" for this script to work. The `-w`, of course, enables warnings.

Line 2: The CGI module is included in the program. The `qw(:standard)` causes the standard set of functions from the CGI module to be imported into your program.

Line 3: `use strict` enforces good coding practices; this is no different for CGI programs.

Line 5: The `header` function is imported from the CGI module. It prints a standard header that the server (and the client) must see to process the output of the CGI program.

Line 6: After the header is printed, any output will appear normally to the browser. In this case, when the CGI is run, the browser will display the words `Hello, World`.

That's it!

Well, not quite. You still have to *install* this CGI program and test it. You've won only half the battle.

Installing the CGI Program on the Server

How you install your CGI program depends heavily on what kind of server you have, whether you have local access to it, whether it's a server that you can only FTP files to, and so on. The following sections describe how to install for the different scenarios.

Local Access to a File System on a Unix Web Server

Use these instructions if you're able to telnet, rlogin, or otherwise log in to the Unix web server:

1.

Put the CGI program `hello.cgi` (or `hello.pl`) onto the Unix server, perhaps with FTP.

You may have actually written it on the server with vi. That's fine, too.

2.

Move the CGI program into the correct directory by using the `mv` or `cp` command.

You should have found the correct directory in the checklist under "Location of CGI program directory."

3.

Under Unix, you need to make the program *executable*. Do so by issuing the following command:

```
chmod 755 hello.cgi
```

Use `hello.pl` instead if it is the program's name. This command makes the file writable by the owner and readable and executable by everyone else (this is correct for a CGI program).

FTP-Only Access to a Unix Web Server

Use these instructions if you have FTP-only access:

1.

Use your FTP client to put the `hello.cgi` (or `hello.pl`) program into the CGI program directory. You should have found the correct directory in the checklist under "Locati

 PREV

< Day Day Up >

NEXT 

What to Do When Your CGI Program Doesn't Work

The following sections provide a general-purpose CGI program debugging guide. Before you read through all this information to find the problem in your first CGI program, go back through the instructions to make sure you didn't skip any steps.

By the time you reach the end of this hour, you should be able to find any problems your CGI program has.

The diagnostics in these sections assume you're debugging a program named `hello.cgi`. If your program has another name, use it instead.

Is It Your CGI Program?

The first potential source of problems to eliminate is the CGI program itself. There's no sense in debugging web server configurations if the CGI program doesn't work.

CGI programs can be run interactively like all Perl programs, and running them this way is very useful for debugging. To run your program, start it from the command prompt like this:

```
perl hello.cgi
```

The Perl interpreter should reply with a line that looks like this:

```
(offline mode: enter name=value pairs on standard input)
```

This prompt indicates that the CGI module is attempting to get your CGI form values; they are covered in [Hour 22](#), "Basic Forms."

To this prompt, you should—for now—respond with just an end-of-file character. Under Unix, it is Ctrl+D; you just press the Ctrl key and type D. In Windows, you press Ctrl+Z. Perl should then print the following:

```
Content-Type: text/html
```

```
<b>Hello, World!</b>
```

The `Content-Type: text/html` message indicates that what follows should be interpreted as text or HTML. The meaning of this message is covered fully in [Hour 24](#). For now, just know that it's important that this message is the *first* thing printed by your program—with the `header` function—and that it's necessary. If anything else is emitted before the `Content-Type` message, the CGI program will fail.

Problem: Perl responds with a syntax error.

Solution: Fix the syntax error.

Problem: Perl responds with `Can't locate CGI.pm in @INC....`

Solution: Your Perl installation is incomplete. The CGI module ships with Perl by default. If you need to reinstall it, see [Appendix A](#), "Installing Modules." Make sure you've spelled CGI correctly.

Server Problems

After you've eliminated your script as the cause of the problems, it's time to check the script's installation and the server's configuration.

Problem: The server responds with the message `Not Found` or `404 Not Found`.

Solution: These messages usually indicate one of the following two problems:

- The URL you used is incorrect. You typed `http://www.server.com/cgi/hello.cgi` when you should have typed `http://www.server.com/cgi-bin/hello.cgi`. Go back to the checklist and verify that the URL of your CGI directory is correct, you've spelled the name of the script right, and you have the right extension.
- You placed the script in the wrong directory on the web server. Verify from the checklist that the CGI directory is correct, and if not, move the script to the correct location.

Problem: The text of your script is displayed.

Solution: The program is displayed because the web server thinks that the program is really a document, not a program to be executed.

- You used the wrong extension on your CGI program. Instead of `.pl`, you used `.cgi` or some other wrong extension. Check the checklist, and make sure you're using the correct extension on your CGI program.
- You put the script in the wrong directory, *and* you're using the wrong URL to access it. Verify that you've put the script in the correct CGI program directory and that you're using the right URL.
- The server is misconfigured. If you're using your own web server, reread the documentation and verify that you've set it up correctly. Sometimes the installation includes a test CGI script; if so, try it. If you're using a commercial web hosting server, verify that you're placing your script in the correct

directory, or contact the web host for help.

Problem: The server responds with the message [Forbidden or 403 Error](#).

Solution: The permissions on the CGI program are not correct. This problem is most likely to occur on a Unix web server.

You can view the permissions on the [hello.cgi](#) program by typing `ls -l hello.cgi` at a command prompt in the proper directory. If you have FTP access to the server, you can view the file permissions by typing [dir](#). The permissions should look something like this:

```
-rwxr-xr-x    1 user            93 Aug 03 23:06 hello.cgi
```

The permissions are the `rwxr-xr-x` characters on the left. If they don't look exactly like this, go back to the installation instructions for details on how to properly set the permissions on the CGI program.

Fixing Internal Server or 500 Errors

If the server replies with an [Internal Server Error](#) or [500 Error](#) message, this means that your CGI program failed somehow. This general-purpose failure message is generated by many different problems.

The most important tool you have in debugging an "Internal Server Error" is the server's log file. As the web server receives requests for web pages, it writes each page request into a file for later analysis. Any errors that the server encounters are also logged, including error messages generated by CGI programs.

Find the location of the server error log file, which you should have noted in the checklist. The log is typically written by appending any new items to the bottom of the log file. To see the last few entries under Unix, at a command prompt, type

```
tail error_log
```

to see just the bottom of the file. Some web servers have a utility—often a CGI program itself—to view the log file. If you have FTP-only access to the server, you might need to download the log file and view it on your local PC to see the error entries.

If you do not have access to the server's error logs, you have a big problem. Finding an "Internal Server Error" will be a hit-and-miss affair. Following the checklist shown here, you should find the problem eventually. (The messages you'll find in the server logs are approximated; the text will vary from server to server.)

Log entry: [No such file or directory: exec of /cgi-bin/hello.cgi failed](#)

Possible causes:

- The `#!` line of the script may not be correct. Make sure that the location of Perl in the `#!` line matches the location of Perl on the web server from the checklist. Verify that Perl is actually installed there by using `ls` or `dir` in FTP or locally.
- If you used FTP to transfer the CGI program to the server, you may not have used ASCII mode for the transfer. Moving a script written in Windows to a Unix server—and vice versa—in binary mode does not work.
- The permissions are not correct on the CGI program (Unix). See the description of [Forbidden](#) in the "Server Problems" section.

Log entry: `Can't locate CGI.pm in @INC....`

Possible causes: The Perl installation is incomplete, corrupt, or very old. Perl is apparently unable to locate the CGI module, which is a standard part of the Perl installation. You might need to reinstall the module or contact your system administrator to have him or her reinstall Perl. See [Appendix A](#).

Log entry: `Syntax error, warning, Global symbol requires , and so on`

Possible causes: Your Perl program apparently has a typo or incorrect syntax. Follow the instructions in the "Is It Your CGI Program?" section to determine the problem.

Log entry: `Premature end of script headers`

Possible causes: This catchall error message describes any situation in which your script is run, and the `Content-Type` header—printed by the CGI module's `header` function—is not the first thing emitted by the script. Sometimes a secondary message appears in the log file, immediately before or after this one. The other message may be more helpful in determining what went wrong. You can try these solutions:

- Make sure you're not printing *anything*—including error messages—before calling the `header` function. Anything printed before the `header` function causes this error.

By the Way

You may see Perl CGI programs printing "`Content-Type: text/html\n\n`" at the beginning of the program instead of calling `header`. Apparently, printing this and calling `header` are supposed to do the same thing, but they don't. The `header` function takes into account that `\n\n` doesn't always mean the same thing on every server and emits the proper sequence for that server.

- A problem called **output buffering** can cause output generated by the `system` function or backticks (````) to appear in the output before the `header` function's output. To make sure that the output from

`header` is always shown first, you might want to try rewriting the beginning of your Perl CGI program to look like this:

```
#!/usr/bin/perl -wT  
  
use strict;  
  
use CGI;  
  
  
$|=1;          # ensures that header's output always prints first  
  
print header;
```



< Day Day Up >



Summary

In this hour, you learned the basics of how CGI programs work. The differences between static web pages and dynamic web content were explained, and will be explored in further hours. You also got to write your first CGI program and, I hope, make it work.

This hour also provided a guide to debugging your CGI programs that will be useful to you in upcoming hours.

Q&A

Q1: *I don't have the CGI module loaded. Do I have to use it?*

A1: Quite frankly, yes. CGI is not an easy protocol to get right. Many programs that are published try to replicate the CGI module's functionality—and do so poorly. They are riddled with security problems and incompatibilities; plus, they do not follow Internet Standards as they should. In Hour 16, "The Perl Community," you learned why "reinventing the wheel" is a bad thing. CGI is an incredibly hard wheel to reinvent, and neither you nor I would get it right the first time or the hundredth time.

The Perl CGI module included with the standard distribution has been tested by hundreds of thousands of programmers and is very robust. Use it.

As I explain in Appendix A, you can install the module for just your use, if necessary. You don't have any excuse not to use it. All the examples in this book require the CGI module, and the explanations assume you have it installed.

Q2: *I have this copy of cgi-lib.pl. Can I use it instead?*

A2: You really shouldn't. All the functionality of `cgi-lib.pl` is in the CGI module. The `cgi-lib.pl` library is quite old and is no longer being maintained.

Q3: *Why does everyone use Perl for CGI? Why not C or TCL?*

A3: Perl has features that make it particularly useful for CGI. The short list is as follows:

- Perl has excellent text-handling capabilities.
- Perl's tainting capabilities—you'll learn about them soon—help make it a secure language for doing CGI programs.
- Perl is an excellent **glue** language, good for sticking together different technologies such as operating system utilities, databases, and CGI.
- Perl is easy to use.

Q4: *If I have questions about Perl and CGI, should I post a message to the comp.lang.perl.misc newsgroup or at PerlMonks.org?*

A4: Probably not. A more appropriate group would be `comp.infosystems.www.authoring.cgi`. First, though, you should check the FAQ at <http://www.w3.org/CGI/>.

Workshop

Quiz

1: CGI programs can be written in

a.

Perl, the Unix shell, or C only

b.

Perl only

c.

Almost any programming language that can be run on the server

2: The Web was invented before Perl.

a.

True

b.

False

Answers

A1: c. Perl is not special in this regard; it simply has some features for making CGI programming much easier and more reliable.

A2: b. Perl was invented in 1987, and the Web wasn't invented at CERN until 1991.

Activities

- Spice up your "Hello, World!" program a little bit. Print the current time (by using `localtime`), and maybe add some color and a table or two with HTML tags. Be creative. Remember that printing HTML in your Perl program will make the program appear in the final web page when the page is loaded.

 PREV

< Day Day Up >

NEXT 

Hour 22. Basic Forms

What You'll Learn in This Hour:

- How to process basic forms in Perl CGI programs
- How to debug CGI forms
- How to write safer CGI programs

While surfing the Web, you've undoubtedly filled out a few HTML forms. HTML forms, such as email web forms, shopping carts, guestbooks, online auction forms, mailing list subscriptions, and order forms, are used to collect information, such as login information, and even web site preference settings from the web browser's user.

When the user clicks the submit button on those forms, what happens? Almost always, the data from that form is passed to a CGI program. In this hour, you'll learn how to take data from a form and manipulate it in your CGI programs.

How Forms Work

You've undoubtedly used forms on the Web, and you might even understand how they're laid out and how they work. Regardless, a review of HTML form basics is in order.

Short Review of HTML Form Elements

Before you jump into how a form works, you should review how HTML represents a form and the roles that all the elements in that form play.

By the Way

The HTML presented in this book should not be taken as a model of great HTML. Enough HTML is presented to demonstrate the CGI features needed, but not much else. There are generally no `<head>` or `<body>` tags shown, nor a `<!DOCTYPE>` in any of the HTML presented in the examples. In addition, the screens are quite plain. Feel free to add your own HTML to make the screens more attractive and complete. Most browsers will display the HTML presented here correctly.

An **HTML form** is a portion of an HTML document that accepts user input. When the browser loads HTML containing a form, various HTML tags create user input areas on the page. The user input is put into various form elements such as check boxes, radio buttons, selection menus, and text entry elements. After the user has manipulated the input elements with the web browser, the form is typically submitted to a CGI program for processing.

A typical HTML form is constructed something like the one shown in Listing 22.1.

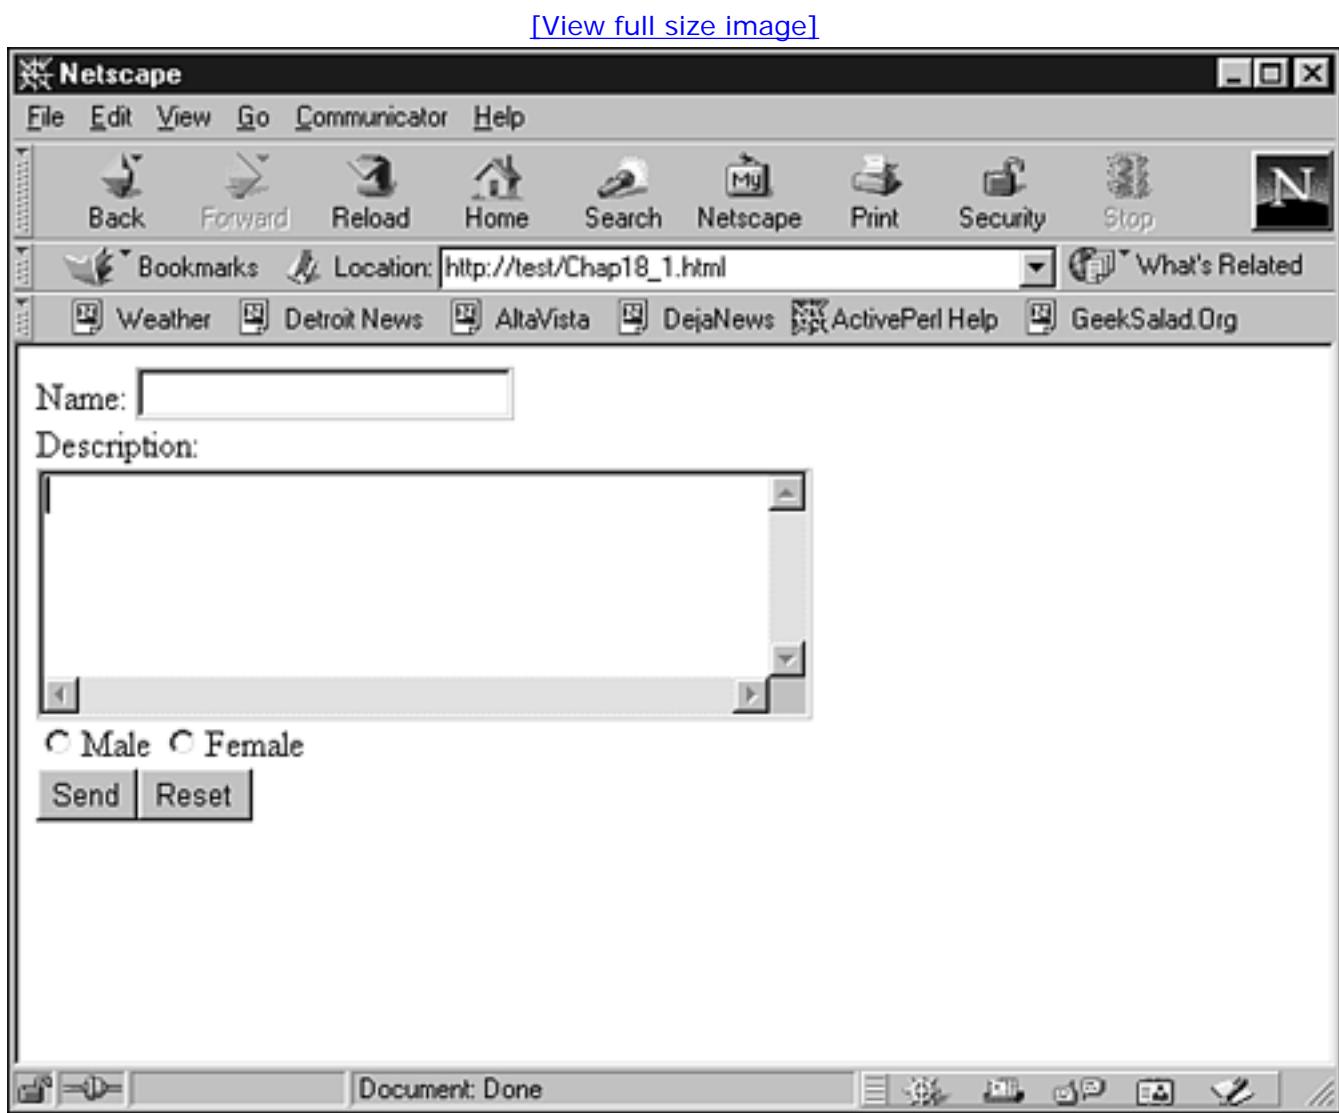
Listing 22.1. A Small HTML Form

```
1:   <form action="http://www.server.com/cgi-bin/submit.cgi" method="get">
2:     Name <input type="text" name="name" /><br/>
3:     Description<br/>
4:     <textarea name="description" rows="5" cols="40">
5:     </textarea>
6:     <input type="radio" name="sex" value="male"/>Male
7:     <input type="radio" name="sex" value="female"/>Female
```

```
8: <br />
9: <input type="submit" value="Send" /><input type="reset" />
10: </form>
```

Figure 22.1 shows the form from Listing 22.1 rendered with the Netscape browser.

Figure 22.1. The form from Listing 22.1 rendered in Netscape.



The `<form>` tag specifies the beginning of the form within the complete HTML document. The `method` attribute specifies whether this form will use the `GET` or `POST` method for form submission. If this attribute is not specified, the browser submits the form to the CGI program using the `GET` method. The difference between the methods is documented later. The `action` attribute specifies the URL of the CGI program that will receive the form data.

The `<input>` tag presents the user with an input field—in this case, a blank text box. The text box is assigned a name—coincidentally "name".

The `<textarea>` tag causes the browser to display a multiline text box to accept input. The important attribute to notice is the `name` attribute; in this case, the name of the field is `description`. Each form element in an HTML form should have a different `name` attribute. When the CGI program is given the form for processing, the `name` attribute is used to tell the fields apart.

Finally, the submit button is displayed. When the user clicks this button, the form values are passed to a CGI program for processing, as explained in the next section.

The HTML 4.0 specification contains quite a few different form element types, so I will not try to document them all fully in this book. For example, many form elements contain attributes to fix certain properties of the form element, such as `rows` and `cols` in the `textarea` in the preceding form. For the rest of this book, whenever HTML form elements are used, only the most basic attributes will be used.

By the Way

You can find the full HTML 4.0 specification, including the valid form elements and their attributes, at <http://www.w3c.org>.

What Happens When You Click Submit?

After a user has filled in a form on his or her web browser, a short chain of events takes place:

- The data on the form is collected by the web browser into name and value pairs (see [Figure 22.2](#)). For example, in the sample form, the field named `body` takes on the value of the text entry field. The field named `sex` takes on the value of the radio buttons. The web browser does all this work before anything else happens.

Figure 22.2. Browser pairing up data and fiel

Passing Information to Your CGI Program

When a CGI program is run as the result of a form submission, the field names and values passed in from the form—called **parameters**—need to be processed by the CGI program. This is done with the `param` function.

Without any arguments, the `param` function returns the names of the fields being passed into the CGI program. If the CGI program was accepting the form in Listing 22.1, the `param` function would return `body`, `sex`, `name`, and `submit`.

With an argument, `param` returns the value of that parameter. For example, `param('sex')` would return the value of the radio buttons—either `male` or `female`—depending on which was selected.

Listing 22.2 contains a short CGI program to print these parameters.

Listing 22.2. CGI Program for Printing Parameters

```
1: #!/usr/bin/perl -w
2: use strict;
3: use CGI qw(:standard);
4:
5: print header;
6: print "<p>The name was", param('name'), "<br/>";
7: print "The sex selected: ", param('sex'), "<br/>";
8: print "The description was:<br/>", param('description'),
9:      "</p>";
```

If the parameter specified by `param` is not used in the form, `param` returns `undef`.

GET and POST Methods

In the form in Listing 22.1, the `<FORM>` tag has an attribute called `method`. The `method` attribute specifies how the web browser should transmit the data to the web server. Currently, two methods are available.

The first method—also the default method if you don't specify one in the `<FORM>` tag—is called `GET`. With the `GET` method, the form values are passed to the CGI program by having their values encoded in the URL. You may have seen URLs like this while surfing the Web:

<http://www.server.com/cgi-bin/sample.pl?name=foo&desc=Basic%20Forms>

The CGI program, when it is run, can decode the remainder of the URL into fields and values. The `param` function essentially does the same when it's called. You should *not* try decoding these values yourself. The `param` function does a thorough job of it, and you don't have any reason to use anything else to extract the values.

The other method, `POST`, produces exactly the same result but through a different means. Instead of encoding all the form values into the URL, the web server is contacted and the HTML form values are sent as input to the CGI program. Again, you don't need to know at this time exactly how this process works; the CGI module takes care of that for you. Simply calling the `param` function takes care of reading the values, decoding them, and passing them to your program.

By the Way

You may have downloaded CGI programs from the Internet or seen examples in other books that try to decode an environment variable called `QUERY_STRING` or use a variable called `REQUEST_METHOD` to figure out whether the form uses a `GET` or `POST` method. Those programs attempt to reproduce the work already done in the standard CGI module—and probably not as well. You should avoid doing this yourself.

So, which method should you choose? Each method has advantages and disadvantages. The `GET` method allows the web browser to bookmark the particular URL that generated the page. For example, the URL <http://www.server.com/cgi-bin/sample.pl?name=foo&desc=Basic%20Forms> could be bookmarked and always returned to by the browser. From the perspective of the CGI program—`sample.pl`—it doesn't know whether you've just come from actually viewing the form. It receives the normal CGI parameters as it always has. Being able to call a CGI program repeatedly using the `GET` method's URL-encoded values is called **idempotency**.

However, you might not particularly want browsers to be able to bookmark in your site to run your CGI programs directly, and the URLs for invoking a CGI program with the `GET` method are, frankly, ugly.

The `POST` method doesn't encode the form data in the URL at all; it relies on the browser to send the data as it's negotiating for the page. However, because the data isn't encoded into the URL, you can't bookmark a page that was generated by a CGI program using the `POST` method.

Web Security 101

Before you put CGI programs on the World Wide Web, you need to know a few things. By putting a CGI program on a web page, you are giving remote users (using web browsers) limited access to your system. Using normal HTML documents, they can retrieve only static documents from your web site. Using CGI programs, however, they're actually able to run programs on your web server.

Knowing how to write safe and secure CGI programs will make you and your web server administrator much happier. Writing such programs is not hard; you just need to follow a few simple precautions.

A Clear Link

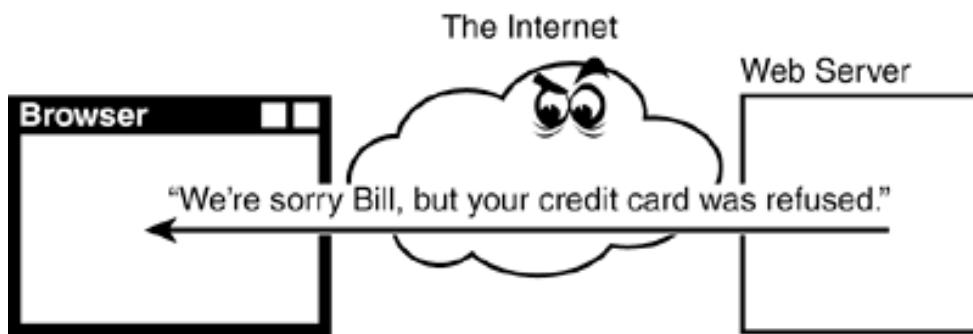
When a web browser retrieves a page from your web server, the HTML is sent over a cleartext channel (see [Figure 22.6](#)). This means that as the data winds its way through the Internet, it's not encrypted, encoded, or otherwise obscured.

Figure 22.6. Plain text being transmitted to server.



The data that a user puts into a form and submits to your CGI program is transmitted with the same protocol as the initial web page. All the fields of the form are free for anyone to see as they go by (see [Figure 22.7](#)).

Figure 22.7. Server responding in plain text.



The problem of data being transmitted in the clear is a real one that you should be concerned with. The Internet is not a secure place, and anyone along the wire between the web browser and the web server can eavesdrop on the information being sent out or in.

Keeping this point in mind, you should never transmit some kinds of data with normal CGI forms:

- Passwords of any kind
- Personal information (Social Security numbers, phone numbers)
- Financial information (account numbers, PINs, credit card numbers)

Keep this rule of thumb in mind: Never send anything on the Internet that you wouldn't put on a postcard. If you need to transmit data that needs to be kept secure, you need to use extraordinary methods to do so: use secure HTTP ([https](https://)) instead of HTTP, make sure that the servers you're sending the data to are secure, and trust the recipient with your information.

By the Way

"But wait!" you say. "I've seen forms on the Internet asking for all those pieces of information and claiming to be secure." Reasonably secure transactions can be performed over the Web with some additional tools. Secure web transactions are actually performed by encrypting the *entire* browser/server conversation. This is done by using a secure version of the http protocol called **https**.

Watching for Insecure Data

Another point to consider in writing secure CGI programs is that you're writing programs that will execute Perl commands based on the input given to you by a web page. The Internet—and possibly your intranet—is full of people who are not nice and would take great glee and pride in harming your web server. Also, benign users could accidentally send invalid data to your CGI program.

Consider the HTML form in [Listing 22.3](#) and the CGI program in [Listing 22.4](#).

Listing 22.3. Directory Listing Web Form

```
1: <form action="/cgi-bin/directory.cgi">
2: <p>What directory to list?
3: <input type==text name==dirname/>
4: <input type=submit name=submit value="Run This"/>
5: </p>
6: </form>
```

Listing 22.4. An Insecure CGI Program Named *directory.cgi*

```
1: #!/usr/bin/perl -w
```

```
2: # Do NOT use this CGI program, it's very insecure
3: use strict;
4: use CGI qw(:all);
5:
6: print header('text/plain');
7: my $directory=param('dirname');
8: print `ls -l $directory`; # Do a directory listing
```

[Listing 22.3](#) presents the user with a short form to collect a directory name and then passes the name to a CGI program called `directory.cgi`. In [Listing 22.4](#), the `directory.cgi` program takes the directory and performs an `ls -l` on it—for DOS/Windows users, this is the equivalent of `dir`—giving the user a directory listing.

A program of this sort allows a remote web surfer to explore your entire directory structure. The CGI program does no validation of what the directory name is, and if the browser wants to explore your sensitive data, it may be able to.

The more important point is that `$directory` might not contain a directory at all! If the web browser had sent back the value `/home; cat /etc/passwd` for `dirname`, the command run by the CGI program would have been as follows:

```
ls -l /home; cat /etc/passwd
```

This command would send a copy of the system's password file back to the web browser. In fact, any Unix shell command or MS-DOS command could be run this way, including commands to delete files on your server! If your web server has not been set up properly, this could allow any user on the Internet to take full control of your server remotely.

Perl has a mechanism to help prevent you from doing silly things like this. The `-T` switch on the `#!` line enables **data tainting**. As data is received from external sources—filehandles, network sockets, the command line, and so on—it's marked as **tainted**. Tainted data cannot be used in backticks, system calls (such as `opening` a file for output), the `system` command, or other places that might compromise your security.

You also cannot use the `open` function, `system` function, or backticks in your Perl program when taint checking is in effect unless you explicitly set your `PATH` environment variable first.

[Listing 22.5](#) shows a more secure version of this program.

[Listing 22.5. A Somewhat More Secure Version of `directory.cgi`](#)

```
1: #!/usr/bin/perl -Tw
2: # tainting is enabled!
3: use strict;
4: use CGI qw(:all);
5:
6: print header('text/plain');
7: # Explicitly set the path to something reasonable
8: $ENV{PATH}='/bin:/usr/bin';
9: my $dir=param('dirname');
10: # Only allow directory listings under /home/projects
11: if ($dir=~m{^(/home/projects/[\w/]*)$} ) {
12:     $dir=$1;    # This "untaints" the data, see "perldoc perlsec"
13:     print 'ls -l $dir';
14: }
```

By the Way

For more information on tainted data, how to untaint data, and Perl programs, see the perlsec manual included with the Perl distribution.

Doing the Impossible

HTML/CGI forms can also be undermined in another way. Consider the HTML in [Listing 22.6](#).

Listing 22.6. A Simple Form

```
1: <form action="/cgi-bin/doit.cgi">
2: <p>Please type in your favorite color:
3: <input type="text" length="15" name="color"/>
4: <input type="submit" value="Submit color"/></p>
```

5: </form>

In this form, the maximum allowable width of the `color` field is 15, right? Almost. The HTML specification says that `length` in a text field will allow at most that many characters. However, the browser could be broken, or a malicious person could put more than 15 characters in that field by simply bypassing your form or creating a new one.

 PREV

NEXT 

A Guestbook

This example allows you to write a customizable **guestbook** for a web site. A guestbook is an HTML form in which the user indicates his or her name and some comments. The guestbook can be used for collecting feedback on a topic, as a simple message board, or for submitting problems to a help desk. The data is kept in a file and can be displayed after the form is filled out; it also can be displayed on a page of its own.

[Listing 22.7](#) provides a short bit of HTML to present a guestbook form for a fictional help desk. You can modify the form to suit your own needs.

Listing 22.7. Help Desk Form

```
1:  <form action="/cgi-bin/helpdesk.cgi" name="helpdesk">
2:  <p>
3:  Problem type:
4:  <input type="radio" name="probtype" value="hardware"/>Hardware
5:  <input type="radio" name="probtype" value="software"/>Software
6:  <br/>
7:  <textarea name="problem" rows="10"
```

Summary

In this hour, you learned how HTML forms and CGI programs interact and how to have your CGI programs interpret the contents of forms with the CGI module's `param` function. Also, you learned what to do to make your CGI programs more secure and how to handle tainted data. You also looked at a simple CGI guestbook application, ready to customize and modify for your own uses.

Q&A

Q1: *I can't get the form submission to work; I keep getting an error.*

A1: Try using the CGI debugging guide from [Hour 21](#), "Introduction to CGI," to find your problem. Just because it's a form doesn't mean that debugging it is any different from debugging normal CGI.

Q2: *I just found this great program on the Internet, but I don't understand why it's trying to use `$ENV{QUERY_STRING}` to get the form parameters.*

A2: The author of that program decided to forgo the CGI module's form-processing capabilities. This fact indicates that it is either a very old Perl program that predates the CGI module or that the author decided to use his or her own form-processing code. Either way, this is a good indication that the program should be looked at skeptically and used cautiously.

Q3: *When I run programs with the `-T` option in the `#!` line, from the command prompt, I get the error `Too late for -T option` and then the program stops. Why?*

A3: The `-T` option must be given to Perl as soon as possible so that it knows to look for tainted data. By the time the `#!` line in your program was processed, it was too late—Perl had already processed your command line options without tainting. To run Perl from a command prompt—for example, in the debugger—you need to specify the `-T` on the command prompt as well, before any other options:

```
perl -T -d foo.cgi
```

Q4: *Will Perl's data-tainting capabilities keep me from making stupid mistakes in my CGI programs? Are they guaranteed to be secure now?*

A4: No CGI program is completely secure. The data-tainting capabilities of Perl go a long way toward keeping you from making silly mistakes, but they do not guarantee secure programs.

Workshop

Quiz

1: The `param` function in an array context with no arguments returns

a.

`undef`

b.

The number of form elements

c.

A list of the form element names

2: The differences between the `POST` and `GET` methods are transparent if you use the CGI module.

a.

True

b.

False

3: The password field input type on an HTML form is secure because it obscures the password before sending it.

a.

False

b.

True

Answers

A1: c. `param`, without arguments, returns a list of the element names from the submitted form.

A2: a. True.

A3: a. False. With normal HTTP and CGI, all form fields are transmitted in cleartext and are not secure. The password field input type just hides the field as you type it.

Activities

- Spice up the help desk form a bit. Add timestamps to each entry and maybe some color on the output.
- Challenge: The `display()` function prints the entries in the help desk form starting with the oldest. Change the `display()` function to print the newest entries first.

 PREV

< Day Day Up >

NEXT 

Hour 23. Complex Forms

What You'll Learn in This Hour:

- How to create multipage forms

The forms on the Web aren't limited to simple one-page forms. Sometimes forms are spread out over several pages. These complex forms appear as surveys, questionnaires, and shopping cart applications.

These more complex forms require some different programming techniques, which you'll learn about this hour.

The Stateless Web

Writing complex multipage forms using CGI programs offers a unique programming challenge. The connection between a Web browser and a Web server is not a long-term connection. The Web browser connects to the server, retrieves the page, and then disconnects from the Web server. No sustained connection is kept between the server and your browser.

To make matters more complicated, each time the browser connects to the Web server, the Web server doesn't recognize the browser as having been to the site beforehand. The server has no easy way of recognizing the browser each time.

An analogy that works well is that of a conversation between a library patron and a librarian with no memory. The patron is allowed to ask only one question at a time. The patron can ask the librarian for a book—say, on Arizona—and the librarian can retrieve it. The librarian can retrieve the book because he or she can fulfill simple requests easily. However, the patron cannot request another book on the same subject. The librarian would have no memory of the prior request and would be helpless. If the request were rephrased as "give me another book on Arizona," the librarian would still be unable to help because he or she might retrieve the same book on Arizona as he or she did for the first request.

The only way to retrieve a second book on the subject would be to say, "I need another book on Arizona; I already have *Settling of the Southwest*." That request has enough information to state the problem and let the librarian know what's not an appropriate response.

Writing multipage forms for the Web uses the same solution: Each question/answer session needs to contain enough information so that the Web server can figure out what needs to be done and what has already been done. You can create such sessions in a few different ways; one of them—using hidden HTML fields—is presented in this hour.

Hidden Fields

The easiest way to have Web forms "remember" information is to embed previous information in them by using hidden fields. Hidden fields are part of the HTML forms specification that allow fields and values to be part of an HTML form but not appear in the form when it's rendered visually. In HTML, they're written as follows:

```
<input type="hidden" name="fullname" value="Pink Floyd" />
```

If the preceding HTML were put into a form, a new name ("fullname") and value ("Pink Floyd") would be part of the form. If the form were submitted to a Perl CGI program, the `param` function would return a key and a value for the hidden field.

The Online Store

For an example of how to use hidden fields, consider an online store that uses a series of Web pages to allow people to shop for items from an online catalog. For now, I'll just walk you through the workings of a complex form. Later in this hour, I'll present another complex form complete with code to create an online survey.

Watch Out!

Do *not* use this example of an online store without implementing some kind of secure Web transaction, such as https. Consult your web server's documentation to set this up. Notice that this example doesn't contain any real personal information such as phone numbers or credit card numbers because hidden fields are just like regular HTML forms. No security is involved at all.

The first page of the online store, illustrated in Figure 23.1, presents a list of featured items from the store.

Figure 23.1. The first page of an online store.



When the user clicks the Go to Store button, the CGI program accepts the values from the form and then presents the full catalog, as shown in [Figure 23.2](#).

Figure 23.2. Catalog of online store shown.

Online Store, full catalog

Fish heads
 Chitterlings
 Head Cheese
 Elephant Ears

Hidden:

Chocolate Bananas

Ship Items

The second page presents the full catalog. When the first page—with the featured items—was submitted, the CGI program accepted the values, and then as it was printing the HTML for the full catalog, it placed the selected quantities for the featured items on the new form as hidden fields.

Each time the CGI program receives the values from the HTML form, the new page contains the old values in hidden fields, as well as the new values in the regular form elements.

In this way, you avoid the "forgetful librarian" problem. When the form for the full catalog is submitted, the hidden fields in the form remind the CGI program which items were selected from the first form, as well as from the current form.

If a third page is needed, the values from the first two pages can be squirreled away on the third page as hidden fields, as shown in [Figure 23.3](#).

Figure 23.3. Shipping info for online store.

Online Store, Shipping

Name

Address

Hidden:

Chocolate Bananas

Fish Heads

Done

You should be aware of some problems with hidden fields on an HTML page. First, the values in the hidden fields can be seen by anyone. To see the values, users can simply view the HTML source for the page; most Web browsers have an option for viewing the source.

Second, the values in hidden fields can be changed by remote users. Remote users can use a doctored Web browser or can submit the form manually using HTTP. For example, the online store should not store the *prices* in hidden fields—only the quantities. The CGI program should look up the prices on the server side as they are needed for display.

By the Way

When you're designing forms, take a look at how other people design theirs. This way, you learn whether they're using hidden fields to preserve information. Most Web browsers have a View Page Source selection. You should use it on any form you find particularly interesting to see how it was put together. Do *not* copy the form, however. Most of the time copying violates the original author's copyright.

A Multipage Survey

Surveys are common places to find forms that span several different Web pages. They're sometimes too long to fit on a single page, and they can usually be broken up into categories.

What follows is a simple multipage Web survey to find out aspects about your personality. This survey presents four different pages but could easily be changed to support as many pages as you would like. The four pages are as follows:

- A series of general questions to find out what kind of personality you have
- Some specific questions about your habits and a question based on an answer from the first survey page
- A page that allows you to enter your name and comments about the survey
- A "thank you" message that prints after the survey is complete

The same CGI program is used to perform all four functions. It decides which page to open next, depending on which page it has just displayed. The core of the survey program is shown in [Listing 23.1](#).

Did you Know?

By including the line

```
use CGI::Carp qw(fatalsToBrowser);
```

your CGI program's `die()` messages, which might normally go to the Web server's log file, are printed as part of the Web page. When you write longer CGI programs, this might help in debugging.

The results of the survey are saved in a text file but are not displayed by this program at all. This program simply collects the survey answers and stores them. You will have to write another CGI program to display the results.

Listing 23.1. First Part of the Survey Program

```
1:  #!/usr/bin/perl -w
2:  use Fcntl qw(:flock);
3:  use CGI qw(:all);
```

```

4:  use CGI::Carp qw(fatalsToBrowser);
5:  use strict;
6:  my $surveyfile="/tmp/survey.txt";
7:  my @survey_answers=qw(petttype daytype clothes
8:                         castaway travel risky ownpet
9:                         realname comments);
10: my $semaphore_file="/tmp/survey.sem";
11: print header;
12: if (! param ) {
13:     page_one();          # Survey just started
14: } elsif (defined param('pageone')) {
15:     page_two();          # Answered one page, print the second
16: } elsif (defined param('pagetwo')) {
17:     page_three();         # Print the last page.
18: } else {
19:     survey_done();        # Print a thank-you note, and save
20: }

```

Lines 7–9: During the survey, each HTML form contains input fields. The names of each of the fields appear in this array. The `save()` function and the `repeat_hidden()` functions will use this array of names later.

Lines 12–13: If this CGI program is passed no parameters—that is, it's not loaded as the result of a form posting—the function `page_one()` is called to print the first page of the survey.

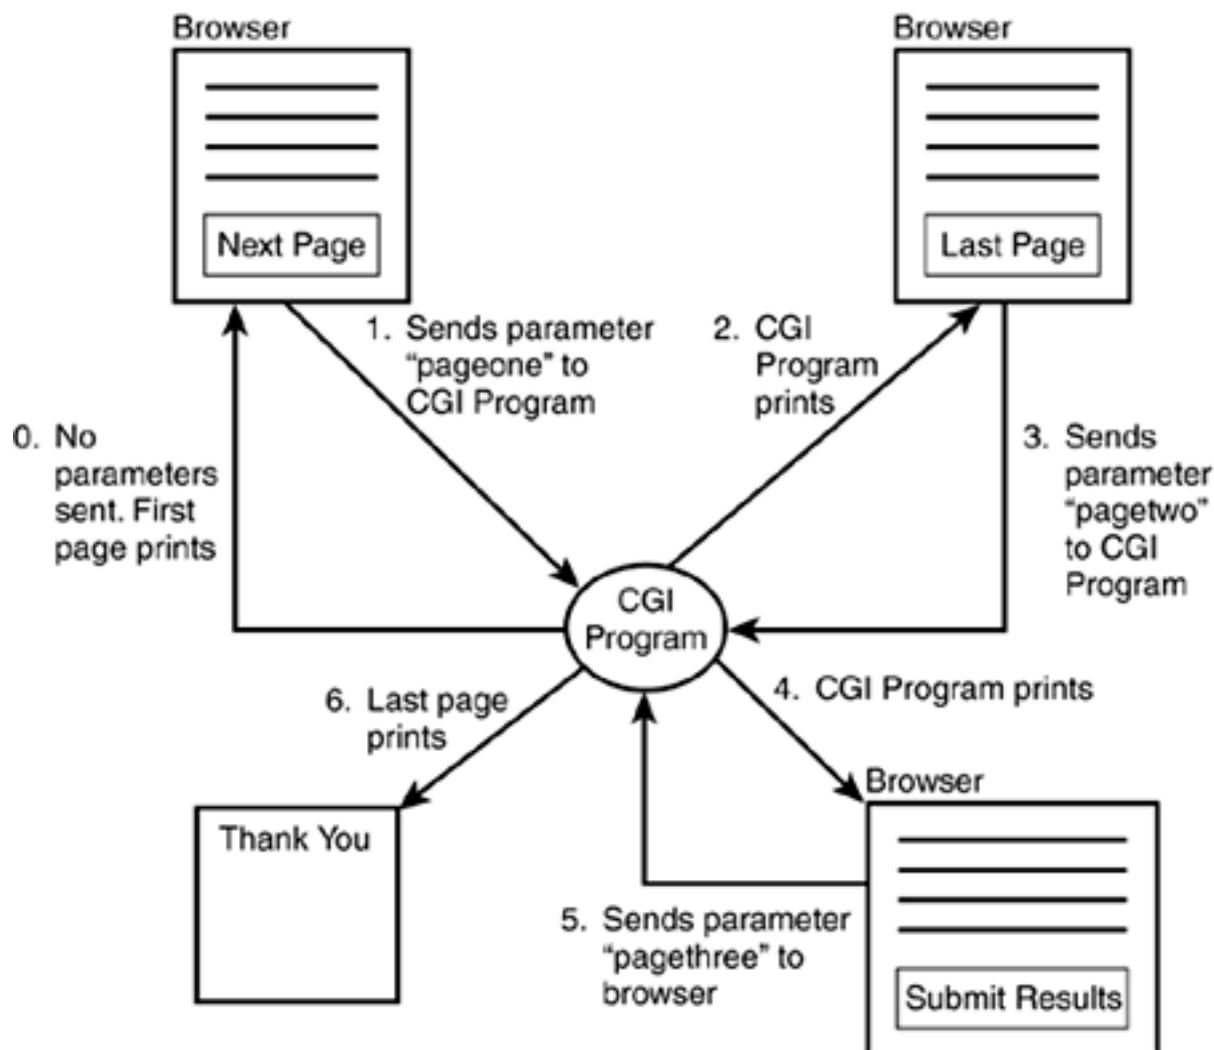
Lines 14–17: If an HTML form parameter called `pageone` is passed to this CGI program, the function `page_two()` is called. If `pagetwo` is passed, then `page_three()` is called.

Line 19: If HTML form parameters are passed to this CGI program, but not `pageone` or

`pagetwo`, the survey is complete, the results are saved, and the "thank you" message is printed in the `survey_done()` function.

The submit buttons on each of the pages provide the clues as to which page should be loaded next, as you can see in [Figure 23.4](#). Because the submit button name is passed into the CGI program as a parameter, it can be used to show which version of the page was just submitted to the program.

Figure 23.4. Diagram of which button causes which action.



[Listing 23.2](#) continues the survey program.

Listing 23.2. Second Part of the Survey Program

```
21: sub page_one {  
22:     print<<END_PAGE_ONE;  
23:     <form>  
24:         <p>Are you a "cat person" or a "dog person"?<br/>
```

```

25: <input type="radio" name="pettype" value="dog"/>Dog<br/>
26: <input type="radio" name="pettype" value="cat"/>Cat<br/>
27: </p><p>
28: Are you more of an early-riser or a night owl?<br/>
29: <input type="radio" name="daytype" value="early"/>Early riser<br/>
30: <input type="radio" name="daytype" value="late"/>Night owl<br/>
31: </p><p>
32: At work, if you had a choice on how to dress....<br/>
33: <input type="radio" name="clothes" value="casual"/>Casual<br/>
34: <input type="radio" name="clothes" value="business"/>Business<br/>
35: </p><p>
36: If stranded on a desert island,
37: who would you rather be stuck with?<br/>
38: <input type="radio" name="castaway" value="ginger"/>Ginger<br/>
39: <input type="radio" name="castaway" value="marya"/>Mary Ann<br/>
40: <input type="radio" name="castaway" value="prof"/>Professor<br/>
41: <input type="radio" name="castaway" value="skipper"/>Skipper<br/>
42: <input type="submit" name="pageone" value="Next page"/></p>
43: </form>
44: END_PAGE_ONE
45: }

```

Lines 22 through 44 are a new Perl construct that you haven't seen yet; it's called a **here document**. A here document allows you to specify a string that spans several lines, includes other quotation marks, and acts like a normal double-quoted string. To start a here document, you use `<<` followed by a word. The quotation continues until the next occurrence of the word at the beginning of the line, as in this example:

```
$a=<<END_OF_QUOTE;
```

This is included as part of the string.

```
END_OF_QUOTE
```

The word that identifies the beginning of the here document—`END_OF_QUOTE` in the previous snippet or `END_PAGE_ONE` in Listing 23.2—must be followed by a semicolon. At the end of the here document, the word must appear starting in the first column and must not have any trailing characters such as spaces or semicolons. Inside the here document, variables expand as they do in normal double-quoted strings (`" "`), so be careful using \$ and @ inside a here document.

Using here documents allows you to embed a large amount of HTML in your Perl program without having to fuss with quotation marks, multiple `print` statements, and so on.

The function in Listing 23.2 simply prints an HTML form. The `<form>` tag does not contain an action or a method. When a `<form> action` attribute is not specified, the current CGI program (the one that produced the form) is reloaded when the form is submitted. When the `method` attribute is not specified, the default method, `GET`, is used.

Notice that the submit button on the form is named `pageone`. When this form is submitted, a parameter called `pageone` will be sent to the CGI program; the value isn't important here. When this parameter is submitted, it is the cue to the CGI program to load the second page.

Listing 23.3 continues the survey CGI program.

Listing 23.3. Third Part of the Survey Program

```
46: # Print out any of the responses so far as hidden fields
47: sub repeat_hidden {
48:     foreach my $answer ( @survey_answers ) {
49:         if (defined param($answer)) {
50:             print "<input type=hidden";
51:             print " name=\"$answer\" ";
52:             print " value=\"", param($answer), "\"/>\n";
53:         }
54:     }
55: }
```

```
54:         }
55:     }
56: sub page_two {
57:     my $pet=param('pettype');
58:     if (! defined $pet) {
59:         $pet="goldfish";
60:     }
61:     print<<END_PAGE_TWO;
62: <form>
63: <p>Would you rather...<br/>
64: <input type="radio" name="travel" value="travel"/>Travel<br/>
65: <input type="radio" name="travel" value="home"/>Stay at home<br/>
66: </p><p>
67: Do you consider yourself...<br/>
68: <input type="radio" name="risky" value="yes"/>A daredevil<br/>
69: <input type="radio" name="risky" value="no"/>Cautious<br/>
70: </p><p>
71: Do you own a $pet?<br/>
72: <input type="radio" name="ownpet" value="$pet"/>Yes<br/>
73: <input type="radio" name="ownpet" value="no"/>No<br/>
74: </p>
75: <input type="submit" name="pagetwo" value="Last Page"/>
76: END_PAGE_TWO
77:     repeat_hidden();
78:     print "</form>";
```

```
79: }
```

Line 47: As the comment in line 46 indicates, the function in this line prints the values of all the form fields so far as hidden fields. The array `@survey_answers` contains all the possible `"name=` values on the HTML forms. When run the first time, most of the fields will not exist because those portions of the survey haven't been filled out yet.

Lines 48–49: Each of the possible parameters in `@survey_answers` is checked, and each parameter is defined. An `<input type="hidden">` HTML tag is printed to hold the value on the current form.

Line 56: This function prints the second page of the survey.

Lines 57–60: This function is called as the second page of the survey. If the first page of the survey was properly filled out, then `param('pettype')` will hold either `dog` or `cat`, and this value will be stored in `$pet`. If the responder skipped that question, and `param('pettype')` is not defined, then `goldfish` is used instead.

Lines 61–76: The rest of the second page of the form is printed, and `$pet` is substituted as a question; thus, this questionnaire is customized based on the response on the first page.

Line 77: Any HTML form parameters from the first page are carried over onto this form as hidden fields.

If you view the survey form at this point—the second page—all the answers from the first page are stored as hidden fields at the end of the second page. The code for the third page is shown in Listing 23.4.

Listing 23.4. Fourth Part of the Survey Program

```
80: sub page_three {
81:     print<<END_PAGE_THREE;
82:     <form><p>
83:     Last page! This information is optional!<br/>
84:     Your name:</p><p>
85:     <input type="text" name="realname" /><br/>
86:     Any comments about this survey:<br/>
```

```

87:   <textarea name="comments" cols="40" rows="10">
88:   </textarea>
89:   </p>
90:   <input type="submit" name="pagethree"
91:         value="Submit survey results"/>
92: END_PAGE_THREE
93:   repeat_hidden();
94:   print "</form>";
95: }

```

The function `page_three()` is fairly straightforward, simply printing a text box and a text area in a form. At the end, it again calls `repeat_hidden()` to put all the hidden fields into the third page of the survey. The conclusion to the survey CGI program is shown in Listing 23.5.

Listing 23.5. Last Part of the Survey Program

```

96: sub survey_done {
97:   save();
98:   print "Thank You!";
99: }
100: #
101: # Save all of the survey results to $surveyfile
102: #
103: sub save {
104:   get_lock();
105:   open(SF, ">>$surveyfile") || die "Cannot open $surveyfile: $!";
106:   foreach my $answer (@survey_answers) {

```

```
107:         if (defined param ($answer) ) {
108:             print SF $answer, " = ", param($answer), "\n";
109:         }
110:     }
111:     close(SF);
112:     release_lock();
113: }
114: #
115: # Locks and unlocks the survey file so that multiple survey-takers
116: # Don't clash and write at the same time.
117: #
118:
119: # Function to lock (waits indefinitely)
120: sub get_lock {
121:     open(SEM, ">$semaphore_file")
122:         || die "Cannot create semaphore: $!";
123:     flock SEM, LOCK_EX;
124: }
125:
126:# Function to unlock
127: sub release_lock {
128:     close(SEM);
129: }
```

Line 96: This function is called simply to print a "thank you" message—always a nice thing to do after having someone go through three pages of survey—and call the `save()` function.

Line 103: The `save()` function here is almost a clone of the `save` function from Hour 22, "Basic Forms." It locks the survey file with `get_lock()`, writes the answers to it using a mechanism similar to that in `repeat_hidden()`, and then unlocks the file with `release_lock()`.

Feel free to modify this survey to suit your own needs. The design is fairly flexible and could be adapted for many purposes.



< Day Day Up >



Summary

In this hour, you learned how multipage Web forms are assembled. While doing so, you learned about a few problems that need to be addressed in writing programs—most seriously, making programs remember from page to page what happened. You also learned how hidden fields are used to remember information on Web pages that can't be remembered by the server, and then you constructed a bare-bones survey form using hidden fields.

Q&A

Q1: *Do HTML forms have to be this ugly?*

A1: The forms in this book are simple, plain, bare-bones, and—some would say—ugly. The purpose of *Sams Teach Yourself Perl in 24 Hours* is to teach you Perl and CGI programming, not HTML. In fact, most of the HTML in this book is not up to standards and is not complete; it lacks `<head>` tags, `<html>` tags, DTD headers, and so on. By providing basic HTML, I hope you'll correct this for your own needs.

As I mentioned earlier, a good method for spicing up your forms is to explore the Web and find a form you like. By viewing the source, you can get an idea of how the pages are put together.

Q2: *I'm getting the Can't find string terminator "XXXX" anywhere before EOF at.... error. What is it?*

A2: This error is caused by having an open quote somewhere in your program without a matching close quote. When you use here documents, this means that the word you've used to mark the end of the here document can't be found. The format is as follows:

```
print <<MARK;  
text  
text  
text  
MARK
```

The here document beginning and ending words—in this case, the word `MARK`—must be exactly the same. The terminating word must appear with nothing in front of it, on a line by itself, and with nothing after it. MS-DOS and Windows text editors sometimes do not put an end-of-line character after the last line of the program. If your here program ends at the end of the file, try putting a blank line after it.

Workshop

Quiz

1: So that your program remembers long, multipage Web transactions, you need to use

a.

Databases and cookies

b.

Hidden HTML form fields

c.

Some combination of hidden HTML form fields, cookies, and databases

2: Using an HTML `<form>` tag without an `action` attribute will

a.

Not work

b.

Cause the submit button to use the CGI program that originally generated the page

c.

Cause the submit button to reload the current page

3: The survey program, as listed, has a small bug. What is it?

a.

`print<<EOP ;` is not valid syntax in Listing 23.2.

b.

The HTML is not complete because it lacks `<head>` tags and such.

c.

The survey doesn't print its results.

Answers

A1: Either b or c. You can use just hidden HTML fields or just cookies. Using just a database won't quite work.

A2: b. Reloading the current page would just erase all the current form answers. Without an `action` attribute, the `<form>` tag uses the URL of the current page for the submission URL.

A3: b. `print<<EOP ;` is certainly valid syntax, called a here document. Choice c is not correct because the program wasn't designed that way (see the "Activities" section).

Activities

- Write a short CGI program to display the results of the survey program. Possibly build a table to display them like this:

Cat/Dog Own One? Nocturnal Clothes Castaway Traveler Risky?

cat	no	yes	casual	Professor	yes	yes
neither	goldfish	no	business	Skipper	no	no
dog	yes	no	casual	Mary Ann	yes	yes

- For an additional challenge, write a CGI program to summarize the survey results like this:

Cat/Dog Personality: Cat 40% Dog 45% Neither 15%

Own that pet: Cat 20% Dog 15% Goldfish 30% None 35%

Night Person? Yes 35% No 40%

 PREV

< Day Day Up >

NEXT 

Hour 24. Manipulating HTTP and Cookies

In this hour, you'll learn a host of interesting things you can do with the Web by using CGI programs to make your web site a little more flexible and easy to manage.

What You'll Learn in This Hour:

- How HTML gets from the server to your browser
- How to redirect the browser to another URL

The HTTP Conversation

In Hour 21, "Introduction to CGI," you learned about the basic conversation between the web browser (Netscape, Internet Explorer, and so on) and the web server (Apache, IIS, and so on). The discussion in that hour was somewhat oversimplified. Now that you're more comfortable with CGI, it's probably a good time to take a closer look at this subject. Later in this hour, you'll learn some techniques for manipulating this conversation to perform some interesting tasks.

The conversation is described by a protocol known as the **Hypertext Transfer Protocol (HTTP)**. The two current versions of this standard are HTTP 1.0 and HTTP 1.1. For purposes of this discussion, either one is applicable.

By the Way

The Internet standards documents that describe the protocols used on the Internet are called **Request For Comments** or **RFCs**. The RFCs are maintained by the Internet Engineering Task Force (IETF) and can be viewed on the Web at <http://www.ietf.org>. The specific documents that describe HTTP are RFC 1945 and RFC 2616. Be forewarned: These documents are highly technical in nature.

When your web browser makes a connection to the web server, the browser sends an initial message to the server that looks something like this:

```
GET /content.html HTTP/1.0
Connection: Keep-Alive
Accept: image/gif, image/x-xbitmap, image/jpeg, image/pjpeg, image/png, */*
Accept-Charset: iso-8859-1,* ,utf-8
Accept-Encoding: gzip
Accept-Language: en, en-GB, de, fr, ja, ko, zh
Host: testserver:80
User-Agent: Mozilla/4.51 C-c32f404p (WinNT; U)
```

The **GET** line indicates the path part of the URL you're trying to receive and what version of the protocol you're accepting. In this case, you're accepting version 1.0 of the HTTP protocol.

The `Connection` line indicates that you would like this connection kept open for multiple page fetches. By default, a browser makes a separate connection for each frame, page, and image on a web page. The directive `Keep-Alive` asks the server to keep the connection open so that multiple items can be fetched using the same connection.

The `Accept` lines indicate what sorts of data you're willing to accept on the connection. The `*/*` at the end of the first `Accept` line indicates that you're willing to accept any kind of data. The next line (`iso-8859-1` and so on) indicates what character encoding can be used for the document. The `Accept-Encoding` line, in this case, says that the browser can accept content compressed with `gzip` (GNU Zip) for a faster transfer. Finally, `Accept-Language` indicates what languages are acceptable to this browser—English, English—Great Britain, German, French, and so on.

`Host` is the hostname from the URL that you're retrieving. Because of virtual hosting, it might be different than the main hostname of the server.

Finally, the browser identifies itself to the web server as `Mozilla/4.51 C-c32f404p (WinNT; U)`. In web terminology, the browser is called a **user agent**.

The server then sends a reply that will look something like this:

`200 OK`

`Date: Thu, 02 Sep 2004 19:54:39 GMT`

`Server: Netscape-Enterprise/3.5.1G`

`Content-Length: 2222`

`Content-Type: text/html`

`Last-Modified: Wed, 01 Sep 2004 17:12:03 GMT`

The reply is followed by the contents of the page you requested.

The status `200` indicates that everything went fine. The server also identifies itself on the `Server` line; in this case, the server is a Netscape-Enterprise/3.5.1G web server.

The `Content-Length` line indicates that 2,222 bytes of content will be sent back to the browser. Using this information, your browser knows that a page is 50 percent complete, 60 percent complete, and so on. The `Content-Type` is the kind of page that is being sent back. For HTML pages, this line is set to `text/html`. For an image, it might be set to `image/jpeg`.

The `Last-Modified` date indicates to the browser whether the page has changed since it was last fetched.

Most web browsers *cache* pages so that if you look at a web page twice, the date can be compared at this point to a saved copy the browser already has. If the page on the server hasn't changed, downloading the entire page again might not be necessary.

Example: Fetching a Page Manually

You can fetch a web page manually. This capability is often useful when you're just testing and want to make sure that a web server is sending correct replies.

To follow this example, you need a program called a **Telnet client**. The Telnet client is a remote-terminal access program used to log in to Unix workstations remotely; however, it's often useful for tasks such as debugging HTTP.

If you have a Unix machine, it's very likely that you have Telnet installed already. If you have a Microsoft Windows machine, Telnet may already be installed as part of your networking utilities. Simply use the Run option from the Start menu to run the Telnet client.

To start communicating with the web server, use Telnet like this at a prompt:

```
$ telnet www.webserver.com 80
```

Here, www.webserver.com is the name of the web server, and [80](#) is the port number you want to connect to (port 80 is typically where web servers are listening). If your Telnet client is a graphical one, you might need to set these values in a dialog box.

When Telnet connects, you might not receive a prompt or a connection message. Don't worry; that's normal. HTTP expects the client to talk first; the server isn't expected to prompt. Under Unix, you get a message that says something like this:

```
Trying www.webserver.com
```

```
Connected to www.webserver.com
```

```
Escape character is '^]'
```

Other kinds of systems, such as Windows, might not see this message.

You now need to carefully and somewhat quickly type the following:

```
GET http://www.webserver.com/ HTTP/1.0
```

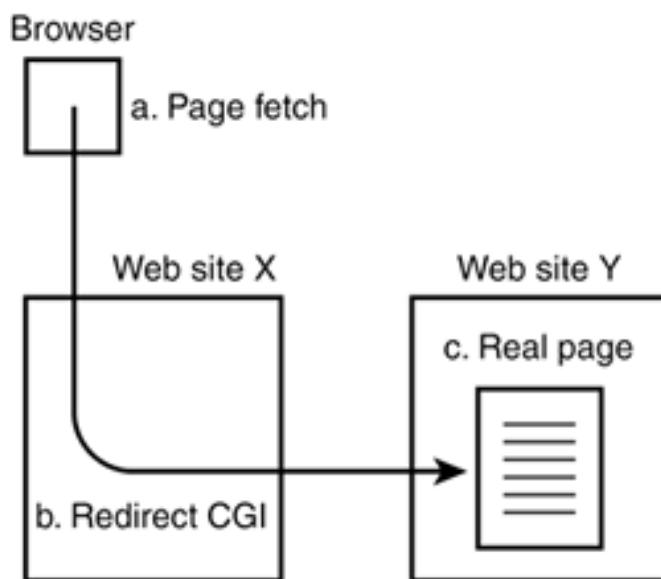
Press the Enter key *twice* after typing this line. The web server should then respond with a normal HTTP header and the top-level page for the web site and then disconnect.

Redirection

One useful trick to use in CGI programs is called **HTTP redirection**. Use redirection when you want a CGI program to load another page based on some computed value.

If you have a series of pages specific to a browser—for example, they contain a plug-in that's available only to Mozilla browsers under Microsoft Windows—you can send all the web site visitors to the same URL and have a CGI program redirect them to the correct page, as illustrated in [Figure 24.1](#).

Figure 24.1. One web site redirecting to another.



Another example might be where your CGI program has produced some kind of output: a PDF file, a ZIP file, an Excel Spreadsheet, or other downloadable file and you've placed it in a file somewhere else on the web server. You can then redirect the browser to pick up the file so that the CGI program doesn't have to be active during the download itself.

To implement a redirection, you need to use the CGI module's `redirect` function. The `redirect` function manipulates the HTTP conversation discussed earlier and causes the browser to load a new page.

[Listing 24.1](#) contains a short program to redirect users of Mozilla under Windows to one page and all other browsers to a different page.

Watch Out!

The `redirect` header has to be printed before anything else is emitted from the CGI script.

Don't print the output of the `header()` function or anything else before calling `redirect()`.

Listing 24.1. Redirection Based on Browser

```
1:  #!/usr/bin/perl -w
2:
3:  use CGI qw(:all);
4:  use strict;
5:  my($browser, $target);
6:
7:  # Fetch the browser's name
8:
9:  $browser=user_agent();
10: $target="http://www.server.com/generic.html";
11:
12: # Test for WinXX and Mozilla
13: if ($browser=~/Mozilla/ and $browser !~ /MSIE/ and $browser=~/Win/) {
14:     $target="http://www.server.com/netscape.html";
15: }
16: print redirect( -uri => $target );
```

Line 9: The browser type is captured in `$browser`.

Line 10: The default URL is put into `$target`. Any non-Mozilla browsers will be sent here.

Lines 13–14: The browser identification, stored in `$browser`, is checked to see if it contains `Mozilla` or `Win` but not `MSIE`. If so, a new target address is assigned.

Line 16: The redirection message is sent to the browser.

Redirection through CGI is seamless, whereas other techniques such as using JavaScript and HTML extensions

have problems. JavaScript is not supported on all platforms, and using a `window.location.href` assignment in JavaScript might not produce the proper results. Using an HTML `<META HTTP-EQUIV="refresh" >` tag for redirection causes a noticeable delay because the browser has to load the page completely before the redirection can take place. JavaScript shares this problem. HTTP redirection happens before any HTML is transmitted and is nearly instantaneous.

By the Way

The `user_agent` name returned by a typical Windows XP Firefox browser is similar to Firefox/1.0.1 (Windows NT 5.1; U; pl-PL). Microsoft Internet Explorer 6 identifies itself as Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1). Coding a script that guesses the browser in a foolproof manner is almost impossible though because a browser can lie about its user agent type, and often will.



< Day Day Up >



More Details on Calling CGI Programs

So far, you've learned two different techniques for starting a CGI program. The first and most obvious is simply to call the CGI program's URL from a link or have the user type it in the browser. A line like this will start and run a CGI program called `time.cgi`:

```
<A HREF="http://server/cgi-bin/time.cgi">Click here for the time</A>
```

When this link is followed, the web server runs the CGI program `time.cgi`, and its output is displayed as a new web page. This example is simple, easy, and straightforward—just like the "Hello, World!" CGI program in Hour 17, "Writing Modules."

The other way to start a CGI program is to make it the target of an HTML form. For example, the following form calls the CGI program `process.cgi` when the submit button is clicked:

```
<form method="get" action="/cgi-bin/process.cgi">  
  
<input type="text" name="STUFF"/><br />  
  
<input type="submit"/>  
  
</form>
```

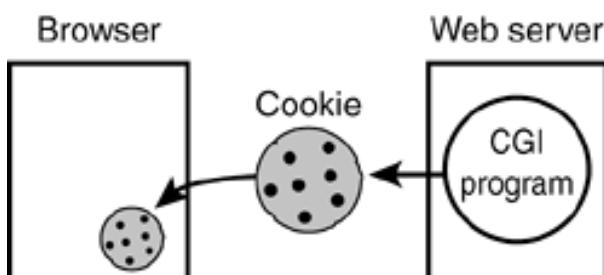
This method of calling a CGI program has an added benefit: You can pass parameters to the CGI program for processing. Well, that's the whole point of an HTML form.

Passing Parameters to CGI Programs

Wouldn't it be nice to pass information into a Perl program from a link? For example, you could have a clickable link in a document that would "run cgi program `foo.cgi`, with value `X=this` and value `Y=that`." This is possible with a little bit of work.

First, you have to use a special kind of URL in your `<A HREF>` tag. The format of this URL is illustrated in Figure 24.2.

Figure 24.2. URL containing parameters.



Each parameter is the name of a value that you want to pass into your CGI program—such as a named HTML form element. The value is the value for that name. For example, to create a link that, when clicked, will run a CGI program with the parameter `sign` set to `Aries` and `year` set to `1969`, you would enter the following:

```
<a href="http://www.server.com/cgi-bin/astrology.cgi?sign=Aries&year=1969">  
Aries, year of the Rooster</a>
```

Inside the CGI program, the parameters are processed with the CGI module's `param` function as normal:

```
#!/usr/bin/perl -w  
  
use CGI qw(:all);  
  
use strict;  
  
print header;  
  
print "The year ", param('year'), " and being born under ",  
      param('sign'), " indicates you are brilliant.\n";
```

You can pass as many parameters as you want. If you want to pass an empty parameter—one with no value—simply leave it off, as in `author` in the following example:

```
<a href="http://www.server.com/cgi-bin/book.cgi?author=&title=Beowulf">Beowulf</a>
```

Special Parameter Considerations

You need to be aware of some special considerations when calling CGI programs with parameters. Some characters are special and cannot be made part of a URL. For example, the `?` (question mark) is special: It marks the separation between the main portion of the URL and the parameters. Other characters that are special are `&`, spaces, and quotation marks.

By the Way

The full list of special characters is enumerated in an Internet standards document called RFC 2396.

To insert one of these special characters into a URL requires that you **escape** the character. In this case, escaping the character means to translate its ASCII value into a two-digit hexadecimal number and then precede it with a percent sign. The encoding for "Hello, World!" is as follows:

```
Hello%2C%20World
```

Obviously, creating a URL-escaped string can be messy. The CGI module provides a function to create such strings automatically for you. The following snippet demonstrates how to print a URL with the proper encoding:

```
#!/usr/bin/perl -w

use strict;

# The 'escape' function must
```

< Day Day Up >

NEXT 

Cookies

In Hour 23, "Complex Forms," you learned how to make your web browser remember information between web pages by using hidden fields in HTML. You must understand this process because you might need to pass information from one instance of a CGI program to another. The only way to do this is to store a piece of information with the browser.

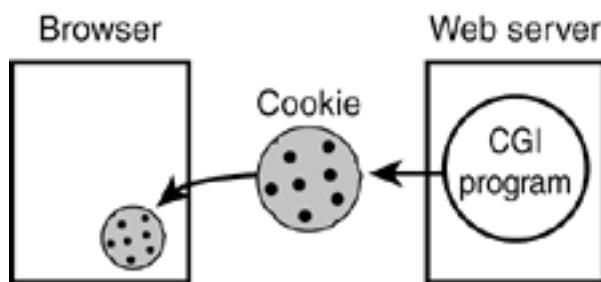
Another way to store information with the browser is to use **HTTP cookies**. HTTP cookies are pieces of information passed between the browser and the CGI program during the HTTP connection. Using cookies can be a much more flexible method of storing information with a browser than using hidden HTML fields.

You can think of a cookie as a movie theater ticket. You can go to the theater ticket booth and obtain a ticket for a subsequent showing. You can then leave the theater, come back, buy some popcorn, and do anything else you'd like. When you're ready to see the movie, you present your ticket to the ticket-taker. The ticket-taker doesn't know how, when, or why you purchased the ticket, but as long as you have it, the ticket-taker will let you into the theater. The ticket entitles the bearer to admission to a subsequent show.

An HTTP cookie is simply an information packet that the CGI program asks the browser to hold. The packet can be reclaimed at any time by another CGI program or by the same program. Cookies are even passed back to the server when regular HTML pages are requested. The cookie can contain any kind of information: information about multipage web forms, visitation information, user preferences, and so on.

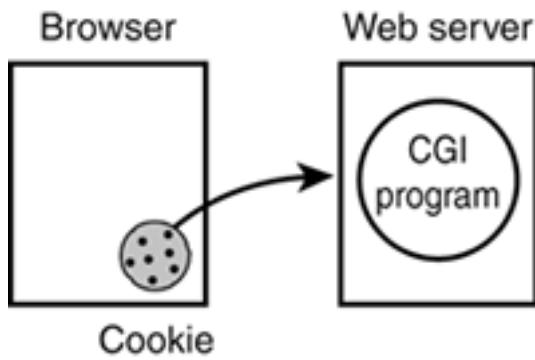
The cookie is transferred from the server to the browser whenever a CGI program requests that a cookie be created (see [Figure 24.3](#)); this process is called **setting a cookie**.

Figure 24.3. Cookie going to the browser from a CGI program.



The cookie can be later reclaimed by a CGI program to retrieve the information stored in the cookie, as illustrated in [Figure 24.4](#).

Figure 24.4. Cookie being returned to the server by the browser.



Why Is a Cookie Called a Cookie?

In computing circles, *cookie* is a very old term. It refers to any piece of information passed between routines or programs that enables the holder of the cookie to perform some operation. Some kinds of cookies are called **magic cookies** because they contain data that's obscure and meaningful only to the sender and the receiver of the cookie. CGI cookies are not magic.

How to Make Cookies

To create a cookie, you can use a CGI function called `cookie`. The syntax of the `cookie` function that you're concerned with is as follows:

```
$cookie_object=cookie( -name => cookie_name,           # Optional
                      -value => cookie_value,
                      -expires => expiration_date, # Optional
                      -path   => path_info,       # Optional
) ;
```

The `cookie` function takes arguments in a way you haven't seen yet, but that is quite common in Perl modules. Each argument in the call to `cookie` has a name, so you don't have to remember the order in which the arguments occur; they simply get named in whatever order you decide to use them.

When called with the first syntax, the `cookie` function returns a cookie, which should be stored in a scalar variable and can then be given to the CGI module's `header` function to be sent to the browser. The only required argument to create a cookie is the `-value` argument. The `-name` argument allows several cookies to be sent to the browser at one time, which can be retrieved individually or as a group. The `-expires` and `-path` arguments are discussed in the next section.

The `header` function in the CGI module takes care of actually transmitting the cookie to the browser. This means you must create the cookie, by using the `cookie` function, and then call the `header` function soon afterward. You shouldn't send any other kind of data to the browser until after the cookie and the header have been sent.

To create a cookie and send it to a browser in a CGI program, you can use a CGI program similar to this:

```
#!/usr/bin/perl -w

use CGI qw(:all);

use strict;

my $cookie=cookie(-name => 'Sample',
                   -value => 'This cookie contains no MSG');

# Transmit the cookie to the browser
print header(-cookie => $cookie);
```

After the preceding snippet has run, a cookie called `Sample` is set on the browser. The cookie contains the information '`This cookie contains no MSG`'.

Watch Out!

Actually, the cookie might not be set. Browsers can refuse to accept a cookie for many reasons. See "[Problems with Cookies](#)" later this hour.

To retrieve cookies from the browser in your CGI program, you use the same function—`cookie`. When called without any arguments, as shown in the following examples, `cookie` returns a list of cookies that the browser has for your server:

```
@cookie_list=cookie();    # Returns names of all cookies set
```

--or--

```
# Returns the value for a particular cookie  
  
$cookie_value=cookie($cookie_name);
```

By default, after the cookie is set on the browser, it is returned to any CGI program that resides on the same server. That is, only the server that set the cookies can fetch those cookies. To view the [Sample](#) cookie created previously, you can use another CGI program:

```
#!/usr/bin/perl -wT  
  
use CGI qw(:all);  
  
use strict;  
  
print header(); # Print out the standard header  
  
print "<p>Sample's value: ", cookie('Sample'), "</p>";
```

The preceding snippet uses [cookie](#) with one argument—the name of the cookie whose value you want to see. That value is retrieved and printed.

The browser should retain the cookie until the browser is terminated. When the browser is restarted, the cookie [Sample](#) will be gone. If you want to make a more permanent cookie, see "[Long Term Cookies](#)" later in this hour.

By the Way

Most browsers have an option to view the cookies as they're being set. In Netscape, you can find the cookie viewing options under the Preferences selection on the Advanced tab. In Internet Explorer, the selection appears on the Advanced tab of the Internet Options dialog, and a radio button controls whether you can view cookies as they're set.

Example: Using Cookies

For this example, you'll create a small program to allow the user using the web browser to set the color of the web page he or she is viewing. The program actually does several things at once:

1.

Checks for a change in the default background color by checking for program parameters.

2.

Sets a cookie on the browser with the correct background color.

3.

Sets the background color of the page to the correct color.

4.

Displays a CGI form, allowing you to change the color.

[Listing 24.2](#) contains the color-changing program.

Listing 24.2. Complete Listing for *ColorChanger*

```
1:  #!/usr/bin/perl -w
2:  use strict;
3:  use CGI qw(:all);
4:  use CGI::Carp qw(fatalsToBrowser);
5:  my($requested_color, $old_color, $color_cookie)=(" ", " ");
6:  $old_color="blue";  # Default value
7:  # Is there a new color requested?
8:  if (defined param('color')) {
9:      $requested_color=param('color');
10: }
11: # What was the old color, if any?
```

```
12: if (defined cookie('bgcolor')) {  
13:     $old_color=cookie('bgcolor');  
14: }  
15: if ($requested_color and ($old_color ne $requested_color)) {  
16:     # Set the cookie in the browser  
17:     $color_cookie=cookie(-name => 'bgcolor',  
18:                           -value => $requested_color);  
19:     print header(-cookie => $color_cookie);  
20: } else {  
21:     # Nothing's changed, no need to set the cookie  
22:     $requested_color=$old_color;  
23:     print header;  
24: }  
25: print<<END_OF_HTML;  
26: <html>  
27: <head>  
28: <title>Set your background color</title>  
29: </head>  
30: <body bgcolor="$requested_color">  
31: <form>  
32: <select name="color">  
33:     <option value='red'>Red</option>  
34:     <option value='blue'>Blue</option>  
35:     <option value='yellow'>Yellow</option>  
36:     <option value='white'>White</option>
```

```
37:  </select>

38:  <input type="submit" value="Set the color"/>

39:  </form>

40:  </body>

41:  </html>

42: END_OF_HTML
```

Lines 7–10: If this program is called as the target of a CGI program, the `param('color')` function returns a defined value—a new color. Otherwise, it returns nothing, and `$requested_color` remains unset.

Lines 12–14: These lines check for a cookie named `bgcolor`; that cookie might or might not exist. If it does exist, it is stored in `$old_color`, which is the last screen color value that was saved to a cookie.

Lines 15–19: If the color has changed (the cookie value doesn't match the new value), a new cookie needs to be set with the new value.

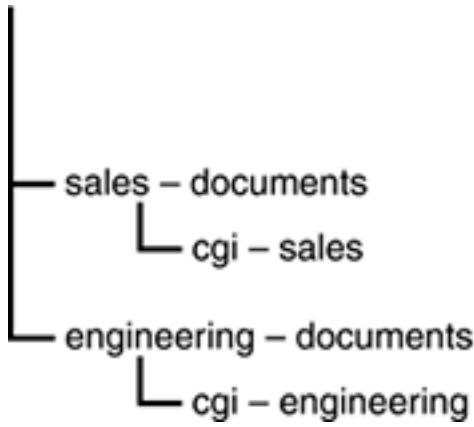
Lines 20–24: Otherwise, a plain header is printed without a cookie. Remember: The browser will retain the previous cookie indefinitely.

Lines 25–42: These lines create a standard HTML form. Note line 30, however; on this line, the color gets substituted into the HTML output.

Restricting Cookies

It's also possible to restrict a cookie's allowed return destinations. By default, when you create a cookie, that cookie is returned to any URL on that web site—including non-CGI URLs. For example, consider an automotive web site organized like the diagram in [Figure 24.5](#).

Figure 24.5. Directory tree of a multiuse web site.



Having the sales CGI programs and the engineering CGI programs reside in different directories might make sense. If a sales CGI program were to set a cookie, the engineering CGI programs would pick it up, and vice versa. This result might not be desirable, and the people writing CGI programs for both sites need to coordinate to make sure they didn't reuse each other's cookie names.

To get around this problem, you can use the `cookie` function's `-path` option. It indicates the pathname—relative to the top of the URL—to which the cookie should be returned. For example, to send a cookie that will be returned only to the sales CGI programs, you could use this snippet:

```
# Cookie only visible to sales CGI programs

$cookie=cookie( -name => 'profile',
                -value => 'sedan,luxury,2-door',
                -path => '/cgi-sales');

print header(-cookie => $cookie);
```

By default, cookies are returned to every site on the server as though you had used the option `-path=> '/'`. To restrict the cookie to return to only one CGI program, you can use the CGI program's URL in the `-path` option:

```
# Return the cookie only to this program

$cookie=cookie( -name => 'profile',
                -value => 'sedan,luxury,2-door',
                -path => script_name() );

print header(-cookie => $cookie);
```

Remember that the `script_name` function in the CGI module returns the partial URL of the current CGI program. This effectively creates a cookie that will be returned only to the program that set the cookie on the browser.

 PREV

< Day Day Up >

NEXT 

Long Term Cookies

So far, the cookies you've set on the browser have been temporary: as soon as the browser is closed, the cookie disappears. When you're using cookies to save values among multiple pages on a form—instead of hidden HTML values—using temporary cookies is entirely appropriate. When a new browser is started, you don't want the cookie returned to the server because the user wouldn't be filling out the form starting in the middle; he or she would be starting at the beginning.

In some cases, you might like the cookie to hang around much longer. Perhaps you want the cookie to persist for days, weeks, or months after the browser is closed and restarted. Creating these cookies is easy with Perl's CGI module.

To set an expiration date for a cookie, you can use the `-expires` option when you create the cookie. The `-expires` option must be followed by a date that you want the cookie to expire. You can specify this date in several formats, as shown in [Table 24.1](#).

Table 24.1. Formats for Cookie Expiration Dates

Format	Sample	Meaning
Number of seconds	<code>+30s</code>	30 seconds from now
Number of minutes	<code>+15m</code>	15 minutes from now
Number of hours	<code>+12h</code>	12 hours from now
Number of months	<code>+6M</code>	6 months from now
Number of years	<code>+1y</code>	1 year from now
	<code>now</code>	Cookie expires immediately
Any negative time	<code>-10m</code>	Cookie expires immediately
A specific time	<code>Saturday, 28-Aug-1999 22:51:05 GMT</code>	

When specifying an expiration time, you must follow the format listed in [Table 24.1](#) exactly. All the other possible values are offsets from the current time, and the fully qualified time will be computed for you and sent to the browser.

This small program sets a cookie on the browser that expires in eight days:

```
#!/usr/bin/perl -w

use CGI qw(:all);

use strict;

my $cookie=cookie(-name => 'Favorite',
                   -value => 'soft oatmeal raisin cookies',
                   -expires => '+8d' );

# Transmit the cookie to the browser

print header(-cookie => $cookie);
```

Persistent cookies do not last indefinitely. That is, if you send a cookie to the browser and expect that the cookie will be around weeks, months, or years from now, you might be disappointed.

As you'll learn later in "Problems with Cookies," browsers don't have to store cookies. In fact, they don't have to accept your cookies at all—and your CGI program is not informed that the cookies aren't kept.

Browsers can flush out their cookies at any time to make room for new cookies received from other sites—or for no reason at all. Some browsers can allow users to edit cookies or to add new ones.

Users can erase the cookies, either accidentally or on purpose. If the users install a new version of a browser or an operating system, the cookies can be wiped out or misplaced. Simply using a different browser can make the cookies seem to "disappear." The cookies are usually stored in a file when the browser isn't active, and that file can be edited by users, erased, or corrupted.

By the Way

If you're curious, most browsers store the cookies in files when they're not active, and usually they're just text files that you can view with an editor. Netscape stores the cookies in a file called `cookies.txt` under the user's home directory (which varies from system to system). Internet Explorer stores cookies under `\Windows\Cookies`.

Storing critical information in an HTTP cookie, therefore, is really not a good idea. Any information that you want to store persistently in a cookie should be easily replaced—user preferences, a replaceable entry key for

a restricted web page, last-visited information, and so on.

 PREV

< Day Day Up >

NEXT 

Problems with Cookies

Before using cookies in an application, you should be aware of the problems associated with cookies. For these reasons—and others that may crop up in the future—you should be careful to design your web pages and CGI programs so that cookies are entirely *optional*.

For example, if you're using cookies to store user preferences, you should use a default set of preferences if cookies are not available. You must code defensively.

Cookies Are Ephemeral

As I've mentioned throughout this hour, cookies are ephemeral. Cookies can be deleted from the user's system, edited by the user, or simply thrown away by the browser for no apparent reason. A browser can accept the cookie, use it for a short time, and then simply forget it for no apparent reason. If you've set up persistent cookies using `-expire`, the browser can still throw away the cookie without telling the user.

Cookies Aren't Always Supported

Not all browsers support HTTP cookies. The Internet standards that apply to HTTP and web traffic don't guarantee that browsers must support cookies.

I'm not suggesting that *most* browsers don't support cookies, because they do. Netscape (since version 1.1), Internet Explorer (all versions), Lynx, Opera, and most popular web browsers do support cookies. In most of them, though, the user has the option to turn off cookie support.

Even if you've used the CGI module's `user_agent` function and determined that the intended browser should support cookies, don't count on it.

Some People Don't Like Cookies

The title of this section might be hard to fathom: Why on earth would someone not like cookies?

Surfing the web is essentially an anonymous activity. As you've seen in previous lessons, when a web page is requested by a browser, the request happens in a vacuum. The server doesn't necessarily know where the browser's been or the last time a page on this site was requested by this particular browser.

By the Way

Remember: One browser does not necessarily represent one user. A browser can be shared by many people in a household, kiosk, Internet café, or public access site, such as a library. Setting (or changing) a cookie for one person might actually set it for several people.

Cookies can be used to track where people have been on a particular site and what items they've clicked. If you're concerned about privacy, this might be alarming. For example, a bookstore might track which books the web surfer clicked on for more information. That information can be used to tailor lists of books to present to the surfer.

On the surface, this feature seems convenient. For people concerned about privacy, it presents two problems. First, now an entity keeps track of the kinds of books in which the web surfer is interested. If that information is ever correlated with the web surfer's name and address—possibly obtained from a fill-out form from another site that is sharing information with [Congo.com](#)—the surfer might start receiving junk mail related to his or her book selections. The more information that's shared by cookie-gathering sites, the more detailed a "profile" that can be assembled about the web surfer.

Beyond the privacy issue, if the first two books examined were categorized as "Computer" books, the web site might stop offering books to the web surfer in the "Romance" and "Cooking" categories. The web site channels the surfer into categories that might be difficult to escape.

By the Way

You might be surprised at how often cookies are used to collect and store information on your browser. Just turn on cookie confirmation in your browser, and visit any popular web site.

People have been known to go to great lengths to avoid cookies. Web browsers that support cookies all offer features to turn off cookies, and some allow you to examine cookies as they're set on the browser. Add-on software packages are available to filter cookies that are sent to and from your browser and allow you to edit them. Web sites are designed to allow you to surf other web sites anonymously without having cookies gathering information about you.

In short, some people consider HTTP cookies as an invasion of privacy, so you should be careful how you use them.



< Day Day Up >



Summary

In this chapter you learned to use the HTTP protocol to achieve a couple of new effects. After examining how web pages are fetched, you learned that parameters could be passed in the URL and what kinds of identification information is sent in the HTTP request. On the server side, by adding a redirection header or a cookie header to the HTTP response, you can redirect the browser elsewhere or send cookies.

Q&A

Q1: *The Telnet example doesn't work.*

A1: If Telnet fails to connect, make sure you're using Telnet against the Web server's name and that you're using the correct port, probably 80. You might need to look at the Telnet client's documentation to set the port number properly.

Another common problem is that you can't see yourself type. Some Telnet clients echo the characters to you, and some do not. Don't worry about this; just type carefully. The characters are still being transmitted carefully. After you type the `GET` line, make sure that you press the Enter key *twice*.

Q2: *How do I put more than one piece of information in an HTTP cookie?*

A2: The easiest way is to combine multiple items in a single cookie separated by a field separator, as in this example:

```
$cookie=cookie(-name => 'preferences',
               -value => 'bgcolor=blue,fgcolor=red,banners=no,java=no');
```

Later, when you retrieve the cookie, you can use `split` to separate the items:

```
$cookie=cookie('preferences');

@options=split(/,,,$cookie);

# Now, make a hash with the option as the key,
# and the option's value as the hash value

foreach $option (@options) {
    ($key,$value)=split(/=/, $option)
    $Options{$key}=$value;
}
```

Q3: *How do I use cookies to track which links a user clicks on a Web page?*

- A3:** Before I give you the answer, you must realize that some people consider tracking an invasion of privacy. Having said that, let me show you the general method:

1.

Write your `<a href>` links so that they go to a CGI program, passing the *real* target URL as a parameter:

```
<a
```

```
    href="http://server/cgi/redirect.pl?target=http://www.congo.com">Congo</a>
```

2.

The `redirect.pl` program in the example should use the CGI module's `param` function to get the real URL (`http://www.congo.com`) from the parameter `target`:

```
$target_url=param('target');
```

3.

A cookie is then created with the target URL in the value, with a name you can search for later, such as:

```
$tracking_cookie=cookie(-name => 'tracker',
                         -value => $target_url,
                         -expires => '+1w');
```

4.

The redirect is then sent to the browser along with the cookie:

```
print redirect(-uri => $target_url,
               -cookie => $tracking_cookie);
```

Later, when the browser returns to your Web site, you can look for the cookie named `TRacker`, which will contain the URL that the user visited when he or she left your site.

There are other methods of tracking users, this is just one.

- Q4:** Can I transmit a cookie while redirecting the browser to another page?

A4: Yes. The CGI module's `redirect` function can also take a `-cookie` argument such as the `header`:

```
my $cookie=cookie(-name => 'target',
                   -value => 'redirected to foo.html');

print redirect(-uri => "http://www.server.com/foo.html",
               -cookie => $cookie);
```

One complication: If you're using IIS as a web server, to send a cookie and a redirect at the same time you will need to enable Non-Parsed-Header scripts (usually done by having `nph-` at the beginning of the script name).

[PREV]

< Day Day Up >

[NEXT]

Workshop

Quiz

- 1:** The following URL will work as expected.

```
<a href="/cgi/foo.pl?name=Ben Franklin&Job=printer">
```

a.

True.

b.

False; you can't pass two parameters to a CGI program in this way.

c.

False; the space in the name Ben Franklin is not allowed.

- 2:** Why does using cookies for long-term storage sometimes fail?

a.

Browsers can "throw away" cookie information.

b.

Cookies can be lost during software updates.

c.

Users may turn off cookie support in their browsers.

3: To make a cookie expire one week from now, use the `-expire` option to the `cookie` function with what?

a.

`+7d`

b.

`+1w`

c.

`+10080m`

Answers

A1: c. You should use `escape` to properly hide spaces and other special characters.

A2: All the above.

A3: Either choice a or c. The `+1w` argument is invalid.

Activities

- Use a Telnet client to connect to one of your favorite Web sites, and try fetching a Web page manually.
- Extend the background color changer to set foreground colors and fonts. Randomly select an image to display on the page by editing the target of an `` tag.

Part IV: Appendixes

APPENDIX A Installing Modules

Appendix A. Installing Modules

Installing modules in Perl isn't hard, and learning how to install them is essential if you ever want to master Perl. This appendix contains information on how to install the modules you need.

By the Way

Detailed documentation on installing modules under almost any architecture is available in Perl's documentation. The document, called "perlmodinstall," even contains instructions for building under obscure operating systems such as OS/2 and VMS.

Picking the Right Module

First, you need to pick the right module. A good starting place is the CPAN itself at <http://www.perl.com/CPAN>. You need to decide which module you're interested in.

CPAN modules are named, roughly, by what they do. For example, `Image::Size` takes an image and reports what size the image is; this module is useful for working with Web pages. Some modules have unusual names though. `LWP` derives its name from a Perl library called `libwww-perl`.

You also can find bundles of modules on CPAN. These bundles contain several related modules—and usually prerequisite modules—all in one large package. For example, the `libnet` bundle installs like a module, but in the process, you get several modules all related to networking. `LWP` is part of the `libnet` bundle, for example.

By the Way

When you install a module, you also automatically get all the documentation needed for that module.

Installing the Modules Under...

For each of the examples in the following sections, you will attempt to install the `Date::Manip` module from CPAN. To install your own module or bundle, simply substitute the name of your module in place of `Date::Manip`.

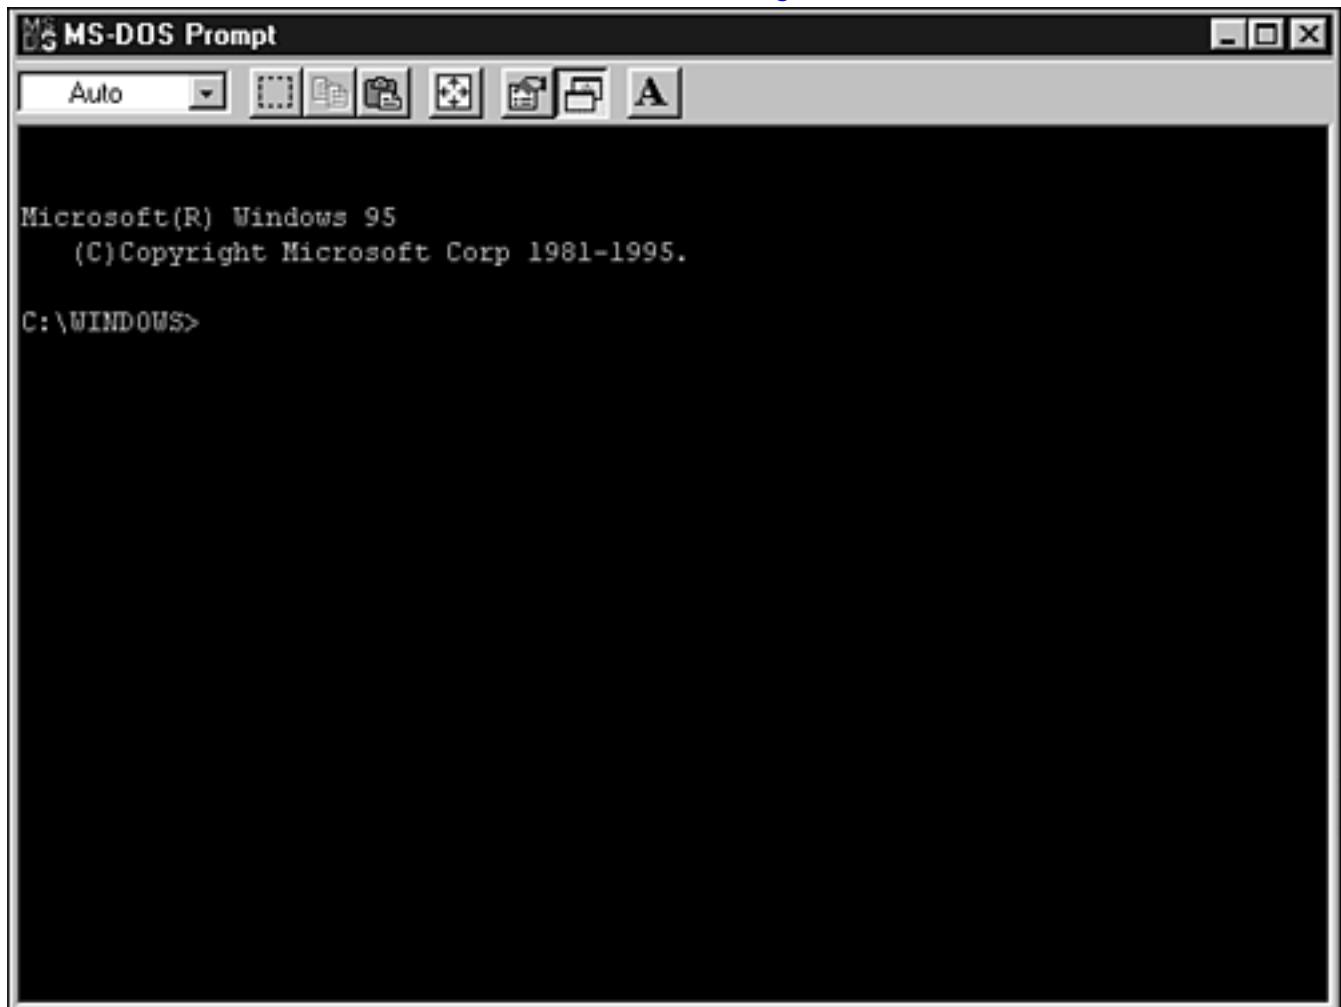
...Windows

With Windows, the easiest solution is to use prepackaged modules available from ActiveState Tool Corp., assuming, of course, that the Perl installation you have is from ActiveState.

To install a prebuilt module under Windows, you first need to start the Perl Package Manager (PPM). This utility simplifies the module-building process for you by providing an interactive interface to module installation. To start the PPM, you need to open a DOS command prompt, as shown in [Figure A.1](#); you should be connected to the Internet.

Figure A.1. You can work from the DOS command prompt.

[View full size image]



At the command prompt, simply type `ppm`, as shown here. The PPM utility should then start; if it doesn't, you need to look for the `ppm.bat` file that was installed with ActiveState Perl and run it with the full pathname.

```
C:\Windows>ppm
```

```
PPM interactive shell (2.1.6) - type 'help' for available commands
```

```
PPM>
```

To search for a particular module, use the `search` command, as shown here. You need to use this command because ActiveState doesn't have prebuilt packages for all the modules in CPAN, only the more requested ones. Also, for installation, you have to be able to spell the name of the module correctly.

```
PPM> search Date
```

```
Packages available from http://www.ActiveState.com/packages:
```

```
Date-Calc
```

```
Date-Manip
```

```
TimeDate
```

```
PPM>
```

I've summarized the results of the search. There are actually more than 40 available date modules. After you've found the module you want—`Date-Manip` for this example—you can install it by using the `install` command, as follows:

```
PPM> install Date-Manip
```

```
Install package 'Date-Manip?' (y/N): y
```

```
Installing C:\Perl\html\lib\Date\Manip.html
```

```
Installing C:\Perl\htmlhelp\pkg-Date-Manip.chm
```

```
Installing C:\Perl\htmlhelp\pkg-Date-Manip.hhc
```

```
Installing C:\Perl\site\lib\Date\Manip.pm
```

```
Writing C:\Perl\site\lib/auto/Date-Manip/.packlist
```

```
PPM>
```

The `Date::Manip` module is now installed!

If you want to download the bundle and install it manually—perhaps the PC doesn't have Internet access or is behind a firewall—you can find instructions for downloading and installing modules by hand and dealing with firewalls on ActiveState's Web site (<http://www.ActiveState.com>). ActiveState maintains a FAQ specific to its distribution of Perl, and you can find the necessary instructions there.

By the Way

Installing modules without PPM—for example, using your own C compiler under Microsoft Windows—is well beyond the scope of this book. The source distribution of Perl contains instructions for building Perl yourself under Windows, but this job is not for beginners. If you can accomplish that, building modules for yourself shouldn't be too difficult because the processes are about the same.

...UNIX, Using CPAN

Installing modules under UNIX can be exciting and full of challenges, or it can be ridiculously easy. You need an ANSI C compiler (the one used to build Perl will be fine) and a compiler license if your vendor requires one. You also need a copy of the GNU compression program gzip/gunzip; some UNIX vendors ship it as a standard utility. You can obtain a copy from <http://www.fsf.org> if you don't have it.

One last thing: You'll probably need root privileges on the installation machine. Normally, Perl is installed as a systemwide utility. Installing modules into the system directories requires that you have adequate permission to do so—root permission.

The Perl installation comes with a module called CPAN that is designed to help you install other modules. To get the installation started, you start Perl using the CPAN module's shell like this:

```
$ perl -MCPAN -e shell
```

The first time you run this command, the CPAN module interrogates you to find out where you want to fetch Perl modules from and how you want them built. Most of the time, the default answers are sufficient. You are asked about the location of temporary directories, which CPAN mirror you want to use (a list is presented), and whether you're accessing the Internet through a proxy.

When CPAN is done interrogating you, you are presented with this prompt:

```
cpan shell -- CPAN exploration and modules installation (v1.3901)

ReadLine support available (try [ag][ag]install Bundle::CPAN'')
```

cpan>

At this prompt, you can use the command `i /pat/` to search for information on a package, where `pat` specifies the pattern to search for. To find the `Date::Manip` module, for example, search like this:

```
cpan> i /Manip/
```

The CPAN module may need to contact a CPAN server and obtain a new copy of the index. This situation happens only when needed, and the process takes just a moment. When the query is done, CPAN replies with something like the following information:

Distribution	SBECK/DateManip-5.35.tar.gz
Module	Date::Manip (SBECK/DateManip-5.35.tar.gz)

To install the module, type the following:

```
cpan> install Date::Manip
```

Now, the CPAN module goes through the process of fetching, compiling, testing, and installing the module for you. The display is quite messy but looks something like this greatly abbreviated example (# comments normally do not appear, they're added for clarity here):

```
Running make for SBECK/DateManip-5.35.tar.gz

Fetching with LWP: # Fetching the module
```

```
ftp://ftp.perl.org/pub/perl/CPAN/authors/id/SBECK/DateManip-5.35.t  
ar.gz
```

```
Writing Makefile for Date::Manip
```

```
mkdir blib          # Building the module
```

```
mkdir blib/lib
```

```
Target "makemakerdflt" is up to date.
```

```
/usr/bin/make -- OK
```

```
Running make test      # Testing to ensure it works
```

```
PERL_DL_NONLAZY=1 /usr/bin/perl -Iblib/arch -Iblib/lib
```

```
-I/usr/local/lib/
```

```
perl5/5.00502/aix -I/usr/local/lib/perl5/5.00502 -e 'use
```

```
Test::Harness qw(&runte
```

```
sts $verbose); $verbose=0; runtests @ARGV;' t/*.t
```

```
t/settime.....ok
```

```
t/unixdate.....ok
```

```
All tests successful.
```

```
Files=30, Tests=826, 178 wallclock secs (168.85 cusr + 5.23
```

```
csys = 174.08 CPU)
```

```
Target "test" is up to date.
```

```
/usr/bin/make test -- OK
```

```
Running make install      # Installing the module
```

```
Target "install" is up to date.
```

```
/usr/bin/make install -- OK
```

Your output will vary considerably from the preceding. The module is now tested and installed. Enjoy!

...UNIX, The Hard Way

Although you can install modules in UNIX without using the CPAN module, most of the time you don't need to install this way. This method is presented for completeness only; you should really use the CPAN module whenever possible.

First, you need to download the module from CPAN. It is a gzip'ed tar bundle. If the module is `Date::Calc`, for example, you would need to obtain the latest version, which has a name something like `Date-Calc-X.Y.tar.gz`. After you've downloaded the bundle, go to that directory and unpack the bundle like this:

```
$ gunzip Date-Calc-4_2.tar.gz
```

```
$ tar xf Date-Calc-4_2.tar
```

Unpacking creates a subdirectory called `Date-Calc-4.2`. Change into that subdirectory by using `cd`, and type the following command:

```
$ perl Makefile.PL
```

Checking if your kit is complete...

Looks good

```
Writing Makefile for Date::Calc
```

Now you have a makefile, which is required for the build process. Next, build the module by using the `make` command as follows:

```
$ make
```

```
mkdir blib
```

```
mkdir blib/lib
```

```
:
```

```
Manifying blib/man3/Date::Calc.3
```

```
Target "makemakerdfilt" is up to date.
```

This process might take awhile.

At the next prompt, you need to test the module to see whether it is built correctly. Just type `make test` as follows:

```
$ make test  
  
PERL_DL_NONLAZY=1 /usr/bin/perl -Iblib/arch -Iblib/lib  
-I/usr/local/lib/  
  
perl5/5.00502/aix -I/usr/local/lib/perl5/5.00502 -e 'use  
Test::Harness qw(&runte  
sts $verbose); $verbose=0; runtests @ARGV;' t/*.t  
  
t/f000.....ok  
t/f001.....ok  
:  
t/f032.....ok  
t/f033.....ok  
  
All tests successful.  
  
Files=34, Tests=1823, 14 wallclock secs ( 9.81 cusr + 1.10 csys  
= 10.91 CPU)  
  
Target "test" is up to date.
```

You should *always* run the `make test` command to ensure the module was built correctly. It can save you (and others) hours of debugging later. When the test is complete, you need to install the module, as shown here. This step is usually done as the root user because the installation must write to system directories.

```
$ su  
Password: *****
```

```
# make install
```

Installing

```
/usr/local/lib/perl5/site_perl/5.005/aix/auto/Date/Calc/Calc.so
```

:

Appending installation info to

```
/usr/local/lib/perl5/5.00502/aix/perllocal.pod
```

Target "install" is up to date.

#

At this point, you're done!

...Mac OS X

As noted earlier in this book, under the hood Mac OS X is essentially Unix. Use the instructions from [Hour 1](#), "Getting Started with Perl," to get to a prompt, and then follow the Unix instructions for installing modules.



< Day Day Up >



What to Do When You're Not Allowed to Install Modules

If you can install programs on a system, you can install modules. Your ability to do so depends on how complex the module is and what kind of trouble you're willing to go through. Sometimes the system administrator does not allow you to install a module because he or she doesn't want it used by others. In some cases, only you, or a group of people, want specific modules, and having them installed systemwide is too much hassle.

In either case, installing private copies of Perl modules in your own directories is not difficult.

First, you need to build the module using the instructions given previously—with a small exception. You can specify to the installation program that you want the modules placed in a particular directory. With PPM under Microsoft Windows, before you install the module, you need to tell PPM that you want to install into a different directory. You do so by using the `set` command like this:

```
PPM> set root c:\myperl
```

```
PPM> set build c:\myperl
```

The module is then assembled and installed in the directory `c:\myperl`.

Under UNIX, when you're using the CPAN module, you can specify the installation directory by using the `makepl_arg` setting like this:

```
cpan> o conf makepl_arg PREFIX="/home/clintp/perl/lib"
```

Or, if you're installing modules by hand using the `make` utility, you can specify the installation directory by using the `PREFIX` argument on the first line during the build:

```
$ perl Makefile.PL PREFIX="/home/clintp/perl/lib"
```

Using either method, the module you're trying to install is installed into `/home/clintp/perl/lib`. You can then move this module to another directory if you need to.

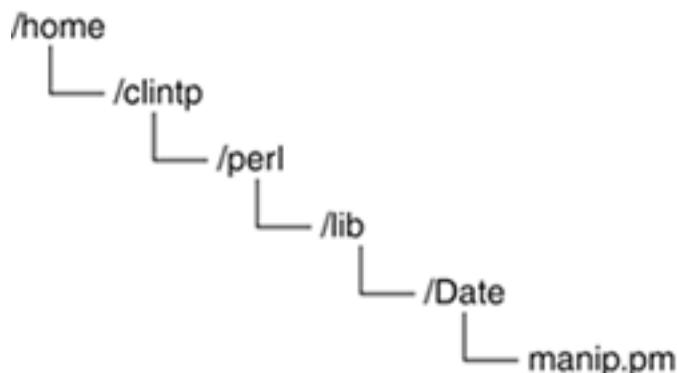
Watch Out!

You should be careful not to move modules between different kinds of machines. A compiled module will only work on one kind of system, like Perl itself. Also, try not to move modules between different versions of Perl; it sometimes does not work. In that case, you have to rebuild the module.

Using Modules Installed in Strange Places

To use modules that are installed in a directory other than the standard directory, you need to use a directive called `use lib`. For example, if you install the module `Date::Manip` in the directory `/home/clintp/perl/lib` using the instructions in the preceding section, you would have a file tree like the one shown in [Figure A.2](#).

Figure A.2. File tree after installation for Date::Manip.



At the beginning of your program, you can simply put the following:

```
use lib '/home/clintp/perl/lib'; # Look for module here  
else  
use Date::Manip;
```

Perl then searches that directory for the module before it searches any of its own directories. You also can use this technique to install newer versions of modules onto the system—for testing—without overwriting older versions and creating incompatibilities.

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

! (logical not) 2nd

! (not)

\$ variable

 warning against storing values 2nd

\$_ (special scalar variable)

% (hash type identifier)

% (modulus operator)

&& (logical and) 2nd

(backslash)

 references

* (multiplication operator)

* (quantifier)

** (exponentiation operator)

+ (addition operator)

+ (quantifier)

++ (autoincrement operator)

- (subtraction operator)

-- (autodecrement operator)

-c (check syntax)

 command line switch

-d (debug)

 command line switch

-e (execute one-line program)

 command line switch

-n (input loop)

 command line switch

-p (input loop with print)

 command line switch

-V (detailed version)

 command line switch

-v (version)

 command line switch

. (concatenation operator)

. (dot)

regular expressions

.. (range operator)

/ (division operator)

< (numeric less than)

<=

= (assignment operator) 2nd

combined with other operators

context

contrast to numeric equality test 2nd

== (numeric equality test)

contrast to assignment operator 2nd

=~ (regular expression binding operator)

=® (comma-arrow operator)

initializing hashes

? (quantifier)

@@@EDIT NOTE

global replace ++ with two colons@@@

@INC array

module searches

[ds]

regular expressions

anchor

[LT][GT] (angle operator)

command-line tools

condition of while

list context

reading files

[LT][GT] operator

globbing

^ [See [caret](#)]

{ } (quantifier)

|| (logical or) 2nd



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

A Small Number Guessing Game (Listing 3.1)

 if statement example 2nd

accelerate subroutine

 Car class example

Accept directive (HTTP)

Accept-Encoding directive (HTTP)

accessor methods

 classes (OOP) 2nd

Acrobat Reader

 downloading

 PDF documents

 reading

ActiveState Tool Corp

ActiveState.com

 module installation resources

Add Descriptions to Parts (Listing 19.4) 2nd 3rd

Adobe.com

 Acrobat Reader

 downloading

advisory locking

airports

 NOAA (National Oceanic and Atmospheric Administration)

 database queries/retrievals 2nd 3rd 4th

 METAR codes

alphanumeric relational operators (if statement) 2nd

anchors

 metacharacters 2nd

and (logical operator &&) 2nd

angle operator ([LT][GT])

 command-line tools

 condition of while

 globbing

 list context

 reading files

anonymous storage

building structures

Apache Web server

applications

weather station

 features application 2nd

 NOAA (National Oceanic and Atmospheric Administration) METAR codes

 NOAA (National Oceanic and Atmospheric Administration), airport queries/retrievals 2nd 3rd 4th

 NOAA (National Oceanic and Atmospheric Administration), weather report queries/retrievals 2nd

 U.S. Census Bureau zip code resources

 U.S. Census Bureau, zip code queries/retrievals 2nd 3rd

arguments

 modifying values

 subroutines

 arrays and hashes

arithmetic mean [See [mean](#)]

arithmetic operators [See [numeric operators](#)]

array

 @INC

 module searches

arrays

 arrays

 assignments

 converting between scalars and

 databases

 text files

 difference

 code listing

 elements

 extracting elements

 finding the end 2nd

 finding unique elements

 code listing

 implementing stacks

 push and pop

 shift and unshift

 intersection

 code listing

 iterating over

 foreach

 manipulating

passing to subroutines

references

dereferencing

reordering

reverse function

sort function

size 2nd

slices

splicing

two-dimensional

traversing

arrays~index

negative

ASCII

text

string literals

assignment

contrast with referencing

operator (=) 2nd 3rd 4th

value of

assignment operator (=)

combined with other operators

assignments

array

context

associative arrays [See [hashes](#)]

autodecrement operator (--)

autoincrement operator (++)

AutoLoader module

AutoSplit module

awk tool (Unix)

replacement by Perl



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

b (set breakpoint)

 debugger command

background colors

 web pages

 setting via cookies 2nd 3rd 4th

backreferences

 regular expressions 2nd

backslash

backtick operator ([ag]...[ag])

 capturing output

backtick operator ([ag][ag])

 qx notation

backward searches

 scalars

 rindex function

bare statement blocks

base-8 number system (Octal)

Benchmark module

binary

 file test

binary data

 WidgetCo Bike Parts Company example 2nd

binary files

binding operator

blocks

 private variables

blocks (statements)

 bare

 format of

 naked

 nested

books

 troubleshooting resources

Perl Cookbook

Boolean context

brackets

 unmatched error message

breakpoints

 debugger

 continuing to

 examining

 removing

 selecting position

 setting

browsers

 retrieving dynamic pages

CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th 20th

21st 22nd 23rd 24th 25th 26th 27th 28th

 retrieving static pages 2nd 3rd 4th 5th

building

 structures

 exercise

Bulk Order Sample Document (Listing 19.4) 2nd

bulk orders

 WidgetCo Bike Parts Company

 Bulk Order Sample Document (Listing 19.10) 2nd

 Reading Orders in Bulk Using XML (Listing 19.11) 2nd 3rd

 Using the Bulk Order Function (Listing 19.12)



< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

c (continue to breakpoint)

 debugger command

calling

 CGI programs 2nd

 modules 2nd

Car class (OOP)

 A Race Using the Car Class (Listing 18.5) 2nd 3rd 4th

 accelerate subroutine

 blessed reference

 design example 2nd 3rd

 First Part of the Class (Listing 18.1) 2nd 3rd

 methods 2nd 3rd

 Methods in the Car Class (Listing 18.3) 2nd

 properties 2nd 3rd

 Property Methods in the Class (Listing 18.2) 2nd

 Remainder of the Car Class (Listing 18.4) 2nd

 vin subroutine

caret (^)

 regular expressions

 anchor

 character classes

Carp module

carriage return

 metacharacter

case

 regular expressions

 insensitivity modifier

case-sensitivity

 naming variables

CGI

 forms

 processing 2nd 3rd 4th 5th 6th 7th

 processing; GET and POST methods 2nd 3rd

 processing; guestbook example 2nd 3rd 4th

processing; listing

processing; security 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

processing; submitting 2nd 3rd

installing programs on Web servers 2nd 3rd

troubleshooting

 Usenet newsgroup

writing programs

 listing

CGI (Common Gateway Interface) 2nd

 CGI module

 dynamic Web pages 2nd 3rd 4th

 debugging programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

 installing programs on Web servers 2nd 3rd 4th

 running programs 2nd

 Web server needs 2nd 3rd 4th 5th

 writing CGI programs 2nd

CGI programs

 calling 2nd

 cookies

 restricting 2nd 3rd

 retrieving to browsers 2nd

 sending to browsers

HTTP redirection

 when to use

 installing on Web servers

 launching 2nd

parameters

 passing 2nd 3rd

 special characters, escaping 2nd 3rd

CGI++* module (CPAN)

character classes 2nd

 metacharacters

 special

characters

 substitutions

 transliteration operator

chdir function

 change current directory 2nd

chmod

 Perl function

 Unix command 2nd 3rd

choosing

modules 2nd

Christiansen, Tom

class methods (OOP) 2nd

classes (OOP)

accessor methods 2nd

Car example

A Race Using the Car Class (Listing 18.5) 2nd 3rd 4th

accelerate subroutine

blessed reference

design example 2nd 3rd

First Part of The Car Class (Listing 18.1) 2nd 3rd

Methods in the Car Class (Listing 18.3) 2nd

Property Methods in the Car Class (Listing 18.2) 2nd

Remainder of the Car Class (Listing 18.4) 2nd

vin subroutine

File Information example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

instances

clients

Web

close function

closing files

writing

closedir function

cmp (string compare operator)

colors

web page backgrounds

setting via cookies 2nd 3rd 4th

comma (,)

as operator

context

comma-arrow (=®) operator

initializing hashes

command line

switches

-c (check syntax)

-d (debug)

-e (execute one-line program)

-n (input loop)

-p (input loop with print)

-V (detailed version)

-v (version)

command prompt

MS-DOS

Unix

Windows

Windows NT

command prompts (MS-DOS)

 directory identification

command-line tools

 angle operator ([LT][GT])

 debugger

 breakpoints

 commands 2nd

 exercise

 modifying variables

 references

 starting

 one-line programs

 quotation mark rules

commands

 debugger

 b (set breakpoint)

 c (continue)

 L (examine breakpoint)

 s (step into subroutine)

comments

Common Gateway Interface [See [CGI](#)]

Common Separated Values (CSV)

spreadsheets 2nd

 converting to native Excel format 2nd 3rd 4th

tables

 example of 2nd

compiler directives

Complete Listing for ColorChanger (24.2) 2nd 3rd

Complete Listing for mygrep (Listing 10.1) 2nd 3rd

Complete Source for Primes (Listing 3.4) 2nd 3rd

complex forms 2nd 3rd 4th

 hidden HTML fields 2nd 3rd 4th 5th 6th 7th 8th 9th

 survey example 2nd

 listings 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

Comprehensive Perl Archive Network [See [CPAN](#)]

concatenation operator (.)

conditional statements

connecting

 web servers

 HTTP connection process 2nd 3rd

 manually (Telnet client) 2nd 3rd

Connection directive (HTTP)

constant module

Content-Length directive (HTTP)

context

 Boolean

 list 2nd 3rd

 scalar 2nd 3rd

Convert CSV to a Native Excel Spreadsheet (Listing 20.6) 2nd 3rd

converting

 CSV tables

 to Excel spreadsheets 2nd 3rd 4th

cookie function 2nd 3rd 4th

 -expires option 2nd 3rd 4th 5th

 -path option 2nd 3rd

cookies

 creating 2nd 3rd 4th

 disabling

 erasing

 file storage

 flushing

information items

 combining (HTTP) 2nd

links

 tracking user activity 2nd

listings

 Complete Listing for ColorChanger (24.2) 2nd 3rd

long-term

 expiration dates 2nd 3rd 4th 5th

persistent

potential problems

 privacy issues 2nd 3rd

 support levels 2nd

 transient nature of

restricting return destinations 2nd 3rd

setting (HTTP)

transmitting while redirecting browsers 2nd

web browsers

 retrieving via CGI programs 2nd

 sending via CGI programs

 viewing options (IE6)

web pages

 background colors, setting 2nd 3rd 4th

cookies (HTTP)

CGI programs

 requesting

 function of 2nd

cooking data

 example of

 Sample Output from Moon (Listing 6.1) 2nd

 Your Moon Weight (Listing 6.2) 2nd 3rd

cos operator

CPAN

 (Comprehensive Perl Archive Network

 contributor motivations 2nd 3rd

modules

 CGI++*

 Data++*

 Date++*

 DBD++*

 DBI++*

 GD++*

 HTML++*

 Image++*

 LWP++*

 Math++*

 Net++*

 String++*

 Term++*

 Time++*

 Tk

 Tree++*

 URI++*

 Win32++*

 WWW++*

Web site

 search functions

Web site

CPAN module

creating

subroutines

current working directories

cwd function

current directory identification 2nd 3rd



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

d (remove breakpoint)

 debugger command

data

 binary

 WidgetCo Bike Parts Company example 2nd

 hierarchical

 WidgetCo Bike Parts Company example 2nd

 permanent storage

 DBM files

 DBM files; exercise

 locking

 random access files

 text files

 references

 references to

 exercise

 representations

 WidgetCo Bike Parts Company example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

table

 WidgetCo Bike Parts Company example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th

16th

 unstructured

 WidgetCo Bike Parts Company example 2nd

Data Munging with Perl

data structures

 anonymous storage

 building

 variations

 references 2nd 3rd

 exercise

 stacks

 implementation with lists

 two-dimensional arrays

 traversing

data types

arrays

text files

hashes

adding elements

applications 2nd

DBM files

each function

initializing

inverting

key-value pairs

list context

removing keys

retrieving elements

sorting

testing for keys

lists

array variables 2nd 3rd 4th 5th 6th

constructing

literals

numeric

string

scalar

comparing numbers

comparing strings

scalars

converting between arrays and

variables

scalar

Data++* module (CPAN)

databases

DBM files

dbmclose function

dbmopen function

exercise

hashes

implementing using hashes

code listing

data listing

output listing

locking

exclusive

race condition

reading and writing

semaphores

shared

Unix

Windows 95/98

Windows NT

random access files

opening

read pointers

text files

adding records

advantages

arrays

disadvantages

removing records

Date++* module (CPAN)

Date++Manip module

installations

UNIX platform 2nd 3rd 4th 5th 6th

Windows 95/98/NT platforms 2nd 3rd

DBD++* module (CPAN)

DBI++* module (CPAN)

DBM files

dbmclose function

exercise

code listing

output listing

hashes

each function

opening

dbmopen function

dbmclose function

dbmopen function

debugger

breakpoints

continuing to

examining

removing

selecting position

setting

commands

d (remove breakpoint)

h (help)

l (list)

n (run next instruction)

q (quit)

R (restart program)

s (step into subroutine)

exercise

code listing

modifying variables

references

repeating previous command

starting

tracing variables

debugging

CGI programs

on Web servers 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

dynamic Web pages

CGI (Common Gateway Interface) 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

logic errors

Perl documentation 2nd

print statement 2nd

programs

diagnostics module

references 2nd

syntax errors

decrement operators

autodecrement (--)

defensive programming

delete function

hash elements

delimiters

nesting

regular expressions 2nd

denial of service

writing CGI programs 2nd

dereferencing

array references

hashes

diagnostics module

 debugging programs

diamond operator

 command-line tools

diamond operator ([LT][GT])

 globbing

die function

 opening files

difference

 arrays

digit

 character class

directories 2nd

 closing

 closedir

 creating

 mkdir function 2nd

 current

 changing (chdir function) 2nd

 identifying 2nd

files

 renaming

 stat function 2nd 3rd 4th

 Unix permissions 2nd 3rd

 identifying (cwd function)

listings

 globbing 2nd 3rd 4th

 grep utility 2nd 3rd 4th

modules

 grouping by name

 separating from programs

moving files between

navigating

reading

 readdir

removing

 rmdir function 2nd

 removing files 2nd

renaming

 Renamer program (Listing 10.3) 2nd 3rd

 renaming files 2nd

en masse 2nd 3rd 4th

searching

File::Find module

directory

file test

DirHandle module

disabling

cookies

DLLs

(dynamically linked libraries)

documentation

modules 2nd

perldoc program

Plain Old Documentation (POD)

Pod::* modules

documentation (perlre)

Mastering Regular Expressions

documents

PDF

graphics coordinates system

image placement

object-oriented

reading

XML

elements

uses

dynamic Web pages

CGI (Common Gateway Interface) 2nd 3rd 4th 5th 6th

debugging programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

installing programs on Web servers 2nd 3rd 4th

running programs 2nd

Web server needs 2nd 3rd 4th 5th

writing programs 2nd

designing

CGI (Common Gateway Interface) 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th

17th 18th 19th

retrieving

CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th 20th

21st 22nd 23rd 24th 25th 26th 27th 28th

dynamically scoped variables

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

each function

editing

scalars

substr function

elements (XML)

 attributes

 root

 tags

email header

 WidgetCo Bike Parts Company 2nd

email order taker

 WidgetCo Bike Parts Company

 Email Reader (Listing 19.1) 2nd 3rd 4th

 Order Printer (Listing 19.2) 2nd

email order verifier

 WidgetCo Bike Parts Company

 Add Descriptions to Parts (Listing 19.4) 2nd 3rd

 Vendor Parts List (Listing 19.3) 2nd 3rd

Email Reader (Listing 19.1) 2nd 3rd 4th

empty file test

encryption

 writing CGI programs 2nd 3rd 4th

end-of-line sequence

English module

 renaming special variables

Env module

eq (string equality test)

erasing

 cookies

errors

 fence-post

 modules

 one-by-one

opening files

die function

escape sequences 2nd

pathnames

escaping

special characters

CGI programs 2nd 3rd

existence

file test

existing

loops

last statement 2nd

exists function

test for existence of key

exit statement 2nd

expiration dates

cookies 2nd 3rd 4th 5th

exponent

Exporter

modules

subroutines, forcing names on 2nd 3rd 4th 5th

Exporter module

expressions

comparison

contrast to assignment 2nd

interchanging numbers and strings

operators

assignment (=) 2nd 3rd 4th 5th

assignment combined with other operators

decrement

increment

named unary

numeric

precedence

string

unary

regular

alternation 2nd 3rd

backreferences 2nd

binding

character classes 2nd 3rd

grep function 2nd 3rd

grouping

interpretation rules 2nd

map function

metacharacters 2nd 3rd 4th

modifiers 2nd 3rd 4th

split function

substitution 2nd

scalar

testing (if statement) 2nd

testing (if-else statement)

testing (if-elsif-else statement) 2nd

undef

eXtensible Markup Language [See [XML](#)]

ExtUtils::* modules

 PREV

< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

factorials

 code listing

 defined

false values (if statement) 2nd 3rd 4th

FAQ [See [Usenet newsgroups; Frequently Asked Questions](#)]

FAQs (Frequently Asked Questions)

fence-post errors

fetching

 web pages

 manually (HTTP) 2nd 3rd

Fetching the Weather (Listing 20.3) 2nd

field specifiers

 printf function 2nd

fields

 hidden HTML 2nd 3rd 4th 5th 6th 7th 8th 9th

File Information class (OOP) 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

file input operator [See [angle operator](#)]

file pointers [See [read pointers](#)]

file systems 2nd 3rd

File::: * modules

File:::Copy module

 moving files

File:::Find module

 finding files

 code listing

 removing files

 code listing

File:::Spec::: * modules

FileCache module

filehandles

 automatic

 close function

 writing

 indirect

- names
- pipes
 - closing
- reading
 - angle operator ([LT][GT])
 - line of input

- STDERR

- STDIN

- STDOUT

- writing

- print function

- FileInfo Module, as a Class (Listing 18.6) 2nd 3rd 4th

- files

- binary

- test

- closing

- close function

- writing

- copying

- File::Copy module

- DBM

- dbmclose function

- dbmopen function

- exercise

- hashes

- defensive programming

- directories 2nd

- creating 2nd

- globbing 2nd 3rd 4th

- listings 2nd 3rd 4th

- navigating

- removing 2nd

- directory

- test

- emptiness

- test

- existence

- test

- identifying

- stat function 2nd 3rd 4th

- information requests

handling (TYPFileInfo module) 2nd 3rd 4th 5th 6th
locking
 exclusive
 race condition
 reading and writing
 semaphores
 shared
 Unix
 Windows 95/98
 Windows NT
metadata
 testing
modification date
 test
moving
 File::Copy module
moving between directories
 rename function
opening
 die function
 error
 filehandles
 modes
 open function
 reading
 warn function
 writing
Perl's capability
permissions
 dbmopen
 displaying 2nd
 setting
 Unix 2nd 3rd
random access
 opening
 read pointers
readability
 test
reading
 angle operator ([LT][GT])
 line of input

removing 2nd

renaming

en masse 2nd 3rd 4th

exercise

rename function 2nd

reversing

searching

File::Find module

semaphore

size

test

testing

operators 2nd

text

databases

record separators

test

writability

test

writing

print function

risks

filtering

cookies

FindBin module

Finding the Nearest Airport (Listing 20.2) 2nd 3rd

First Part of the Car Class (Listing 18.1) 2nd 3rd

flags

floating-point numbers

flock function

unavailable

Windows 95/98

flow control (programs)

conditional statements

exit statement 2nd

for loop 2nd 3rd

if statement

alphanumeric relational operators 2nd

expression testing 2nd

logical operators 2nd 3rd 4th 5th

numeric relational operators 2nd 3rd

syntax variations 2nd

true/false values 2nd 3rd 4th

if-else statement

expression testing

if-elsif-else statement

expression testing 2nd

labels 2nd 3rd

last statement 2nd

next statement

redo statement

statement blocks 2nd

bare

format of

naked

nested

switch statement 2nd

unmatched brackets

while loop 2nd

flushing

cookies

for loop 2nd 3rd

foreach loops

formfeed

metacharacter

forms

HTML 2nd 3rd 4th 5th

listing

processing with CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th

18th 19th

submitting 2nd 3rd

multipage 2nd 3rd 4th

hidden HTML fields 2nd 3rd 4th 5th 6th 7th 8th 9th

survey example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th

frequency distributions

determining with hashes

Frequently Asked Questions (FAQ)

Perl documentation

Usenet newsgroups

FTP

NT Web server access

installing CGI programs 2nd

Unix Web server access

installing CGI programs 2nd 3rd

function

cwd

current directory identification 2nd

function libraries [See [modules](#)]

functions

chdir

change current directory 2nd

chmod

setting file permissions

close

writing

closedir

context

cookie 2nd 3rd 4th

dbmclose

dbmopen

delete

die

opening files

each

editing scalars

substr

exists

flock

unavailable

Windows 95/98

grep

regular expressions 2nd 3rd

index

join

converting arrays to scalars

localtime

translating time stamps 2nd

map

regular expressions

mkdir

creating directories 2nd

permissions

open

- error
- pipes
- param
 - CGI form processing
- pop
- print
 - context
 - files
- printf
 - field specifiers 2nd
- programs
- pure
- push
- readdir
- recursive
 - termination condition
- rename
 - moving files
 - renaming files 2nd
- reverse
 - inverting hash
 - reordering arrays
- rindex
 - searching scalars
- rmdir
 - removing directories 2nd
- scope
- seek
- shift
- sort
 - reordering arrays
- splice
- split
 - converting scalars to arrays
 - regular expressions in
- sprintf
- stat
 - identifying files 2nd 3rd 4th
 - permissions, displaying 2nd
 - return values 2nd 3rd
 - time stamps 2nd

subroutines

arguments

calling

creating

return values

scope

system

capturing output

command interpreter

tell

unlink

removing files from directories 2nd

unshift

warn

opening files



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

GD++ * module (CPAN)

ge (string greater than or equal)

GET directive (HTTP)

GET method

CGI form processing 2nd 3rd

Getopt::* modules

glob operator

C shell-style versus Bourne or Korn shell

global matching

regular expressions

list context

global variables

globbing

getting directory listings 2nd 3rd 4th

patterns 2nd 3rd

glue language

GNU (GNU's Not Unix)

General Public License

Web site

GNU General Public License 2nd 3rd

Google

Usenet archive

grep function

regular expressions 2nd 3rd

grep utility

Complete Listing for mygrep (Listing 10.1) 2nd 3rd

implementing in Perl 2nd 3rd 4th

Output for mygrep (Listing 10.2)

grouping

modules by name

subroutines in modules

gt (string greater than)

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

h (help)

 debugger command

hackable (insecure) data

 writing CGI programs 2nd 3rd 4th 5th 6th

hacking

 forms

 writing CGI programs 2nd

handles

 directory

 globbing 2nd 3rd 4th

 file

 pipes

Hangman

 output

 program

hashes

 adding elements

 applications

 computing intersection and difference of arrays

 database

 DBM files 2nd

 dbmclose function

 dbmopen function

 functions

 iterating

 finding unique elements in arrays

 initializing

 comma-arrow operator

 key-value pairs

 DBM files

keys

 testing for existence

list context

passing to subroutines

references

dereferencing

removing keys

retrieving elements

sorting

type identifier><\$!type identifiers

hash (%)

help

finding resources

documentation 2nd

requesting

Usenet newsgroups 2nd 3rd 4th 5th

troubleshooting resources 2nd 3rd 4th 5th 6th

books 2nd

PerlMonks website 2nd 3rd

hidden fields

HTML 2nd 3rd 4th 5th 6th 7th 8th 9th

hierarchical data

WidgetCo Bike Parts Company example 2nd

Hobbe's Internet Timeline website

Host directive (HTTP)

host names

URLs

HTML

forms 2nd 3rd 4th 5th

listing

processing with CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th

18th 19th

submitting 2nd 3rd

hidden fields 2nd 3rd 4th 5th 6th 7th 8th 9th

HTML++* module (CPAN)

HTTP

(Hypertext Transfer Protocol)

cookies

CGI program requests

function of 2nd

items, combining 2nd

directives

Accept

Accept-Encoding

Connection

Content-Length

GET

Host

Last-Modified

redirection

implementing 2nd 3rd 4th

Redirection Based on Browser (Listing 24.1) 2nd 3rd

RFC 1945

RFC 2616

server/browser connection process 2nd 3rd

versions

web pages

fetching manually 2nd 3rd

Hypertext Transfer Protocol [See [HTTP](#)]

 PREV

< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

I18N::Collate module

if statement

 alphanumeric relational operators 2nd

 expression testing 2nd

 listings

 A Small Number Guessing Game (3.1) 2nd

 logical operators 2nd 3rd 4th 5th

 numeric relational operators 2nd 3rd

 syntax variations 2nd

 true/false values 2nd 3rd 4th

if-else statement

 expression testing

if-elsif-else statement

 expression testing 2nd

Image++* module (CPAN)

images

 PDF documents

 placement of

increment operators

 autoincrement (++)

increments

 assignment operator

index

 array

 negative

index function

 searching scalars

indirect filehandles

input

 connections

 pipes

 line

insecure (hackable) data

 writing CGI programs 2nd 3rd 4th 5th

installing

modules

choosing modules 2nd

in non-standard directory 2nd

private copies 2nd

UNIX 2nd 3rd 4th 5th 6th

Windows 95/98/NT 2nd 3rd

instances (OOP)

int operator

integer module

Internet

early history of 2nd

history

Web sites

Internet Engineering Task Force (IETF)

RFCs (Requests for Comments)

Internet Explorer

cookies

viewing options

interpolation

regular expressions

variable

interpreter

angle operator ([LT][GT])

debugger

breakpoints

commands 2nd

exercise

modifying variables

references

one-line programs

quotation mark rules

interpreters

intersection

arrays

inverting hashes

invoking subroutines [See [subroutines](#), [calling](#)]

IPC::* modules

iterating

arrays

iterator

foreach



< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

join function

 converting arrays to scalars

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

key

 hashes

key-value pairs

 hashes

 DBM files

keys

 hashes

 DBM files

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

L (examine breakpoint)

 debugger command

I (list)

 debugger command

labels

 loops

 labeling 2nd 3rd

last statement

 Example of the last Statement (Listing 3.3)

 loops

 exiting 2nd

 with label

Last-Modified directive (HTTP)

launching

 CGI programs 2nd

Ic operator

le (string less than or equal)

length operator

lexical variables 2nd

lexically scoped variables

licenses

 Perl interpreter

 GNU General Public License 2nd 3rd

 Perl Artistic License 2nd 3rd

line of input

links

 cookies

 tracking user activity 2nd

listings

 A Race Using the Car Class (18.5) 2nd 3rd 4th

 A Small Number Guessing Game (3.1) 2nd

 Car class (OOP)

 First Part of the Class (18.1) 2nd 3rd

 Methods in the Car Class (18.3) 2nd

Property Methods in the Class (18.2) 2nd

Remainder of the Car Class (18.4) 2nd

CGI programs

writing

Complete Listing for mygrep (10.1) 2nd 3rd

Complete Listing for Renamer (10.3) 2nd 3rd

Complete Source for Primes (3.4) 2nd 3rd

computing intersection of arrays

Convert CSV to a Native Excel Spreadsheet (20.6) 2nd 3rd

cookies

Complete Listing for ColorChanger (24.2) 2nd 3rd

customer database

code

output

sample data

DBM memo pad

code

output

debugger exercise

directory

globbing 2nd 3rd 4th

Even More Convoluted XML Parse in Regular Expressions (19.8) 2nd 3rd

Example of the last Statement (3.3)

factorial

file I/O with locking

File Information Module (17.7) 2nd 3rd

File::Find

finding file

removing temporary files

FileInfo Module, as a Class (18.6) 2nd 3rd 4th

formatted printing

code

output

general-purpose locking functions

Hangman

output

program

hashes

computing difference between arrays

finding unique elements with

HTML

form processing 2nd 3rd 4th

forms

Maze example

code

output

modules

Program to Use Statistics Module (Listing 17.2) 2nd

Statistics Module (Listing 17.1) 2nd

multipage forms

survey example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

Output for mygrep (10.2)

Reading an Excel Spreadsheet (20.7) 2nd 3rd

Reading the XML File, Almost (19.5) 2nd

Redirection Based on Browser (24.1) 2nd 3rd

reversing file

Reworking of Listing 17.9 with a Class (18.7) 2nd

Sample Output from Moon (6.1) 2nd

sample Perl program

Sample Usage of the File Information Module (17.8) 2nd 3rd

Sample Using XML++Simple to Parse Your Album Document (19.9) 2nd 3rd

Sample while Loop (3.2) 2nd

scoping

Small Program Using Package and Lexical Variables (17.4) 2nd

Small Scoping Example (17.3) 2nd

statistics example

code

output

Statistics Program with Function Name Export (17.6) 2nd 3rd 4th 5th

Still Flawed XML Reading (19.6) 2nd

The Album XML Document, re-arranged (19.7) 2nd

The out Statement (17.5) 2nd

Totaler

code

Using Multiple FileInfo Objects at Once (18.8) 2nd 3rd 4th

weather station application

Fetching the Weather (20.3) 2nd

Finding the Nearest Airport (20.2) 2nd 3rd

Retrieving the ZIP Information (20.1) 2nd

Setting Up the PDF Document (20.4) 2nd 3rd 4th 5th 6th 7th

WidgetCo Bike Parts Company

Add Descriptions to Parts (19.4) 2nd 3rd

Bulk Order Sample Document (19.10) 2nd

Email Reader (19.1) 2nd 3rd 4th

Order Printer (19.2) 2nd

Reading Orders in Bulk Using XML (19.11) 2nd 3rd

Using the Bulk Order Function (19.12)

Vendor Parts List (19.3) 2nd 3rd

Your Moon Weight (6.2) 2nd 3rd

lists

array variables

assignments

converting between scalars and

elements

extracting elements

finding the end 2nd

iterating over

manipulating

reordering

size 2nd

slices

constructing

qw operator

range operator

hashes as

implementing stacks

push and pop

shift and unshift

lists

splicing arrays

literals

numeric

decimal

underscores

string

local operator

local variables

contrast with private variables

locale module

localtime function

translating time stamps 2nd

locking

advisory

databases

exclusive locks

general-purpose functions

 code listing

mandatory

race condition

reading and writing files

 code listing

semaphore files

shared locks

Windows 95/98

logic errors

 troubleshooting

logical operators

 if statement 2nd 3rd 4th 5th

 short-circuiting

loops

 Example of the last Statement (Listing 3.3)

 exiting (last statement) 2nd

 exiting with last

 with label

 for 2nd 3rd

 foreach

 labeling 2nd 3rd

 passing (next statement)

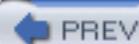
 Sample while Loop (Listing 3.2) 2nd

 skipping body with next

 while 2nd

lt (string less than)

LWP++ module (CPAN)

 PREV

< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

m// (pattern match operator)

 modifiers 2nd 3rd 4th

Macintosh

 installing Perl

 mapping

 pathnames

Perl

 preinstalled

 Perl documentation

man

mandatory locking

mantissa

map function

 regular expressions

Mastering Regular Expressions

Math++* module (CPAN)

Math::* modules

mean

median

memo pad

 exercise

 code listing

 output listing

metacharacters

 anchors 2nd

 carriage return

 character classes

dot (.)

 regular expressions

formfeed

newline

 quantifiers 2nd 3rd 4th

*

+

?

{ }

regular expressions 2nd 3rd

tab

unprintables 2nd

metadata

testing

METAR codes

NOAA (National Oceanic and Atmospheric Administration)

methods (OOP) 2nd 3rd

Methods in the Car Class (Listing 18.3) 2nd

Microsoft Word

RTF documents

mkdir function

creating directories 2nd

permissions

modification time

file test

modules

(function libraries)

AutoLoader

AutoSplit

Benchmark

building 2nd

calling 2nd

Carp

CGI

constant

copying files

File::Copy

CPAN 2nd

CGI++*

Data++*

Date++*

DBD++*

DBI++*

GD++*

HTML++*

Image++*

LWP++*

Math++*

Net++^{*}
String++^{*}
Term++^{*}
Time++^{*}
Tk
Tree++^{*}
URI++^{*}
Web site
Win32++^{*}
WWW++^{*}
debugging
 diagnostics
directories
 separating from programs
DirHandle
documentation 2nd
 perldoc program
Env
errors
Exporter
ExtUtils::^{*}
file information requests
 File Information Module (Listing 17.7) 2nd 3rd
 Sample Usage of the File Information Module (Listing 17.8) 2nd 3rd
File::^{*}
File::Find
 file removal code listing
 file-finding code listing
FileCache
FileSpec::^{*}
FindBin
Getopt::^{*}
grouping by name
I18N::Collate
installing
 choosing modules 2nd
 in non-standard directory 2nd
 private copies 2nd
 UNIX 2nd 3rd 4th 5th 6th
 Windows 95/98/NT 2nd 3rd
integer

IPC::*

listings

A Reworking of Listing 17.9 with a Class (18.7) 2nd

FileInfo Module, as a Class (18.6) 2nd 3rd 4th

Program to Use Statistics Module (17.2) 2nd

Statistics Module (17.1) 2nd

Using Multiple FileInfo Objects at Once (18.8) 2nd 3rd 4th

locale

Math::*

moving files

File::Copy

namespace pollution 2nd 3rd

namespaces 2nd

naming special variables

English

Net::*

OLE++Win32

PDF++API2

PDF file writing 2nd 3rd 4th 5th 6th 7th

pm file extension

Pod::*

reasons for

reasons for using

Spreadsheet++ParseExcel

spreadsheet reading 2nd 3rd

Spreadsheet++WriteExcel

CSV table conversion to spreadsheet 2nd 3rd 4th

standard

list

Statistics Program with Function Name Export (Listing 17.6) 2nd 3rd 4th 5th

subroutines

forcing names on

forcing names on (Exporter) 2nd 3rd 4th 5th

grouping

Symbol

Sys::Hostname

Sys::Syslog

Term::*

testing systems

Net::Ping

Text::Abbrev

Text::ParseWords

Text::Soundex

Tie::*
Time::*
TYPFileInfo.pm

file information requests 2nd 3rd 4th 5th 6th

TYPStats.pm 2nd 3rd 4th

mean subroutine 2nd

median subroutine 2nd

use directive 2nd

using

XML++Simple

weather station application, zip code retrieval 2nd 3rd

MS-DOS

command prompt

pathnames

multipage forms 2nd 3rd 4th

hidden HTML fields 2nd 3rd 4th 5th 6th 7th 8th 9th

survey example 2nd

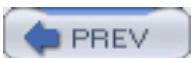
listings 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

my operator

private variables

blocks

subroutines



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

n (next instruction)

 debugger command

naked statement blocks

namespaces

 modules 2nd

 pollution practices 2nd 3rd

navigating

 directories

ne (string not equal)

nested statement blocks

Net++ module (CPAN)

Net::* modules

Net::Ping module

networks

 security

 defensive programming

 testing systems

 Net::Ping module

newline

 metacharacter

newsgroups [See [Usenet newsgroups](#)]

newsgroups (Usenet)

 flaming responses

newsreaders

 Usenet newsgroups

 obtaining

next statement

 loops

 passing

 with label

NOAA (National Oceanic and Atmospheric Administration)

 airport queries/retrievals 2nd 3rd 4th

 METAR codes

 weather report queries/retrievals 2nd

non-standard directories

installing modules in 2nd

not (logical operator !) 2nd

numeric operators

numeric relational operators (if statement) 2nd 3rd



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

Object Linking and Embedding (OLE)

Object Oriented Perl

object-oriented programming [See [OOP](#)]

objects

 properties

 accessing 2nd

objects (OOP)

Octal

 base-8 numbering system

OLE (Object Linking and Embedding)

OLE++Win32 module

one-by-one errors

one-line programs

 command-line tools

 quotation mark rules

one-liners [See [one-line programs](#)]

OOP

 (object-oriented programming) 2nd

 class methods 2nd

 classes

 Car example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th

 File Information example 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

 instances

 classes accessor methods 2nd

 development of 2nd

 methods 2nd 3rd

Object Oriented Perl

 objects

 accessing properties 2nd

 Perl object documentation

 properties

 versus conventional programming 2nd

open function

 error

pipes

Open Source

essays

Web site

Open Source software

Apache Web Server development

concept

development benefits

Firefox browser development

FreeBSD development

Linux development

Perl interpreter

licenses 2nd 3rd

Raymond, Eric S

opening

files

die function

operating system

calling external programs

capturing output

pipes

system function

operating systems

portability

example

operator

qq

qw

operators

! (logical not) 2nd

! (not)

&& (logical and) 2nd

< (numeric less than)

== (numeric equality test)

contrast to assignment operator 2nd

=~ (regular expression binding)

[LT][GT]

globbing

alphanumeric relational

if statements 2nd

angle ([LT][GT])

command-line tools

assignment (=)

contrast to numeric equality test 2nd

assignment(=) 2nd

combined with other operators

context

backtick ([ag]...[ag])

capturing output

backtick ([ag][ag])

qx notation

combined numeric

comma (,)

context

comma-arrow (=®)

decrement

autodecrement (--)

documentation

eq (string equality test)

file test 2nd

-B (binary)

-d (is directory)

-e (existence)

-f (is file)

-M (time since last modified)

-r (readable)

-s (size)

-T (text)

-w (writable)

-z (empty file)

ge (string greater than or equal)

gt (string greater than)

increment

autoincrement (++)

le (string less than or equal)

local

logical

and (&&)

not (!)

or (||)

precedence

short-circuiting

logical (if statement) 2nd 3rd 4th 5th

lt (string less than)

my

blocks

subroutines

ne (string not equal)

numeric

precedence

numeric relational

if statements 2nd 3rd

pattern match (m//)

modifiers 2nd 3rd 4th

pattern substitution (s///) 2nd

modifiers 2nd 3rd 4th

q

range(..)

scalar

<=

cmp (string compare)

string

concatenation (.)

repetition (x)

transliteration

character substitutions

unary

named

|| (logical or) 2nd

or (logical operator || 2nd

order information

WidgetCo Bike Parts Company 2nd

Order Printer (Listing 19.2) 2nd

our statement

The out Statement (Listing 17.5) 2nd

output

connections

pipes

external programs

capturing

Output for mygrep (Listing 10.2)

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

package variables 2nd

param function

 CGI form processing

parameters

 CGI programs

 passing 2nd 3rd

 processing forms with CGI 2nd

 GET and POST methods 2nd 3rd

 security 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

passing

 CGI program parameters 2nd 3rd

loops

 next statement

pathname

pathnames

 Macintosh

 reading

 Universal Naming Convention (UNC)

 Unix

 VMS

 Windows and MS-DOS

writing

 print function

pattern

 splitting scalars

pattern matching 2nd

 example

payment block

 WidgetCo Bike Parts Company 2nd

PDF

 (Portable Document Format)

documents

 graphics coordinates system

 image placement

object-oriented

reading

PDF++API2 module

PDF file writing 2nd 3rd 4th 5th 6th 7th

Perl

announcements

Usenet newsgroup

as replacement for awk and sed tools (Unix)

camel symbol

reason for

characteristics

discussions

Usenet newsgroups

documentation

FAQs

Frequently Asked Questions

Macintosh

man

not found

perldoc

reinstalling

troubleshooting resources 2nd

Unix 2nd

Web site

Windows 2nd

first program

function

history 2nd 3rd

development process 2nd

version 2 release

version 4 release

version 5 release

Web sites

installing

Macintosh

Unix

Windows

interpreter

-c switch

-d switch

-e switch

- n switch
- p switch
- v switch
- V switch
- debugger

- licenses 2nd 3rd 4th

- Macintosh

- documentation
- installing
- preinstalled
- running programs

- name

- original description in manual 2nd

- preinstalled

- Macintosh

- Unix

- Windows

- programs

- comments

- errors

- running

- scripts

- similarities with other languages 2nd

- statement

- statement separator

- style

- syntax

- Unix

- documentation 2nd

- installing

- path to interpreter

- preinstalled

- running programs

- version numbers

- versions

- numbering scheme 2nd

- Web site

- Windows

- documentation 2nd

- installing

- preinstalled

running programs

Perl Artistic License 2nd 3rd 4th

Perl Community

book and magazine resources 2nd

CPAN (Comprehensive Perl Archive Network)

contributor motivations 2nd 3rd

modules

culture 2nd 3rd

Open Source software

development benefits

Usenet newsgroups

help resources 2nd 3rd 4th

requesting help 2nd 3rd 4th 5th

Wall, Larry

Perl Cookbook

Perl Data Structures Cookbook

Perl Developer's *Dictionary*

Perl interpreter

path

Perl Journal, The

Perl Package Manager (PPM)

module installations

launching 2nd 3rd

Perl Review Magazine

perlboot

perlbot

perldoc

finding

perldoc command

troubleshooting 2nd

perldoc program

module documentation

PerlMonks website

help resources 2nd 3rd

perlobj

perlre documentation

perltoot

perltootc

permissions

files

dbmopen

displaying 2nd

Unix 2nd 3rd 4th

persistent cookies

Pierce, Clinton

pipes

filehandles

input/output connections

standard input and output

Totaler

code listing

Plain Old Documentation (POD)

Pod:: * modules

platforms

portability

example

Pod:: * modules

pollution (namespaces)

modules 2nd 3rd

pop function

port numbers

URLs

portability

example

general rules

POST method

CGI form processing 2nd 3rd

prime numbers

Complete Source for Primes (Listing 3.4) 2nd 3rd

print

print function

files

print statement

troubleshooting 2nd

printf function

field specifiers 2nd

printing

formatted

code listing

field specifiers

output listing

weather station reports

private modules

installing 2nd

private variables

contrast with local variables

processing

forms (HTML) in CGI programs 2nd 3rd 4th 5th 6th 7th

GET and POST methods 2nd 3rd

guestbook example 2nd 3rd 4th

listing

security 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

submitting 2nd 3rd

program flow control

block only

brackets

unmatched errors

conditional statements 2nd

exiting Perl 2nd

for loop 2nd 3rd

if statement

alphanumeric relational operators 2nd

expression testing 2nd

logical operators 2nd 3rd 4th 5th

numeric relational operators 2nd 3rd

syntax variations 2nd

true/false values 2nd 3rd 4th

if-else statement

expression testing

if-elsif-else statement

expression testing 2nd

labels 2nd 3rd

last statement 2nd

looping

foreach

last

next

next statement

redo statement

statement blocks 2nd

bare

format of

naked

nested

switch statement 2nd

using short-circuiting in logical expressions

while loop 2nd

Program to Use Statistics Module (Listing 17.2) 2nd

programming

defensive

programming language

free-form

general-purpose

glue

Programming Perl

programs

CGI

processing HTML forms 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th

19th 20th 21st 22nd 23rd 24th 25th 26th 27th 28th

debugging

diagnostics module

editing

software

interpreting

running 2nd

structure

style 2nd

writing

functions

properties (OOP)

Property Methods in the Car Class (Listing 18.2) 2nd

protocols

URLs

pure functions

push function

pwd command (Unix)

directory identification



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

q (quit)

 debugger command

quantifiers

 regular expressions 2nd 3rd 4th

quotation marks

 command-line rules

 one-line programs

quoting

 backticks

 double quotes

 q operator

 qq operator

 qw operator

 single quotes

qx operator

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

R (restart program)

 debugger command

race condition

 defined

Race Using the Car Class (Listing 18.5) 2nd 3rd 4th

rand operator

random access files

 opening

 read pointers

 seek function

 tell function

range operator (...)

Raymond, Eric S

 Open Source software essays 2nd

read pointers

 random access files

 seek function

 tell function

readability

 file test

readdir function

reading

 spreadsheets 2nd 3rd

 XML files

 via regular expressions 2nd 3rd 4th 5th 6th 7th 8th 9th

 via XML++Simple 2nd 3rd

Reading an Excel Spreadsheet (Listing 20.7) 2nd 3rd

Reading Orders in Bulk Using XML (Listing 19.11) 2nd 3rd

record separators

recursion

 termination condition

recursive subroutines

 termination condition

redirecting

web browsers

cookie transmission 2nd

redirection (HTTP)

implementing 2nd 3rd 4th

redo statement

loops

repeating without evaluation

references

anonymous storage

arrays

contrast with assignment

creating

data structures 2nd

two-dimensional array

variations

debugging

defined

dereferencing

arrays

hashes

levels

exercise 2nd

hashes

Maze example

code listing

output listing

modifying original value 2nd

references

regexp [See [regular expressions](#)]

regular expressions 2nd

(regexp)

alternation 2nd 3rd

backreferences 2nd

binding operator (= ~)

character classes 2nd

special

creating

delimiters 2nd

Even More Convoluted XML Parse in Regular Expressions (Listing 19.8) 2nd 3rd

grep function 2nd 3rd

grouping

list context

interpretation rules 2nd

map function

Mastering Regular Expressions

metacharacters 2nd 3rd

anchors 2nd

character classes

dot (.)

quantifiers 2nd 3rd 4th

unprintables 2nd

modifiers 2nd 3rd 4th

case insensitivity

combining

global matching

list context

pattern matching example

Reading the XML File, Almost (Listing 19.5) 2nd

scalars

split function

Still Flawed XML Reading (Listing 19.6) 2nd

substitution 2nd

The Album XML Document, re-arranged (Listing 19.7) 2nd

use in multiple languages

variables

 matching against other than \$_ 2nd

XML files

 reading 2nd 3rd 4th 5th 6th 7th 8th 9th

Remainder of the Car Class (Listing 18.4) 2nd

remote execution

 capturing output

 input/output

 pipes

 system function

 command interpreter

rename function

 moving files

 renaming files 2nd

renaming

 directories

 en masse (Renamer program) 2nd 3rd

files

exercise

repetition operator (`x`)

list context

resources

finding help

print statement 2nd

Retrieving the ZIP Information (Listing 20.1)

weather station application 2nd

return values

subroutines

reverse function

inverting hash

reordering arrays

Reworking of Listing 17.9 with a Class (Listing 18.7) 2nd

RFCs (Requests for Comments)

HTTP documentation

Internet Engineering Task Force (IETF)

rindex function

searching scalars

rmdir function

removing directories 2nd

root element (XML)



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

s (step into subroutine)

 debugger command

s/// (pattern substitution operator) 2nd

 modifiers 2nd 3rd 4th

Sample Output from Moon (Listing 6.1) 2nd

Sams Teach Yourself Unix in 24 Hours

scalar variables

 use in regular expressions

scalars

 assignment

 converting between arrays and

 defined

 editing

 substr function

 expressions

 interchanging numbers and strings

 literals

 numeric

 string

 numbers

 comparing

 string searches

 index function

 rindex function

 strings

 comparing

 type identifier

 undef

 variables

 default values

 special 2nd

scientific notation

scope

 defined

functions 2nd

global variables

local variables

private variables

blocks

subroutines

variables

dynamically scoped

lexically scoped

scoping

variables 2nd 3rd 4th

Small Program Using Package and Lexical Variables (Listing 17.4) 2nd

Small Scoping Example (Listing 17.3) 2nd

scripts

searching

scalars

index function

rindex function

security

defensive programming

writing CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th

denial of service 2nd

encryption 2nd 3rd 4th

hacking forms 2nd

insecure (hackable) data 2nd 3rd 4th 5th

sed tool (Unix)

replacement by Perl

seek function

random access file pointers

semaphore files

code listing

servers

denial of service

CGI form security 2nd

HTTP

connection process 2nd 3rd

Web

Apache

debugging CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

installing CGI programs 2nd 3rd 4th

running CGI programs 2nd

writing CGI programs 2nd 3rd 4th 5th 6th 7th

Setting Up the Document (Listing 20.4) 2nd 3rd 4th 5th 6th 7th

shift function

short-circuiting

logical operators

size

array 2nd

file test

slice

array

Small Program Using Package and Lexical Variables (Listing 17.4) 2nd

Small Scoping Example (Listing 17.3) 2nd

sort function

reordering arrays

sorting

hashes

spaceship operator (<=

special characters

CGI programs

escaping 2nd 3rd

splice function

split function

converting arrays to scalars

Spreadsheet++ParseExcel module

spreadsheet reading 2nd 3rd

Spreadsheet++WriteExcel module

CSV table conversion to spreadsheet 2nd 3rd 4th

spreadsheets

Common Separated Values (CSV) format 2nd

CSV (Common Separated Values)

converting to native Excel format 2nd 3rd 4th

listings

Convert CSV to a Native Excel Spreadsheet (20.6) 2nd 3rd

Reading an Excel Spreadsheet (20.7) 2nd 3rd

reading 2nd 3rd

writing 2nd

sprintf function

stacks

lists

push and pop

shift and unshift

standard deviation

stat function

identifying files 2nd 3rd 4th

permissions

printing 2nd

return values 2nd 3rd

time stamps 2nd

statement

statement blocks 2nd

bare

format of

naked

nested

statement separator

statements

exit 2nd

foreach

if

alphanumeric relational operators 2nd

expression testing 2nd

logical operators 2nd 3rd 4th 5th

numeric relational operators 2nd 3rd

syntax variations 2nd

true/false values 2nd 3rd 4th

if-else

expression testing

if-elsif-else

expression testing 2nd

last 2nd

with label

next

with label

our

The out Statement (Listing 17.5) 2nd

print 2nd

redo

static Web pages

retrieving 2nd 3rd 4th 5th

statistics

code listing

mean

median

output listing

standard deviation

Statistics Module (Listing 17.1) 2nd

Statistics Program with Function Name Export (Listing 17.6) 2nd 3rd 4th 5th

STDERR filehandle

STDIN filehandle

STDOUT filehandle

storage

anonymous

String++* module (CPAN)

strings

.. (range operator)

backslash

comparing

cmp

concatenation

by interpolation

using . operator

delimiters

double quotes

escape sequences 2nd

interchanging with numbers

interpolation

q operator

qq operator

quoting

qw operator

repetition

searching scalars

index function

rindex function

single quotes

structures

exercise

subroutines

arguments

arrays and hashes

calling

creating

debugging

stepping into

modules

forcing names on

forcing names on (Exporter) 2nd 3rd 4th 5th

grouping

namespaces 2nd

pure functions

recursive

termination condition

return values

scope

substitution

regular expressions 2nd

modifier 2nd 3rd 4th

substr function

editing scalars

surveys

writing multipage forms 2nd

listings 2nd 3rd 4th 5th 6th 7th 8th 9th 10th

switch statement

program flow control 2nd

Symbol module

syntax errors

troubleshooting

Sys::Hostname module

Sys::Syslog module

system function

capturing output

command interpreter



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

tab

metacharacter

table data

WidgetCo Bike Parts Company example 2nd

email header 2nd

email order taker 2nd 3rd 4th 5th 6th

email order verifier 2nd 3rd 4th 5th 6th

order information 2nd

payment block 2nd

tables

CSV (Common Separated Values)

example of 2nd

tell function

random access file pointers

Telnet

web servers

connecting 2nd 3rd

Term++* module (CPAN)

Term::* modules

testing

files

operators

text

file test

text editor

text files

databases

adding records

advantages

arrays

disadvantages

removing records

Text::Abbrev module

Text::ParseWords module

Text::Soundex module

The Perl Journal

Web site

The Perl Review

Tie::* modules

time stamps

printing (localtime function) 2nd

Time++* module (CPAN)

Time::* modules

Tk module (CPAN)

tools

command-line

debugger

Torkington, Nathan

tr/// (transliteration operator)

tracking

link clicks by users via cookies 2nd

transliteration operator (tr///)

character substitutions

Tree++* module (CPAN)

troubleshooting

help resources 2nd 3rd 4th

logic errors

modules

programs

diagnostics module

requesting help on Usenet 2nd 3rd 4th 5th

resources

books 2nd

documentation 2nd

print statement 2nd

syntax errors

unmatched brackets in programs

XML++Simple

document errors

returning hashes

true values (if statement) 2nd 3rd 4th

type identifier

TYPOFileInfo.pm module

file information requests

processing 2nd 3rd 4th 5th 6th

Reworking of Listing 17.9 with a Class (Listing 18.7) 2nd

The FileInfo Module, as a Class (Listing 18.6) 2nd 3rd 4th

Using Multiple FileInfo Objects at Once (Listing 18.8) 2nd 3rd 4th

TYPStats.pm module 2nd 3rd 4th

mean subroutine 2nd

median subroutine 2nd



< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

U.S. Census Bureau

 zip code queries/retrievals 2nd 3rd

 zip code resources

U.S. Gazetteer Place (zip code database) 2nd 3rd

uc operator

unary operators

 named

UNC (Universal Naming Convention)

 pathnames

underscores

 numeric literals

Uniform Resource Locators [See [protocols; URLs](#)]

Unix

 awk tool

 replacement by Perl

 command prompt

 end of line

 file locking

 file permissions 2nd 3rd

 dbmopen

 implementing in Perl 2nd

 grep utility 2nd 3rd 4th

 Complete Listing for mygrep utility (Listing 10.1) 2nd 3rd

 Output for mygrep utility (Listing 10.2)

UNIX

 installing modules under 2nd 3rd 4th 5th 6th

Unix

 installing Perl

 man

 path to Perl interpreter

 pathnames

 Perl documentation

 Perl preinstalled

Sams Teach Yourself Unix in 24 Hours

sed tool

replacement by Perl

Web servers

installing FTP-only CGI program access 2nd 3rd

installing local CGI program access 2nd

unlink function

removing files from directories 2nd

unmatched brackets

program flow control

unshift function

unstructured data

WidgetCo Bike Parts Company example 2nd

URI++* module (CPAN)

URLs (Uniform Resource Locators)

host names

port numbers

protocols

requested items

use directive

modules 2nd

use strict

Usenet newsgroups

archive

Google

CGI-specific questions

flaming responses

Frequently Asked Questions

reading

troubleshooting resources

question procedure 2nd 3rd 4th 5th

user-defined functions [See [subroutines](#)]

Using Multiple FileInfo Objects at Once (Listing 18.8) 2nd 3rd 4th

Using the Bulk Order Function (Listing 19.12)

utilities

grep

getting directory listings 2nd 3rd 4th



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

value

 hashes

values

 hashes

 DBM files

variables

 array

 assignments

 converting between scalars and

 elements

 extracting elements

 finding the end 2nd

 iterating over

 reordering

 size 2nd

 slices

 splicing

 arrays

 passing to subroutines

 two-dimensional

data structures

 building

 references

debugging

 modifying as program runs

 output

filehandles

 closing 2nd

 indirect

 names

 reading

 writing to files

global

hashes

adding elements

applications 2nd

DBM files

each function

initializing

inverting

key-value pairs

list context

passing to subroutines

removing keys

retrieving elements

testing for keys

interpolation

lexical 2nd

local

names

package 2nd

private

blocks

subroutines

references

arrays

creating

dereferencing

hashes

modifying original value 2nd

regular expressions

matching against other than \$ 2nd

renaming

English module

scalar

converting between arrays and

default values

regular expressions

scalars

assignment

string searches

scope

defined

dynamic

lexical

scoping 2nd 3rd 4th

Small Program Using Package and Lexical Variables (Listing 17.4) 2nd

Small Scoping Example (Listing 17.3) 2nd

sorting

special

\$_

type identifiers

Vendor Parts List (Listing 19.3) 2nd 3rd

version numbers

Perl

 module errors

vin subroutine

 Car class example

VMS

 pathnames



< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

Wall, Larry

 Perl history

warn function

 opening files

weather reports

 NOAA (National Oceanic and Atmospheric Administration)

 queries/retrievals 2nd

weather station application

 features requirements 2nd

 information elements

 assembling 2nd

listings

 Fetching the Weather (20.3) 2nd

 Finding the Nearest Airport (20.2) 2nd 3rd

 Retrieving the ZIP Information (20.1) 2nd

 Setting Up the Document (20.4) 2nd 3rd 4th 5th 6th 7th

NOAA (National Oceanic and Atmospheric Administration)

 airport queries/retrievals 2nd 3rd 4th

 METAR codes

 weather report queries/retrievals 2nd

reports

 printing in PDF format 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th

U.S. Census Bureau

 zip code queries/retrievals 2nd 3rd

 zip code resources

Web browsers

web browsers

 cookies

 disabling

 erasing

 expiration dates 2nd 3rd 4th 5th

 flashing

 long-term 2nd 3rd 4th 5th

 privacy issues 2nd 3rd

restricting return destinations 2nd 3rd

retrieving via CGI programs 2nd

sending via CGI programs

storage file

support levels 2nd

transient nature of

viewing options (IE6)

HTTP

server connection process 2nd 3rd

HTTP cookies

CGI program requests

function of 2nd

HTTP redirection 2nd 3rd

redirecting while transmitting cookies 2nd

Web browsers

retrieving dynamic pages

CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th 20th

21st 22nd 23rd 24th 25th 26th 27th 28th

retrieving static pages 2nd 3rd 4th 5th

Web clients

web pages

background colors

setting via cookies 2nd 3rd 4th

Web pages

CGI (Common Gateway Interface) 2nd 3rd 4th 5th 6th

debugging programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

installing programs on Web servers 2nd 3rd 4th

running programs 2nd

Web server needs 2nd 3rd 4th 5th

writing programs 2nd

dynamic

retrieving 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th 20th 21st

22nd 23rd 24th 25th 26th 27th 28th

web pages

fetching manually (HTTP) 2nd 3rd

Web pages

HTML

forms 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th 19th 20th 21st 22nd

23rd 24th 25th 26th 27th 28th

web pages

links

tracking user activity via cookies 2nd

Web pages

static

retrieving 2nd 3rd 4th 5th

Web servers

debugging CGI programs 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th

installing CGI programs 2nd 3rd 4th

FTP-only access on NT servers 2nd

FTP-only access on Unix servers 2nd 3rd

local access on NT servers

local access on Unix servers 2nd

web servers

manual connections

Telnet client program 2nd 3rd

Web servers

running CGI programs 2nd

writing CGI programs 2nd 3rd 4th 5th 6th 7th

Apache

Web sites

Apache Web server

CPAN

searching

CPAN (Comprehensive Perl Archive Network)

modules

GNU

history of Perl

Internet history

online documentation

Open Source

essays

Perl Online Documentation

The Perl Journal 2nd

Usenet archive on Google

Yahoo!

websites

ActiveState

module installation resources

cookies

anonymous surfing

privacy issues 2nd 3rd

restricting return destinations 2nd 3rd

CPAN 2nd

Hobbe's Internet Timeline

HTTP redirection

CGI programs

implementing 2nd 3rd 4th

Redirection Based on Browser (Listing 24.1) 2nd 3rd

NOAA (National Oceanic and Atmospheric Administration)

airport queries 2nd 3rd 4th

METAR codes

weather report queries 2nd

PerlMonks

help resources 2nd 3rd

U.S. Census Bureau

zip code queries 2nd 3rd

weights

Earth to Moon conversion 2nd 3rd 4th 5th

while loop 2nd

Sample while Loop (Listing 3.2) 2nd

whitespace

character class

WidgetCo Bike Parts Company

data representation 2nd

binary data 2nd

hierarchical data 2nd

table data 2nd

unstructured data 2nd

table data

email header 2nd

email order taker 2nd 3rd 4th 5th 6th

email order verifier 2nd 3rd 4th 5th 6th

order information 2nd

payment block 2nd

XML data

Bulk Order Sample Document (Listing 19.10) 2nd

Reading Orders in Bulk Using XML (Listing 19.11) 2nd 3rd

Using the Bulk Order Function (Listing 19.12)

wildcards 2nd 3rd 4th

Win32++ * module (CPAN)

Windows

end of line

pathnames

Universal Naming Convention (UNC)

Windows 95/98

command prompt

locking

Perl

installing

preinstalled

Perl documentation 2nd

Windows 95/98/NT

installing modules under 2nd 3rd

Windows NT

command prompt

file locking

Perl

installing

Perl documentation 2nd

Perl preinstalled

Web servers

installing FTP-only CGI program access 2nd

installing local CGI program access

Windows NT Resource Kit

word

character class

word processor

for programming

World Wide Web [See [WWW](#)]

World Wide Web Consortium (W3C)

XML

format development

specifications

writeability

file test

writing

CGI programs

security

programs

functions

WWW

Web browsers

retrieving dynamic pages 2nd 3rd 4th 5th 6th 7th 8th 9th 10th 11th 12th 13th 14th 15th 16th 17th 18th

19th 20th 21st 22nd 23rd 24th 25th 26th 27th 28th

retrieving static pages 2nd 3rd 4th 5th

WWW (World Wide Web

WWW++* module (CPAN)



< Day Day Up >

Index

[SYMBOL] [A] [B] [C] [D] [E] [F] [G] [H] [I] [J] [K] [L] [M] [N] [O] [P] [Q] [R] [S] [T] [U] [V] [W] [X] [Y] [Z]

x (repetition operator)

list context

XML

(extensible Markup Language)

Data Munging with Perl

documents

elements

uses

Even More Convoluted XML Parse in Regular Expressions (Listing 19.8) 2nd 3rd

files

reading via regular expressions 2nd 3rd 4th 5th 6th 7th 8th 9th

reading via XML++Simple 2nd 3rd

Reading the XML File, Almost (Listing 19.5) 2nd

Sample Using XML++Simple to Parse Your Album Document (Listing 19.9) 2nd 3rd

Still Flawed XML Reading (Listing 19.6) 2nd

The Album XML Document, re-arranged (Listing 19.7) 2nd

WidgetCo Bike Parts Company example

Bulk Order Sample Document (Listing 19.10) 2nd

Reading Orders in Bulk Using XML (Listing 19.11) 2nd 3rd

Using the Bulk Order Function (Listing 19.12)

World Wide Web Consortium (W3C)

format development

specifications

XML++Simple

document errors

troubleshooting

returning hashes

troubleshooting

Sample Using XML++Simple to Parse Your Album Document (Listing 19.9) 2nd 3rd

XML files

reading 2nd 3rd

XML++Simple module

weather station application

zip code retrieval 2nd 3rd

 PREV

< Day Day Up >

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

y/// [See [tr///](#)]

Yahoo!

Web site

Your Moon Weight (Listing 6.2) 2nd 3rd

Index

[[SYMBOL](#)] [[A](#)] [[B](#)] [[C](#)] [[D](#)] [[E](#)] [[F](#)] [[G](#)] [[H](#)] [[I](#)] [[J](#)] [[K](#)] [[L](#)] [[M](#)] [[N](#)] [[O](#)] [[P](#)] [[Q](#)] [[R](#)] [[S](#)] [[T](#)] [[U](#)] [[V](#)] [[W](#)] [[X](#)] [[Y](#)] [[Z](#)]

zip codes

 U.S. Census Bureau

 database queries/retrievals 2nd 3rd

 weather station application resources